

SPECIALE

**FRA HADES
OF
DOOM**



BJARKE MØNSTED

PRETENTIOUS
QUOTE.

- FAMOUS PERSON, BORN-DIED

FRA HADES OF DOOM

Author	My Name
Advisor	His Name
Co-Advisor	Her Name



Ting, som KU siger der skal stå her

CMOL

Center for Models of Life

Submitted to the University of Great Justice
June 14, 2014

ACKNOWLEDGEMENTS

Thank you! Thank you all!

CONTENTS

I	Social Fabric Project	1
1	Data Extraction	3
1.1	Temporal Correlations in Activity	4
1.1.1	Influence of phone calls	5
1.1.2	Influence of GPS activity	8
1.1.3	Influence of Bluetooth signal	11
1.2	Extraction of Input Data	15
1.2.1	Simple Call/Text Data	15
1.2.2	Location Data	16
1.2.3	Time Series Analysis	21
1.2.4	Facebook Data	21
1.2.5	Bluetooth Data	21
1.3	Output Data	23
1.4	Linear Discriminant Analysis	23
1.4.1	A measure of separation for projected Gaussians . .	24
1.4.2	Optimizing separation	26
2	Psychological Profiling	31
2.1	SVM	31
2.1.1	Obtaining the Maximally Separating Hyperplane .	33
2.1.2	Generalizing to the non-linear case	35
2.1.3	Implementing a custom weighted radial basis func- tion kernel	41
2.1.4	Statistical subtleties	44
2.2	Decision Trees & Random Forests	45
2.3	Nearest Neighbour-classifiers	48
2.4	Results	49
2.4.1	Big Five Personality Traits	49
2.4.2	Miscellaneous Traits	50

II Wikipedia-based Explicit Semantic Analysis	53
3 ESA	55
3.1 Methods	55
3.1.1 Bag-of-Words	56
3.1.2 Semantic Analysis	56
3.2 Constructing a Semantic Analyser	58
3.2.1 XML Parsing	59
3.2.2 Index Generation	60
3.2.3 Matrix Construction	61
3.3 Applications & Results	62
3.3.1 Trend Discovery and Monitoring	65
3.3.2 Measuring Social Media Impact	67
A Appendix	71
A.1 Social Fabric-related Code	72
A.1.1 Phonetools	72
A.1.2 Code Communication Dynamics	77
A.1.3 Preprocessing	83
A.1.4 Social Fabric Code	96
A.1.5 Lloyd's Algorithm	113
A.1.6 Smallest Enclosing Circle	116
A.2 Source Code for Explicit Semantic Analysis	118
A.2.1 Parser	118
A.2.2 Index Generator	123
A.2.3 Matrix Builder	126
A.2.4 Library for Computational Linguistics	130
A.2.5 Wikicleaner	134
A.2.6 Application Examples	140
Bibliography	147

ENGLISH ABSTRACT

WORDS! SOOOOO MANY WORDS!

DANSK SAMMENFATNING

ORD! SAAAAAAAAAAAA MANGE ORD

Part I

SOCIAL FABRIC PROJECT

PHONE ACTIVITY AND QUANTITATIVE DATA EXTRACTION

THE main objective of part I of this thesis is to investigate behavioural patterns in the phone activities of the participants in the Social Fabric Project, and to predict various traits of the users based only on their phone logs. Throughout the part, I'll provide brief examples of usage for the software I've written to process the large amount of data available and to apply various prediction schemes to it, while the source code itself is included in the appendix.

My first objective was to investigate how various phone activities correlate with each other temporally, i.e. how a given user's probability for e.g. receiving a call increases or decreases around other activities such as moving around physically. This is the topic of section 1.1.

Next, I set out to replicate some recent research results claiming that people's phone activities predict certain psychological traits. In the most general terms, then, the task consists of predicting a collection of numbers or labels denoted Y based on a set of corresponding data points X . The topic of section 1.2 is the extraction of the many-dimensional data points or *feature vectors* X from the phone logs of the participants, while section 1.3 gives a brief description of the psychological traits Y . Finally, section 1.4 derives an often used linear classification method known as Linear Discriminant Analysis or Fisher's Discriminant and provides a discussion of why it fails for the present dataset, which serves to motivate the more sophisticated prediction schemes introduced in chapter 2.

kilder!!!

1.1 Temporal Correlations in Activity

One category of interesting quantities is the predictability of mobile phone behaviour from recorded behaviour at different times, i.e. the influence of certain events deduced from a user's Bluetooth, GPS or call log data on the tendency of some event to happen in the near past or future. A simple example would be to determine how much placing or receiving a call increases or decreases the probability of a user placing or receiving another call in the period following the first call.

This analysis was performed by comparing an 'activity' signal with a 'background' signal in the following fashion: For each user, the time period around each call is sliced into bins and the each of the remaining calls placed in the bin corresponding to the time between the two calls. Once divided by the total number of calls for the user, this is the activity signal. The background is obtained in a similar fashion but comparing each call in a given user's file with calls in the remaining users' files.

This involves repeatedly binning the time around certain events and then determining in which bin to place other events; a situation in which confusion may arise easily and errors may be hard to identify. To accommodate this, I started out by writing a custom array class designed to greatly simplify the binning procedure. This class called features the following:

- Methods to bin the time around a given event and determine determine which bin a given event falls into. This is useful to implement in the class itself as one then avoids having to continually worry about which bin an event fits into, and as it ensures that bin placement errors can only arise in one small piece of code which can then be tested rigorously.
- Attributes that keep track of the number of events that didn't fit into any bin, and of the current centre of the array, which can then be manipulated to move the array and a method to use this to return a normalized list of the bins.

In short, the binarray can be visualized as a collection of enumerated buckets that can be moved so as to center it on some event and then let other events 'drip' into the buckets. The code for this class is included in [A.1.1](#). In general, objects can be converted to byte streams and stored using Python's pickle module, but as that tends to be both slow and insecure, I generally used json to save my objects. This poses a slight problem as some data types, such as tuples, and custom classes in general are not json

serializable. I got around this by writing some recursive helper methods to help store the relevant information about arbitrary nested combinations of some such objects and to help reconstruct said objects again. These are also included in section A.1.1. As an example of usage, the following code constructs a Binarrray, centers it around the present time, and generates a number of timestamps which are then placed in the event. It is then saved to a file using the helper method previously described.

```

from time import time
from random import randint
#Create Binarrray with interval +/- one hour and bin size ten minutes.
ba = Binarrray(interval = 60*60, bin_size = 10*60)
#Center it on the present
now = int(time())
ba.center = now
#Generate some timestamps around the present
new_times = [now + randint(-60*60, 60*60) for _ in xrange(100)]
for tt in new_times:
    ba.place_event(tt)

#Save it
with open('filename.sig', 'w') as f:
    dump(ba, f)

```

This data will be visualized by plotting the relative signal from the activity of some event, such as in- or outgoing calls or texts, over the background (simply A/B) around another type of event hypothesized to trigger the activity.

1.1.1 Influence of phone calls


I first investigated the effects of incoming and outgoing calls as triggers for other phone activities. The call logs were stored in a format where each line represents a hashmap with quantities such as call time or call duration mapping to their corresponding value. Below is an example of one such line, where any personal data have been replaced by a random number or hexadecimal string of equal length.

```


{
  "timestamp": 6335212287,
  "number": "c4bdd708b1d7b82e349780ee1e7875caa600c579",
  "user": "ea42a1dbe422f83b0178d158f154f4",
  "duration": 483,
  "type": 2,
  "id": 45687
}

```

the text logs are similar except for the missing duration entry. Computing the relative signal in Binarrrays centered on each incoming and outgoing




(a) Relative activity of events triggered by incoming calls.



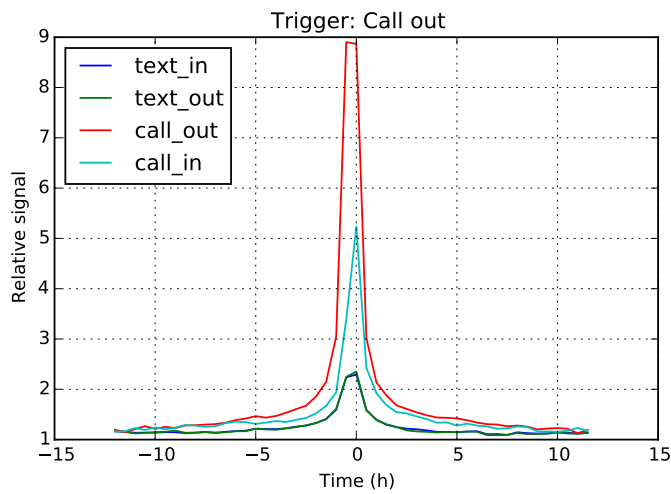
(b) Relative activity of events triggered by outgoing calls.

Figure 1.1: Comparison of the increased activity caused by incoming and outgoing calls over an interval of ± 3 hours around an event with bins of three minutes.

call using bin sizes of three and thirty minutes resulted in the plots shown in figures 1.1 and 1.2, respectively. As the figures clearly show, the all four activities increase significantly for the average user around incoming and outgoing calls.



(a) Relative activity of events triggered by incoming calls.



(b) Relative activity of events triggered by outgoing calls.

Figure 1.2: Comparison of the increased activity caused by incoming and outgoing calls over an interval of ± 12 hours around an event with bins of thirty minutes.

1.1.2 Influence of GPS activity

The raw format of the users' GPS logs looks similar to those of the call and text logs:

```
{
  "timestamp": 8058876274,
  "lon": 6.45051654,
  "user": "0c28e8f4ad9619bca1e5ea4167e10a",
  "provider": "gps",
  "lat": 28.20527041,
  "id": 6429902,
  "accuracy": 39.4
}
```

An analysis similar to that of described in section 1.1.1 was carried out using GPS phone data as triggers. I chose to define a user as being 'active' if they travelled at an average speed of 0.5 m/s between two consecutive GPS log entries, while discarding measurements closely following each other. The reason for this is that the uncertainty on the location measurements could yield false measurements of high average speeds when the measurements are not temporally separated. A lot of the measurements turned out to be grouped somewhat tightly - for instance, approximately 80% of the time intervals were below 100 s. This occurs because the Social Fabric data not only actively records its users' locations with some set interval, but also passively records the location when another app requests it, so when users spend time on apps that need to continually update their position such as Google Maps, a location log entry is written every second. The distribution of intervals between consecutive GPS measurements is shown in figure 1.3. A typical uncertainty on civilian GPS devices is at most 100 m[21], so because I choose to consider a user active if they travel at a mean speed of 0.5 m/s , and based on the time spacings shown in figure 1.3, I chose to discard measurements separated by less than 500 s.

An analysis like that of section 1.1.1 reveals that a user's phone activity is significantly increased around times when they are on the move, as shown in figure 1.4. Note the asymmetry of the signal, especially visible in figure 1.4(a). After a measurement of a user being active, the signal dies off about two and a half hours into the future, whereas it persists much longer into the past. Concretely, this means that people's phone activity (tendency to call or text) becomes uncorrelated with their physical activity after roughly two and a half hours, whereas their tendency to move around is increased for much longer time after calling or texting.

The relative signal in figure 1.4(b) appears to be increasing at around $\pm 24 \text{ h}$, which would seem reasonable assuming people have slightly

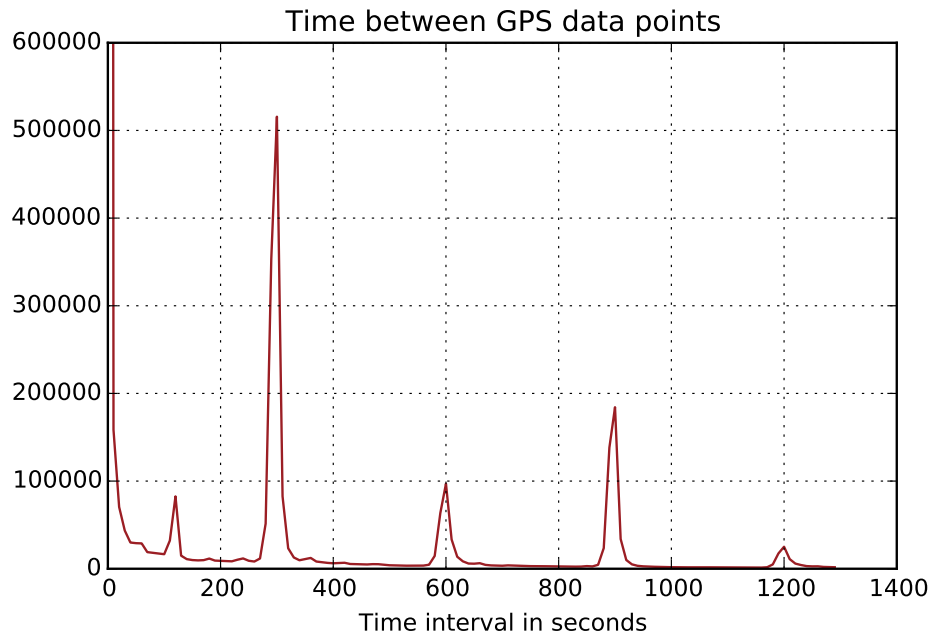



Figure 1.3: Plot of typical temporal spacings between consecutive GPS measurements.

different sleep schedules - if a person is on the move and hence more likely to place a call at time $t = 0$, they're slightly more likely than the general user to be on the move around $t = \pm 24$ h. Figure 1.5 shows the same signal extended to ± 36 h where slight bumps are visible 24 hours before and after activity.



(a) Interval: 5 hours. Bin size: 5 minutes.



(b) Interval: 24 hours. Bin size: 15 minutes.

Figure 1.4: Relative increase of activities triggered by GPS activity.




Figure 1.5: GPS-triggered activity increase over an interval of 36 hours using a bin size of 10 minutes.

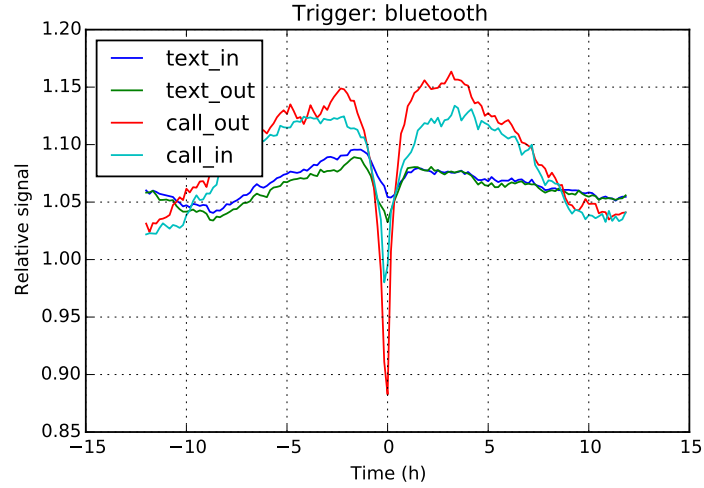
1.1.3 Influence of Bluetooth signal

The following is a randomized entry in a user's bluetooth log.

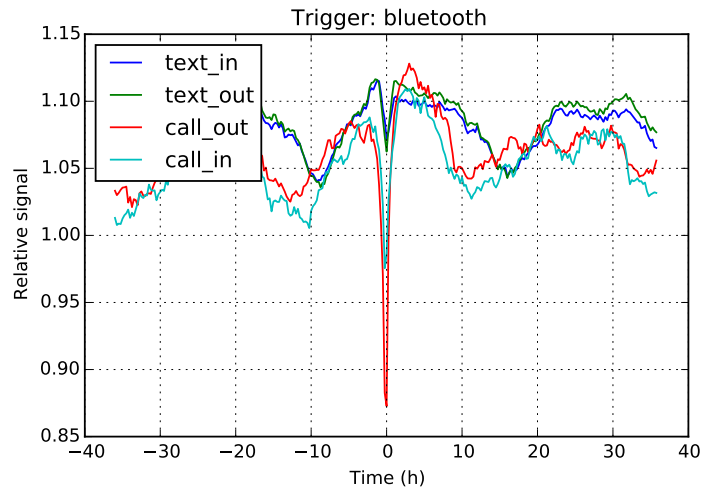
```
{
  "name": "d5306a3672b7a0b8f9696d294ec4b731",
  "timestamp": 6870156680,
  "bt_mac": "1f158ae269d69efa5bb4794ee2a0b2dd68bd3a9badfeaf70f258ad3c74b0c09b",
  "class": 1317046,
  "user": "41cdb7ecaaec3d33391ed063e7fa2",
  "rssi": -76,
  "id": 4139043
}
```

The 'bt_mac' entry is the MAC-address of the device which the Bluetooth receiver in the user's phone has registered, so it is reasonable to assume several different MAC addresses occur at several consecutive timestamps. I call the number of repeated MAC addresses needed for a user to be considered social the 'social threshold'. Figures 1.6, 1.7 and 1.8 show the increased activity around times when users were considered social with a threshold of 1, 2 and 4 repeated pings.

Contrary to the previous analyses, phone activities decreased somewhat when users were social. As stated, each of these analyses were fairly similar, I've only explicitly included the code used to extract and save




(a) Interval: 12 hours, Bin size: 10 minutes.




(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 1.6: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as one repeated signal.




(a) Interval: 12 hours, Bin size: 10 minutes.




(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 1.7: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as two repeated signals.



(a) Interval: 12 hours, Bin size: 10 minutes.



(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 1.8: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as four repeated signals.

Bluetooth data, as well as the code used to load the data and generate figures 1.6 through 1.8. This code is included in section A.1.2.

1.2 Extraction of Input Data

The predictive powers of mobile phone behaviour on the user's psychological profile is currently an area of active research. As part of my thesis work, I have tried to predict the psychological profiles of the SFP participants using various machine learning methods on the available phone logs.

1000 kilder!!!

The software I've written first preprocesses the phone logs to extract various relevant parameters, then collects the parameters and psychological profile scores for each user to serve as input and output, respectively, for the various learning methods. Many of the parameters are chosen following a recent article by de Montjoye et al[6]. The following contains an outline and brief explanation of the extracted parameters.

This section contains a list of the extracted parameters used for psychological profiling along with a brief description of the extraction process when necessary. The preprocessing code is included in section A.1.3.

1.2.1 Simple Call/Text Data

The most straightforward data type is the timestamps from a given user's call/text logs. Six of the parameters used were simply the standard deviation and median of the times between events in the logs for each user's call log, text log, and the combination thereof, excluding time gaps of more than three says on the assumption that it would indicate a user being on vacation or otherwise having a period of telephone inactivity. The entropy S_u of each of the three was also included simply by computing the sum

$$S_u = - \sum_c \frac{n_c}{n_t} \ln_2 \frac{n_c}{n_t}, \quad (1.1)$$

where c denotes a given contact and n_t the total number of interactions, and n_c the number of interactions with the given contact. The number of contacts, i.e. the number of unique phone numbers a given user had contacted by means of calls, texts, and the combination thereof, was also extracted along with the total number of the various kinds of interactions and the contact to interaction ratios. The response rates, defined as the rate of missed calls and incoming texts, respectively, that a given user replied to within an hour, were also determined along with the text

latency defined as the median test response time. Finally the percentage calls and texts that were outgoing was determined as well as the fraction of call interactions that took places during the night, defined as between 22-08.

1.2.2 Location Data

A number of parameters based on the spacial dynamics of the user were also extracted. Among these is the radius of gyration, meaning simply the radius of the smallest enclosing circle enclosing all the registered locations of the user on the given day, and the distance travelled per day. I chose to extract the median and standard deviation of each, filtering out the radii that exceeded 500km so as to keep information about long distance travels in the distance parameter and information about travel within a given region in the radius of gyration parameter.

Cluster Analysis

One parameter which has strong links[6] to psychological traits is the number of locations in which the user typically spends time, and the entropy of their visits to that location. Hence, the task at hand is to identify dense clusters of GPS coordinates for each user. This is a typical example of a task which is very intuitive and quickly approximated by humans, but is extremely computationally expensive to solve exactly. Concretely, the problem of finding the optimal division of n data points into K clusters is formulated as minimizing the 'score' defined as

$$S = \sum_K \sum_{x_n \in C_k} |x_n - c_k|^2, \quad (1.2)$$


where c_k denotes the centroid of the cluster C_k . Each point x_n is assigned to the cluster corresponding to the nearest centroid. The usual way of approaching this problem is to use Lloyd's algorithm, which consists of initializing the centroids randomly, assigning each point to the cluster corresponding to the centroid which is nearest, then moving each centroid to the center of its points and repeating the last two steps until convergence. As this isn't guaranteed to converge on the global minimum of (1.2), the process can be repeated a number of times, keeping only the result with the lowest value of S . I accomplished this by writing a small Python module to perform various variations of Lloyd's algorithm and to produce plots of the resulting clusters. The code is included in section A.1.5.

This allows one to implement Lloyd's algorithm and visualize its result easily, as the code allows automatic plotting of the result from the

algorithm while automatically selecting different colors for the various clusters. As an example, the following code snippet generates 1000 random points, runs Lloyd's algorithm to determine clusters and saves a plot of the results.

```
points = [[random.uniform(-10,10), random.uniform(-10,10)] for _ in xrange(10**3)]
clusters = lloyds(X = points, K = 6, runs = 1)
draw_clusters(clusters = clusters, filename = 'lloyds_example.pdf')
```

This results in the following visualization:



I chose to modify the algorithm slightly on the following basis: Usually, the algorithm takes as its initial centroids a random sample of the data. I'll call this 'sample' initialization. This leads to a greater number of clusters being initialized in the areas with an increased density of data points, meaning that centroids will be highly cluttered at first, 'fighting' over the dense regions of data points then slowly spreading out. A few such iterations are shown in figure 1.9. However, this method is dangerous: The goal is to identify locations in which a user spends much of their time, i.e. in which more than some threshold of their GPS pings originated, and this initialization is likely to 'cut' the more popular locations into several clusters, neither of which contains more data points than the threshold. One example might be the DTU campus, which is a risk of being divided into several locations with too few data points in each, giving the false impression that user doesn't visit the campus that often. To avoid this effect, I implemented another initialization, 'scatter', in which the clusters start out on points select randomly from the entire range of x, y -values in the user's dataset. This turned out to not only solve the problem described above, but also converge much quicker and reach a slightly lower score as define in (1.2). A few such iterations are shown in figure 1.10. The difference in end results for the two methods is exemplified in figure 1.11. While this works great for users who stay in or around Copenhagen, it will cause problems for people who travel a lot. A user who has visited Australia, for instance, will have their initial clusters spread out across




Figure 1.9: A few iterations of Lloyd's algorithm using 'sample' initialization. The axes denote the distance in km to some typical location for the user. Note that clusters are initially cluttered, then slowly creep away from the denser regions.

the globe, and it's highly likely that one them will end up representing all of Denmark. I ended up simply running both versions and keeping the result yielding the highest amount of locations.






Figure 1.10: A few iterations of Lloyd's algorithm using 'scatter' initialization. The axes denote the distance in km to some typical location for the user. Note that clusters are initially randomly spread across the entire range of x, y -values and converge quickly to a local minimum for (1.2).



(a) Sample initialization.



(b) Scattered initialization.

Figure 1.11: Comparison of the final results of the two initialization methods using 100 initial clusters, a threshold of 5% of the data points before a cluster is considered a popular location and running the algorithm 10 times and keeping the best result. Clusters containing more than 5% of the total amount of data points are in color, whereas the remaining points are black dots.

1.2.3 Time Series Analysis

Another interesting aspect to include is what one somewhat qualitatively might call behavioural regularity - some measure of the degree in which a user's phone activities follow a regular pattern. Quantifying this turns out to take a bit of work. First of all, any user's activity would be expected to closely follow the time of day, so the timestamps of each user's outgoing texts and calls are first converted into 'clock times' meaning simply the time a regular clock in Copenhagen's time zone would display at the given time. This process is fairly painless when using e.g. the UTC time standard, which does not observe daylight saving time (DST), but some subtleties arise in countries that do use DST, as this makes the map from Unix/epoch time to clock time 'almost bijective' - when changing *away* from DST, two consecutive hours of unix time map to the same clock time period (02:00 to 03:00), whereas that same clock period is skipped when changing *to* DST. The most commonly used Python libraries for datetime arithmetic accommodate this by including a dst boolean in their datetime objects when ambiguity might arise, however I simply mapped the timestamps to clock times and ignored the fact twice a year, one time bin will artificially contain contributions from one hour too many or few. One resulting histogram is shown in figure 1.12.

Tilføj lidt om
AR-serier når
du har bo-
gen!!!

1.2.4 Facebook Data

Unfortunately, the only available Facebook data was a list of each user's friends, so the only contribution of each user's Facebook log was the number of friends the user had.

1.2.5 Bluetooth Data

I extracted a number of different features from each user's Bluetooth log file. First, I set a threshold for when a given user is considered social, as described in section 1.1.3. I chose to use a threshold of two. I then tried to estimate how much time each user spends in the physical company of others in the following way: for each time stamp in the user's Bluetooth log, I checked if the user was social or not and assumed that this status was the same until the following log entry, unless the delay was more than two hours. The rationale behind this is to avoid skewing the measurements if a user turns off their phone for extended periods of time. Otherwise, e.g. studying with a few friends at DTU, turning off your phone and going on vacation for two weeks would give the false impression that the user




Figure 1.12: Histogram of a user's outgoing calls and texts with a bin size of six hours.

were highly social for a long period of time. I then recorded the fraction of times the user was estimated as being social in this fashion.

Finally, I also wanted some measure of the degrees to which a user's social behaviour follows a pattern. I looked for temporal patterns by fitting AR-series and computing autocorrelation coefficients for each user's social behaviour as described in section 1.2.3. I also chose to compute a 'social entropy' much like (1.1), but weighted by the time the user spends with each acquaintance:

$$E = - \sum_i f_i \ln_2(f_i), \quad (1.3)$$

$$f_i = \frac{\text{time spent with } i}{\sum_j \text{time spent with } j}. \quad (1.4)$$

Note that the denominator of (1.4) is not equal to the total amount of time spent being social, as the contribution from each log entry is multiplied by the number of people present.

1.3 Output Data

The main emphasis of this part of the thesis is on predicting so-called *Big Five* personality traits. This section contains a brief description of those, following[7]. **Extraversion** signifies how extroverted and sociable a person is. People with high extraversion scores are supposed to be more eager to seek the company of others. **Agreeableness** is supposed to be a measure of how sympathetic or cooperative a person is, whereas **conscientiousness** denotes constraint, self discipline, level of organization etc.. **Neuroticism** signify the tendency to experience mood swings, and is complementary to emotional stability. Finally, **Openness**, also called 'openness to experience', or 'inquiring intellect' in earlier works, signifies thoughtfulness, imagination and so on. These five are collectively referred to as the 'big five' or 'OCEAN' after their initials.

In addition to the above, I also had access to a range self-explanatory traits about the participants such as their gender, whether they smoke etc.

1.4 Linear Discriminant Analysis & Gender Prediction

Linear discriminant analysis is basically a dimensionality reduction technique developed by Fisher in 1936 [8] for separating data points into two or more classes. The general idea is to project a collection of data points in n -dimensional variable space, onto the line or hyperplane which maximizes the separation between classes. Representing data points in n -space by vectors denoted x , the objective is to find a vector ω such that separation between the projected data points on it

$$y = \omega^T x \quad (1.5)$$

is maximized.

To break down the derivation of this method, I will first define a convenient distance measure used to optimize the separation between classes, then solve the resulting optimization problem. For clarity, I'll only describe the case of projection of two classes onto one dimension (i.e. using 'line' rather than 'hyperplane' and so on), although the method generalizes easily.

1.4.1 A measure of separation for projected Gaussians

If the projected data points for two classes a and b follow distributions \mathcal{N}_a and \mathcal{N}_b , which are standard Gaussians, $\mathcal{N}_i(x) = \mathcal{N}(x; \mu_i, \sigma_i^2)$, the joint probability distribution for the distance between the projections will be the convolution

$$P(x) = \int_{-\infty}^{\infty} \mathcal{N}_a(y) \cdot \mathcal{N}_b(x - y) dy. \quad (1.6)$$

Computing this for a Gaussian distribution,

$$\mathcal{N}(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}, \quad (1.7)$$

becomes easier with the convolution theorem, which I'll derive in the following.

Denoting convolution by $*$ and Fourier transforms by

$$\mathcal{F}(f) = \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(x) \cdot e^{-i\omega x} dx, \quad (1.8)$$

the convolution theorem is derived as follows:

$$\mathcal{F}(f * g) = \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} \int_{\mathbb{R}^n} f(y) \cdot g(x - y) dy e^{-i\omega x} d\omega, \quad (1.9)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) \int_{\mathbb{R}^n} g(x - y) e^{-i\omega x} dy d\omega, \quad (1.10)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) \int_{\mathbb{R}^n} g(z) e^{-i\omega(z+y)} dz d\omega, \quad (1.11)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) e^{-i\omega y} \int_{\mathbb{R}^n} g(z) e^{-i\omega z} dz d\omega, \quad (1.12)$$

$$\boxed{\mathcal{F}(f * g) = (2\pi)^{n/2} \mathcal{F}(f) \cdot \mathcal{F}(g)}, \quad (1.13)$$

where the factor in front of the usual form of the theorem $\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$ stems from the convention of using angular frequency in Fourier transforms, as in (1.8), rather than

$$\mathcal{F}(f) = \int_{\mathbb{R}^n} f(x) \cdot e^{-2\pi i v x} dx. \quad (1.14)$$

Using this, the convolution of two Gaussians can be calculated as

$$\mathcal{N}_a * \mathcal{N}_b = (2\pi)^{n/2} \mathcal{F}^{-1}(\mathcal{F}(\mathcal{N}_a) \cdot \mathcal{F}(\mathcal{N}_b)). \quad (1.15)$$

The required Fourier transform can be massaged into a nicer form by displacing the coordinate system and cancelling out terms with odd parity:

$$\begin{aligned}
 \mathcal{F}(\mathcal{N}(x)) &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \cdot e^{-i\omega x} dx, \\
 &= \frac{1}{2\pi\sigma} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} e^{-i\omega(x+\mu)} dx, \\
 &= \frac{1}{2\pi\sigma} e^{-i\omega\mu} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} (\cos(\omega x) + i \sin(\omega x)) dx, \\
 &= \underbrace{\frac{1}{2\pi\sigma} e^{-i\omega\mu}}_a \underbrace{\int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} \cos(\omega x) dx}_{I(\omega)}. \tag{1.16}
 \end{aligned}$$

Noting that $I(\omega)$ reduces to an ordinary Gaussian integral at $\omega = 0$ so $I(0) = \sqrt{2\pi}\sigma$, this can be solved with a cute application of Feynman's trick:

$$\begin{aligned}
 \frac{\partial I}{\partial \omega} &= - \int_{-\infty}^{\infty} x e^{-\frac{x^2}{2\sigma^2}} \sin(\omega x) dx, \\
 &= \int_{-\infty}^{\infty} \sigma^2 \frac{\partial}{\partial x} \left(e^{-\frac{x^2}{2\sigma^2}} \right) \sin(\omega x) dx, \\
 &= \sigma^2 e^{-\frac{x^2}{2\sigma^2}} \sin(\omega x) \Big|_{-\infty}^{\infty} - \omega \int_{-\infty}^{\infty} \sigma^2 e^{-\frac{x^2}{2\sigma^2}} \cos(\omega x) dx, \\
 &= -\omega \sigma^2 I(\omega) \Leftrightarrow \\
 I(\omega) &= C e^{-\sigma^2 \omega^2 / 2}, \\
 I(0) &= C = \sqrt{2\pi}\sigma, \\
 I(\omega) &= \sqrt{2\pi}\sigma e^{-\sigma^2 \omega^2 / 2}.
 \end{aligned}$$

Plugging this into (1.16) gives the result

$$\mathcal{F}(\mathcal{N}) = \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu} e^{-\sigma^2 \omega^2 / 2}. \tag{1.17}$$

This can be used in conjunction with (1.13) to obtain

$$\mathcal{F}(\mathcal{N}_a * \mathcal{N}_b) = \sqrt{2\pi} \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu_a} e^{-\sigma_a^2 \omega^2 / 2} \cdot \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu_b} e^{-\sigma_b^2 \omega^2 / 2}, \tag{1.18}$$

$$= \frac{1}{\sqrt{2\pi}} e^{-i\omega(\mu_a + \mu_b)} e^{-(\sigma_a^2 + \sigma_b^2) \omega^2 / 2}, \tag{1.19}$$

which is recognized as the transform of another Gaussian describing the separation with $\mu_s = \mu_a - \mu_b$ and $\sigma_s^2 = \sigma_a^2 + \sigma_b^2$, so taking the inverse Fourier transformation gives the convolution

$$\mathcal{N}_a * \mathcal{N}_b = \frac{1}{\sqrt{2\pi}\sigma_s} e^{-\frac{(x-\mu_s)^2}{2\sigma_s^2}}. \quad (1.20)$$

Hence, a reasonable measure of the separation of two projected distributions is

$$d = \frac{(\mu_a - \mu_b)^2}{\sigma_a^2 + \sigma_b^2}. \quad (1.21)$$

1.4.2 Optimizing separation

To maximize the separation, the numerator and denominator, respectively, of (1.21) can be rewritten in terms of w in the following way (using $\tilde{\mu}_i$ to denote projected means) and simplified by introducing scattering matrices:

$$(\tilde{\mu}_a - \tilde{\mu}_b)^2 = (w^T (\mu_a - \mu_b))^2, \quad (1.22)$$

$$= w^T (\mu_a - \mu_b) (\mu_a - \mu_b)^T w, \quad (1.23)$$

$$= w^T S_B w, \quad (1.24)$$

and

$$\tilde{\sigma}_i^2 = \sum_{y \in i} \frac{1}{N} (y - \tilde{\mu}_i)^2, \quad (1.25)$$

$$= w^T \sum_{y \in i} (x - \mu_i) (x - \mu_i)^T w, \quad (1.26)$$

$$= w^T S_i w, \quad (1.27)$$

$$\tilde{\sigma}_a^2 + \tilde{\sigma}_b^2 = w^T S_W w, \quad (1.28)$$

having introduced the between-class and within-class scatter matrices S_B and S_W by

$$S_B = (\mu_a - \mu_b) (\mu_a - \mu_b)^T, \quad (1.29)$$

$$S_i = \sum_{y \in i} (x - \mu_i) (x - \mu_i)^T, \quad (1.30)$$

$$S_W = S_a + S_b. \quad (1.31)$$

Hence, the objective is to solve

$$\frac{d}{dw} J(w) = \frac{d}{dw} \left(\frac{w^T S_B w}{w^T S_W w} \right) = 0, \quad (1.32)$$

$$\frac{\frac{d[w^T S_B w]}{dw} w^T S_W w - w^T S_B w \frac{d[w^T S_W w]}{dw}}{(w^T S_W w)^2} = 0, \quad (1.33)$$

$$2S_B w \cdot w^T S_W w - w^T S_B w \cdot 2S_W w = 0, \quad (1.34)$$

$$S_B w - \frac{w^T S_B w \cdot S_W w}{w^T S_W w} = 0, \quad (1.35)$$

$$S_B w - S_W w J(w) = 0, \quad (1.36)$$

$$S_B w = S_W w J(w), \quad (1.37)$$

$$S_W^{-1} S_B w = J(w) w. \quad (1.38)$$

The optimal projection vector w^* which satisfies this is

$$w^* = S_W^{-1} (\mu_a - \mu_b). \quad (1.39)$$

Vær lige
sikker på at
du forstår det
her.

Figure 1.13 shows a visualization of this that I generated by drawing (x, y) points from two random distributions to simulate two distinct classes of points. If the distributions are independent and Gaussian, the projections will also form Gaussian distributions, and the probability of a new point belonging to e.g. class a given its coordinates d can be estimated using Bayesian probability

$$P(a|d) = \frac{P(d|a)P(a)}{P(d|a)P(a) + P(d|b)P(b)}, \quad (1.40)$$

where $P(a)$ and $P(b)$ are simply the prior probabilities for encountering the respective classes, and the conditional probabilities, e.g. $P(d|a)$ are simply given by the value of the projected Gaussian $\mathcal{N}(x'; \tilde{\mu}_a, \tilde{\sigma}_a)$ at the projected coordinate x' . In practise, even when the points are not independent or Gaussian, so that (1.40) is not a precise estimate of the probability of the point representing a given class, the class with the highest posteriori according to (1.40) still often turns out to be a good guess.

This method accurately predicted the gender of 79.8% of the participants, which is not particularly impressive as 77.3% of participants were male, so a classifier that assumes that every participant is male would have a comparable success rate. An immediate source of concern is the assumption of linearity: It is possible that the data is ordered in such a way that it is possible to separate data points fairly well based on




Figure 1.13: Two collections of points drawn from independent Gaussian distributions, representing **class a** and **class b**. If the points are projected onto the **straight line**, which is given by (1.39), the separation between the peaks representing the two classes is maximized.

gender or some psychological trait, just not using a linear classifier. As an extreme example of this, figure 1.14 shows a situation where the points representing one class are grouped together in an 'island' in the middle, isolating them from points representing the remaining class. While it is clear that there's a pattern here, a linear classifier fails to predict classes more precisely than their ratio. Support Vector Machines, or SVMs are another linear classification technique which can be generalized to detect patterns like that in figure 1.14. This is described in section 2.1




Figure 1.14: An example of data points representing **class a** are clearly discernible from those of **class b**, yet a linear Fisher classifier fails to predict the classes more precisely than the ratio of **b** to **a**.

PSYCHOLOGICAL PROFILING & MACHINE LEARNING



MACHINE learning is currently a strong candidate for prediction of psychological profiles from phone data. This chapter describes the application of the quantitative data described in section 1.2 and various machine learning schemes, starting with support vector machines (SVMs).

1000 kilder!!!

Uddyb når
der er flere
modeller.

2.1 Support Vector Machines

The purpose of this section is to introduce SVMs and attempt to apply them to the data obtained in 1.2. The introduction is mainly based on introductory texts by Marti Hearst [11] and Christopher Burges [5]. SVMs in their simplest form (*simplest* meaning using a linear kernel, which I'll explain shortly) can be thought of as a slight variation on the linear classifier described in section 1.4. However, where LDA finds a line such that the distribution of the points representing various classes projected onto the line is maximized, the aim of SVMs is to establish the hyperplane that represents the best possible slicing of the feature space into regions containing only points corresponding to the different classes. A simple example of this is shown in figure 2.1. Using labels ± 1 to denote classes, the problem may be stated as trying to guess the mapping from an N-dimensional data space to classes $f : \mathbb{R}^N \rightarrow \{\pm 1\}$ based on a set of training data in $\mathbb{R}^N \otimes \{\pm 1\}$. I'll describe separately the properties of




Figure 2.1: The same points as those shown in figure 1.13, except points in class a and class b are now pictured along with their maximally separating hyperplane.

this maximally separating hyperplane, how it is obtained, and how the method is generalized to non-linear classification problems as the 'island' illustrated in figure 1.14.

The well-known equation for a plane is obtained by requiring that its normal vector \mathbf{w} be orthogonal to the vector from some point in the plane \mathbf{p} to any point \mathbf{x} contained in it:

$$\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}) = 0. \quad (2.1)$$

The left hand side of (2.1) gives zero for points in the plane and positive or negative values when the point is displaced in the same or opposite direction as the normal vector, respectively. Hence, $\text{sign}(\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}))$ may be taken as the decision function. It is clear from (2.1) that the normal vector may be scaled without changing the actual plane (of course the decision function is inverted if a negative value is chosen), so \mathbf{w} is usually rescaled such that

$$\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}) = \mathbf{w} \cdot \mathbf{x} + b = \pm 1, \quad (2.2)$$

for the points that are closest to the separating plane. Those points located on the margin are encircled in figure 2.1. In general then, the meaning of

the sign and magnitude of

$$\mathbf{w} \cdot \mathbf{x} + b \quad (2.3)$$

will be the predicted class and a measure of prediction confidence, respectively, for new data points. Finally, note that \mathbf{w} can be expanded in terms of the data points that are on the margin in figure 2.1 as

$$\mathbf{w} = \sum_i v_i \mathbf{x}_i, \quad (2.4)$$

these \mathbf{x}_i , the position vectors of the margin points in data space, are the ‘support vectors’ that lend their name to the method.

2.1.1 Obtaining the Maximally Separating Hyperplane

Assuming first that it is possible to slice the data space into two regions that contain only points corresponding to one class each, and that the plane’s normal vector has already been rescaled according to (2.2), the following inequalities hold:

$$\begin{aligned} \mathbf{x}_i \cdot \mathbf{w} + b &\geq 1, y_i = +1, \\ \mathbf{x}_i \cdot \mathbf{w} + b &\leq -1, y_i = -1. \end{aligned} \quad (2.5)$$

Multiplying by y_i , both simply become

$$y_i (\mathbf{x}_i \cdot \mathbf{w} + b) - 1 \geq 0. \quad (2.6)$$

The distance between the separating plane and each of the margins in figure 2.1 is $1/|\mathbf{w}|$, so in order to maximize the separation, $|\mathbf{w}|$ must be minimized. For mathematical convenience, $\frac{1}{2}|\mathbf{w}|^2$, rather than $|\mathbf{w}|$ is included in the Lagrangian, which then becomes

$$L = \frac{1}{2}|\mathbf{w}|^2 - \sum_i \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{w} + b - 1), \quad (2.7)$$

must be minimized with the constraints

$$\alpha_i \geq 0, \quad (2.8)$$

$$\frac{\partial L}{\partial \alpha_i} = 0. \quad (2.9)$$

A result from convex optimization theory known as Wolfe Duality[20] states that one may instead maximize the above Lagrangian subject to

$$\nabla_w L = \frac{\partial L}{\partial b} = 0, \quad (2.10)$$

which gives conditions

$$\mathbf{w} = \sum_j \alpha_j y_j \mathbf{x}_j, \quad (2.11)$$

$$\sum_j \alpha_j y_j = 0. \quad (2.12)$$

These can be plugged back into (2.7) to obtain

$$L_D = \frac{1}{2} \sum_i \sum_j \alpha_i y_i \alpha_j y_j \mathbf{x}_i \cdot \mathbf{x}_j - \sum_i \alpha_i y_i \left(\mathbf{x}_i \cdot \sum_j \alpha_j y_j \mathbf{x}_j + b \right) + \sum_i \alpha_i, \quad (2.13)$$

$$L_D = -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_i \alpha_i. \quad (2.14)$$

A problem with this is that eqs 2.5 can only be satisfied in the completely separable case, although it is easy to imagine an example in which a classifier performs well but not flawlessly on the training set. For instance, if two points, one from each class, in figure 2.1 were permuted, the classifier shown in the plot would still do a very good job, but eqs. 2.5 would not be satisfiable, causing the method to fail. This is remedied by introducing slack variables[3]

$$\begin{aligned} \mathbf{x}_i \cdot \mathbf{w} + b &\geq 1 - \xi_i, & y_i &= +1, \\ \mathbf{x}_i \cdot \mathbf{w} + b &\leq -(1 - \xi_i), & y_i &= -1, \\ \xi_i &\geq 0, \end{aligned} \quad (2.15)$$

which allows the algorithm to misclassify. This should come without a cost to the overall Lagrangian, or one would just end up classifying randomly, so a 'cost term', $C \cdot \sum_i \xi_i$ is added as well. The value of C (as well as a few similar parameters which have yet to be introduced) is usually determined experimentally by simply varying it across some space of possible values and choosing the value resulting in the best performance - see for instance figure 2.3. In the case of the misclassification cost parameter C , low values will result in low performance, whereas too large values will result in overfitting. The above can be rewritten exactly as previously, except another set of non-negative Lagrange multipliers μ_i are added to (2.7) to ensure positivity of the ξ_i , resulting in

$$L = \frac{1}{2} |\mathbf{w}|^2 + C \cdot \sum_i \xi_i - \sum_i \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{w} + b - 1 + \xi_i) - \sum_i \mu_i \xi_i. \quad (2.16)$$

This results in the same dual Lagrangian L_D as before, but with an upper bound on the α_i :

$$0 \leq \alpha_i \leq C. \quad (2.17)$$

The methods outlined above can also be used to solve regression, rather than classification, problems[19]. The training data will then be in \mathbb{R}^{N+1} rather than in $\mathbb{R}^N \otimes \{\pm 1\}$, and the value of the decision function in (2.3) is predicted instead of using only its sign. The criterion of correct classification from (2.6) is replaced by the demand that predictions be within some tolerated margin ϵ of the true value of the training point y_i , so (2.5) becomes

$$-\epsilon \leq \mathbf{x}_i \cdot \mathbf{w} + b - y_i \leq \epsilon \quad (2.18)$$

so when slack variables ξ_i and ξ_i^* (for the lower and upper bound, respectively) like in (2.15) are introduced, the Lagrangian from (2.16) becomes

$$L = \frac{1}{2}|\mathbf{w}|^2 + C \sum_i (\xi_i + \xi_i^*), \quad (2.19)$$

with constraints

$$\mathbf{x}_i \cdot \mathbf{w} + b - y_i \geq -(\epsilon + \xi_i), \quad (2.20)$$

$$\mathbf{x}_i \cdot \mathbf{w} + b - y_i \leq \epsilon + \xi_i^*, \quad (2.21)$$

$$\xi_i, \xi_i^* \geq 0. \quad (2.22)$$

The main point to be emphasized here is that the training data \mathbf{x}_i only enter into the dual Lagrangian of (2.14) as inner products. This is essential when extending the SVM model to nonlinear cases, which is the subject of the following section.

2.1.2 Generalizing to the non-linear case

The fact that the data \mathbf{x}_i only occur as inner products in (2.14) makes one way of generalizing to non-linearly separable datasets straightforward: Referring back to figure 1.14, one might imagine bending the plane containing the data points by curling the edges outwards in a third dimension after which a two-dimensional plane could separate the points very well. In general, this means applying some mapping

$$\Phi : \mathbb{R}^l \rightarrow \mathbb{R}^h, \quad h > l, \quad (2.23)$$

to the \mathbf{x}_i (l and h are for low and high, respectively). For example, one could look for a mapping such that the new inner product becomes

$$\Phi(\mathbf{x}_i) \cdot \Phi(\mathbf{x}_j) = (\mathbf{x}_i \cdot \mathbf{x}_j)^2. \quad (2.24)$$

I'll describe the components of each vector separately, so I'm going to change notation to let the subscripts denote coordinates and using \mathbf{x} and \mathbf{y} as two arbitrary feature vectors, where the latter shouldn't be confused with the class labels used earlier. As an example, in two dimensions the above becomes

$$(\mathbf{x} \cdot \mathbf{y})^2 = \left(\sum_{i=1}^2 x_i y_i \right)^2 = x_1^2 y_1^2 + 2x_1 y_1 x_2 y_2 + x_2^2 y_2^2, \quad (2.25)$$

meaning that one possibility for Φ is

$$\Phi : \mathbf{x} \mapsto \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1 x_2 \\ x_2^2 \end{pmatrix} \quad (2.26)$$

This can be generalized to d -dimensional feature vectors and to taking the n 'th power rather than the square using the multinomial theorem:

$$\left(\sum_{i=1}^d x_i \right)^n = \sum_{\sum_{i=1}^d k_i = n} \frac{n!}{\prod_{l=1}^d k_l!} \prod_{j=1}^d x_j^{k_j}, \quad (2.27)$$

where the subscript $\sum_{i=1}^d k_i = n$ simply means that the sum goes over any combination of d non-negative integers k_i that sum to n . I wish to rewrite this slightly for two reasons: to simplify the notation in order to make a later proof more manageable, and to help quantify how quickly the number of dimensions in the output space grows to motivate a trick to avoid these explicit mappings.

As stated, the sum on the RHS of (2.27) runs over all combinations of d integers which sum to n . This can be simplified by introducing a function K , which simply maps

$$K : n, d \mapsto \left\{ \{k\} \in \mathbb{N}^d \mid \sum_{i=1}^d k_i = n \right\}, \quad (2.28)$$

and denoting each of those collections $\{k\}_i$ so each of the coefficients in (2.27) can be written

$$\frac{n!}{\prod_{i=1}^d k_i!} = C_{\{k\}}. \quad (2.29)$$

Then, (2.27) becomes

$$\left(\sum_{i=1}^d x_i \right)^n = \sum_{K(n,d)} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} \quad (2.30)$$

To show how quickly the dimensions of the required embedding space grows, note that the dimension is equal to the number of terms in the sum above, i.e.

$$\dim(\mathbb{R}^h) = |K(n, d)| = \left| \left\{ \{k\} \in \mathbb{N}^d \left| \sum_{i=1}^d k_i = n \right. \right\} \right|, \quad (2.31)$$

which can be computed using a nice trick known from enumerative combinatorics.

Consider the case where $n = 5$ and $d = 3$. $K(5, 3)$ then contains all sets of 3 integers summing to 5, such as 1, 3, 1 or 0, 1, 4. Each of these can be uniquely visualized as 5 unit values distributed into 3 partitions in the following fashion:

$$\begin{array}{c} \circ \mid \circ \circ \circ \mid \circ, \\ \mid \circ \mid \circ \circ \circ \circ, \end{array}$$

and so on. It should be clear that you need n \circ -symbols and $d - 1$ separators. The number of possible such combinations, and hence the dimensionality of the embedding space, is then

$$\binom{n + d - 1}{n} = \frac{(n + d - 1)!}{n!(d - 1)!}. \quad (2.32)$$

This number quickly grows to be computationally infeasible, which motivates one to look for a way to compute the inner product in the embedded space without performing the explicit mapping itself. This is the point of the so-called 'kernel trick', which I'll introduce in the following.

The idea of the kernel trick is that since only the inner products between feature vectors in the embedded space are required, one might as well look for some function K of the original feature vectors which gives the same scalar as the inner product in the embedded space, i.e.

$$K(\mathbf{x}, \mathbf{y}) = \Phi(\mathbf{x}) \cdot \Phi(\mathbf{y}). \quad (2.33)$$

In the polynomial case treated above, the correspondence between the kernel function $K(\mathbf{x}, \mathbf{y})$ and the explicit mapping Φ is straightforward:

$$K(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y})^n, \quad (2.34)$$

$$\Phi(\mathbf{x}) = \sum_{K(n,d)} \sqrt{C_{\{k\}}} \prod_{j=1}^d x_j^{k_j}, \quad (2.35)$$

so that (2.33) is true by the multinomial theorem and the above considerations. However, situations arise in which the explicit mapping Φ isn't directly obtainable, and the correspondence of the kernel function to inner products in higher dimensional spaces is harder to demonstrate. This is the subject of the following section.

Radial Basis Functions

One commonly used kernel function is the RBF, or radial basis function, kernel:

$$K(\mathbf{x}, \mathbf{y}) = e^{|\mathbf{x}-\mathbf{y}|^2/2\sigma}. \quad (2.36)$$

Burges [5] shows that the polynomial kernel is valid, so I'll show how the argument extends to the RBF kernel in the following.

Mercer's condition[18] states that for a kernel function $K(\mathbf{x}, \mathbf{y})$, there exists a corresponding Hilbert space \mathcal{H} and a mapping Φ as specified earlier, iff any L^2 -normalizable function $g(\mathbf{x})$ satisfies

$$\int K(\mathbf{x}, \mathbf{y}) g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \geq 0. \quad (2.37)$$

This can be shown by rewriting (2.36) as

$$K(\mathbf{x}, \mathbf{y}) = e^{(\mathbf{x}-\mathbf{y}) \cdot (\mathbf{x}-\mathbf{y})/2\sigma} = e^{|\mathbf{x}|^2/2\sigma} e^{|\mathbf{y}|^2/2\sigma} e^{-\mathbf{x} \cdot \mathbf{y}/\sigma}, \quad (2.38)$$

and expanding the last term in $(\mathbf{x} \cdot \mathbf{y})$ as

$$e^{-\mathbf{x} \cdot \mathbf{y}/\sigma} = \sum_{i=0}^{\infty} \frac{(-1)^i}{i! \sigma^i} (\mathbf{x} \cdot \mathbf{y})^i, \quad (2.39)$$

but using (2.30) on the dot product gives

$$(\mathbf{x} \cdot \mathbf{y})^i = \left(\sum_{j=1}^d x_j y_j \right)^i = \sum_{K(i,d)} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} y_j^{k_j} \quad (2.40)$$

so the Taylor expansion becomes

$$e^{-\mathbf{x} \cdot \mathbf{y}/\sigma} = \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} y_j^{k_j}, \quad (2.41)$$

which can be plugged back into (2.38) to yield

$$K(\mathbf{x}, \mathbf{y}) = \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} e^{|\mathbf{x}|^2/2\sigma} e^{|\mathbf{y}|^2/2\sigma} \prod_{j=1}^d x_j^{k_j} y_j^{k_j}. \quad (2.42)$$

The underlying reason for these algebraic shenanigans is that (2.42) is clearly separable so that the integral in (2.37) from Mercer's condition becomes

$$\int K(\mathbf{x}, \mathbf{y}) g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \quad (2.43)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \int_{\mathbb{R}^{2d}} e^{|\mathbf{x}|^2/2\sigma} e^{|\mathbf{y}|^2/2\sigma} \prod_{j=1}^d x_j^{k_j} y_j^{k_j} g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \quad (2.44)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \left(\int_{\mathbb{R}^d} e^{|\mathbf{x}|^2/2\sigma} \prod_{j=1}^d x_j^{k_j} g(\mathbf{x}) d\mathbf{x} \right) \cdot \left(\int_{\mathbb{R}^d} e^{|\mathbf{y}|^2/2\sigma} \prod_{j=1}^d y_j^{k_j} g(\mathbf{y}) d\mathbf{y} \right) \quad (2.45)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \left(\int_{\mathbb{R}^d} e^{|\mathbf{x}|^2/2\sigma} \prod_{j=1}^d x_j^{k_j} g(\mathbf{x}) d\mathbf{x} \right)^2 \quad (2.46)$$

$$\geq 0. \quad (2.47)$$

Hence, radial basis functions satisfy Mercer's condition and the kernel described above can be plugged into the dual Lagrangian from (2.14) to obtain

$$L_D = -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j e^{|\mathbf{x}_i - \mathbf{x}_j|^2/2\sigma} + \sum_i \alpha_i, \quad (2.48)$$

which must be maximized subject to the same constraints as earlier. The concrete optimization procedure is complicated and already implemented in most machine learning libraries, so I choose not to go into details with that, but instead to demonstrate the effectiveness of the RBF kernel approach on the non-linear-separable points that were generated earlier. figure 2.2 shows the points again, along with the *decision frontier* i.e. the curve which separates regions in which points are classified into separate classes. The danger of overfitting should be clear from figure 2.2. If the cost of misclassification C and the sharpness of the RBFs, usually denoted by $\gamma = 2/\sigma$ are set sufficiently high, the algorithm will simply end up with a tiny decision boundary around every training point of class a, resulting in flawless classification on the training set, but utter failure on new data. The typical way of evaluating this is to perform k-fold validation, meaning that the available data is k equal parts and the SVM is consecutively trained on $k - 1$ parts and tested on the remaining. A variant of this, which my code uses, is stratified k-fold validation, which only differs in that the data is partitioned so as to keep the ratio between the different classes in each parts as close to equal as possible.




Figure 2.2: The 'island' scenario of figure 1.14 revisited. The points representing **class a** and **class b** have been mapped to a higher-dimensional space in which it is possible to construct a separating hyperplane whose **decision frontier** is also shown.

The γ parameter is often fixed by performing a grid search similar to that discussed earlier. Figure 2.3 shows the resulting heat map from a grid search.




Figure 2.3: Result of a grid search for the optimal combination of values for the cost parameter C and the sharpness γ of the Gaussian kernel function giving optimal values of $C = 0.52$ and $\gamma = 0.12$.

2.1.3 Implementing a custom weighted radial basis function kernel

In general, the features extracted as described in section 1.2 are too numerous for efficient implementations of most machine learning schemes. [6] work around this by determining the linear correlation between each feature and the target values, and then discarding features whose correlation is below some threshold and letting the rest contribute equally in the SVM. They never explicate this threshold but state that they include features 'significantly related' to the target values, which convention would suggest means a threshold of 0.05. [15] have demonstrated significant improvements in SVM performance by assigning variable importances to each feature, so I implemented a modified kernel function where I don't just discard features with linear correlations below some threshold but also assign to each remaining feature a normalized weight given by its correlation with the output variable.

I do this by defining a diagonal matrix \hat{M} where the i, i th element is the linear correlation coefficient between feature i and the target variable, and using the matrix as a metric in the exponent of the usual radial basis

functions so the kernel $K(\mathbf{x}, \mathbf{y})$ becomes

$$e^{\gamma \|\mathbf{x}-\mathbf{y}\|^2} \rightarrow e^{\gamma (\mathbf{x}-\mathbf{y})^T \hat{M} (\mathbf{x}-\mathbf{y})} \quad (2.49)$$

Note that this does not change the validity of the proof I gave in section 2.1.2, so this weighted RBF kernel also satisfies Mercer's condition and may hence be used as a kernel function. It turned out to be nontrivial to write an implementation of this which had a syntax consistent with that of the default methods available in the library I used, and which provided a similarly simple way to grid search over its parameters (C and γ as in the usual RBF kernel SVM along with a threshold parameter). I ended up solving this by writing a metamethod to provide a kernel matrix based on an input list of variable importances as well as a default sharpness parameter denoted γ .

```
def make_kernel(importances, gamma = 1.0):
    '''Returns a weighted radial basis function (WRBF) kernel which can be
    passed to an SVM or SVR from the sklearn module.

    Parameters:
    -----
    importances : list
        The importance of each input feature. The value of element i can mean
        e.g. the linear correlation between feature i and target variable y.
        None means feature will be weighted equally.

    gamma : float
        The usual gamma parameter denoting inverse width of the gaussian used.
    '''
    def kernel(x,y, *args, **kwargs):
        d = len(importances) #number of features
        impsum = sum([imp**2 for imp in importances])
        if not impsum == 0:
            normfactor = 1.0/np.sqrt(impsum)
        else:
            normfactor = 0.0
        #Metric to compute distance between points
        metric = dok_matrix((d,d), dtype = np.float64)
        for i in xrange(d):
            metric[i,i] = importances[i]*normfactor
        #
        result = np.zeros(shape = (len(x), len(y)))
        for i in xrange(len(x)):
            for j in xrange(len(y)):
                diff = x[i] - y[j]
                dist = diff.T.dot(metric*diff)
                result[i,j] = np.exp(-gamma*dist)
        return result
    return kernel
```

With that in place, it was a simple matter to implement classifiers and regressors by inheriting from the default classes and overriding the constructor methods like this:




Figure 2.4: Heat map resulting from a grid search over the parameter space of the linear correlation threshold and the default sharpness γ of the radial basis functions showing optimal values of 0.25 and 0.15 for γ and the threshold, respectively.

```
class WRBFR(svm.SVR):
    '''Weighted radial basis function support vector regressor.'''
    def __init__(self, importances, C = 1.0, epsilon = 0.1,
                  gamma = 0.0):
        kernel = make_kernel(importances = importances, gamma = gamma)
        super(WRBFR, self).__init__(C = C, epsilon = epsilon, kernel = kernel)
```

Figures 2.4 and 2.5 show heat maps resulting from grid searches over the threshold parameter versus γ and C , respectively. Interestingly, the ideal correlation threshold seems to be well above the value used in [6].

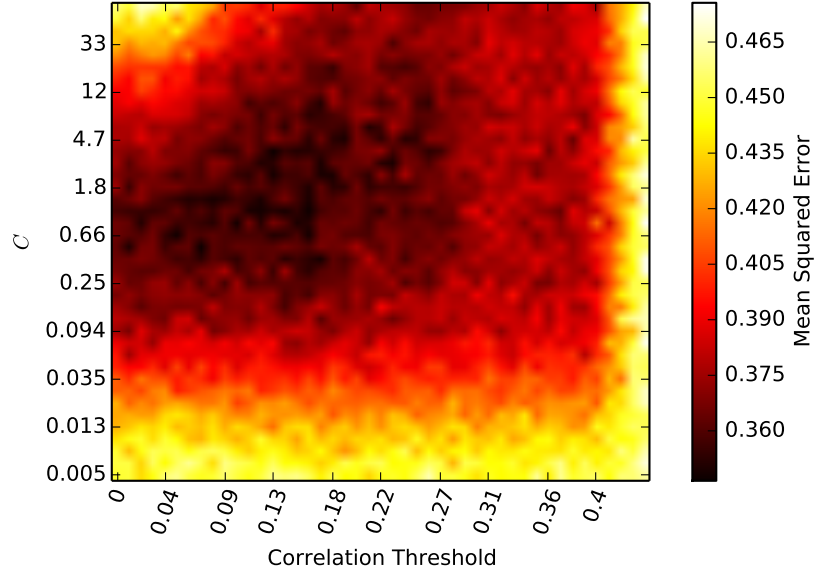


Figure 2.5: Heat map resulting from a grid search over the parameter space of the linear correlation threshold and the cost parameter C of the radial basis functions showing optimal values of 3.6 and 0.13 for γ and the threshold, respectively.

2.1.4 Statistical subtleties

An important note should be made here about some often neglected subtleties relating to uncertainties. Physicists often deal with measurements that can assumed to be independently drawn from a normal distribution $\mathcal{N}(x_i; \mu, \sigma^2)$ due to the central limit theorem. With a large number of measurements n , the standard deviation of a sample

$$\sigma^2 = \frac{1}{N} \sum_i^N (x_i - \mu)^2, \quad (2.50)$$

converges as $N \rightarrow \infty$ to the maximum likelihood, minimum variance unbiased estimator for the true variance of the underlying distribution with unknown mean

$$\hat{\sigma}^2 = \frac{1}{N-1} \sum_i^N (x_i - \mu)^2. \quad (2.51)$$

The standard deviation σ and the width of the underlying gaussian $\hat{\sigma}^2$ can then often be used interchangeably. This tempts some people into the

questionable habit of always assuming that the sample standard deviance can be used as the 68% confidence interval of their results.

When using a K-fold validation scheme, the performance scores for the various folds cannot be assumed to be independently drawn from an underlying distribution, as the test set of one fold is used in the training sets of the remaining folds. In fact, it has been shown [1] that there is no unbiased estimator for the variance of the performance estimated using K-fold validation. However, as K-fold validation is more effective than keeping the test, and training data separate, which can be shown using Jensen's inequality along with some basic properties of expectation values [2], I'll mostly use K-fold regardless. As the standard deviation still provides a qualitative measure of the consistency of the model's performance, I'll still use the sample STD in a usual fashion, such as error bars, unless otherwise is specified, but the reader should keep in mind that these do not indicate precise uncertainties whenever K-fold validation has been involved.

2.2 Decision Trees & Random Forests

Another popular machine learning scheme is that of random forests, which consist of an ensemble of decision trees. A decision tree is a very intuitive method for classification problems which can be visualized as a kind of flow chart in the following fashion. As usual, the problem consists of a set of feature vectors \mathbf{x}_i and a set of corresponding class labels y_i . A decision tree then resembles a flowchart starting at the root of the tree, at each node splitting into branches and finally branching into leaves at which all class labels should be identical. At each node, part of the feature vector is used to split the dataset into parts. This resembles the 'twenty questions' game, in which one participant thinks of a famous person and another attempts to guess who it is by asking a series of yes/no-questions, each one splitting the set of candidates in two parts. In this riddle game and in decision tree learning, there are good and bad questions (asking whether the person was born on March 14th, 1879 is a very bad first question, for instance). There are several ways of quantifying how 'good' a yes/no-question, corresponding to a partitioning of the dataset, is.

One metric for this is the Gini Impurity Index I_G , which is computed by summing over each class label:

$$I_G = \sum_i f_i (1 - f_i) = 1 - \sum_i f_i^2, \quad (2.52)$$

where f_i denotes the fraction of the set the consists of class y_i . Using this as a metric, the best partitioning is the one which results in the largest drop in total Gini impurity following a branching. Another metric is the information gain measured by comparing the entropy before a split with a weighted average of the entropy in the groups resulting from the split. Denoting the fractions of the various classes in the parent group, i.e. before splitting, by f_i and the two child groups by a_i and b_i , the information gain is

$$I_E = - \sum_i f_i \log_2 f_i + \frac{n_a}{N} \sum_i a_i \log_2 a_i + \frac{n_b}{N} \sum_i b_i \log_2 b_i. \quad (2.53)$$

However, if too many such nodes are added to a decision tree, overfitting, i.e. extreme accuracies on training data but poor performance on new data, becomes a problem. This can be remedied by instead predicting with a majority vote, or averaging in the case of regression problems, from an ensemble of randomized decision trees called a random forest. The main merits of random forests are their accuracy and ease of use, and their applications as auxiliary methods in other machine learning schemes, which I'll elaborate on shortly.

The individual trees in a random forest are grown using a randomly selected subset of the training data for each tree. The data used to construct a given tree is referred to as 'in bag', whereas the remaining training data is referred to as 'out of bag' (OOB) for the given tree. At each node, a set number of features is randomly selected and the best possible branching, cf. the above considerations, is determined. The only parameters that must be tweaked manually are the number of trees in the forest, number of features to include in each branching, and the maximum tree depth. While other variables such as the metric for determining branching quality as described above, may be customized, those aren't essential to achieve a decent predictor, which is robust in regard to both overfitting and irrelevant parameters.[4]

There doesn't seem to be a single universally accepted way of adjusting these parameters, so I chose a somewhat pragmatic approach of simply checking how well various choices for each parameter performed on a randomly selected trait. For instance, figure 2.6 shows how well a random forest predicted the tertiles of participants' extroversion as a function of the fraction of available features each tree was allowed to include in each branching. This was done using a large number of trees ($n = 1000$) and using each of the two metrics described earlier. The number of features used pr split doesn't seem to have any significant effect on performance, and as the entropy metric seems to perform as well or slightly better than Gini impurity, I decided to stick to that. A similar plot of the performance




Figure 2.6: Performance of a random forest with 1000 decision trees using various fractions of the available features in each branching using both the **entropy** and the **Gini impurity** metric to determine the optimal branching. The number of features seems not to play a major role, and the **entropy** metric seems to perform slightly better in general.

of various numbers of decision trees in the forest is shown in figure 2.7. The performance seems to stagnate around 100 trees, and remain constant after that, so I usually used at least 500 trees to make sure to get the optimal performance, as runtime wasn't an issue.

The robustness to irrelevant features and overfitting described earlier also plays a role in the application of random forests in conjunction with other schemes. SVMs as described in section 2.1 can be sensitive to irrelevant data[15]. There exist off-the-shelf methods, such as recursive feature elimination (RFE)[10], for use with linear SVMs, but to my knowledge, there is no 'standard' way to eliminate irrelevant features when using a non-linear kernel. However, it is possible to use a random forest approach to obtain the relative importance of the various features and then use only the most important ones in another machine learning scheme which is less tolerant to the inclusion of irrelevant data. The relative importance of feature j can be estimated by first constructing a random forest and evaluating its performance s , then randomly permuting the values of feature j across the training sample and measure the damage it does to

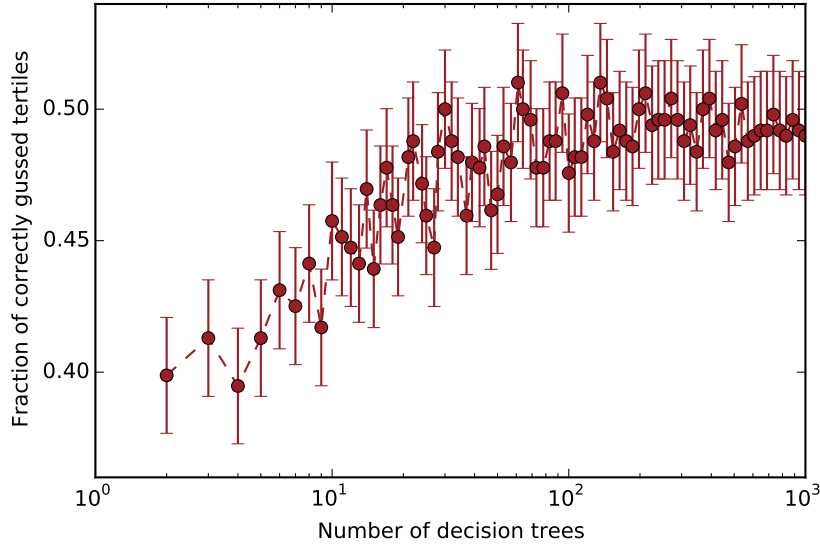


Figure 2.7: Example of a random forest performance versus number of decision trees. Performance seems to increase steadily until about 100 trees, then stagnate.

the performance of the forest by comparing with the permuted score s_p . The ratio of the mean to standard deviation of those differences:

$$w_j = \frac{\langle s - s_p \rangle}{\text{std}(s - s_p)} \quad (2.54)$$

Random forests also provide a natural measure of similarity between data points. Given data points i and j these can be plugged into all their OOB decision trees, or a random subset thereof, and the fraction of the attempts in which both end up at the same leaf can be taken as a measure of similarity. This can be used to generate a proximity matrix for the data points, and it can be used as a metric for determining the nearest neighbours of a new point in conjunction with a simple nearest neighbour classifier.

2.3 Nearest Neighbour-classifiers

Do iiit!

— Skriv en masse om den smart random forest NN-model.

2.4 Results

2.4.1 Big Five Personality Traits

HARJ

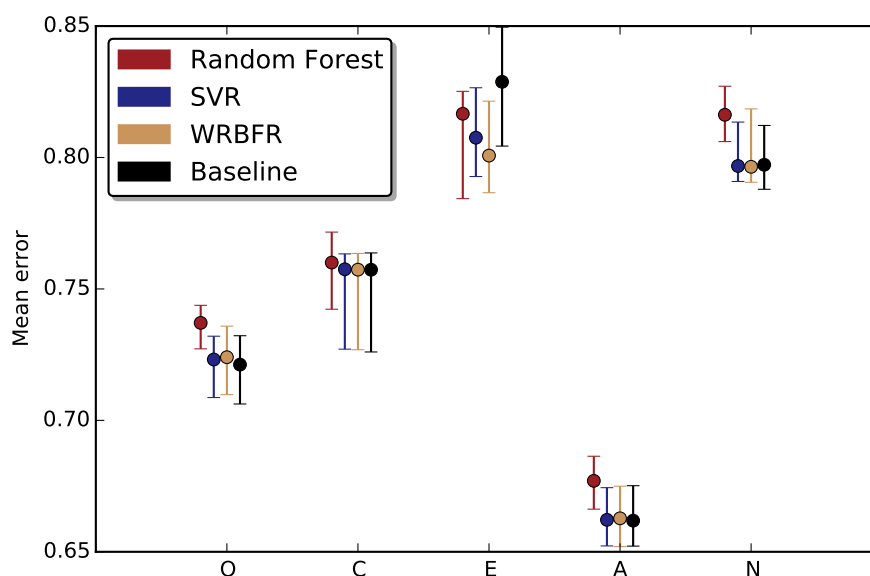


Figure 2.8: Comparison of performance of models using **random forest**, **support vector regression** and **weighted radial basis function regressor** with a **baseline model** which always predicts the mean of the training sample. The y axis shows the mean error of each model and the error bars show the 95-percentile around the median scores obtained by running on 1000 bootstrap samples.

2.4.2 Miscellaneous Traits

Opdater med
ML

Her står der ting om figur 2.9.

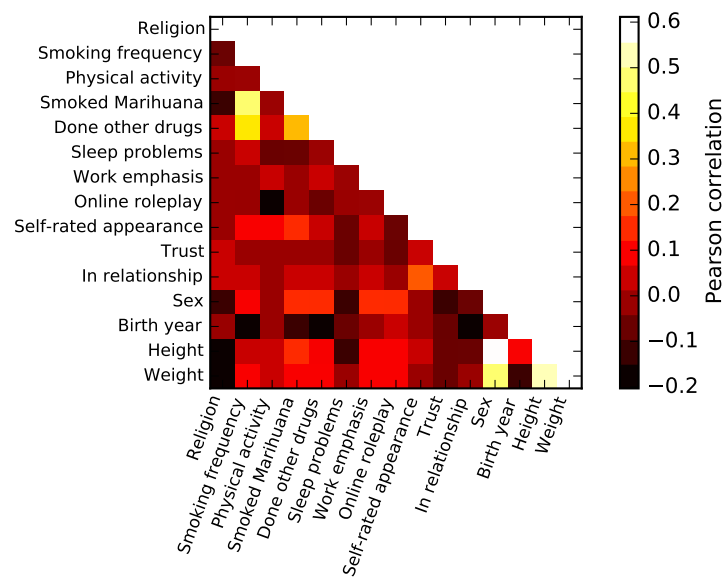


Figure 2.9: Triangle of Pearson correlation coefficients for various miscellaneous traits.

Part II

WIKIPEDIA-BASED EXPLICIT SEMANTIC ANALYSIS

WIKIPEDIA-BASED EXPLICIT SEMANTIC ANALYSIS



NATURAL language processing has long been both a subject of interest and a source of great challenges in the field of artificial intelligence. The difficulty varies greatly depending with the different language processing tasks; certain problems, such as text categorization, are relatively straightforward to convert to a purely mathematical problem, which in turn can be solved by a computer, whereas other problems, such as computing semantic relatedness, necessitates a deeper understanding of a given text, and thus poses a greater problem. This sections aims firstly to give a brief introduction to some of the most prominent techniques used in language processing in order to explain my chosen method of explicit semantic analysis (ESA), and secondly to explain in detail my practical implementation of an ESA-based text interpretation scheme.

ref

3.1 Methods

This section outlines a few methods used in natural language processing, going into some detail on ESA while touching briefly upon related techniques.

3.1.1 *Bag-of-Words*

ref

An example of a categorization problem is the ‘bag of words’ approach, which has seen use in spam filters. Here, text fragments are treated as unordered collections of words drawn from various bags, which in the case of spam filters would be undesired mails (spam) and desired mails (ham). By analysing large amounts of regular mail and spam, the probability of drawing each of the words constituting a given text from each bag can be computed, and the probability of the given text fragment representing draws from each bag can be computed using Bayesian statistics.

More formally, the text T is represented as a collection of words $T = \{w_1, w_2, \dots, w_n\}$, and the probability of T actually representing draws from bag j is hence

$$P(B_j|T) = \frac{P(T|B_j)P(B_j)}{P(T)}, \quad (3.1)$$

$$= \frac{\prod_i P(w_i|B_j)P(B_j)}{\sum_j \prod_i P(w_i|B_j)P(B_j)}, \quad (3.2)$$

ref

for an arbitrary number of bags labelled by j . This method is simple and powerful whenever a text is expected to fall in one of several discrete categories (such as spam filters or language detection). However, for more complex tasks it proves lucrative to attempt instead to assign some kind of meaning to text fragments rather than to consider them analogous to lottery numbers or marbles. This notion of meaning will be elaborated on shortly, as it varies depending on the method of choice, but the overall idea is to ascribe to words a meaning which depends not only on the word itself, but also on the connection between the word and existing repository of knowledge. The reader may think of this as mimicking the reading comprehension of humans. In itself, the word ‘dog’ for instance, contains a mere 24 bits of information if stored with a standard encoding, yet a human reader immediately associates a rich amount of existing knowledge to the word, such as dogs being mammals, related to wolves, being a common household pet, etc. The objective of both explicit and latent semantic analysis is to establish a high-dimensional ‘concept space’ in which words and text fragments are represented as vectors. The difference between explicit and latent semantic analysis is the method used to obtain said concepts, as explained in the following sections.

3.1.2 Semantic Analysis

Salton et al proposed in their 1975 paper *A Vector Space Model for Automatic Indexing*[16] an approach where words and text fragments are mapped

with a linear transformation to vectors in a high-dimensional concept space,

$$T \rightarrow |V\rangle = \sum_i v_i |i\rangle, v_i \in \mathbb{R}, \quad (3.3)$$

where a similarity measure of two texts can be defined as the inner product of two normalized such vectors,

$$S(V, W) = \langle \hat{V} | \hat{W} \rangle = \frac{\sum_i v_i w_i}{\left(\sum_i v_i^2\right) \left(\sum_i w_i^2\right)}, \quad (3.4)$$

and the cosine quasi-distance can be considered as a measure of semantic distance between texts:

$$D(V, W) = 1 - S(V, W). \quad (3.5)$$

This approach has later seen use in the methods of Latent Semantic Analysis (LSA) and Explicit Semantic Analysis (ESA). Both methods can be said to mimic human cognition in the sense that the transformation from (3.3) is viewed as a mapping of a text fragment to a predefined *concept space* and thus, processing of texts relies heavily on external repositories of knowledge.

The difference between LSA and ESA is how the concept space is established. Although I have used solely ESA for this project, I will give an extremely brief overview of LSA for completeness following (Landauer 1998 [12]). LSA constructs its concept space by first extracting every unique word encountered in a large collection of text corpora and essentially uses the leading eigenvectors (i.e. corresponding to the largest eigenvalues) of the word-word covariance matrix as the basis vectors of its conceptual space. This is the sense in which the concept are latent - rather than interpret text in terms of explicit concepts, such as 'healthcare', LSA would discover correlations between words such as 'doctor', 'surgery' etc. and consider that a latent concept. Owing to the tradeoff between performance and computational complexity, only about 400 such vectors are kept[14]. In psychology, LSA has been proposed as a possible model of fundamental human language acquisition as it provides computers a way of estimating e.g. word-word relatedness (a task which LSA does decently) using nothing but patterns discovered in the language it encounters[12].

In contrast, the concepts in ESA correspond directly to certain parts of the external text corpora one has employed to construct a semantic analyser. Concretely, the matrix playing the role of the reduced covariance matrix in LSA has columns corresponding to each text corpus used and

rows corresponding to individual terms or words, with the value of each matrix element denoting some measure of relatedness between the designated word and concept. I have used the English Wikipedia, so naturally each concept consists of an article, although the process could easily be tuned to be more or less fine-grained and associate instead each concept with e.g. a subsection or a category, respectively. Of course, a wholly different collection of texts could also be used - for instance a version of ESA more suited to compare the style or period of literary works could be constructed using a large collection of literature such as the Gutenberg Project. However, with no prior knowledge of the subject matter of the text to be analysed, Wikipedia seems like a good all-round solution considering its versatility and the massive numbers of volunteers constantly keeping it up to date.

I wish to point out two advantages of ESA over LSA. First, it is more successful, at least using the currently available text corpora and implementation techniques. A standard way of evaluating the performance of a natural language processing program is to measure the correlation between relatedness scores assigned to pairs of words or text fragments by the computer and by human judges. In both disciplines, ESA has outperformed LSA since it was first implemented ([9], p 457).

Second, the concepts employed in ESA are directly understandable by a human reader, whereas the concepts in LSA correspond to the leading moments of the covariance matrix. For example, to test whether the first semantic analyser built by my program behaved reasonably, I fed it a snippet of a news article on CNN with the headline *"In Jerusalem, the 'auto intifada' is far from an uprising"*. This returned an ordered list of top scoring concepts as follows: "Hamas, Second Intifada, Palestinian National Authority, Shuafat, Gaza War (2008-09), Jerusalem, Gaza Strip, Arab Peace Initiative, Yasser Arafat, Israel, West Bank, Temple Mount, Western Wall, Mahmoud Abbas", which seems a very reasonable output.

3.2 Constructing a Semantic Analyser

The process of applying ESA to a certain problem may be considered as the two separate subtasks of first a very computationally intensive construction of the machinery required to perform ESA, followed by the application of said machinery to some collection of texts. For clarity, I'll limit the present section to the details of the former subtask while describing its application and results in section 3.3.

The construction itself is divided into three steps which are run in

succession to create the desired machinery. The following is a very brief overview of these steps, each of which is elaborated upon in the following subsections.

1. First, a full Wikipedia XML-dump¹ is parsed to a collection of files each of which contains the relevant information on a number of articles. This includes the article contents in plaintext along with some metadata such as designated categories, inter-article link data etc.
2. Then, the information on each concept (article) is evaluated according to some predefined criteria, and concepts thus deemed inferior are purged from the files. Furthermore, two lists are generated and saved, which map unique concepts and words, respectively, to an integer, the combination of which is to designate the relevant row and column in the final matrix. For example, the concept 'Horse' corresponded to column 699221 in my matrix, while the word 'horse' corresponded to row 11533476.
3. Finally, a large sparse matrix containing relevance scores for each word-concept pair is built and, optionally, pruned (a process used to remove 'background noise' from common words as explained in section 3.2.3)

These steps are elaborated upon in the following.

3.2.1 XML Parsing

The Wikipedia dump comes in a rather large (~50GB unpacked for the version I used) XML file which must be parsed to extract each article's contents and relevant information. This file is essentially a very long list of nested fields where the data type in each field is denoted by an XML tag, such as `<text> blabla </text>`. A very simplified example of a field for one Wikipedia article is shown in 3.1. The content of each field has already been sanitised by Wikipedia so that if for instance the symbol '<' is entered into an article, it is instead represented as '<' in the XML file. To this end, I wrote a SAX parser, which processes the dump sequentially to accommodate its large size. When running, the parser walks through the file and sends the various elements it encounters to a suitable processing function depending on the currently open XML tag. For example, when a 'title' tag is encountered, a callback method is triggered which assigns

¹These are periodically released at <http://dumps.wikimedia.org/enwiki/>

```

<page>
  <title>Horse</title>
  <ns>0</ns>
  <id>12</id>
  <revision>
    <id>619093743</id>
    <parentid>618899706</parentid>
    <timestamp>2014-07-30T07:26:05Z</timestamp>
    <contributor>
      <username>Eduen</username>
      <id>7527773</id>
    </contributor>
    <text xml:space="preserve">
      ===Section title===
      [[Image:name_of_image_file|Image caption]]
      Lots of educational text containing, among other things links to [[other article|text to display]].
    </text>
    <sha1>n57mnhtthxpxq1nanak3zhmmmc622</sha1>
    <model>wikitext</model>
    <format>text/x-wiki</format>
  </revision>
</page>

```

Snippet 3.1: A simplified snippet, of a Wikipedia XML dump.

a column number to the article title and adds it to the list of processed articles. For the callback method for processing the 'text' fields, I used a bit of code from the pre-existing Wikiextractor project to remove Wiki markup language (such as links to other articles being displayed with square brackets and the like) from the content. This code is included in section A.2.5. The remainder of the parser is my work, and is included in section A.2.1. Throughout this process, the parser keeps a list of unique words encountered as well as outgoing link information for each article. These lists, along with the article contents, are saved to files each time a set number of articles have been processed. The link information is also kept in a hashmap with target articles as keys and a set articles linking to the target as values. The point of this is to reduce the computational complexity of the link processing as detailed in the following section.

3.2.2 Index Generation

The next step reads in the link information previously saved and adds it as ingoing link information in the respective article content files. The point of this approach is that link information is initially saved as a hashmap so the link going to a given article can be found quickly, rather than having to search for outgoing links in every other article to determine the ingoing

links to each article, which would be of $O(n^2)$ complexity.

Following that, articles with sufficiently few words and/or ingoing/-outgoing links are discarded and index lists for the remaining articles and words are generated to associate a unique row/column number with each word/concept pair. The code performing the step described here is included in section A.2.2.

3.2.3 Matrix Construction

This final step converts the information compiled in the previous steps into a very large sparse matrix. The program allows for this to be done in 'chunks' in order to avoid insane RAM usage. Similarly, the matrix is stored in segments with each file containing a set number of rows in order to avoid loading the entire matrix to interpret a short text.

The full matrix is initially constructed using a DOK (dictionary of keys) sparse format in which the i, j th element simply counts the number of occurrences of word i in the article corresponding to concept j . This is denoted $\text{count}(w_i, c_j)$. The DOK format as a hashmap using tuples (i, j) as keys and the corresponding matrix elements as values and is the fastest format available for element-wise construction. The matrix is subsequently converted to CSR (compressed sparse row) format, which allows faster operations on rows which performs much quicker when computing TF-IDF (term frequency - inverse document frequency) scores and extracting concept vectors from words, i.e. when accessing separate rows corresponding to certain words.

Each non-zero entry is then converted to a TF-IDF score according to

$$T_{ij} = \left(1 + \ln\left(\text{count}(w_i, c_j)\right)\right) \ln\left(\frac{n_c}{df_i}\right), \quad (3.6)$$

where n_c is the total number of concepts and

$$df_i = |\{c_k, w_i \in c_k\}| \quad (3.7)$$

is the number of concepts whose corresponding article contains the i th word. Thus, the first part of (3.6), $1 + \ln(\text{count}(w_i, c_j))$ is the *text frequency* term, as it increases with the frequency of word i in document j . Similarly, $\ln(\frac{n_c}{df_i})$ in (3.6) is the *inverse document frequency* term as it decreases with the frequency of documents containing word i . Thus, the TF-IDF score as somewhat complement to entropy in that it goes to zero as the fraction of documents containing word i goes to 1, and takes its highest values if

Giver det
mening?

word i occurs with high frequency in only a few documents[13]. While (3.6) is not the only expression to have those properties, empirically it tends to achieve superior results in information retrieval[17].

Each row is then L^2 normalized (divided by their Euclidean norm):

$$T_{ij} \rightarrow \frac{T_{ij}}{\sqrt{\sum_i T_{ij}^2}}. \quad (3.8)$$

Finally, each row is pruned to reduce spurious associations between concepts and articles with a somewhat uniform occurrence rate. This was done in practice by following the pragmatic approach of Gabrilovich [9] of sorting the entries of each row, move a sliding window across the entries, truncating when the falloff drops below a set threshold and finally reversing the sorting. The result of this step is the matrix which computes the interpretation vectors as described in 3.1.2. The code is included in section A.2.3.

3.3 Applications & Results

Having constructed a necessary machinery, I wrote a small Python module to provide an easy-to-use interface with the output from the computations described earlier. The code for this is included in section A.2.4. The module consists mainly of a `SemanticAnalyser` class, which loads in the previously mentioned index lists and provides methods for various computations such as estimating the most relevant concepts for a text, determining semantic distance etc. For example, the following code will create a semantic analyser instance and use it to guess the topic of the input string:

```
sa = SemanticAnalyser()
sa.interpret_text("Physicist from Austria known for the theory of relativity")
```

This returns a sorted list of the basis concepts best matching the input string, where the first element is of course 'Albert Einstein'. The `SemanticAnalyser` class contains equally simple methods to interpret a target text file or keyboard input, to calculate the semantic similarity or cosine distance between texts, and to compute interpretations vectors from a text.

The same module contains a `TweetHarvester` class which I wrote in order to obtain a large number of tweets to test the semantic analyser on, as tweets are both numerous and timestamped, which allows investigations of the temporal evolution of tweets matching a given search term. The

Hearst nævner noget offentligt tilgængeligt Reuters-data som folk øver tekstklassifikation på. Det kunne være ret sjovt. Det kunne være sjovt at lave 'semantisk nearest neighbor'

TweetHarvester class provides an equally simply interface - for instances, the 100 most recent tweets regarding a list of companies can be mined and printed by typing

```
terms = ['google', 'carlsberg', 'starbucks']
th = TweetHarvester()
th.mine(terms, 100)
print th.harvested_tweets
```

in addition to actively 'mining' for tweets matching a given query, the class can also passively 'listen' for tweets while automatically saving its held tweets to a generated date-specific filename once a set limit of held tweets is exceeded:

```
th = TweetHarvester(tweets_pr_file = 100)
th.listen(terms)
```

The downloaded tweets are stored as tweet objects which contain a built in method to convert to a JSON-serializable hashmap, an example of which is provided in 3.2. As can be seen in the example, the tweet object contains not only the tweets textual content but also a wide range of metadata such as the hashtags contained in the tweet, users mentioned, time of creation, language etc. Using this, I wrote a script that harvested tweets mentioning some selected brand names for about 5 months which resulted in a massive dataset of 370 million tweets. This section demonstrates some applications of the methods outlined earlier to this dataset.

```

{u'contributors': None,
 u'coordinates': None,
 u'created_at': u'Wed Jun 03 12:16:23 +0000 2015',
 u'entities': {u'hashtags': [],
               u'symbols': [],
               u'urls': [],
               u'user_mentions': [{u'id': 630908729,
                                   u'id_str': u'630908729',
                                   u'indices': [0, 12],
                                   u'name': u'Alexandra White ',
                                   u'screen_name': u'lexywhite86'}]}},
 u'favorite_count': 0,
 u'favorited': False,
 u'geo': None,
 u'id': 606071930282745857L,
 u'id_str': u'606071930282745857',
 u'in_reply_to_screen_name': u'lexywhite86',
 u'in_reply_to_status_id': 605991263129714688L,
 u'in_reply_to_status_id_str': u'605991263129714688',
 u'in_reply_to_user_id': 630908729,
 u'in_reply_to_user_id_str': u'630908729',
 u'is_quote_status': False,
 u'lang': u'en',
 u'metadata': {u'iso_language_code': u'en', u'result_type': u'recent'},
 u'place': None,
 u'retweet_count': 0,
 u'retweeted': False,
 u'source': u'<a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone
 </a>',
 u'text': u'@lexywhite86 if carlsberg did mornings \U0001f602',
 u'truncated': False,
 u'user': {u'contributors_enabled': False,
           u'created_at': u'Thu Apr 05 13:48:00 +0000 2012',
           u'description': u'',
           u'favourites_count': 169,
           u'follow_request_sent': False,
           u'followers_count': 110,
           u'friends_count': 268,
           u'geo_enabled': False,
           u'id': 545986280,
           u'id_str': u'545986280',
           u'lang': u'en',
           u'listed_count': 0,
           u'location': u'',
           u'name': u'Robert Murphy',
           u'notifications': False,
           u'protected': False,
           u'screen_name': u'7Robmurphy',
           u'statuses_count': 1650,
           u'time_zone': None,
           u'url': None,
           u'utc_offset': None,
           u'verified': False}}

```

Snippet 3.2: Example of a downloaded tweet.

3.3.1 Trend Discovery and Monitoring

A simple application of the methods outlined above is a purely exploratory analysis. As an example, using code included in section A.2.6, I extracted the 10 concepts most closely related to every tweet mentioning Carlsberg for each week over a period of a few months. This gave a bar chart like figure 3.1 for each of those weeks. Figure 3.1 implies some significant relation between Carlsberg and Raidió Teilifís Éireann blablabla!!!

Erstat med det med den irske radio

Another application is to manually select a few concepts of interest and then monitor how the number of tweets strongly related to them develops over time. For example I did an exploratory analysis as described above and occasionally saw the concept 'Kim Little' surface. This turned out to be a young football player signed to a team sponsored by Carlsberg, and when I produced a series of bar charts like that shown in figure 3.2 and combined them into an animation, the bar representing Little tended to peak around dates on which news stories featuring her or her team could be found. Of course these method are only effective insofar the concepts of interest actually have an associated Wikipedia article. However, these concepts are merely the basis vectors of the semantic space in which an arbitrary text can be represented, so this procedure can be extended fairly elegantly to estimate the impact of e.g. a press release on social media.

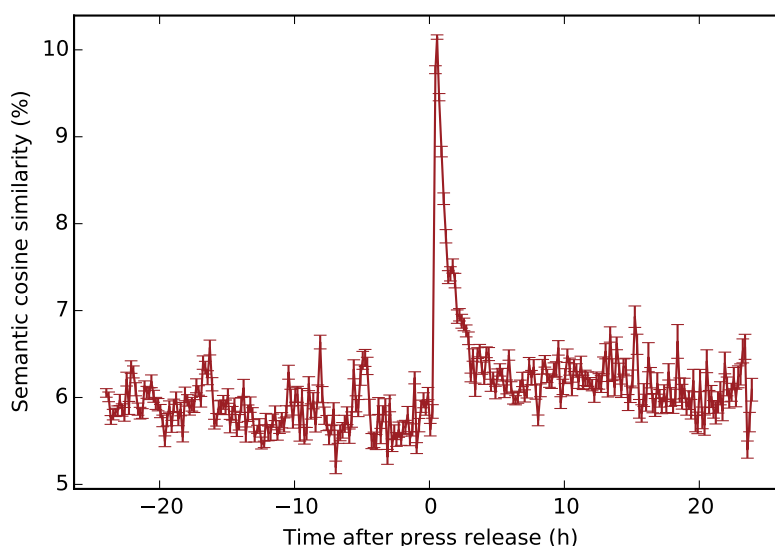


Figure 3.1: Sæt det rigtige ind når du kommer hjem!!!

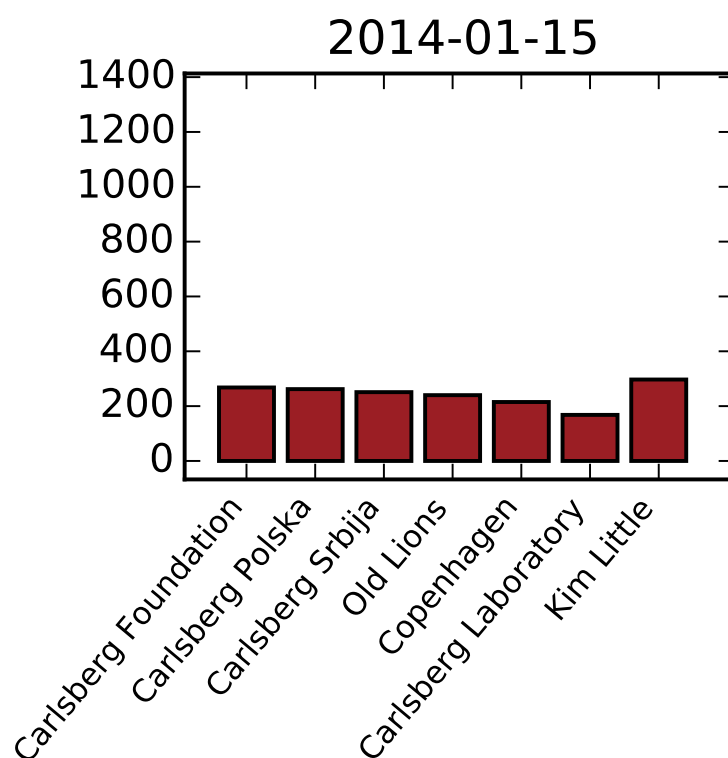


Figure 3.2: Bar chart showing the concepts most related to tweets about Carlsberg for January 15th 2014.

This is the subject of section [3.3.2](#).

3.3.2 Measuring Social Media Impact

One possible application of the software I wrote to perform ESA is to provide a quantitative measure of the impact of some event on social media. For instance, some corporation or organization might be interested in learning precisely effectively a campaign or press release has reached the general public as it is expressed by social media. While my tweet harvesting script was running, Google posted a blog entry² on their experiments with producing psychedelic images with deep neural network called *Deep Dreams* which received widespread attention on social media.

Using the text from the Deep Dreams blog post as a reference text, I converted each English tweet about Google from a period around the blog post to a semantic vector using (3.3) and computed their semantic cosine similarity with the reference text as described by (3.4). Figure 3.3 shows this behaviour in the time around the release of the blog post. A measure of impact should of course take into account the rate at which new tweets occur. Just as one would expect, not only semantic relatedness, but also tweet frequency increase drastically around an interesting event. To make the signal from figure 3.3 independent of the normal tweet frequency, yet sensitive to changes in it, I first find the difference in semantic relatedness between the signal and a reference signal obtained long before the event to be investigated, then modulate the signal by multiplying each bin with the activity A given by the ratio between the current and baseline tweet frequency. As normal Twitter activity tends to vary greatly over a 24 hour period, I obtained a baseline tweet frequency by computing the average tweet rate, weighted by the activity, for each discreet time interval of the day, and then repeating that signal into the period after the event of interest. The baseline along with the observed tweet rate is shown in figure 3.4(a). Modulating in this fashion and computing the difference between the observed signal and the baseline gives an expression of the 'impact rate' of a given time bin, which can be integrated to obtain a measure of the total impact which is independent of the average tweet rate before publication, which is practical if one wishes to compare e.g. the success of media campaigns for companies or organizations of varying sizes. This is shown in figure 3.5. If one wishes to include the number of tweets posted in the impact measure, the result can be computed and visualized nicely by again obtaining a baseline cosine similarity from a period prior to publication and then considering e.g. the cumulative deviation from that following the release, as shown in figure 3.6.

²The original blog post is available at <http://googleresearch.blogspot.dk/2015/06/inceptionism-going-deeper-into-neural.html>




Figure 3.3: Graph of the mean semantic cosine similarity of tweets around the Deep Dreams press release. There is a clear peak around $t = 0$ and the similarity appears increased in the period following the release. The error bars represent the true standard deviation of the sample mean σ / \sqrt{N} for each time bin, each representing a 10-minute interval.




Figure 3.4: Tweet activity around the time the Deep Dreams blog entry was posted. The **signal tweet rate** increases with a factor of about 5 relative to the **baseline rate** around the post. Figure 3.4(b) shows the resulting **modulation factor** to the signal from in figure 3.3.




Figure 3.5: A measure of the total impact on Twitter by a press release can be obtained by integrating over the difference between a signal corresponding of the average tweet cosine similarity modulated by the relative activity, and a corresponding baseline. This measure does not depend on the typical Twitter activity and so can be used to compare the effectiveness of campaigns of varying size.




Figure 3.6: The cumulative deviation from the mean cosine similarity between the Deep Dreams post and tweets in the period around its release.



APPENDIX

A.1 Social Fabric-related Code

This section contains the code referred to in part I of the thesis.

A.1.1 Phonetools

```

1  # -*- coding: utf-8 -*-
2  """
3  Created on Tue Dec 02 15:12:29 2014
4
5  @author: Bjarke
6  """
7
8  import datetime
9  from pkg_resources import resource_filename
10 from ast import literal_eval as LE
11 import numpy as np
12 import json
13 import os
14
15 def _global_path(path):
16     '''Helper method to ensure that data files belonging to the social_fabric
17     module are available both when importing the module and when running
18     individual parts from it for testing.
19     usage: Always use _global_path(somefile) rather than somefile'''
20     if __name__ == '__main__':
21         return path
22     else:
23         return resource_filename('social_fabric', path)
24
25 def make_filename(prefix, interval, bin_size, ext=''):
26     '''Returns a filename like "call_out-intervalsize-binsize_N", where N is
27     an int so files aren't accidentally overwritten.
28     I've been known to do that.'''
29     stem = "%s-int%s-bin%s" % (prefix, interval, bin_size)
30     n = 0
31     ext = '.' + ext
32     attempt = stem + ext
33     while os.path.isfile(attempt):
34         n += 1
35         attempt = stem + "_" + str(n) + ext
36     return attempt
37
38 def unix2str(unixtime):
39     '''Converts timestamp to datetime object'''
40     dt = datetime.datetime.fromtimestamp(unixtime)
41     return str(dt)
42
43 #Converts bluetooth MAC-addresses to users
44 with open(_global_path('user_mappings/bt_to_user.txt'), 'r') as f:
45     bt2user = LE(f.read())
46
47 #Converts phone 'number' code to users
48 with open(_global_path('user_mappings/phonenumbers.txt'), 'r') as f:
49     number2user = LE(f.read())
50
51 #Converts IDs from psychological profile to users

```

```

52 with open(_global_path('user_mappings/user_mapping.txt'), 'r') as f:
53     psych2user = {}
54     #This is in tab-separated values, for some reason.
55     for line in f.read().splitlines():
56         (ID, usercode) = line.split('\t')
57         assert ID.startswith('user_')
58         psych2user[ID] = usercode
59
60 #Converts users to info on their psych profiles
61 with open(_global_path('user_mappings/user2profile.json'), 'r') as f:
62     user2profile = json.load(f)
63
64 def is_valid_call(call_dict):
65     '''Determine whether an entry is 'valid', i.e. make sure user isn't
66     calling/texting themselves, which people apparently do...'''
67     caller = call_dict['user']
68     try:
69         receiver = number2user[call_dict['number']]
70     except KeyError:
71         receiver = None
72     return caller != receiver
73
74 def readncheck(path):
75     '''Reads in all valid call info from the file at path'''
76     try:
77         with open(path, 'r') as f:
78             raw = [LE(line) for line in f.readlines()]
79     except IOError:
80         return [] #no file :(
81     #File read. Return proper calls
82     return [call for call in raw if is_valid_call(call)]
83
84 class Binarray(list):
85     '''Custom array type to automatically bin the time around a set center
86     and place elements in each bin.
87     The array can be centered using Binarray.center = <some time>.
88     After centering, timestamps can be placed in bins around the center with
89     Binarray.place_event(<some other time>).'''
90
91     def __init__(self, interval = 3*60*2, bin_size = 3*60, center = None,
92                 initial_values = None):
93         '''
94         Args:
95         -----
96         interval : int
97             Total number of seconds covered by the Binarray.
98
99         bin_size : int
100             Width of each bin measured in seconds. The total interval must be
101             an integer multiplum of the bin size.
102
103         center : int
104             Where to center the Binarray. If the array is centered at time t,
105             any event placed in it will placed in a bin depending on how long
106             before or after t the event occurred.
107
108         initial_values : list
109             List of values to start the Binarray with. Default is zeroes.'''
110
111         #Make sure interval is an integer multiplum of bins
112         if not interval % bin_size == 0:

```

```

114         suggest = interval - interval % bin_size
115         error = "Interval isn't an integer multiple of bin size. \
116             Consider changing interval to %s." % suggest
117         raise ValueError(error)
118
119     #Set parameters
120     self.bin_size = bin_size
121     self.interval = interval
122     self.size = 2*int(interval/bin_size)
123     self.centerindex = int(self.size//2)
124     self.center = center
125     #Keep track of how many events missed the bins completely
126     self.misses = 0
127     #Call parent constructor
128     if not initial_values:
129         startlist = [0]*self.size
130     else:
131         if not len(initial_values) == self.size:
132             msg = '''Array of start value must have length %d. Tried to
133                 instantiate with length of %d.''' % (self.size,
134                                                         len(initial_values))
135             raise ValueError(msg)
136         startlist = initial_values
137     super(Binarray, self).__init__(startlist)
138
139
140     def place_event(self, position):
141         '''Places one count in the appropriate bin if event falls within
142         <interval> of <center>. Returns True on success.'''
143         if self.center == None:
144             raise TypeError('Center must be set!')
145         delta = position - self.center
146         #Check if event is outside current interval
147         if np.abs(delta) >= self.interval:
148             self.misses += 1
149             return False
150         #Woo, we're in the correct interval
151         index = int(delta//self.bin_size) #relative to middle of array
152         self[self.centerindex + index] += 1
153         return True
154
155     def normalized(self):
156         '''Returns a normalized copy of the array's contents.'''
157         events = sum(self) + self.misses
158         #Use numpy vectorized function for increased speed.
159         f = np.vectorize(lambda x: x*1.0/events if events else 0)
160         return f(self)
161
162     def _todict(self):
163         '''Helper method to allow dumping to JSON format.'''
164         attrs = ['misses', 'interval', 'bin_size']
165         d = {att: self.__getattr__(att) for att in attrs}
166         d['values'] = list(self)
167         d['type'] = 'binarray'
168         return d
169
170
171     def _dumphelper(obj):
172         '''Evil recursive helper method to convert various nested objects to
173         a JSON-serializable format.
174         This should only be called by the dump method!'''
175         if isinstance(obj, Binarray):

```

```

176         d = obj._todict()
177         return _dumpHelper(d)
178     elif isinstance(obj, tuple):
179         hurrayimhelping = [_dumpHelper(elem) for elem in obj]
180         return {'type': 'tuple', 'values': hurrayimhelping}
181     elif isinstance(obj, dict):
182         temp = {'type': 'dict'}
183         contents = [{'key': _dumpHelper(key), 'value': _dumpHelper(value)}
184                     for key, value in obj.iteritems()]
185         temp['contents'] = contents
186         return temp
187     #Do nothing if obj is an unrecognized type. Let JSON raise errors.
188     else:
189         return obj
190
191 def _hook(obj):
192     '''Evil recursive object hook method to reconstruct various nested
193     objects from a JSON dump.
194     This should only be called by the load method!'''
195     if isinstance(obj, (unicode, str)):
196         try:
197             return _hook(LE(obj))
198         except ValueError: #happens for simple strings that don't need eval
199             return obj
200     elif isinstance(obj, dict):
201         if not 'type' in obj:
202             raise KeyError('Missing type info')
203         if obj['type'] == 'dict':
204             contents = obj['contents']
205             d = {_hook(e['key']): _hook(e['value']) for e in contents}
206             #Make sure we also catch nested expressions
207             if 'type' in d:
208                 return _hook(d)
209             else:
210                 return d
211         elif obj['type'] == 'binarray':
212             instance = Binarray(initial_values = obj['values'],
213                                bin_size = obj['bin_size'],
214                                interval = obj['interval'])
215             instance.misses = obj['misses']
216             for key, val in obj.iteritems():
217                 if key == 'values':
218                     continue
219                 instance.__setattr__(key, val)
220             return instance
221         elif obj['type'] == 'tuple':
222             #Hook elements individually, then convert back to tuple
223             restored = [_hook(elem) for elem in obj['values']]
224             return tuple(restored)
225         else:
226             temp = {}
227             for k, v in obj.iteritems():
228                 k = _hook(k)
229                 temp[k] = _hook(v)
230             return temp
231     #
232     #Do nothing if obj is an unrecognized type
233     else:
234         return obj
235
236 def load(file_handle):

```

```
238     '''Reads in json serialized nested combinations of dicts, binarrays
239     and tuples.'''
240     temp = json.load(file_handle, encoding='utf-8')
241     return _hook(temp)
242
243 def dump(obj, file_handle):
244     '''json serializes nested combinations of dicts, binarrays
245     and tuples.'''
246     json.dump(_dumphelper(obj), file_handle, indent=4, encoding='utf-8')
247
248
249 if __name__ == '__main__':
250     from time import time
251     from random import randint
252     #Create Binarray with interval +/- one hour and bin size ten minutes.
253     ba = Binarray(interval = 60*60, bin_size = 10*60)
254     #Center it on the present
255     now = int(time())
256     ba.center = now
257     #Generate some timestamps around the present
258     new_times = [now + randint(-60*60, 60*60) for _ in xrange(100)]
259     for tt in new_times:
260         ba.place_event(tt)
261
262     #Save it
263     with open('filename.sig', 'w') as f:
264         dump(ba, f)
265
266     print ba
```


A.1.2 Code Communication Dynamics

Code for extracting Bluetooth signals

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import os
5  import glob
6  import itertools
7  import random
8  from ast import literal_eval as LE
9  from social_fabric.phonetools import readncheck, Binarray, dump, make_filename
10
11  #=====
12  # Parameters
13  #=====
14
15  #Grap all the files we need to read
16  userfile_path = "userfiles/" #linux
17
18
19  #Interval of interest and bin size (seconds)
20  interval = 12*60*2
21  bin_size = 10*60
22
23  #How many other users contribute to background for a given user
24  number_of_background_samples = 1
25
26  #Number of users to analyse. Use None or 0 to include everyone.
27  max_users = 2
28
29  #Number of repeated signals required to be considered social
30  social_threshold = 2
31
32  #=====
33  # Functions
34  #=====
35
36  def call_type(n):
37      if n == 1:
38          return 'call_in'
39      elif n == 2:
40          return 'call_out'
41      else:
42          return None
43
44  def text_type(n):
45      if n == 1:
46          return 'text_in'
47      elif n == 2:
48          return 'text_out'
49      else:
50          return None
51
52
53  def user2social_times(user):
54      '''Converts user code to a list of times when the user was social i.e. had
55      two or more repeated Bluetooth signals.'''

```

```

56 with open(userfile_path+user+"/bluetooth_log.txt", 'r') as f:
57     raw = [LE(line) for line in f.readlines()]
58     #List of times when user was social
59     social_times = []
60     #Temporary variables
61     current_time = 0
62     previous_time = 0
63     current_users = []
64     previous_users = []
65     for ind in xrange(len(raw)-1):
66         signal = raw[ind]
67         new_time = signal['timestamp']
68         #Check if line represents a new signal and if so, update values
69         if new_time != current_time:
70             #Determine if previous signal was social and append to results
71             overlap = set(previous_users).intersection(set(current_users))
72             if len(overlap) >= social_threshold:
73                 social_times.append(previous_time)
74             #Update variables
75             previous_time = current_time
76             previous_users = current_users
77             current_users = []
78             current_time = new_time
79             new_user = signal['name']
80             if new_user == '-1' or not new_user:
81                 continue
82             else:
83                 current_users.append(new_user)
84             #
85     return social_times
86
87 #=====
88 # Time to crunch some numbers
89 #=====
90
91 user_folders = [f for f in glob.glob(userfile_path+"/*")
92                 if os.path.isfile(f+"/call_log.txt")
93                 and os.path.isfile(f+"/sms_log.txt")
94                 and os.path.isfile(f+"/bluetooth_log.txt")]
95 users = [folder.split(userfile_path)[1] for folder in user_folders]
96
97 if not users:
98     raise IOError('Found no users. Check userfile path.')
99
100 if max_users:
101     users = users[:max_users]
102
103 trigger = 'bluetooth'
104 events = ['call_out', 'call_in', 'text_in', 'text_out']
105 pairs = [p for p in itertools.product([trigger], events)]
106
107 activity = {p: Binarray(interval, bin_size) for p in pairs}
108 background = {p: Binarray(interval, bin_size) for p in pairs}
109
110
111 #Read in data
112 call_data = {user: readncheck(userfile_path+user+"/call_log.txt")
113              for user in users}
114 text_data = {user: readncheck(userfile_path+user+"/sms_log.txt")
115              for user in users}
116
117 count = 0

```

```

118
119 for user in users:
120     count += 1
121     print "Analyzing user %s out of %s. Code: %s" % (count, len(users), user)
122
123     #Get user data
124     user_calls = call_data[user]
125     user_texts = text_data[user]
126     user_social_times = user2social_times(user)
127
128     if not user_social_times:
129         continue
130
131     #Get background data
132     others = []
133     other_social_times = []
134     while len(others) < number_of_background_samples:
135         temp = random.choice(users)
136         if not (temp in others or temp == user):
137             newstuff = user2social_times(temp)
138             if not newstuff:
139                 continue
140             others.append(temp)
141             other_social_times += newstuff
142
143     #Determine the interval in which we have data on the current user
144     first = min(user_social_times)
145     last = max(user_social_times)
146
147     #=====
148     # #Establish activity signal
149     #=====
150     for time in user_social_times:
151         for e in events:
152             activity[(trigger, e)].center = time
153
154         for user_call in user_calls:
155             event = call_type(user_call['type'])
156             if not event:
157                 continue
158             time = user_call['timestamp']
159             activity[(trigger, event)].place_event(time)
160
161         for user_text in user_texts:
162             event = text_type(user_text['type'])
163             if not event:
164                 continue
165             time = user_text['timestamp']
166             activity[(trigger, event)].place_event(time)
167
168     #=====
169     # #Establish background signal
170     #=====
171     for other_time in other_social_times:
172         #Reposition the relevant binarrays
173         if not first <= other_time <= last:
174             continue
175         for e in events:
176             background[(trigger, e)].center = other_time
177
178     #Determine call background
179     for user_call in user_calls or []:

```

```

180         event = call_type(user_call['type'])
181         if not event:
182             continue
183         time = user_call['timestamp']
184         background[(trigger,event)].place_event(time)
185
186         #Determine text background
187         for user_text in user_texts or []:
188             event = text_type(user_text['type'])
189             if not event:
190                 continue
191             time = user_text['timestamp']
192             background[(trigger,event)].place_event(time)
193         #
194     #
195 #
196
197 #=====
198 # Done. Save signals
199 #=====
200
201 #Make a filename for the output file
202 filename = make_filename(prefix = trigger, interval=interval,
203                          bin_size=bin_size, ext = 'json')
204 with open(filename, 'w') as f:
205     dump((activity, background), f)
206
207 print "Saved to "+filename

```

Code for loading and plotting Bluetooth data

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import glob
5  import numpy as np
6  import matplotlib.pyplot as plt
7
8  from social_fabric.phonetools import (load, make_filename)
9
10 #=====
11 display = True
12 #These are just default - they're updated when data is read
13 interval = 12*60*2
14 bin_size = 10*60
15 #=====
16
17
18 #Numpy-compliant method to convert activity and background to relative signal
19 get_signal = np.vectorize(lambda act, back : act * 1.0/back if back else 0)
20
21 def read_data(filename):
22     '''Reads in data to plot. Updates the interval and bin sizes parameters.
23     Returns data to be plotted as a hashmap with the following structure:
24     {trigger : {event : signal}}.'''
25     data = {}
26     with open(filename, 'r') as f:
27         (a,b) = load(f)
28         first = True
29         for key in a.keys():
30             #Update interval and bin info
31             if first:
32                 global interval
33                 global bin_size
34                 interval = a[key].interval
35                 bin_size = a[key].bin_size
36                 first = False
37
38             (trigger, event) = key
39             act = a[key].normalized()
40             back = b[key].normalized()
41             signal = get_signal(list(act), list(back))
42             if not trigger in data:
43                 data[trigger] = {}
44             data[trigger][event] = signal
45         return data
46
47
48 def make_plot(trigger, signals):
49     '''Generate a plot of relative user activity signal.'''
50     legendstuff = []
51     for event, vals in signals.iteritems():
52         t = np.arange(-interval, interval, bin_size)
53         t = map(lambda x: x*1.0/3600, t)
54         legendstuff.append(event)
55         s = list(vals)
56         plt.plot(t, s)
57     plt.legend(legendstuff, loc='upper left')
58     plt.xlabel('Time (h)')
59     plt.ylabel('Relative signal')
60     plt.title("Trigger: " + trigger)
61     plt.grid(True)

```

```
62     saveas = make_filename(trigger, interval, bin_size, ext = 'pdf')
63     plt.savefig(saveas)
64     print "Saved to "+saveas
65     if display:
66         plt.show()
67
68 if __name__ == '__main__':
69     # Read in data
70     filenames = glob.glob('bluetooth-int43200-bin600.json')
71     for filename in filenames:
72         data = read_data(filename)
73         for trigger, signals in data.iteritems():
74             make_plot(trigger, signals)
75     #
```

A.1.3 Preprocessing

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import os
5  import glob
6  import numpy as np
7  from ast import literal_eval as LE
8  from social_fabric.phonetools import readncheck, user2profile
9  from social_fabric.secrets import pushme
10 from social_fabric.smallestenclosingcircle import make_circle
11 from social_fabric import lloyds
12 from collections import Counter
13 import math
14 from datetime import datetime, timedelta
15 import pytz
16 import json
17 from statsmodels.tsa import ar_model
18
19 #=====
20 # Parameters
21 #=====
22
23 #Grap all the files we need to read
24 #userfile_path = "c:\\ userfiles\\" #@Windows
25 userfile_path = "/lscr_paper/amoellga/Data/Telefon/userfiles/" #linux
26
27 #Filename to save output to.
28 output_filename = 'data_single_thread_numedgps.json'
29
30 #Number of users to analyse. Use None or 0 to include everyone.
31 max_users = None
32
33 #Specific user codes to analyze for debugging purposes. Empty means include all
34 exclusive_users = []
35
36 #Conversion factors from degrees to meters (accurate around Copenhagen)
37 longitude2meters = 111319
38 latitude2meters = 110946
39
40 #In meters. These improve convergence of the stochastic SEC algorithm.
41 x_offset = 1389425.2238223257
42 y_offset = 6181209.0059678229
43
44 required = ['bluetooth_log.txt', 'call_log.txt', 'facebook_log.txt',
45            'gps_log.txt', 'sms_log.txt']
46
47 user_folders = [f for f in glob.glob(userfile_path+"*")
48                if all(os.path.isfile(f+"/"+stuff) for stuff in required)]
49 users = [folder.split(userfile_path)[1] for folder in user_folders]
50
51 if not users:
52     raise IOError('Found no users. Check userfile path.')
53
54 if exclusive_users:
55     users = list(set(exclusive_users).intersection(users))
56
57 if max_users:

```

```

58 |     users = users[:max_users]
59 |
60 | # Number of wallclock hours pr bin when fitting autoregressive series
61 | hours_pr_ar_bin = 6
62 | assert 24%hours_pr_ar_bin == 0
63 |
64 | # Number of hours pr bin when compute daily rythm entropy
65 | hours_pr_daily_rythm_bin = 1
66 | assert 24%hours_pr_daily_rythm_bin == 0
67 |
68 | # Time zone information
69 | cph_tz = pytz.timezone('Europe/Copenhagen')
70 |
71 | #Which data kinds to include
72 | include_calls = True
73 | include_ar = True
74 | include_gps = True
75 | include_network = False #Allow geolocation from network data
76 | include_bluetooth = True
77 | include_facebook = True
78 |
79 | #Whether to include not-a-number values in final output
80 | allow_nan = True
81 |
82 | #Threshold values to discard users with insufficient data
83 | minimum_number_of_texts = 10
84 | minimum_number_of_calls = 5
85 | minimum_number_of_gps_points = 100
86 | minimum_number_of_facebook_friends = 1
87 |
88 | #N_pings required to be considered social
89 | bluetooth_social_threshold = 2
90 |
91 | #Whether to output plots of cluster analysis
92 | plot_clusters = False
93 |
94 | #=====
95 | # Define helper methods
96 | #=====
97 |
98 | def get_distance(p, q):
99 |     return math.sqrt(sum([(p[i]-q[i])**2 for i in xrange(len(p))]))
100 |
101 | def is_sorted(l):
102 |     return all([l[i+1] >= l[i] for i in xrange(len(l)-1)])
103 |
104 | def get_entropy(event_list):
105 |     '''Takes a list of contacts from in/ or outgoing call/text events and
106 |     computes its entropy. event_list must be simply a list of user codes
107 |     corresponding to events.'''
108 |     n = len(event_list)
109 |     counts = Counter(event_list)
110 |     ent = sum([-v/n * math.log(v/n, 2) for v in counts.values()])
111 |     return ent
112 |
113 | def next_time(dt, deltahours):
114 |     '''Accepts a datetime object and returns the next datetime object at which
115 |     the 'hour' count modulo deltahours is zero.
116 |     For example, deltahours = 6 gives the next time clock time is
117 |     0, 6, 12 or 18.
118 |     Sounds simple but is pretty annoying due to daylight saving time and so on,
119 |     so take care not to mess with this.'''

```



```

120     base = datetime(dt.year, dt.month, dt.day, dt.hour)
121     interval = timedelta(hours = deltahours - dt.hour%deltahours)
122     naive_guess = base + interval
123     return cph_tz.localize(naive_guess)
124
125 def next_midnight(dt):
126     '''Takes a datetime object and returns dt object of following midnight'''
127     base = datetime(dt.year, dt.month, dt.day)
128     naive = base + timedelta(days = 1)
129     return cph_tz.localize(naive)
130
131 def epoch2dt(timestamp):
132     '''Converts unix timestamp into a pytz timezone-aware datetimeobject.'''
133     utc_time = datetime.utcfromtimestamp(timestamp)
134     smart_time = cph_tz.fromutc(utc_time)
135     return smart_time
136
137 def sort_dicts_by_key(dictlist, key):
138     '''Takes a list of dicts and returns the same list sorted by its
139     key-entries.'''
140     decorated = [(d[key], d) for d in dictlist]
141     decorated.sort()
142     return [d for (k, d) in decorated]
143
144 def get_autocovar_coefficient(X, lag):
145     '''Returns the autocovariance coefficient for the input series at the
146     input lag.'''
147     mu = np.mean(X)
148     temp = sum((X[i] - mu)*(X[i+lag]-mu) for i in xrange(len(X)-lag))
149     return temp*1.0/len(X)
150
151 def get_autocorrelation_coefficients(series, lags):
152     '''Determines the autocorrelation coefficients of input series at
153     each of the input lags. Uses .r_k = c_k/c_0.
154     Accepts a list of lags or an int in which case it returns lags up to
155     and including the input.'''
156     if isinstance(lags, int):
157         lags = range(lags+1)
158     c0 = get_autocovar_coefficient(series, 0)
159     inv = 1.0/c0
160     return [inv*get_autocovar_coefficient(series, lag) for lag in lags]
161
162 def make_time_series(dts, hours_pr_ar_bin):
163     '''Takes a sorted list of datetime objects and converts to a time series
164     where each entry denotes the number of events in the corresponding bin.'''
165     first_time = next_time(dts[0], hours_pr_ar_bin)
166     for i in xrange(len(dts)):
167         if dts[i] >= first_time:
168             dts = dts[i:]
169             break
170     #
171     last_time = next_time(first_time, hours_pr_ar_bin)
172     time_series = []
173
174     summer = 0
175     for dt in dts:
176         while not first_time <= dt < last_time:
177             time_series.append(summer)
178             summer = 0
179             first_time = last_time
180             last_time = next_time(last_time, hours_pr_ar_bin)
181     summer += 1

```

```

182     return time_series
183
184
185 def timestamps2daily_entropy(timestamps, hours_pr_bin):
186     '''Constructs a histogram of hour-values of the input timestamps
187     and computes its entropy.'''
188     if not 24%hours_pr_bin == 0:
189         raise ValueError("24 must be divisible by hours_pr_bin.")
190     bins = {}
191     for timestamp in timestamps:
192         hour = epoch2dt(timestamp).hour
193         _bin = int(hour/hours_pr_bin)
194         try:
195             bins[_bin] += 1
196         except KeyError:
197             bins[_bin] = 1
198     #
199     total = sum(bins.values())
200     entropy = sum([-v/total*math.log(v/total, 2) for v in bins.values()])
201     return entropy
202
203 #Make sure we have a blank file to write to.
204 open(output_filename, 'w').close()
205 assert os.stat(output_filename).st_size == 0 #Check that it worked.
206
207 user_counter = 0
208
209 for user in users:
210     user_counter += 1
211     msg = "Processing user %s of %s: %s" % (user_counter, len(users), user)
212     print msg
213     if user_counter % 100 == 0:
214         pushme(msg)
215
216     #Dict to hold all data extracted on current user
217     data = {}
218
219     #Try to load user psych profile- Discard user if they're not in file
220     try:
221         profile = user2profile[user]
222     except KeyError:
223         print "No psychological data on user."
224         continue # No questionnaire data. Shouldn't happen.
225
226     #Read in calls/texts
227     calls = readncheck(userfile_path+user+"/call_log.txt")
228
229     if len(calls) < minimum_number_of_calls:
230         print "too few calls"
231         continue
232     texts = readncheck(userfile_path+user+"/sms_log.txt")
233     if len(texts) < minimum_number_of_texts:
234         print "too few texts"
235         continue
236
237     #Extract list of times for calls, texts and combination
238     call_times = sorted([d['timestamp'] for d in calls])
239     text_times = sorted([d['timestamp'] for d in texts])
240     call_text_times = sorted(call_times + text_times)
241
242     #Get calls from the first three months
243     tmin = call_times[0]

```

```

244 tmax = call_times[0] + 60*60*24*30*3
245 early_calls = [c['number'] for c in calls if tmin<=c['timestamp']<=tmax]
246
247 #Repeat for texts
248 tmin = text_times[0]
249 tmax = text_times[0] + 60*60*24*30*3
250 early_texts = [t['body'] for t in texts if tmin<=t['timestamp']<=tmax]
251
252 #Get number of unique contacts for first 3 months and append to data.
253 uniques = len(set(early_calls + early_texts))
254 data['n_contacts_first_three_months'] = uniques
255
256 #Compute daily entropy for calls and texts
257 data['call_daily_entropy'] = timestamps2daily_entropy(call_times,
258                                                         hours_pr_daily_rhythm_bin)
259 data['text_daily_entropy'] = timestamps2daily_entropy(text_times,
260                                                         hours_pr_daily_rhythm_bin)
261
262 #Compute median and std for call durations
263 call_durations = [c['duration'] for c in calls if not c['duration'] == 0]
264 data['call_duration_med'] = np.median(call_durations)
265 data['call_duration_std'] = np.std(call_durations)
266
267 #=====
268 # Crunch time-series info
269 #=====
270 if include_ar:
271     #Grab a sorted list of only times of events caused by user
272     outgoing_stuff = sorted([d['timestamp'] for d in calls
273                             if d['type'] == 2] + [d['timestamp']
274                                                    for d in texts if d['type'] == 2])
275
276     n_params = int(24*7/hours_pr_ar_bin + 1) #1 week plus 1 extra bin
277
278     #Fit time series and extract parameters
279     try:
280         #Convert into timezone aware datetime objects
281         dts = [epoch2dt(timestamp) for timestamp in outgoing_stuff]
282
283         time_series = make_time_series(dts, hours_pr_ar_bin)
284
285         model = ar_model.AR(time_series)
286         result = model.fit(n_params)
287         #Grab parameters from fitted model
288         params = result.params[1:]
289         while len(params) < n_params:
290             params.append(float('nan'))
291     except:
292         if not allow_nan:
293             continue
294         else:
295             params = [float('nan') for _ in xrange(n_params)]
296
297     #Append AR-coefficients to user data
298     count = 0
299     for par in params:
300         count += 1
301         name = "outgoing_activity_AR_coeff_" + str(count)
302         data[name] = par
303
304     #Get autocorrelation coefficients as well
305     try:

```

```

306         accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
307     except:
308         if not allow_nan:
309             continue
310         accs = [float('nan') for _ in xrange(n_params + 1)]
311
312     #Append autocorrelation coefficients to user data
313     for i in xrange(len(accs)):
314         name = "outgoing_activity_acc_"+str(i)
315         data[name] = accs[i]
316
317     #Repeat with incoming signals. Might be interesting.
318     incoming_stuff = sorted([d['timestamp'] for d in calls
319                             if d['type'] == 1] + [d['timestamp']
320                                                  for d in texts if d['type'] == 1])
321
322     try:
323         # Convert into timezone aware datetime objects
324         dts = [epoch2dt(timestamp) for timestamp in incoming_stuff]
325
326         time_series = make_time_series(dts, hours_pr_ar_bin)
327         model = ar_model.AR(time_series)
328         result = model.fit(n_params)
329         params = result.params[1:]
330         while len(params) < n_params:
331             params.append(float('nan'))
332     except:
333         if not allow_nan:
334             continue
335         else:
336             params = [float('nan') for _ in xrange(n_params)]
337
338     # Name each of them and append to user data
339     count = 0
340     for par in params:
341         count += 1
342         name = "incoming_activity_AR_coeff_"+str(count)
343         data[name] = par
344
345     #Get autocorrelation coefficients as well
346     try:
347         accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
348     except:
349         if not allow_nan:
350             continue
351         accs = [float('nan') for _ in xrange(n_params + 1)]
352
353     for i in xrange(len(accs)):
354         name = "incoming_activity_acc_"+str(i)
355         data[name] = accs[i]
356
357
358     #=====
359     #    Crunch call/text info
360     #=====
361
362     if include_calls:
363         #Add values to temporary data map.
364         d = {'call': call_times, 'text': text_times, 'ct': call_text_times}
365         for label, times in d.iteritems():
366             timegaps = [times[i+1] - times[i] for i in xrange(len(times)-1)]
367             timegaps = filter(lambda x: x < 259200, timegaps) #3 days, tops

```

```

368     data[label+'_iet_med'] = np.median(timegaps)
369     data[label+'_iet_std'] = np.std(timegaps)
370
371     #Generate lists of the contact for each text/call event
372     call_numbers = [call['number'] for call in calls]
373     text_numbers = [text['address'] for text in texts]
374     ct_numbers = call_numbers + text_numbers
375
376     #Compute entropy and add to data
377     data['call_entropy'] = get_entropy(call_numbers)
378     data['text_entropy'] = get_entropy(text_numbers)
379     data['ct_entropy'] = get_entropy(ct_numbers)
380
381     #Compute contact list info
382     call_contacts = Counter([c['number'] for c in calls
383                             if c['type'] == 2]).keys()
384     text_contacts = Counter([t['address'] for t in texts
385                             if t['type'] == 2]).keys()
386
387     #Grap number of contacts
388     n_call_contacts = len(call_contacts)
389     n_text_contacts = len(text_contacts)
390     n_ct_contacts = len(set(call_contacts).union(set(text_contacts)))
391
392     #Add to data map
393     data['n_call_contacts'] = n_call_contacts
394     data['n_text_contacts'] = n_text_contacts
395     data['n_ct_contacts'] = n_ct_contacts
396
397     #Compute and add contact/interaction ratio (cir)
398     data['call_cir'] = n_call_contacts/len(calls)
399     data['text_cir'] = n_text_contacts/len(texts)
400     data['ct_cir'] = n_ct_contacts/(len(calls) + len(texts))
401
402     #Add data on number of interactions
403     data['n_calls'] = len(calls)
404     data['n_texts'] = len(texts)
405     data['n_ct'] = len(calls + texts)
406
407     #Determine percentage of calls/texts that were initiated by user.
408     initiated_calls = len([c for c in calls if c['type'] == 2])
409     data['call_percent_initiated'] = initiated_calls/len(calls)
410     initiated_texts = len([t for t in texts if t['type'] == 2])
411     data['text_percent_initiated'] = initiated_texts/len(texts)
412
413     #Determine call response rate.
414     with open(userfile_path+user+"/call_log.txt", 'r') as f:
415         all_calls = [LE(line) for line in f.readlines()]
416     #Make sure the call data is sorted
417     if not is_sorted([c['timestamp'] for c in all_calls]):
418         all_calls = sort_dicts_by_key(all_calls, 'timestamp')
419
420     '''Check for unanswered called that are replied to within an hour.
421     This is performed in the following fashion: iterate through all the
422     calls. If a call is unanswered, add it to "holding" list. If a call
423     from holding matches the current call, it counts as a reply.
424     If the time of the current call is more than hour after a held call,
425     it is discarded'''
426     missed = 0
427     replied = 0
428     holding = []
429     for call in all_calls:
430         if call['type'] == 3 or call['type'] == 1 and call['duration'] == 0:

```

```

430         holding.append(call)
431         missed += 1
432     else:
433         for held_call in holding:
434             #Drop calls that have been held for too long
435             if call['timestamp'] - held_call['timestamp'] > 3600:
436                 holding.remove(held_call)
437             #Check if given call is a response
438             elif (call['type'] == 2
439                   and call['number'] == held_call['number']):
440                 holding.remove(held_call)
441                 replied += 1
442             #
443         #
444     #
445     data['call_response_rate'] = replied/(missed+replied) if replied else 0
446
447     #Determine text response rate
448     missed = 0
449     replied = 0
450     holding = []
451     response_times = []
452     if not is_sorted([t['timestamp'] for t in texts]):
453         texts = sort_dicts_by_key(texts, 'timestamp')
454
455     for text in texts:
456         #Make sure incoming text is not from a user already held
457         if text['type'] == 1:
458             if not holding or all([text['address'] !=
459                                   t['address'] for t in holding]):
460                 #It's good - append it
461                 holding.append(text)
462                 missed += 1
463             #
464         else:
465             for held_text in holding:
466                 if text['timestamp'] - held_text['timestamp'] > 3600:
467                     holding.remove(held_text)
468                 #Check if text counts as reply
469                 elif (text['type'] == 2
470                       and text['address'] == held_text['address']):
471                     holding.remove(held_text)
472                     replied += 1
473                     dt = text['timestamp'] - held_text['timestamp']
474                     response_times.append(dt)
475             #
476         #
477     #
478     data['text_response_rate'] = replied/(missed+replied) if replied else 0
479     data['text_latency'] = np.median(response_times)
480
481     #Check percentage of calls taken place in during the night
482     count = 0
483     for call in calls:
484         hour = epoch2dt(call['timestamp']).hour
485         if not (8 <= hour < 22):
486             count += 1
487     #
488     data['call_night_activity'] = count/len(calls)
489
490     #Compute % of calls/texts outgoing from user. This works because true=1
491     data['call_outgoing'] = sum([c['type'] == 2 for c in calls])/len(calls)

```

```

492     data['text_outgoing'] = sum([t['type'] == 2 for t in texts])/len(texts)
493
494     #=====
495     # Crunch location data
496     #=====
497
498     if include_gps:
499         with open(userfile_path+user+"/gps_log.txt", 'r') as f:
500             raw = [LE(line) for line in f.readlines()]
501
502             if not is_sorted([l['timestamp'] for l in raw]):
503                 raw = sort_dicts_by_key(raw, 'timestamp')
504
505             #We only want measurements taken at least 500s apart.
506             prev = 0
507             gps_data = []
508             allowed_providers = ['gps', 'network'] if include_network else ['gps']
509             for line in raw:
510                 now = line['timestamp']
511                 if line['provider'] in allowed_providers and now - prev >= 500:
512                     #Convert coordinates to km and note it down
513                     x = (longitude2meters*line['lon'] - x_offset)*0.001
514                     y = (latitude2meters*line['lat'] - y_offset)*0.001
515                     gps_data.append({'point': (x,y), 'timestamp': now,
516                                     'smarttime': epoch2dt(now)})
517                     prev = now
518             #
519             #ignore user if there aren't enough data
520             if not len(gps_data) >= minimum_number_of_gps_points:
521                 continue
522
523             # We want to investigate each day separately so start at midnight.
524             first_midnight = next_midnight(gps_data[0]['smarttime'])
525             for i in xrange(len(gps_data)):
526                 if gps_data[i]['smarttime'] >= first_midnight:
527                     gps_data = gps_data[i:]
528                     break
529
530
531             #Generate list of radii of smallest enclosing circle, SEC, for each day
532             current_points = []
533             prev = gps_data[0]['timestamp']
534             radii = []
535
536             distances = []
537             early_day = gps_data[0]['smarttime']
538             late_day = next_midnight(early_day)
539             for datum in gps_data:
540                 now = datum['smarttime']
541                 while not early_day <= now < late_day:
542                     if len(current_points) > 2:
543                         crds = [p['point'] for p in current_points]
544                         circle = make_circle(crds)
545                         r = circle[2] if circle else 0
546                         if circle and r > 0:
547                             if r <= 500:
548                                 radii.append(r)
549                                 distances.append(sum([get_distance(crds[i], crds[i+1])
550                                                         for i in xrange(len(crds)-1)]))
551                     # Reset counters and update bins
552                     current_points = []
553                     early_day = late_day

```

```

554         late_day = next_midnight(late_day)
555         current_points.append(datum)
556
557
558     data['radius_of_gyration_med'] = np.median(radii)
559     data['radius_of_gyration_std'] = np.std(radii)
560     data['travel_med'] = np.median(distances)
561     data['travel_std'] = np.std(distances)
562
563     #Run Lloyd's algorithm to identify clusters
564     #Determine which points are stationary – less movement than 100m
565     stationary_data = []
566     try:
567         for i in xrange(1,len(gps_data)-1):
568             a,b,c = tuple([gps_data[ind]['point'] for ind in
569                           [i-1, i, i+1]])
570             if (get_distance(a, b) < 0.1 and get_distance(b, c) < 0.1):
571                 stationary_data.append(gps_data[i])
572         #
573         initial_clusters = 50
574         threshold_percent = 0.05
575         points = [elem['point'] for elem in stationary_data]
576
577         minimum_points = int(threshold_percent*len(points))
578
579         clusters_scatter = lloyds.lloyds(points, initial_clusters, runs=3,
580                                         init='scatter')
581
582         clusters_sample = lloyds.lloyds(points, initial_clusters, runs=3,
583                                         init='sample')
584
585         #Determine most succesful method
586         locs = lambda c: [p for p in c.values() if len(p)>=minimum_points]
587         if len(locs(clusters_scatter)) >= len(locs(clusters_sample)):
588             method = 'scatter'
589             best = clusters_scatter
590         else:
591             method = 'sample'
592             best = clusters_sample
593
594         # Number of places which survive cutoff (true = 1, so just sum)
595         n_places = sum([len(pl) >= minimum_points for pl in best.values()])
596
597         #Output plots of clustering
598         if plot_clusters:
599             if not os.path.isdir('pics'):
600                 os.mkdir('pics')
601             for ext in ['.pdf', '.png']:
602                 filename = 'pics/'+user+"_"+method+ext
603                 lloyds.draw_clusters(clusters = best,
604                                     threshold = minimum_points,
605                                     show = False,
606                                     filename = filename)
607
608     except:
609         if not allow_nan:
610             continue
611         n_places = float('nan')
612     data['n_places'] = n_places
613
614     #Compute location entropy and add to data
615     try:

```



```

616     n_points = sum(len(location) for location in best.values())
617     data['location_entropy'] = sum([-len(p)/n_points*math.log(len(p)/n_points)
618                                   for p in best.values()])
619 except ValueError:
620     data['location_entropy'] = float('nan')
621
622 '''Guess where people live. Probably where they spend weeknights...
623 It's important to avoid selection bias here (people probably turn off
624 their phone when sleeping at home but not while partying at DTU, which
625 means fewer data points at their actual home).
626 This is rectified by excluding points that aren't logged monday to
627 thursday and only recording one 'late' or 'early' data point pr date.
628 These points are labelled 'weird' and are used to determine the user's
629 home.'''
630 try:
631     weirdpoints = []
632     latedays = []
633     earlydays = []
634     for datum in stationary_data:
635         now = datum['smarttime']
636         #ignore weekends
637         if now.weekday() > 3:
638             continue
639         thisdate = (now.year, now.month, now.day)
640         if now.hour >= 20 and not thisdate in latedays:
641             weirdpoints.append(datum['point'])
642             latedays.append(thisdate)
643         elif now.hour <= 7 and not thisdate in earlydays:
644             weirdpoints.append(datum['point'])
645             earlydays.append(thisdate)
646
647     best_score = 0
648     home = None
649     for key, val in best.iteritems():
650         score = len(set(val).intersection(weirdpoints))
651         if score > best_score:
652             home = key
653             best_score = score
654     #
655     # Estimate how much user spends at home
656     ordered_gps = sort_dicts_by_key(gps_data, 'timestamp')
657     is_home = lambda p: get_distance(home, p) <= 0.200
658     time_home = 0
659     time_away = 0
660     for i in xrange(len(ordered_gps)-1):
661         a = ordered_gps[i]
662         b = ordered_gps[i+1]
663         dt = b['timestamp'] - a['timestamp']
664         if dt > 7200:
665             continue
666         elif is_home(a['point']) and is_home(b['point']):
667             time_home += dt
668         elif (not is_home(a['point'])) and (not is_home(b['point'])):
669             time_away += dt
670     #
671     data['home_away_time_ratio'] = time_home/time_away
672 except:
673     if not allow_nan:
674         continue
675     data['home_away_time_ratio'] = float('nan')
676
677

```

```

678 #=====
679 # Facebook data
680 #=====
681
682 if include_facebook:
683     with open(userfile_path+user+'/facebook_log.txt', 'r') as f:
684         n = len(f.readlines())
685         if n < minimum_number_of_facebook_friends:
686             continue
687         data['number_of_facebook_friends'] = n
688
689 #=====
690 # Bluetooth data
691 #=====
692
693 if include_bluetooth:
694     with open(userfile_path+user+"/bluetooth_log.txt", 'r') as f:
695         raw = [LE(line) for line in f.readlines()]
696         #Make sure data is sorted chronologically
697         if not is_sorted([entry['timestamp'] for entry in raw]):
698             raw = sort_dicts_by_key(raw, 'timestamp')
699         #List of times when user was social
700         social_times = []
701         total_social_time = 0
702         total_time = 0
703         #maps from each other user encountered to time spend with said user
704         friend2time_spent = {}
705         #Temporary variables
706         current_time = 0
707         previous_time = 0
708         current_users = []
709         previous_users = []
710         for signal in raw:
711             new_time = signal['timestamp']
712             #Check if line represents a new signal and if so, update values
713             if new_time != current_time:
714                 dt = new_time - current_time
715                 #Determine number of pings
716                 overlap = set(previous_users).intersection(set(current_users))
717                 if len(overlap) >= bluetooth_social_threshold:
718                     social_times.append(previous_time)
719                     if dt <= 7200:
720                         total_social_time += dt
721                         total_time += dt
722                         for friend in overlap:
723                             try:
724                                 friend2time_spent[friend] += dt
725                             except KeyError:
726                                 friend2time_spent[friend] = dt
727
728                     #
729                     elif dt <= 7200:
730                         total_time += dt
731                 #Update variables
732                 previous_time = current_time
733                 previous_users = current_users
734                 current_users = []
735                 current_time = new_time
736             new_user = signal['name']
737             if new_user == '-1' or not new_user:
738                 continue
739             else:
740                 current_users.append(new_user)

```

```

740     #
741     # Add fraction of time spent social to output
742     data['fraction_social_time'] = total_social_time/total_time
743     # Compute social entropy
744     normfac = 1.0/sum(friend2time_spent.values())
745     ent = sum(-t*normfac*math.log(t*normfac)
746              for t in friend2time_spent.values())
747     data['social_entropy'] = ent
748
749     data['bluetooth_daily_entropy'] = timestamps2daily_entropy(social_times,
750                                                                hours_pr_daily_rythm_bin)
751
752     #Ensure time span is suficcient to make a time series
753     if not (social_times[-1] - social_times[0] > 24*3600*7
754            +1+3600*hours_pr_ar_bin):
755         continue
756
757     #Fit AR-series and append parameters to output
758     try:
759         dts = [epoch2dt(timestamp) for timestamp in social_times]
760         time_series = make_time_series(dts, hours_pr_ar_bin)
761         model = ar_model.AR(time_series)
762         n_params = int(24*7/hours_pr_ar_bin + 1) #1 week plus 1 extra bin
763         result = model.fit(maxlag = None, ic = None)
764         params = result.params#[1:]
765         while len(params) < n_params:
766             params.append(float('nan'))
767     except:
768         if not allow_nan:
769             continue
770         params = [float('nan') for _ in xrange(n_params)]
771
772     count = 0
773     for par in params:
774         count += 1
775         name = "bluetooth_activity_AR_coeff_" + str(count)
776         data[name] = par
777
778     # Compute autocorrelation coeffs and append to output
779     try:
780         accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
781     except:
782         if not allow_nan:
783             continue
784         accs = [float('nan') for _ in xrange(n_params + 1)]
785     for i in xrange(len(accs)):
786         name = "bluetooth_activity_acc_" + str(i)
787         data[name] = accs[i]
788
789     #=====
790     # Wrap up user
791     #=====
792
793     # Double check thata doesn't containing nan values
794     if any(np.isnan(value) for value in data.values()) and not allow_nan:
795         continue #Discard user due to insufficient data
796
797     #Collect results
798     final = {'user': user, 'data': data, 'profile': profile}
799
800     with open(output_filename, 'a') as f:
801         json.dump(final, f)

```

```

802         f.write("\n")
803
804     #Done.
805     pushme("Data extraction done.")

```

A.1.4 Social Fabric Code

```

1  # -*- coding: utf-8 -*-
2  """This module aims to allow sharing of some common methods and settings
3  when testing and tweaking various machine learning schemes.
4  Always import settings and the like from here!"""
5
6  from __future__ import division
7  import abc
8  from collections import Counter
9  import itertools
10 import json
11 import math
12 import matplotlib.colors as mcolors
13 import matplotlib.pyplot as plt
14 import matplotlib.patches as mpatches
15 import multiprocessing
16 import numpy as np
17 import os
18 import random
19 from scipy.sparse import dok_matrix
20 from sklearn import svm
21 from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
22 from sklearn.cross_validation import (cross_val_score, LeaveOneOut, KFold,
23                                       StratifiedKFold)
24
25 import sys
26 import traceback
27
28 BASE = os.path.dirname(__file__)
29
30 oldhat = (35/256,39/256,135/256)
31 nude = (203/256,150/256,93/256)
32 wine = (110/256,14/256,14/256)
33 moerkeroed = (156/256,30/256,36/256)
34
35 #Empirically determined optimal hyperparameters for SVRs and RFs
36 with open(BASE+'svr_parameters.json', 'r') as f:
37     svr_parameters = json.load(f)
38
39 wrbf_parameters = svr_parameters #updater!!!
40
41 with open(BASE+'rf_parameters.json', 'r') as f:
42     rf_parameters = json.load(f)
43
44 def _make_colormap(seq):
45     """Return a LinearSegmentedColormap
46     seq: a sequence of floats and RGB-tuples. The floats should be increasing
47     and in the interval (0,1).
48     """
49     seq = [(None,) * 3, 0.0] + list(seq) + [1.0, (None,) * 3]
50     cdict = {'red': [], 'green': [], 'blue': []}

```

```

51     for i, item in enumerate(seq):
52         if isinstance(item, float):
53             r1, g1, b1 = seq[i - 1]
54             r2, g2, b2 = seq[i + 1]
55             cdict['red'].append([item, r1, r2])
56             cdict['green'].append([item, g1, g2])
57             cdict['blue'].append([item, b1, b2])
58     return mcolors.LinearSegmentedColormap('CustomMap', cdict)
59
60 color_map = _make_colormap([oldhat, moerkeroed, 0.33, moerkeroed, nude, 0.67, nude])
61
62 big_five = ['openness', 'conscientiousness', 'extraversion', 'agreeableness',
63            'neuroticism']
64
65 #_default_features = ["n_texts",
66 #                    "ct_iet_std ",
67 #                    "call_cir ",
68 #                    "call_entropy",
69 #                    "text_cir ",
70 #                    "n_calls",
71 #                    "text_latency",
72 #                    "call_outgoing",
73 #                    "fraction_social_time",
74 #                    "text_outgoing",
75 #                    "call_iet_std ",
76 #                    "n_text_contacts",
77 #                    "call_night_activity ",
78 #                    "call_iet_med",
79 #                    "outgoing_activity_AR_coeff_2",
80 #                    "text_entropy",
81 #                    "ct_cir ",
82 #                    "text_response_rate",
83 #                    "n_ct_contacts",
84 #                    "social_entropy",
85 #                    "n_call_contacts",
86 #                    "n_ct",
87 #                    "text_iet_std ",
88 #                    "ct_iet_med",
89 #                    "ct_entropy",
90 #                    "text_iet_med",
91 #                    "call_response_rate",
92 #                    "number_of_facebook_friends"]
93
94 _default_features = ['call_iet_med', 'text_iet_med', 'social_entropy',
95                    'call_entropy', 'travel_med', 'n_places', 'text_latency',
96                    'call_night_activity']
97
98 def split_ntiles(values, n):
99     '''Determines the values that separate the input list into n equal parts.
100     this is a generalization of the notion of median (in the case n = 2) or
101     quartiles (n=4).
102     Usage: ntiles([5,6,7], 2) gives [6] for instance.'''
103     result = []
104     for i in xrange(1,n):
105         percentile = 100/n * i
106         result.append(np.percentile(values, percentile,
107                                   interpolation='linear'))
108     return result
109
110 def determine_ntile(value, ntiles):
111     '''Determines which n-tilde the input value belongs to.
112     Usage: determine_ntile([7,9,13], 10) gives 2 (third quartile).

```

```

113 | This uses zero indexing so data split into e.g. quartiles will give results
114 | like 0,1,2,3 - NOT 1,2,3,4.'''
115 | #Check if value is outside either extreme, meaning n-tile 1 or n.
116 | if value >= ntiles[-1]:
117 |     return len(ntiles) #Remember the length is n-1
118 | elif value < ntiles[0]:
119 |     return 0 #Values was in the first n-tile
120 | # Define possible region and search for where value is between two elements
121 | left = 0
122 | right = len(ntiles)-2
123 | #Keep checking th middle of the region and updating region
124 | ind = (right + left)//2
125 | while not ntiles[ind] <= value < ntiles[ind + 1]:
126 |     #Check if lower bound tile is on the left
127 |     if value < ntiles[ind]:
128 |         right = ind - 1
129 |     else:
130 |         left = ind + 1
131 |     ind = (right + left)//2
132 | # Being between ntiles 0 and 1, means second n-tile and so on.
133 | return ind + 1
134 |
135 | def assign_labels(Y,n):
136 |     '''Accepts a list and an int n and returns a list of discrete labels
137 |     corresponding to the ntile each original y-value was in.'''
138 |     ntiles = split_ntiles(Y,n)
139 |     labels = [determine_ntile(y,ntiles) for y in Y]
140 |     return labels
141 |
142 | def normalize_data(list_):
143 |     '''Normalizes input data to the range [-1, 1]'''
144 |     lo, hi = min(list_), max(list_)
145 |     if lo == hi:
146 |         z = len(list_)*[0]
147 |         return z
148 |     else:
149 |         return [2*(val-lo)/(hi - lo) - 1 for val in list_]
150 |
151 |
152 | def read_data(filename, trait, n_classes = None, normalize = True,
153 |               features='default', interpolate = True):
154 |     '''This reads in a preprocessed datafile, splits psych profile data into
155 |     n classes if specified, filters desired psychological traits and
156 |     features and returns as a tuple (X,Y, indexdict), which can be fed to a
157 |     number of off-the-shelf ML schemes.
158 |     If trait=='Sex', female and male are converted to 0 and 1, respectively.
159 |     indexdict maps each element of the feature vectors to their label, as in
160 |     {42 : 'distance_travelled_pr_day'} etc.
161 |
162 |     Args:
163 |         filename : str
164 |             Name of the file containing the data.
165 |         trait : str
166 |             The psychological trait to extract data on.
167 |         n_classes : int
168 |             Number of classes to split data into. Default is None,
169 |             i.e. just keep the decimal values. Ignored if trait == 'Sex', as data
170 |             only has two discreet values.
171 |         normalize : bool
172 |             Whether to hard normalize data to [-1, 1].
173 |         features : str/list
174 |             Which features to read in. Can also be 'all'

```

```

175         or 'default', meaning the ones I've pragmatically found to be
176         reasonable.
177     interpolate : bool:
178         Whether to replace NaN's with the median value of
179         the feature in question.'''
180     if trait == 'sex':
181         n_classes = None
182     #Read in the raw data
183     with open(filename, 'r') as f:
184         raw = [json.loads(line) for line in f.readlines()]
185     #Get list of features to be included – everything if nothing's specified
186     included_features = []
187     if features == 'default':
188         included_features = _default_features
189     elif features == 'all':
190         included_features = raw[0]['data'].keys() if features == 'all' else features
191     else:
192         included_features = features
193
194     #Remove any features that only have NaN values.
195     for i in xrange(len(included_features)-1,-1,-1):
196         feat = included_features[i]
197         if all(math.isnan(line['data'][feat]) for line in raw):
198             del included_features[i]
199
200
201     # ----- Handle feature vectors -----
202     # Dict mapping indices to features
203     indexdict = {ind: feat for ind, feat in enumerate(included_features)}
204     # N_users x N_features array to hold data
205     rows = len(raw)
206     cols = len(included_features)
207     X = np.ndarray(shape = (rows, cols))
208     for i in xrange(rows):
209         line = raw[i]
210         #Construct data matrix
211         for j, feat in indexdict.iteritems():
212             val = line['data'][feat]
213             X[i, j] = val
214         #
215     #Replace NaNs with median values
216     if interpolate:
217         for j in xrange(cols):
218             #Get median of feature j
219             med = np.median([v for v in X[:, j] if not math.isnan(v)])
220             if math.isnan(med):
221                 raise ValueError('Feature %s contains only NaN's and should
222                                have been removed.' % indexdict[j])
223             for i in xrange(rows):
224                 if math.isnan(X[i, j]):
225                     X[i, j] = med
226
227     if normalize:
228         for j in xrange(cols):
229             col = X[:, j]
230             X[:, j] = normalize_data(col)
231
232     # ----- Handle class info -----
233     trait_values = []
234     for line in raw:
235         #Add value of psychological trait
236         psych_trait = line['profile'][trait]

```

```

237         if trait == 'Sex':
238             if psych_trait == 'Female':
239                 psych_trait = 0
240             elif psych_trait == 'Male':
241                 psych_trait = 1
242             else:
243                 raise ValueError('My code is binary gender normative, sorry.')
244         trait_values.append(psych_trait)
245     Y = []
246     if n_classes == None:
247         Y = trait_values
248     else:
249         ntiles = split_ntiles(trait_values, n_classes)
250         Y = [determine_ntile(tr, ntiles) for tr in trait_values]
251
252     return (X, Y, indexdict)
253
254 def get_strong_indices(X, Y, threshold):
255     '''Returns the indices of the features vectors X whose linear correlation
256     with Y is above the input threshold.'''
257     correlations = get_correlations(X, Y)
258     indices = [i for i in xrange(len(correlations))
259               if correlations[i] >= threshold]
260     return indices
261
262 def reduce_data(X, indices):
263     '''Return feature vectors including only the feature numbers in the input
264     list of indices.'''
265     reduced_X = np.array([X[i] for i in indices] for x in X)
266     return reduced_X
267
268 def plot_stuff(input_filename, output_filename=None, color='moerkeroed'):
269     with open(input_filename, 'r') as f:
270         d = json.load(f)
271         x = d['x']
272         y = d['y']
273         yerr = d['mean_stds']
274         plt.plot(x, y, color=color, linestyle='dashed',
275                 marker='o')
276         plt.errorbar(x, y, yerr=yerr, linestyle="None", marker="None",
277                     color=color)
278     if output_filename:
279         plt.savefig(output_filename)
280
281 def get_TPRs_and_FPRs(X, Y, forest = None, verbose = False):
282     '''Accepts a list of feature vectors and a list of labels and returns a
283     tuple of true positive and false positive rates (TPRs and FPRs,
284     respectively) for various confidence thresholds.'''
285     kf = LeaveOneOut(n=len(Y))
286
287     results = []
288     thresholds = []
289
290     counter = 0
291     for train, test in kf:
292         counter += 1
293         if counter % 10 == 0 and verbose:
294             print "Testing on user %s of %s..." % (counter, len(Y))
295
296         result = {}
297         train_data = [X[i] for i in train]
298         train_labels = [Y[i] for i in train]

```



```

299     test_data = [X[i] for i in test]
300     test_labels = [Y[i] for i in test]
301
302     if not forest:
303         forest = RandomForestClassifier(
304             n_estimators = 1000,
305             n_jobs=-1,
306             criterion='entropy')
307
308     forest.fit(train_data, train_labels)
309     result['prediction'] = forest.predict(test_data)[0]
310     result['true'] = test_labels[0]
311     confidences = forest.predict_proba(test_data)[0]
312     result['confidences'] = confidences
313     thresholds.append(max(confidences))
314
315     results.append(result)
316
317 #ROC curve stuff – false and true positive rates
318 TPRs = []
319 FPRs = []
320
321 unique_thresholds = sorted(list(set(thresholds)), reverse=True)
322
323 for threshold in unique_thresholds:
324     tn = 0
325     fn = 0
326     tp = 0
327     fp = 0
328     for result in results:
329         temp = result['prediction']
330         if temp == 1 and result['confidences'][1] >= threshold:
331             pred = 1
332         else:
333             pred = 0
334         if pred == 1:
335             if result['true'] == 1:
336                 tp += 1
337             else:
338                 fp += 1
339         #
340         elif pred == 0:
341             if result['true'] == 0:
342                 tn += 1
343             else:
344                 fn += 1
345         #
346     #
347     TPRs.append(tp/(tp + fn))
348     FPRs.append(fp/(fp + tn))
349     return (TPRs, FPRs)
350
351 def make_roc_curve(TPRs, FPRs, output_filename = None):
352     '''Accepts a list of true and false positive rates (TPRs and FPRs,
353     respectively) and generates a ROC-curve.'''
354     predcol = moerkeroed
355     basecol = oldhat
356     fillcol = nude
357
358     fig = plt.figure()
359     ax = fig.add_subplot(1,1,1)
360

```

```

361 TPRs = [0] + TPRs + [1]
362 FPRs = [0] + FPRs + [1]
363
364 area = 0.0
365 for i in xrange(len(TPRs)-1):
366     dx = FPRs[i+1] - FPRs[i]
367     y = 0.5*(TPRs[i] + TPRs[i+1])
368     under_curve = dx*y
369     baseline = dx*0.5*(FPRs[i] + FPRs[i+1])
370     area += under_curve - baseline
371
372 baseline = FPRs
373 ax.fill_between(x = FPRs, y1 = TPRs, y2 = baseline, color = fillcol,
374                interpolate = True, alpha=0.8)
375 ax.plot(baseline, baseline, color = basecol, linestyle = 'dashed',
376        linewidth = 1.0, label = 'Baseline')
377 ax.plot(FPRs, TPRs, color=predcol, linewidth = 1,
378        label = 'Prediction')
379
380 plt.xlabel('False positive rate.')
381 plt.ylabel('True positive rate.')
382
383 handles, labels = ax.get_legend_handles_labels()
384 hest = mpatches.Patch(color=fillcol)
385
386 labels += ['Area = %.3f' % area]
387 handles += [hest]
388 ax.legend(handles, labels, loc = 'lower right')
389 # plt.legend(handles = [tp_line, base])
390 if output_filename:
391     plt.savefig(output_filename)
392 plt.show()
393
394 def rank_features(X, Y, forest, indexdict, limit = None):
395     '''Ranks the features of a given dataset and classifier.
396     indexdict should be a map from indices to feature names like
397     {0 : 'average_weight'} etc.
398     if limit is specified, this method returns only the top n ranking features.
399     Returns a dict like {'feature name' : (mean importances, std)}.'''
400     importances = forest.feature_importances_
401     stds=np.std([tree.feature_importances_ for tree in forest.estimators_],
402                axis=0)
403     indices = np.argsort(importances)[::-1]
404     if limit:
405         indices = indices[:limit]
406     d = {indexdict[i]: (importances[i], stds[i]) for i in indices}
407     return d
408
409 def check_performance(X, Y, clf, strata = None):
410     '''Checks forest performance compared to baseline.'''
411     N_samples = len(Y)
412     #Set up validation indices
413     if not strata: #Do leave-one-out validation
414         skf = KFold(N_samples, n_folds=N_samples, shuffle = False)
415     else: #Do stratified K-fold
416         skf = StratifiedKFold(Y, n_folds=strata)
417
418     #Evaluate classifier performance
419     scores = []
420     for train, test in skf:
421         train_data = [X[ind] for ind in train]
422         train_labels = [Y[ind] for ind in train]

```

```

423     test_data = [X[ind] for ind in test]
424     test_labels = [Y[ind] for ind in test]
425
426     #Check performance of input forest
427     clf.fit(train_data, train_labels)
428     score = clf.score(test_data, test_labels)
429     scores.append(score)
430
431     #Compute baseline
432     most_common_label = max(Counter(Y).values())
433     baseline = float(most_common_label)/N_samples
434
435     #Compare results with prediction baseline
436     score_mean = np.mean(scores)/baseline
437     score_std = np.std(scores)/baseline
438
439     return (score_mean, score_std)
440
441 def check_regressor(X, Y, reg, strata = None):
442     '''Checks the performance of a regressor against mean value baseline.'''
443     N_samples = len(Y)
444     #Set up validation indices
445     if not strata: #Do leave-one-out validation
446         skf = KFold(N_samples, n_folds=N_samples,
447                     shuffle = False)
448     else: #Do stratified K-fold
449         skf = StratifiedKFold(Y, n_folds=strata)
450
451     #Evaluate performance
452     model_abs_errors = []
453     baseline_abs_errors = []
454     for train, test in skf:
455         train_data = [X[ind] for ind in train]
456         train_labels = [Y[ind] for ind in train]
457         test_data = [X[ind] for ind in test]
458         test_labels = [Y[ind] for ind in test]
459
460         #Check performance of input forest
461         reg.fit(train_data, train_labels)
462         base = np.mean(train_labels)
463         for i in xrange(len(test_data)):
464             pred = reg.predict(test_data[i])
465             true = test_labels[i]
466             model_abs_errors.append(np.abs(pred - true))
467             baseline_abs_errors.append(np.abs(base - true))
468         #
469     return (np.mean(model_abs_errors), np.mean(baseline_abs_errors))
470
471 class _RFNN(object):
472     __metaclass__ = abc.ABCMeta
473     '''Abstract class for random forest nearest neighbor predictors.
474     This should never be instantiated.'''
475
476     def __init__(self, forest, n_neighbors):
477         self.forest = forest
478         self.n_neighbors = n_neighbors
479         self.X = None
480         self.Y = None
481
482     def fit(self, X, Y):
483         '''Fits model to training data.
484

```

```

485     Args
486     -----
487     X : List
488         List of training feature vectors.
489
490     Y : List
491         List of training labels or values to be predicted.'''
492     if not len(X) == len(Y):
493         raise ValueError("Training input and output lists must have "
494             "same length.")
495     if not self.n_neighbors <= len(X):
496         raise ValueError("Fewer data points than neighbors.")
497
498     self.forest.fit(X, Y)
499     self.X = X
500     self.Y = Y
501
502     def _rf_similarity(self, a, b):
503         '''Computes a similarity measure for two points using a trained random
504         forest classifier.'''
505         if self.X == None or self.Y == None:
506             raise NotImplementedError("Model has not been fitted to data yet.")
507
508         #Feature vectors must be single precision.
509         a = np.array([a], dtype = np.float32)
510         b = np.array([b], dtype = np.float32)
511         hits = 0
512         tries = 0
513         for estimator in self.forest.estimators_:
514             tries += 1
515             tree = estimator.tree_
516             # Check whether the points end up on the same leaf for this tree
517             if tree.apply(a) == tree.apply(b):
518                 hits += 1
519             #
520         return hits/tries
521
522     def find_neighbors(self, point):
523         '''Determine the n nearest neighbors for the given point.
524         Returns a list of n tuples like (yval, similarity).
525         The tuples are sorted descending by similarity.'''
526         if self.X == None or self.Y == None:
527             raise NotImplementedError("Model has not been fitted to data yet.")
528
529         #Get list of tuples like (y, similarity) for the n 'nearest' points
530         nearest = [(None, float('-infinity')) for _ in xrange(self.n_neighbors)]
531         for i in xrange(len(self.X)):
532             similarity = self._rf_similarity(self.X[i], point)
533             # update top n list if more similar than the furthest neighbor
534             if similarity > nearest[-1][1]:
535                 nearest.append((self.Y[i], similarity))
536                 nearest.sort(key = lambda x: x[1], reverse = True)
537                 del nearest[-1]
538             #
539         return nearest
540
541     #Mandatory methods – must be overridden
542     @abc.abstractmethod
543     def predict(self, point):
544         pass
545
546     @abc.abstractmethod

```

```

547     def score(self, X, Y):
548         pass
549
550     def _reservoir_sampler(start = 1):
551         '''Generator of the probabilities need to do reservoir sampling. The point
552         it that this can be used to iterate through a list, discarding each element
553         for the following element with probability P_n and ending up with a random
554         element from the list.'''
555         n = start
556         while True:
557             p = 1/n
558             r = random.uniform(0,1)
559             if r < p:
560                 yield True
561             else:
562                 yield False
563             n += 1
564
565     class RFNNClassifier(_RFNN):
566         '''Random Forest Nearest Neighbor Classifier.
567
568         Parameters
569         ----
570         n_neighbors : int
571             Number of neighbors to consider.
572
573         forest : RandomForestClassifier
574             The forest which will provide a
575             distance measure on which determine nearest neighbors.
576
577         weighting : str
578             How to weigh the votes of different neighbors.
579             'equal' means each neighbor has an equivalent vote.
580             'linear' mean votes are weighed by their similarity to the input point.
581         '''
582     def predict(self, point):
583         '''Predicts the label of a given point.'''
584         neighbortuples = self.find_neighbors(point)
585         if self.weighting == 'equal':
586             #Simple majority vote. Select randomly if it's a tie.
587             predictions = [t[0] for t in neighbortuples]
588             best = 0
589             winner = None
590             switch = _reservoir_sampler(start = 2)
591             for label, votes in Counter(predictions).iteritems():
592                 if votes > best:
593                     best = votes
594                     winner = label
595                     switch = _reservoir_sampler(start = 2)
596                 elif votes == best:
597                     if switch.next():
598                         winner = label
599                     else:
600                         pass
601                 else:
602                     pass
603             #
604             return winner
605
606         #Weigh votes by their similarity to the input point
607         elif self.weighting == 'linear':
608             #The votes are weighted by their similarity

```

```

609         d = {}
610         for yval, similarity in neighbortuples:
611             try:
612                 d[yval] += similarity
613             except KeyError:
614                 d[yval] = similarity
615         best = float('-infinity')
616         winner = None
617         for k, v in d.iteritems():
618             if v > best:
619                 best = v
620                 winner = k
621             else:
622                 pass
623         return winner
624
625     def score(self, X, Y):
626         if not len(X) == len(Y):
627             raise ValueError("Training data and labels must have same length.")
628
629         hits = 0
630         n = len(X)
631         for i in xrange(n):
632             pred = self.predict(X[i])
633             if pred == Y[i]:
634                 hits += 1
635         #
636         return hits/n
637
638
639     def __init__(self, forest = None, n_neighbors = 3, weighting = 'equal',
640                 n_jobs = 1):
641         #Make sure we have a forest * classifier *
642         if forest == None:
643             forest = RandomForestClassifier(n_estimators = 1000,
644                                           criterion = 'entropy',
645                                           n_jobs = n_jobs)
646         if not isinstance(forest, RandomForestClassifier):
647             raise TypeError("Forest must be a classifier")
648
649         self.weighting = weighting
650
651         #Call parent constructor
652         super(RFNNClassifier, self).__init__(forest, n_neighbors)
653
654     class RFNNRegressor(_RFNN):
655         '''Random Forest Nearest Neighbor Regressor.
656
657         Parameters
658         ----
659         n_neighbors : int
660             Number of neighbors to consider.
661
662         forest : RandomForestRegressor
663             The forest which will provide a
664             distance measure on which determine nearest neighbors.
665
666         weighting : str
667             How to weigh the votes of different neighbors.
668             'equal' means each neighbor has an equivalent weight.
669             'linear' mean votes are weighed by their similarity to the input point.
670         '''

```

```

671 def predict(self, point):
672     # lists of the y vaues and similarities of nearest neighbors
673     neighbortuples = self.find_neighbors(point)
674     yvals, similarities = zip(*neighbortuples)
675
676     # Weigh each neighbor y value equally is that's how we roll
677     if self.weighting == 'equal':
678         weight = 1.0/len(yvals)
679         result = 0.0
680         for y in yvals:
681             result += y*weight
682         return result
683     #
684     # Otherwise, weigh neighbors by similarity
685     elif self.weighting == 'linear':
686         weight = 1.0/(sum(similarities))
687         result = 0.0
688         for i in xrange(len(yvals)):
689             y = yvals[i]
690             similarity = similarities[i]
691             result += y*similarity*weight
692         return result
693
694
695 def score(self, X, Y):
696     if not len(X) == len(Y):
697         raise ValueError("X and Y must be same length.")
698     errors = [Y[i] - self.predict(X[i]) for i in xrange(len(X))]
699     return np.std(errors)
700
701
702 def __init__(self, forest = None, n_neighbors = 3, weighting = 'equal'):
703     #Check forest type.
704     if forest == None:
705         forest = RandomForestRegressor(n_estimators = 1000, n_jobs = -1)
706     if not isinstance(forest, RandomForestRegressor):
707         raise TypeError("Must use Random Forest Regressor to initialize.")
708     # Set params
709     self.weighting = weighting
710     # Done. Call parent constructor
711     super(RFNNRegressor, self).__init__(forest, n_neighbors)
712
713
714 class _BaselineRegressor(object):
715     '''Always predicts the mean of the training set.'''
716     def __init__(self, guess=None):
717         self.guess = guess
718     def fit(self, xtrain, ytrain):
719         '''Find the average of input list of target values and guess on that
720         from now on.'''
721         self.guess = np.mean(ytrain)
722     def predict(self, x):
723         return self.guess
724
725
726 class _BaselineClassifier(object):
727     '''Always predicts the most common label in the training set'''
728     def __init__(self, guess=None):
729         self.guess = guess
730     def fit(self, xtrain, ytrain):
731         '''Find the most common label and guess on that from now on.'''
732         countmap = Counter(ytrain)

```

```

733     best = 0
734     for label, count in countmap.iteritems():
735         if count > best:
736             best = count
737             self.guess = int(label)
738         #
739     #
740     def predict(self, x):
741         return self.guess
742
743
744     def _worker(X, Y, score_type, train_percentage, classifier, clf_args, n_groups,
745               replace, threshold):
746         '''Worker method for parallelizing bootstrap evaluations.'''
747         #Create bootstrap sample
748         try:
749             rand = np.random.RandomState() #Ensures PRNG works in children
750             indices = rand.choice(xrange(len(X)), size = len(X), replace = replace)
751             xsample = [X[i] for i in indices]
752             ysample = [Y[i] for i in indices]
753
754             #Generate training and testing set
755             cut = int(train_percentage*len(X))
756             xtrain = xsample[:cut]
757             xtest = xsample[cut:]
758             ytrain = ysample[:cut]
759             ytest = ysample[cut:]
760
761             #Discard features with too low correlation with output vectors
762             inds = get_strong_indices(X=xtrain, Y=ytrain, threshold = threshold)
763             xtrain = reduce_data(xtrain, inds)
764             xtest = reduce_data(xtest, inds)
765
766             #Create regressor if we're doing regression
767             if classifier == 'RandomForestRegressor':
768                 clf = RandomForestRegressor(**clf_args)
769             elif classifier == 'SVR':
770                 clf = svm.SVR(**clf_args)
771             elif classifier == 'baseline_mean':
772                 clf = _BaselineRegressor()
773             elif classifier == 'WRBFR':
774                 importances = get_correlations(xtrain, ytrain)
775                 clf_args['importances'] = importances
776                 clf = WRBFR(**clf_args)
777
778             #Create classifier and split dataset into labels
779             elif classifier == 'RandomForestClassifier':
780                 clf = RandomForestClassifier(**clf_args)
781                 ysample = assign_labels(ysample, n_groups)
782                 ytrain = ysample[:cut]
783                 ytest = ysample[cut:]
784             elif classifier == 'SVC':
785                 clf = svm.SVC(**clf_args)
786                 ysample = assign_labels(ysample, n_groups)
787                 ytrain = ysample[:cut]
788                 ytest = ysample[cut:]
789             elif classifier == 'baseline_most_common_label':
790                 clf = _BaselineClassifier()
791                 ysample = assign_labels(ysample, n_groups)
792                 ytrain = ysample[:cut]
793                 ytest = ysample[cut:]
794             elif classifier == 'WRBFC':

```



```

795     importances = get_correlations(xtrain, ytrain)
796     clf_args['importances'] = importances
797     clf = WRBFC(**clf_args)
798     ysample = assign_labels(ysample, n_groups)
799     ytrain = ysample[:cut]
800     ytest = ysample[cut:]
801
802     #Fail if none of the above classifiers were specified
803     else:
804         raise ValueError('Regressor or classifier not defined.')
805
806     #Fit the classifier or regressor
807     clf.fit(xtrain, ytrain)
808
809     #Compute score and append to output list
810     if score_type == 'mse':
811         scores = [(ytest[i] - clf.predict(xtest[i]))**2
812                  for i in xrange(len(xtest))]
813         return np.mean(scores)
814
815     elif score_type == 'fraction_correct':
816         n_correct = sum([ytest[i] == int(clf.predict(xtest[i]))
817                        for i in xrange(len(ytest))])
818         score = n_correct/len(ytest)
819         return score
820     elif score_type == 'over_baseline':
821         #Get score
822         score = sum([ytest[i] == int(clf.predict(xtest[i]))
823                    for i in xrange(len(ytest))])
824         #Get baseline
825         baselineclf = _BaselineClassifier()
826         baselineclf.fit(xtrain, ytrain)
827         baseline = sum([ytest[i] == baselineclf.predict(xtest[i])
828                       for i in xrange(len(ytest))])
829         return score/baseline
830
831     #Fail if none of the above performance metrics were specified
832     else:
833         raise ValueError('Score type not defined.')
834
835     #Job's done!
836     return None
837 except:
838     raise Exception(''.join(traceback.format_exception(*sys.exc_info())))
839
840
841 def bootstrap(X, Y, classifier, score_type = 'mse', train_percentage = 0.8,
842             clf_args = {}, iterations = 1000, n_groups = 3, n_jobs = 1,
843             replace = False, threshold = 0.0):
844     '''Performs bootstrap resampling to evaluate the performance of some
845     classifier or regressor. Note that this takes the *complete dataset* as
846     arguments as well as arguments specifying which predictor to use and which
847     function to estimate the distribution of.
848     This seems to be the most straightforward generalizable implementation
849     which can be parallelized, as passing e.g. the scoring function directly
850     clashed with the mechanisms implemented to work around the GIL for
851     multiprocessing for obscure reasons.
852
853     Parameters:
854     -----
855     X : list
856         All feature vectors in the complete dataset.

```

```

857
858 Y : list
859     All 'true' labels or output values in the complete dataset.
860
861 classifier : str
862     Which classifier to use to predict the test set. Allowed values:
863     'RandomForestRegressor', 'baseline_mean', 'SVR',
864     'RandomForestClassifier', 'SVC', 'baseline_most_common_label, WRBFR,
865     WRBFC'
866
867 score_type : str
868     String signifying which function to estimate the distribution of.
869     Allowed values: 'mse', 'fraction_correct', 'over_baseline'
870
871 train_percentage : float
872     The percentage [0:1] of each bootstrap sample to be used for training.
873
874 clf_args : dict
875     optional arguments to the constructor method of the regressor/classifier.
876
877 iterations : int
878     Number of bootstrap samples to run.
879
880 n_jobs : int
881     How many cores (maximum) to use.
882
883 replace : bool
884     Whether to sample with replacement when obtaining the bootstrap samples.
885
886 threshold : float
887     features whose linear correlations with vectors in the output space
888     are below this value are discarded.
889     ,''
890
891 if not len(X) == len(Y):
892     raise ValueError("X and Y must have equal length.")
893
894 #Arguments to pass to worker processes
895 d = {'X' : X, 'Y' : Y, 'train_percentage' : train_percentage,
896     'classifier' : classifier, 'clf_args' : clf_args,
897     'score_type' : score_type, 'n_groups' : n_groups,
898     'replace' : replace, 'threshold' : threshold}
899
900 #Make job queue
901 pool = multiprocessing.Pool(processes = n_jobs)
902 jobs = [pool.apply_async(_worker, kwds = d) for _ in xrange(iterations)]
903 pool.close() #run
904 pool.join() #Wait for remaining jobs
905
906 #Make sure no children died too early
907 if not all(job.successful() for job in jobs):
908     raise RuntimeError('Some jobs failed.')
909
910 return [j.get() for j in jobs]
911
912 def get_correlations(X, Y, absolute = True):
913     '''Given a list of feature vectors X and labels or values Y, returns a list
914     of correlation coefficients for each dimension of the feature vectors.'''
915     n_feats = len(X[0])
916     correlations = []
917     for i in xrange(n_feats):
918         temp = np.corrcoef([x[i] for x in X], Y)

```

```

919     correlation = temp[0,1]
920     if math.isnan(correlation):
921         correlation = 0
922     correlations.append(correlation)
923 if absolute:
924     correlations = [np.abs(c) for c in correlations]
925 return correlations
926
927
928 def make_kernel(importances, gamma = 1.0):
929     '''Returns a weighted radial basis function (WRBF) kernel which can be
930     passed to an SVM or SVR from the sklearn module.
931
932     Parameters:
933     -----
934     importances : list
935         The importance of each input feature. The value of element i can mean
936         e.g. the linear correlation between feature i and target variable y.
937         None means feature will be weighted equally.
938
939     gamma : float
940         The usual gamma parameter denoting inverse width of the gaussian used.
941     '''
942     def kernel(x,y, *args, **kwargs):
943         d = len(importances) #number of features
944         impsum = sum([imp**2 for imp in importances])
945         if not impsum == 0:
946             normfactor = 1.0/np.sqrt(impsum)
947         else:
948             normfactor = 0.0
949         #Metric to compute distance between points
950         metric = dok_matrix((d,d), dtype = np.float64)
951         for i in xrange(d):
952             metric[i,i] = importances[i]*normfactor
953         #
954         result = np.zeros(shape = (len(x), len(y)))
955         for i in xrange(len(x)):
956             for j in xrange(len(y)):
957                 diff = x[i] - y[j]
958                 dist = diff.T.dot(metric*diff)
959                 result[i,j] = np.exp(-gamma*dist)
960         return result
961     return kernel
962
963 class WRBFR(svm.SVR):
964     '''Weighted radial basis function support vector regressor.'''
965     def __init__(self, importances, C = 1.0, epsilon = 0.1,
966                 gamma = 0.0):
967         kernel = make_kernel(importances = importances, gamma = gamma)
968         super(WRBFR, self).__init__(C = C, epsilon = epsilon, kernel = kernel)
969
970 class WRBFC(svm.SVC):
971     '''Weighted radial basis function support vector classifier.'''
972     def __init__(self, importances, C = 1.0, gamma = 0.0):
973         kernel = make_kernel(importances = importances, gamma = gamma)
974         super(WRBFC, self).__init__(C = C, kernel = kernel)
975
976 if __name__ == '__main__':
977
978     random.seed(42)
979     X, Y, ind_dict = read_data('../data.json', trait = 'extraversion',

```

```

980 #         features = ['call_iet_med', 'text_iet_med', 'social_entropy', 'call_entropy', 'travel_med', 'n_places', 'text_latency', 'call_night_activity']
981         features = 'all'
982     )
983     corrs = get_correlations(X, Y)
984
985     # for i in xrange(len(corrs)):
986     #     print corrs[i], ind_dict[i]
987
988     cut = int(0.6*len(X))
989     xtrain = X[:cut]
990     ytrain = Y[:cut]
991     xtest = X[cut:]
992     ytest = Y[cut:]
993
994     inds = get_strong_indices(xtrain, ytrain, threshold = 0.08)
995     ind_dict = {i : ind_dict[inds[i]] for i in xrange(len(inds))}
996     xtrain = reduce_data(xtrain, inds)
997     xtest = reduce_data(xtest, inds)
998
999     importances = get_correlations(xtrain, ytrain)
1000     clf = WRBFR(importances = importances, C = 26.5, epsilon=0.11, gamma = 0.23)
1001     clf_baseline = _BaselineRegressor()
1002     #     clf = svm.SVR(C = 26.5, epsilon=0.11, gamma = 0.23)
1003     clf.fit(xtrain, ytrain)
1004     clf_baseline.fit(xtrain, ytrain)
1005     scores_baseline = []
1006     scores = []
1007     for i in xrange(len(xtest)):
1008         pred = clf.predict(xtest[i])
1009         pred_baseline = clf_baseline.predict(xtest[i])
1010         print pred
1011         scores.append((pred - ytest[i])**2)
1012         scores_baseline.append((pred_baseline - ytest[i])**2)
1013     print np.mean(scores)**0.5
1014     print np.mean(scores_baseline)**0.5
1015
1016     # X = [[1,2,7,0],[3,1,6,0.01],[6,8,1,0],[10,8,2,0.01]]
1017     # Y = [1,1,0,0]
1018     #
1019     #
1020     #
1021     # corrs = get_correlations(X, Y)
1022     # print corrs
1023     #
1024     # C = 70
1025     # gamma = 3.75
1026     #
1027     # kernel = make_kernel(corrs, 0.0)
1028     #
1029     # clf = svm.SVC(kernel = kernel)
1030     # clf = svm.SVC()
1031     # clf.fit(xtrain, ytrain)
1032     #
1033     # hits = 0
1034     #
1035     # for i in xrange(len(xtest)):
1036     #     if clf.predict(xtest[i]) == ytest[i]:
1037     #         hits += 1
1038     #
1039     # print 100.0*hits/len(ytest)
1040     #

```

```

1041 # for i in xrange(len(corrs)):
1042 #     print ind_dict[i], corrs[i]
1043 #
1044 #     test = WRBFC(threshold=0.1, gamma = 0.2, importances=corrs)
1045 #     test.fit(xtrain, ytrain)
1046 #
1047 #     preds = []
1048 #     for i in xrange(len(ytest)):
1049 #         pred = test.predict(xtest[i])[0]
1050 #         print pred, ytest[i]
1051 #         preds.append(pred)
1052 #
1053 #     print len(Counter(preds).keys())
1054 #
1055 #     print len(X[0])
1056 ##     print i
1057 #     print [el for el in i.values() if 'init' in el]

```

A.1.5 Lloyd's Algorithm

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import numpy as np
5  import matplotlib
6  matplotlib.use('Agg') #ugly hack to allow plotting from terminal
7  import matplotlib.pyplot as plt
8  import random
9  from copy import deepcopy
10
11 def _dist(p,q):
12     return sum([(p[i]-q[i])**2 for i in xrange(len(p))])
13
14 def _lloyds_single_run(X,K,max_iterations,init):
15     # Initialize with a subset of the data
16     if init == 'sample':
17         initials = random.sample(X,K)
18     # Or initialize with random points across the same range as data
19     elif init == 'scatter':
20         vals = zip(*X)
21         xmin = min(vals[0])
22         xmax = max(vals[0])
23         ymin = min(vals[1])
24         ymax = max(vals[1])
25         initials = [(random.uniform(xmin, xmax),
26                        random.uniform(ymin, ymax)) for _ in xrange(K)]
27     # Or yell RTFM at user
28     else:
29         raise ValueError('Invalid initialization mode!')
30
31     #Construct hashmap mapping integers up to K to centroids
32     centroids = dict(enumerate(initials))
33     converged = False
34     iterations = 0
35
36     while not converged and iterations < max_iterations:
37         clusters = {i: [] for i in xrange(K)}

```

```

38     #Make sure clusters and centroids have identical keys, or we're doomed.
39     assert set(clusters.keys()) == set(centroids.keys())
40     prev_centroids = deepcopy(centroids)
41
42     ### STEP ONE -update clusters
43     for x in X:
44         #Check distances to all centroids
45         bestind = -1
46         bestdist = float('inf')
47         for ind, centroid in centroids.iteritems():
48             dist = _dist(x, centroid)
49             if dist < bestdist:
50                 bestdist = dist
51                 bestind = ind
52         #
53         clusters[bestind].append(x)
54
55     ### STEP TWO -update centroids
56     for ind, points in clusters.iteritems():
57         if not points:
58             pass #Cluster's empty - nothing to update
59         else:
60             centroids[ind] = np.mean(points, axis = 0)
61
62     ### We're converged when all old centroids = new centroids.
63     converged = all([_dist(prev_centroids[k], centroids[k]) == 0
64                     for k in xrange(K)])
65     iterations += 1
66     #
67     return {tuple(centroids[i]): clusters[i] for i in xrange(K)}
68
69 def lloyds(X, K, runs = 1, max_iterations = float('inf'), init = 'sample'):
70     '''Runs Lloyd's algorithm to identify K clusters in the dataset X.
71     X is a list of points like [[x1,y1],[x2,y2]---].
72     Returns a hash of centroids mapping to points in the corresponding cluster.
73     The objective is to minimize the sum of distances from each centroid to
74     the points in the corresponding cluster. It might only converge on a local
75     minimum, so the configuration with the lowest score (sum of distances) is
76     returned.
77     init denotes initialization mode, which can be 'sample', using a randomly
78     select subset of the input data, or 'scatter', using random points selected
79     from the same range as the data as initial centroids.
80
81     Parameters
82     -----
83     X : array_like
84         list of points. 2D example: [[3,4],[3.4, 7.2], ...]
85
86     K : int
87         Number of centroids
88
89     runs : int
90         Number of times to run the entire algorithm. The result with the lowest
91         score will be returned.
92
93     max_iterations : int or float
94         Number of steps to allow each run. Default if infinit, i.e. the algorithm
95         runs until it's fully converged.
96
97     init : str
98         Initialization mode. 'sample' means use a random subset of the data as
99         starting centroids. 'scatter' means place starting centroids randomly in

```

```

100     the entire x-y range of the dataset.
101
102     Returns
103     -----
104     result : dict
105         A dictionary in which each key is a tuple of coordinated corresponding to
106         a centroid, and each value is a list of points belonging to that cluster.
107     '''
108
109     record = float('inf')
110     result = None
111     for _ in xrange(runs):
112         clusters = _lloyds_single_run(X, K, max_iterations = max_iterations,
113                                     init = init)
114         #Determine how good the clusters came out
115         score = 0
116         for centroid, points in clusters.iteritems():
117             score += sum([_dist(centroid, p) for p in points or []])
118         if score < record:
119             result = clusters
120             record = score
121     #
122     return result
123
124
125 def _makecolor():
126     i = 0
127     cols = ['b', 'g', 'r', 'c', 'm', 'y']
128     while True:
129         yield cols[i]
130         i = (i+1)%len(cols)
131
132
133 def draw_clusters(clusters, threshold = 0, show = True, filename = None):
134     '''Accepts a dict mapping cluster centroids to cluster points and makes
135     a color-coded plot of them. Clusters containing fewer points than the
136     threshold are plottet in black.'''
137     colors = _makecolor()
138     plt.figure()
139     for centroid, points in clusters.iteritems():
140         if not points:
141             continue
142         if len(points) < threshold:
143             style = ['k,']
144         else:
145             color = colors.next()
146             style = [color+'+']
147             #Plot centroids
148             x,y = centroid
149             plt.plot(x,y, color = color, marker = 'd', markersize = 12)
150             #plot points
151             plt.plot(*(zip(*points))+style))
152     if filename:
153         plt.savefig(filename, bbox_inches = 'tight')
154     if show:
155         plt.show()
156
157
158 if __name__ == '__main__':
159     points = [[random.uniform(-10,10), random.uniform(-10,10)] for _ in xrange(10*3)]
160     clusters = lloyds(X = points, K = 6, runs = 1)
161     draw_clusters(clusters = clusters, filename = 'lloyds_example.pdf')

```

A.1.6 Smallest Enclosing Circle

```

1  # -*- coding: utf-8 -*-
2  #
3  # Smallest enclosing circle
4  #
5  # Copyright (c) 2014 Project Nayuki
6  # http://www.nayuki.io/page/smallest-enclosing-circle
7  #
8  # This program is free software: you can redistribute it and/or modify
9  # it under the terms of the GNU General Public License as published by
10 # the Free Software Foundation, either version 3 of the License, or
11 # (at your option) any later version.
12 #
13 # This program is distributed in the hope that it will be useful,
14 # but WITHOUT ANY WARRANTY; without even the implied warranty of
15 # MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
16 # GNU General Public License for more details.
17 #
18 # You should have received a copy of the GNU General Public License
19 # along with this program (see COPYING.txt).
20 # If not, see <http://www.gnu.org/licenses/>.
21 #
22
23 import math, random
24
25
26 # Data conventions: A point is a pair of floats (x, y). A circle is a triple of floats (center x, center y, radius).
27
28 #
29 # Returns the smallest circle that encloses all the given points. Runs in expected O(n) time, randomized.
30 # Input: A sequence of pairs of floats or ints, e.g. [(0,5), (3.1,-2.7)].
31 # Output: A triple of floats representing a circle.
32 # Note: If 0 points are given, None is returned. If 1 point is given, a circle of radius 0 is returned.
33 #
34 def make_circle(points):
35     '''Accepts list of points as tuples and returns (x, y, r).'''
36     # Convert to float and randomize order
37     shuffled = [(float(p[0]), float(p[1])) for p in points]
38     random.shuffle(shuffled)
39
40     # Progressively add points to circle or recompute circle
41     c = None
42     for (i, p) in enumerate(shuffled):
43         if c is None or not _is_in_circle(c, p):
44             c = _make_circle_one_point(shuffled[0:i+1], p)
45     return c
46
47
48 # One boundary point known
49 def _make_circle_one_point(points, p):
50     c = (p[0], p[1], 0.0)
51     for (i, q) in enumerate(points):
52         if not _is_in_circle(c, q):
53             if c[2] == 0.0:

```



```

54         c = _make_diameter(p, q)
55     else:
56         c = _make_circle_two_points(points[0:i + 1], p, q)
57     return c
58
59 # Two boundary points known
60 def _make_circle_two_points(points, p, q):
61     diameter = _make_diameter(p, q)
62     if all(_is_in_circle(diameter, r) for r in points):
63         return diameter
64
65     left = None
66     right = None
67     for r in points:
68         cross = _cross_product(p[0], p[1], q[0], q[1], r[0], r[1])
69         c = _make_circumcircle(p, q, r)
70         if c is None:
71             continue
72         elif cross > 0.0 and (left is None or _cross_product(p[0], p[1], q[0], q[1], c[0], c[1]) > _cross_product(←
p[0], p[1], q[0], q[1], left[0], left[1])):
73             left = c
74         elif cross < 0.0 and (right is None or _cross_product(p[0], p[1], q[0], q[1], c[0], c[1]) < _cross_product(←
p[0], p[1], q[0], q[1], right[0], right[1])):
75             right = c
76     return left if (right is None or (left is not None and left[2] <= right[2])) else right
77
78
79 def _make_circumcircle(p0, p1, p2):
80     # Mathematical algorithm from Wikipedia: Circumscribed circle
81     ax = p0[0]; ay = p0[1]
82     bx = p1[0]; by = p1[1]
83     cx = p2[0]; cy = p2[1]
84     d = (ax * (by - cy) + bx * (cy - ay) + cx * (ay - by)) * 2.0
85     if d == 0.0:
86         return None
87     x = ((ax * ax + ay * ay) * (by - cy) + (bx * bx + by * by) * (cy - ay) + (cx * cx + cy * cy) * (ay - by)) / d
88     y = ((ax * ax + ay * ay) * (cx - bx) + (bx * bx + by * by) * (ax - cx) + (cx * cx + cy * cy) * (bx - ax)) / d
89     return (x, y, math.hypot(x - ax, y - ay))
90
91
92 def _make_diameter(p0, p1):
93     return ((p0[0] + p1[0]) / 2.0, (p0[1] + p1[1]) / 2.0, math.hypot(p0[0] - p1[0], p0[1] - p1[1]) / 2.0)
94
95
96 _EPSILON = 1e-12
97
98 def _is_in_circle(c, p):
99     return c is not None and math.hypot(p[0] - c[0], p[1] - c[1]) < c[2] + _EPSILON
100
101
102 # Returns twice the signed area of the triangle defined by (x0, y0), (x1, y1), (x2, y2)
103 def _cross_product(x0, y0, x1, y1, x2, y2):
104     return (x1 - x0) * (y2 - y0) - (y1 - y0) * (x2 - x0)
105
106
107 #pts = [(0.0, 0.0), (6.0, 8.0)]
108 #
109 #test = make_circle(pts)
110 #
111 #print test

```

A.2 Source Code for Explicit Semantic Analysis

This section contains code pertaining to part II of the thesis.

A.2.1 Parser

```

1  # -*- coding: utf-8 -*-
2  '''Parses a full Wikipedia XML-dump and saves to files containing
3  a maximum of 1000 articles.
4  In the end, each file is saved as a JSON file containing entries like:
5  {{
6      'concept':
7      {
8          'text': <article contents>,
9          'links_in' : Set of links TO the article in question,
10         'links_out' : Set of links FROM the article in question,
11     }
12 }
13 Although links_in is added by the generate_indices script.
14 Also saved are dicts for keeping track of word and concept indices when
15 building a large sparse matrix for the semantic interpreter.
16 The file structure is like {'word blah' : index blah}'''
17
18 import re
19 import xml.sax as SAX
20 import wikicleaner
21 import os
22 import glob
23 import shared
24 import sys
25
26 DEFAULT_FILENAME = 'medium_wiki.xml'
27
28 def canonize_title(title):
29     # remove leading whitespace and underscores
30     title = title.strip(' _')
31     # replace sequences of whitespace and underscore chars with a single space
32     title = re.compile(r'[\s_]+').sub(' ', title)
33     # remove forbidden characters
34     title = re.sub('[?/\*\\"\'', '', title)
35     return title.title()
36
37 # Import shared parameters
38 from shared import extensions, temp_dir
39
40 # Cleanup
41 for ext in extensions.values():
42     for f in glob.glob(temp_dir + '*' + ext):
43         os.remove(f)
44
45 def filename_generator(folder):
46     '''Generator for output filenames'''
47     if not os.path.exists(folder):
48         os.makedirs(folder)
49     count = 0
50     while True:
51         filename = folder + "content" + str(count)

```

```

52         count += 1
53         yield filename
54
55     make_filename = filename_generator(temp_dir)
56
57     #Format {right title : redirected title }, e.g. {because: ([cuz, cus])}
58     redirects = {}
59
60     #Minimum number of links/words required to keep an article.
61     from shared import min_links_out, min_words
62
63     #Open log file for writing and import logging function
64     logfile = open(os.path.basename(__file__) + '.log', 'w')
65     log = shared.logmaker(logfile)
66
67     class WikiHandler(SAX.ContentHandler):
68         '''ContentHandler class to process XML and deal with the WikiText.
69         It works basically like this:
70         It traverses the XML file, keeping track of the type of data being read and
71         adding any text to its input buffer. When event handlers register a page
72         end, the page content is processed, the processed content is placed in the
73         output buffer, and the input buffer is flushed.
74         Whenever a set number of articles have been processed, the output buffer is
75         written to a file. The point of this approach is to
76         limit memory consumption.'''
77
78     def __init__(self):
79         SAX.ContentHandler.__init__(self)
80         self.current_data = None
81         self.title = ''
82         self.input_buffer = []
83         self.output_buffer = {}
84         self.article_counter = 0
85         self.links = []
86         self.categories = []
87         self.redirect = None
88         self.verbose = False
89         #Harvest unique words here
90         self.words = set([])
91         #keeps track of ingoing article links. format {to : set([from])}
92         self.linkhash = {}
93
94     def flush_input_buffer(self):
95         '''Deletes info on the currently processed article.
96         This is called when a page end event is registered.'''
97         self.input_buffer = []
98         self.current_data = None
99         self.title = ''
100         self.links = []
101         self.categories = []
102         self.redirect = None
103
104     def flush_output_buffer(self):
105         '''Flushes data gathered so far to a file and resets.'''
106         self.output_buffer = {}
107         self.words = set([])
108         self.linkhash = {}
109
110     def startElement(self, tag, attrs):
111         '''EventHandler for element start - keeps track of current datatype.'''
112         self.current_data = tag
113         #Informs the parser of the redirect destination of the article

```

```

114         if tag == "redirect":
115             self.redirect = attrs['title']
116             return None
117
118     def endElement(self, name):
119         '''Eventhandler for element end. This causes the parser to process
120         its input buffer when a pageend is encountered.'''
121         #Process content after each page
122         if name == 'page':
123             self.process()
124         #Write remaining data at EOF.
125         elif name == 'mediawiki':
126             self.writeout()
127
128     def characters(self, content):
129         '''Character event handler. This simply passes any raw text from an
130         article field to the input buffer and updates title info.'''
131         if self.current_data == 'text':
132             self.input_buffer.append(content)
133         elif self.current_data == 'title' and not content.isspace():
134             self.title = content
135
136     def process(self):
137         '''Process input buffer contents. This converts wikilanguage to
138         plaintext, registers link information and checks if content has
139         sufficient words and outgoing links (ingoing links can't be checked
140         until the full XML file is processed).'''
141
142         #Ignore everything else if article redirects
143         if self.redirect:
144             self.flush_input_buffer()
145             return None
146         global redirects
147         try:
148             redirects[self.title].add(self.redirect)
149         except KeyError:
150             redirects[self.title] = set([self.redirect])
151         self.flush_input_buffer()
152         return None
153
154         #Redirects handled – commence processing
155         print "processing: " + self.title.encode('utf8')
156         #Combine buffer content to a single string
157         text = ''.join(self.input_buffer).lower()
158
159         #Find and process link information
160         link_regexp = re.compile(r'\[([^\[\]\(\)]*)\]')
161         links = re.findall(link_regexp, text) #grap stuff like [[<something>]]
162         #Add links to the parsers link hash
163         for link in links:
164             #Check if link matches a namespace, e.g. 'file:something.png'
165             if any([ns+':' in link for ns in wikicleaner.namespaces]):
166                 continue #Proceed to next link
167             #Namespaces done, so remove any colons:
168             link = link.replace(':', '')
169             if not link:
170                 continue #Some noob could've written an empty link...
171             #remove chapter designations/displaytext – keep article title
172             raw = re.match(r'([^\|\\#]*)', link).group(0)
173             title = canonize_title(raw)
174             #note down that current article has outgoing link to 'title'
175             self.links.append(title)

```

```

176         #also note that 'title' has incoming link from here
177         try:
178             self.linkhash[title].add(self.title) #maps target->sources
179         except KeyError:
180             self.linkhash[title] = set([self.title])
181
182     #Disregard current article if it contains too few links
183     if len(self.links) < min_links_out:
184         self.flush_input_buffer()
185         return None
186
187     #Cleanup text
188     text = wikicleaner.clean(text)
189     article_words = text.split()
190
191     #Disregard article if it contains too few words
192     if len(article_words) < min_words:
193         self.flush_input_buffer()
194         return None
195
196     #Update global list of unique words
197     self.words.update(set(article_words))
198
199     #Add content to output buffer
200     output = {
201         'text': text,
202         #Don't use category info for now
203         #'categories': self.categories,
204         'links_out': self.links
205     }
206     self.output_buffer[self.title] = output
207     self.article_counter += 1
208
209     #Flush output buffer to file
210     if self.article_counter%1000 == 0:
211         self.writeout()
212
213     #Done, flushing buffer
214     self.flush_input_buffer()
215     return None
216
217 def writeout(self):
218     '''Writes output buffer contents to file'''
219     #Generate filename and write to file
220     filename = make_filename.next()
221     #Write article contents to file
222     with open(filename+extensions['content'], 'w') as f:
223         shared.dump(self.output_buffer, f)
224
225     #Store wordlist as files
226     with open(filename+extensions['words'], 'w') as f:
227         shared.dump(self.words, f)
228
229     #Store linkhash in files
230     with open(filename+extensions['links'], 'w') as f:
231         shared.dump(self.linkhash, f)
232
233     if self.verbose:
234         log("wrote "+filename)
235
236     #Empty output buffer
237     self.flush_output_buffer()

```

```
238         return None
239
240     if __name__ == "__main__":
241         if len(sys.argv) == 2:
242             file_to_parse = sys.argv[1]
243         else:
244             file_to_parse = DEFAULT_FILENAME
245
246         #Create and configure content handler
247         test = WikiHandler()
248         test.verbose = True
249
250         #Create a parser and set handler
251         ATST = SAX.make_parser()
252         ATST.setContentHandler(test)
253
254         #Let the parser walk the file
255         log("Parsing started...")
256         ATST.parse(file_to_parse)
257         log("...Parsing done!")
258
259         #Attempt to send notification that job is done
260         if shared.notify:
261             try:
262                 shared.pushme(sys.argv[0]+' completed.')
263             except:
264                 log("Job's done. Push failed.")
265
266     logfile.close()
```

A.2.2 Index Generator

```

1  # -*- coding: utf-8 -*-
2  '''This finishes preprocessing of the output from the XML parser.
3  This script reads in link data and removes from the content files those
4  concepts that have too few incoming links. Information on incoming links
5  is saved to each content file.
6  Finally, index maps for words and approved concepts are generated and saved.'''
7
8  from __future__ import division
9  import glob
10 import gc
11 import shared
12 import os
13 import sys
14
15 logfile = open(os.path.basename(__file__) + '.log', 'w')
16 log = shared.logmaker(logfile)
17
18 #Import shared parameters
19 from shared import extensions, temp_dir, min_links_in, matrix_dir
20
21 def listchopper(l):
22     '''Generator to chop lists into chunks of a predefined length'''
23     n = shared.link_chunk_size
24     ind = 0
25     while ind < len(l):
26         yield l[ind:ind+n]
27         ind += n
28
29 def main():
30     #Import shared parameters and verify output dir exists
31     if not os.path.exists(temp_dir):
32         raise IOError
33
34     #=====
35     #   Read in link data and update content files accordingly
36     #=====
37
38     #Get list of files containing link info and chop it up
39     linkfiles = glob.glob(temp_dir + '*' + extensions['links'])
40     linkchunks = listchopper(linkfiles)
41
42     linkfiles_read = 0
43     for linkchunk in linkchunks:
44         #Hash mapping each article to a set of articles linking to it
45         linkhash = {}
46
47         for filename in linkchunk:
48             with open(filename, 'r') as f:
49                 newstuff = shared.load(f)
50                 #Add link info to linkhash
51                 for target, sources in newstuff.iteritems():
52                     try:
53                         linkhash[target].update(set(sources))
54                     except KeyError:
55                         linkhash[target] = set(sources)
56
57     #Log status

```

```

58         linkfiles_read += 1
59         log("Read " + filename + " - " +
60             str(100*linkfiles_read/len(linkfiles))[4] + " % of link data.")
61
62     log("Chunk finished - updating content files")
63     #Update concept with newly read link data
64     contentfiles = glob.glob(temp_dir + '*' + extensions['content'])
65     contentfiles_read = 0
66     for filename in contentfiles:
67         #Read file. Content is like {'article title': {'text': blah}}
68         with open(filename, 'r') as f:
69             content = shared.load(f)
70
71         #Search linkhash for links going TO concept
72         for concept in content.keys():
73             try:
74                 sources = linkhash[concept]
75             except KeyError:
76                 sources = set([]) #Missing key => zero incoming links
77
78             #Update link info for concept
79             try:
80                 content[concept]['links_in'] = set(content[concept]['links_in'])
81                 content[concept]['links_in'].update(sources)
82             except KeyError:
83                 content[concept]['links_in'] = sources
84
85             #Save updated content
86             with open(filename, 'w') as f:
87                 shared.dump(content, f)
88
89         contentfiles_read += 1
90         if contentfiles_read % 100 == 0:
91             log("Fixed " + str(100*contentfiles_read/len(contentfiles))[4]
92                 + "% of content files")
93     pass #Proceed to next link chunk
94
95     #=====
96     # Finished link processing
97     # Remove unworthy concepts and combine concept/word lists.
98     #=====
99
100    #What, you think memory grows on trees?
101    del linkhash
102    gc.collect()
103
104    #Set of all approved concepts
105    concept_list = set([])
106
107    #Purge inferior concepts (with insufficient incoming links)
108    for filename in contentfiles:
109        #Read in content file
110        with open(filename, 'r') as f:
111            content = shared.load(f)
112
113        for concept in content.keys():
114            entry = content[concept]
115            if 'links_in' in entry and len(entry['links_in']) >= min_links_in:
116                concept_list.add(concept)
117            else:
118                del content[concept]
119

```



```
120         with open(filename, 'w') as f:
121             shared.dump(content, f)
122
123     log("Links done - saving index files")
124
125     #Make sure output dir exists
126     if not os.path.exists(matrix_dir):
127         os.makedirs(matrix_dir)
128
129     #Generate and save a concept index map. Structure: {concept : index}
130     concept_indices = {n: m for m, n in enumerate(concept_list)}
131     with open(matrix_dir+'concept2index.ind', 'w') as f:
132         shared.dump(concept_indices, f)
133
134     #Read in all wordlists and combine them.
135     words = set([])
136     for filename in glob.glob(temp_dir + '*' + extensions['words']):
137         with open(filename, 'r') as f:
138             words.update(shared.load(f))
139
140     #Generate and save a word index map. Structure: {word : index}
141     word_indices = {n: m for m, n in enumerate(words)}
142     with open(matrix_dir+'word2index.ind', 'w') as f:
143         shared.dump(word_indices, f)
144
145     log("Wrapping up.")
146     #Attempt to notify that job is done
147     if shared.notify:
148         try:
149             shared.pushme(sys.argv[0]+' completed.')
150         except:
151             log("Job's done. Push failed.")
152
153     logfile.close()
154
155 if __name__ == '__main__':
156     main()
```

A.2.3 Matrix Builder

```

1  # -*- coding: utf-8 -*-
2  '''Builds a huge sparse matrix of Term frequency/Inverse Document Frequency
3  (TFIDF) of the previously extracted words and concepts.
4  First a matrix containing simply the number of occurrences of word i in the
5  article corresponding to concept j is build (in DOK format as that is faster
6  for iterative construction), then the matrix is converted to sparse row format
7  (CSR), TFIDF values are computed, each row is normalized and finally pruned.'''
8
9  from __future__ import division
10 import scipy.sparse as sps
11 import numpy as np
12 from collections import Counter
13 import glob
14 import shared
15 import sys
16 import os
17
18 def percentof(small, large):
19     return str(100*small/large) + "%"
20
21 logfile = open(os.path.basename(__file__)+'.log', 'w')
22 log = shared.logmaker(logfile)
23
24 #import shared parameters
25 from shared import (extensions, matrix_dir, prune, temp_dir, column_chunk_size,
26                     row_chunk_size, datatype)
27
28 def main():
29     #Cleanup
30     for f in glob.glob(matrix_dir + '/*'+extensions['matrix']):
31         os.remove(f)
32
33     #Set pruning parameters
34     window_size = shared.window_size
35     cutoff = shared.cutoff
36
37     #Read in dicts mapping words and concepts to their respective indices
38     log("Reading in word/index data")
39     word2index = shared.load(open(matrix_dir+'word2index.ind', 'r'))
40     concept2index = shared.load(open(matrix_dir+'concept2index.ind', 'r'))
41     log("...Done!")
42
43     #=====
44     #     Construct count matrix in small chunks
45     #=====
46
47     #Count words and concepts
48     n_words = len(word2index)
49     n_concepts = len(concept2index)
50
51     #Determine matrix dimensions
52     matrix_shape = (n_words, n_concepts)
53
54     #Allocate sparse matrix. Dict-of-keys should be faster for iterative
55     #construction. Convert to csr for fast row operations later.
56     mtx = sps.dok_matrix(matrix_shape, dtype = datatype)
57

```

```

58 def matrix_chopper(matrix, dim):
59     '''Generator to split a huge matrix into small submatrices, which can
60     then be stored in individual files.
61     This is handy both when constructing the matrix (building the whole
62     matrix without saving to files in the process takes about 50 gigs RAM),
63     and when applying it, as this allows one to load only the submatrix
64     relevant to a given word.'''
65     ind = 0
66     counter = 0
67     rows = matrix.get_shape()[0]
68     while ind < rows:
69         end = min(ind+dim, rows)
70         #Return pair of submatrix number and the submatrix itself
71         yield counter, sps.vstack([matrix.getrow(i)\
72                                 for i in xrange(ind, end)], format = 'csr')
73         counter += 1
74         ind += dim
75
76 def writeout():
77     '''Saves the matrix as small submatrices in separate files.'''
78     for n, submatrix in matrix_chopper(mtx, row_chunk_size):
79         filename = matrix_dir+str(n)+extensions['matrix']
80         #Update submatrix if it's already partially calculated
81         log("Writing out chunk %s" % n)
82         try:
83             with open(filename, 'r') as f:
84                 submatrix = submatrix + shared.mload(f)
85             #
86         except IOError:
87             pass #File doesn't exist yet, so no need to change mtx
88
89         #Dump the submatrix to file
90         with open(filename, 'w') as f:
91             shared.mdump(submatrix, f)
92     return None
93
94 log("Constructing matrix.")
95 filelist = glob.glob(temp_dir + '*' + extensions['content'])
96 files_read = 0
97 for filename in filelist:
98     with open(filename, 'r') as f:
99         content = shared.load(f)
100
101     #Loop over concepts (columns) as so we don't waste time with rare words
102     for concept, entry, in content.iteritems():
103         #This is the column index (concept w. index j)
104         j = concept2index[concept]
105
106         #Convert concept 'countmap' like so: {word : n}
107         wordmap = Counter(entry['text'].split()).iteritems()
108
109         #Add them all to the matrix
110         for word, count in wordmap:
111             #Find row index of the current word
112             i = word2index[word]
113
114             #Add the number of times word i occurs in concept j to the matrix
115             mtx[i,j] = count
116         #
117     #Update file count
118     files_read += 1
119     log("Processed content file no. %s of %s - %s"

```

```

120         %(files_read, len(filelist)-1, percentof(files_read, len(filelist))))
121
122     if files_read % column_chunk_size == 0:
123         mtx = mtx.tocsr()
124         writeout()
125         mtx = sps.dok_matrix(matrix_shape)
126     #
127
128     #Convert matrix to CSR format and write to files.
129     mtx = mtx.tocsr()
130     writeout()
131
132     #=====
133     # Count matrix/matrices constructed – computing TF-IDF
134     #=====
135
136     log("Done - computing TF-IDF")
137
138     #Grap list of matrix files (containing the submatrices from before)
139     matrixfiles = glob.glob(matrix_dir + "*" + extensions['matrix'])
140     words_processed = 0 #for logging purposes
141
142     for filename in matrixfiles:
143         with open(filename, 'r') as f:
144             mtx = shared.mload(f)
145
146         #Number of words in a submatrix
147         n_rows = mtx.get_shape()[0]
148
149         for w in xrange(n_rows):
150             #Grap non-zero elements from the row corresponding to word w
151             row = mtx.data[mtx.indptr[w]:mtx.indptr[w+1]]
152             if len(row) == 0:
153                 continue
154
155             #Make a vectorized function to convert a full row to TF-IDF
156             f = np.vectorize(lambda m_ij: (1+np.log(m_ij))*
157                             np.log(n_concepts/len(row)))
158
159             #Map all elements to TF-IDF and update matrix
160             row = f(row)
161
162             #Normalize the row
163             assert row.dtype.kind == 'f' #Non floats round to zero w/o warning
164             normfact = 1.0/np.linalg.norm(row)
165             row *= normfact
166
167             #Start inverted index pruning
168             if prune:
169                 #Number of documents containing w
170                 n_docs = len(row)
171
172                 #Don't prune if the windows exceeds the array bounds (duh)
173                 if window_size < n_docs:
174
175                     #Obtain list of indices such that row[index] is sorted
176                     indices = np.argsort(row)[::-1]
177
178                     #Generate a sorted row
179                     sorted_row = [row[index] for index in indices]
180
181                     #Go through sorted row and truncate when pruning condition is met

```

```

182         for i in xrange(n_docs-window_size):
183             if sorted_row[i+window_size] >= cutoff*sorted_row[i]:
184                 #Truncate, i.e. set the remaining entries to zero
185                 sorted_row[i:] = [0]*(n_docs-i)
186                 break
187             else:
188                 pass
189
190         #Unsort to original positions
191         for i in xrange(n_docs):
192             row[indices[i]] = sorted_row[i]
193
194         #Update matrix
195         mtx.data[mtx.indptr[w] : mtx.indptr[w+1]] = row
196
197         #Log it
198         words_processed += 1
199         if words_processed % 10**3 == 0:
200             log("Processing word %s of %s - %s" %
201                 (words_processed, n_words,
202                  percentof(words_processed, n_words)))
203
204         #Keep it sparse – no need to store zeroes
205         mtx.eliminate_zeros()
206         with open(filename, 'w') as f:
207             shared.mdump(mtx, f)
208
209         log("Done!")
210
211         #Notify that the job is done
212         if shared.notify:
213             try:
214                 shared.pushme(sys.argv[0]+' completed.')
215             except:
216                 log("Job's done. Push failed.")
217
218         logfile.close()
219         return None
220
221 if __name__ == '__main__':
222     main()

```

A.2.4 Library for Computational Linguistics

```

1  # -*- coding: utf-8 -*-
2  '''Small module for computational linguistics applied to Twitter.
3  The main classes are a TweetHarvester, which gathers data from Twitters' API,
4  and a SemanticAnalyser, which relies on the previously constructed TFIDF
5  matrices.'''
6
7  from __future__ import division
8  from scipy import sparse as sps
9  from collections import Counter
10 from numpy.linalg import norm
11 import re
12 import shared
13 import tweepy
14 from datetime import date
15 import json
16 import time
17 import sys
18 import codecs
19 import os
20 from pprint import pprint
21 sys.stdout = codecs.getwriter('utf8')(sys.stdout)
22 sys.stderr = codecs.getwriter('utf8')(sys.stderr)
23
24 #=====
25 # This stuff defines a twitter 'harvester' for downloading Tweets
26 #=====
27
28 #Import credentials for accessing Twitter API
29 from supersecretstuff import consumer_key, consumer_secret, access_token, access_token_secret
30 auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
31 auth.set_access_token(access_token, access_token_secret)
32
33 class listener(tweepy.StreamListener):
34     '''Listener class to access Twitter stream.'''
35     #What to do with a tweet (override later)
36     def process(self, content):
37         print content
38         return None
39
40     def on_status(self, status):
41         self.process(status)
42         return True
43
44     def on_error(self, status):
45         print status
46
47 # Exception to be raised when the Twitter API messes up. Happens occasionally.
48 class IncompleteRead(Exception):
49     pass
50
51 class TweetHarvester(object):
52     '''Simple class to handle tweet harvest.
53     Harvest can be performed actively or passively, i.e. using the 'mine'
54     method to gather a fixed number of tweets or using the 'listen' method
55     to stream tweets matching a given search term.
56     Harvested tweets are sent to the process method which by default simply
57     stores them inside the object.'''

```

```

58
59 def __init__(self, max_tweets=-1, verbose = False, tweets_pr_file = 10*5):
60     #Set parameters
61     self.max_tweets = max_tweets #-1 for unlimited stream
62     self.verbose = verbose
63     self.tweets_pr_file = tweets_pr_file
64
65     #Internal parameters to keep track of harvest status
66     self.files_saved = 0
67     self.harvested_tweets = []
68     self.current_filenameumber = 0
69     self.current_date = date.today()
70
71 def filename_maker(self):
72     #Update counter and date if neccessary
73     if not self.current_date == date.today():
74         self.current_date = date.today()
75         self.current_filenameumber = 0
76     else:
77         pass #Date hasn't changed. Proceed.
78     filename = str(self.current_date) + "-data%s.json" % self.current_filenameumber
79     self.current_filenameumber += 1
80     return filename
81
82 #Simple logging function
83 def log(self, text):
84     string = text+ " at "+time.asctime()+"\n"
85     if self.verbose:
86         print string
87     with open('tweetlog.log', 'a') as logfile:
88         logfile.write(string)
89     #Must return true so I can log errors without breaking the stream.
90     return True
91
92 def listen(self, search_term):
93     #Make a listener
94     listener = tweepy.StreamListener()
95     #Override relevant methods.
96     listener.on_status = self.process
97     listener.on_error = lambda status_code: self.log("Error: "+status_code)
98     listener.on_timeout = lambda: self.log("Timeout.")
99
100     twitterStream = tweepy.Stream(auth, listener)
101     twitterStream.filter(track=search_term)
102
103 def mine(self, search_term, n = None):
104     '''Mine a predefined number of tweets using input search word'''
105     if n == None:
106         n = self.max_tweets
107
108     api = tweepy.API(auth)
109     tweets = tweepy.Cursor(api.search, q=search_term).items(n)
110     for tweet in tweets:
111         self.process(tweet)
112
113 def process(self, tweet):
114     self.harvested_tweets.append(tweet)
115     if self.verbose:
116         print "Holding %s tweets." % len(self.harvested_tweets)
117
118     #Write to file if buffer is full
119     if len(self.harvested_tweets) == self.tweets_pr_file:

```

```

120         self.writeout()
121
122         #Check if limit has been reached (returning false cuts off listener)
123         return not (len(self.harvested_tweets) == self.max_tweets)
124
125     def writeout(self):
126         filename = self.filename_maker()
127         with open(filename, 'w') as outfile:
128             outfile.writelines([json.dumps(t._json) + "\n"
129                                 for t in self.harvested_tweets])
130
131         self.harvested_tweets = []
132         self.files_saved += 1
133         #Log event
134         s = "Saved %s files" % self.files_saved
135         self.log(s)
136
137
138     #=====
139     # Defines stuff to analyse text using an already constructed interpretation
140     # matrix.
141     #=====
142
143     from shared import matrix_dir, row_chunk_size, extensions
144
145     class SemanticAnalyser(object):
146         '''Analyser class using Explicit Semantic Analysis (ESA) to process
147         text fragments. It can compute semantic (pseudo) distance and similarity,
148         as well'''
149         def __init__(self, matrix_filename = 'matrix.mtx'):
150             #Hashes for word and concept indices
151             with open(matrix_dir + 'word2index.ind', 'r') as f:
152                 self.word2index = shared.load(f)
153             with open(matrix_dir + 'concept2index.ind', 'r') as f:
154                 self.concept2index = shared.load(f)
155             self.index2concept = {i : c for c, i in self.concept2index.iteritems()}
156
157             #Count number of words and concepts
158             self.n_words = len(self.word2index)
159             self.n_concepts = len(self.concept2index)
160
161         def clean(self, text):
162             text = re.sub('[^\w\s\d\'\-\_]', '', text)
163             text = text.lower()
164
165             return text
166
167         def interpretation_vector(self, text):
168             '''Converts a text fragment string into a row vector where the i'th
169             entry corresponds to the total TF-IDF score of the text fragment
170             for concept i'''
171
172             #Remove mess (quotes, parentheses etc) from text
173             text = self.clean(text)
174
175             #Convert string to hash like {'word': no. of occurrences}
176             countmap = Counter(text.split()).iteritems()
177
178             #Interpretation vector to be returned
179             result = sps.csr_matrix((1, self.n_concepts), dtype = float)
180
181             #Add word count in the correct position of the vector

```



```

182     for word, count in countmap:
183         try:
184             ind = self.word2index[word]
185             #Which file to look in
186             file_number = int(ind/row_chunk_size)
187             filename = matrix_dir+str(file_number)+extensions['matrix']
188
189             #And which row to extract
190             row_number = ind % row_chunk_size
191
192             #Do it! Do it naw!
193             with open(filename, 'r') as f:
194                 temp = shared.mload(f)
195                 result = result + count*temp.getrow(row_number)
196         except KeyError:
197             pass #No data on this word -> discard
198
199     #Done. Return row vector as a 1x#concepts CSR matrix
200     return result
201
202 def interpret_text(self, text, display_concepts = 10):
203     '''Attempts to guess the core concepts of the given text fragment'''
204     #Compute the interpretation vector for the text fragment
205     vec = self.interpretation_vector(text)
206
207     #Magic, don't touch
208     top_n = vec.data.argsort()[::-1-display_concepts:-1]
209
210     #List top scoring concepts and their TD-IDF
211     concepts = [self.index2concept[vec.indices[i]] for i in top_n]
212     return concepts
213 #
214 #     scores = [vec.data[i] for i in top_n]
215 #     #Return as dict {concept : score}
216 #     return dict(zip(concepts, scores))
217
218 def interpret_file(self, filename):
219     with open(filename, 'r') as f:
220         data = self.clean(f.read())
221     return self.interpret_text(data)
222
223 def interpret_input(self):
224     text = raw_input("Enter text fragment: ")
225     topics = self.interpret_text(text)
226     print "Based on your input, the most probable topics of your text are:"
227     print topics[:self.display_concepts]
228
229 def scalar(self, v1, v2):
230     #Compute their inner product and make sure it's a scalar
231     dot = v1.dot(v2.transpose())
232     assert dot.shape == (1,1)
233
234     if dot.data:
235         scal = dot.data[0]
236     else:
237         scal = 0 #Empty sparse matrix means zero
238
239     #Normalize and return
240     sim = scal/(norm(v1.data)*norm(v2.data))
241     return sim
242
243 def cosine_similarity(self, text1, text2):
244     '''Determines cosine similarity between input texts.

```

```

244         Returns float in [0,1]'''
245
246     #Determine interpretation vectors
247     v1 = self.interpretation_vector(text1)
248     v2 = self.interpretation_vector(text2)
249
250     #Compute the normalized dot product and return
251     return self.scalar(v1, v2)
252
253
254     def cosine_distance(self, text1, text2):
255         return 1-self.cosine_similarity(text1, text2)
256
257 if __name__ == '__main__':
258     th = TweetHarvester(verbose=True, max_tweets=10)
259     th.mine('carlsberg', n=10)
260     temp = [t._json for t in th.harvested_tweets if t._json['lang'] == 'en']
261     js = temp[4]
262     with open('tweet_example.json', 'w') as f:
263         pprint(js, stream=f)
264
265 # if len(sys.argv) > 1:
266 #     fn = sys.argv[1]
267 # else:
268 #     fn = 'interpret_me.txt'
269 #     with open(fn, 'r') as f:
270 #         data = f.read()
271 #     #
272 #     data = sa.clean(data)
273 #     guesses = sa.interpret_text(data)
274 #
275 # if len(sys.argv) > 2:
276 #     output_filename = sys.argv[2]
277 # else:
278 #     output_filename = 'guesses.txt'
279 #     with open(output_filename, 'w') as f:
280 #         for line in guesses:
281 #             f.write(line.encode('utf8'))
282 #         f.write('\n')

```

A.2.5 Wikicleaner

```

1  # -*- coding: utf-8 -*-
2  import re
3  from htmlentitydefs import name2codepoint
4
5  namespaces = set(['help', 'file talk', 'module', 'topic', 'mediawiki',
6  'wikipedia talk', 'file', 'user talk', 'special', 'category talk', 'category',
7  'media', 'wikipedia', 'book', 'draft', 'book talk', 'template', 'help talk',
8  'timedtext', 'mediawiki talk', 'portal talk', 'portal', 'user', 'module talk',
9  'template talk', 'education program talk', 'education program',
10 'timedtext talk', 'draft talk', 'talk'])
11
12 def dropNested(text, openDelim, closeDelim):
13     '''Helper function to match nested expressions which may cause problems
14     example: {{something something {{something else}} and something third}}
15     cannot be easily matched with a regexp to remove all occurrences.

```

```

16 Copied from the WikiExtractor project.'''
17 openRE = re.compile(openDelim)
18 closeRE = re.compile(closeDelim)
19 # partition text in separate blocks { } { }
20 matches = [] # pairs (s, e) for each partition
21 nest = 0 # nesting level
22 start = openRE.search(text, 0)
23 if not start:
24     return text
25 end = closeRE.search(text, start.end())
26 next = start
27 while end:
28     next = openRE.search(text, next.end())
29     if not next: # termination
30         while nest: # close all pending
31             nest -= 1
32             end0 = closeRE.search(text, end.end())
33             if end0:
34                 end = end0
35             else:
36                 break
37         matches.append((start.start(), end.end()))
38         break
39     while end.end() < next.start():
40         # { } {
41         if nest:
42             nest -= 1
43             # try closing more
44             last = end.end()
45             end = closeRE.search(text, end.end())
46             if not end: # unbalanced
47                 if matches:
48                     span = (matches[0][0], last)
49                 else:
50                     span = (start.start(), last)
51                 matches = [span]
52                 break
53             else:
54                 matches.append((start.start(), end.end()))
55                 # advance start, find next close
56                 start = next
57                 end = closeRE.search(text, next.end())
58                 break # { }
59         if next != start:
60             # { { }
61             nest += 1
62 # collect text outside partitions
63 res = ''
64 start = 0
65 for s, e in matches:
66     res += text[start:s]
67     start = e
68 res += text[start:]
69 return res
70
71 def unescape(text):
72     '''Removes HTML or XML character references and entities
73     from a text string.
74     @return nice text'''
75     def fixup(m):
76         text = m.group(0)
77         code = m.group(1)

```

```

78     return text
79     try:
80         if text[1] == "#": # character reference
81             if text[2] == "x":
82                 return unichr(int(code[1:], 16))
83             else:
84                 return unichr(int(code))
85         else: # named entity
86             return unichr(name2codepoint[code])
87     except UnicodeDecodeError:
88         return text # leave as is
89
90     return re.sub("&#?(\w+);", fixup, text)
91
92 def drop_spans(matches, text):
93     """Drop from text the blocks identified in matches"""
94     matches.sort()
95     res = ''
96     start = 0
97     for s, e in matches:
98         res += text[start:s]
99         start = e
100     res += text[start:]
101     return res
102
103 ###Compile regexps for text cleanup:
104 #Construct patterns for elements to be discarded:
105 discard_elements = set([
106     'gallery', 'timeline', 'noinclude', 'pre',
107     'table', 'tr', 'td', 'th', 'caption',
108     'form', 'input', 'select', 'option', 'textarea',
109     'ul', 'li', 'ol', 'dl', 'dt', 'dd', 'menu', 'dir',
110     'ref', 'references', 'img', 'imagemap', 'source'
111 ])
112 discard_element_patterns = []
113 for tag in discard_elements:
114     pattern = re.compile(r'<s*s\b[^>]*>.*?<s*/\s*s*>' % (tag, tag), re.DOTALL | re.IGNORECASE)
115     discard_element_patterns.append(pattern)
116
117 #Construct patterns to recognize HTML tags
118 selfclosing_tags = set(['br', 'hr', 'nobr', 'ref', 'references' ])
119 selfclosing_tag_patterns = []
120 for tag in selfclosing_tags:
121     pattern = re.compile(r'<s*s\b[/]?*\s*>' % tag, re.DOTALL | re.IGNORECASE)
122     selfclosing_tag_patterns.append(pattern)
123
124 #Construct patterns for tags to be ignored
125 ignored_tags = set([
126     'a', 'b', 'big', 'blockquote', 'center', 'cite', 'div', 'em',
127     'font', 'h1', 'h2', 'h3', 'h4', 'hier', 'i', 'kbd', 'nowiki',
128     'p', 'plaintext', 's', 'small', 'span', 'strike', 'strong',
129     'sub', 'sup', 'tt', 'u', 'var',
130 ])
131 ignored_tag_patterns = []
132 for tag in ignored_tags:
133     left = re.compile(r'<s*s\b[^>]*>' % tag, re.IGNORECASE)
134     right = re.compile(r'<s*/\s*s*>' % tag, re.IGNORECASE)
135     ignored_tag_patterns.append((left, right))
136
137 #Construct patterns to recognize math and code
138 placeholder_tags = {'math': 'formula', 'code': 'codice'}
139 placeholder_tag_patterns = []

```

```

140 for tag, repl in placeholder_tags.items():
141     pattern = re.compile(r'<s*%s\s*| \^>+>.*?<s*/s*%s\s*>' % (tag, tag), re.DOTALL | re.IGNORECASE)
142     placeholder_tag_patterns.append((pattern, repl))
143
144 #HTML comments
145 comment = re.compile(r'<!--.*?-->', re.DOTALL)
146
147 #Wikilinks
148 wiki_link = re.compile(r'\[[^\[\]\?]*\](?:\|([^\[\]\?)*\])?(\w*)')
149 parametrized_link = re.compile(r'\[[^\[\]\?]*\]')
150
151 #External links
152 externalLink = re.compile(r'\[w+.*? (.*)\]')
153 externalLinkNoAnchor = re.compile(r'\[w+[\&]]*\]')
154
155 #Bold/italic text
156 bold_italic = re.compile(r''''''([^\']*?)''''')
157 bold = re.compile(r''''(.*)''''')
158 italic_quote = re.compile(r''''\(.*)\''''')
159 italic = re.compile(r''''([^\']*?)''''')
160 quote_quote = re.compile(r''''(.*)''''')
161
162 #Spaces
163 spaces = re.compile(r' {2,}')
164
165 #Dots
166 dots = re.compile(r'\.{4,}')
167
168 #Sections
169 section = re.compile(r'(=+)\s*(.*)\s*\1')
170
171 # Match preformatted lines
172 preformatted = re.compile(r'^.*?$', re.MULTILINE)
173
174 #Wikilinks
175 def make_anchor_tag(match):
176     '''Recognizes links and returns only their anchor. Example:
177     <a href="www.something.org">Link text</a> -> Link text'''
178     link = match.group(1)
179     colon = link.find(':')
180     if colon > 0 and link[colon] not in namespaces:
181         return ''
182     trail = match.group(3)
183     anchor = match.group(2)
184     if not anchor:
185         if link[colon] in namespaces:
186             return '' #Don't keep stuff like "category: shellfish "
187         anchor = link
188     anchor += trail
189     return anchor
190
191 def clean(text):
192     '''Outputs an article in plaintext from its format in the raw xml dump.'''
193     # Drop transclusions (template, parser functions)
194     # See: http://www.mediawiki.org/wiki/Help:Templates
195     text = dropNested(text, r'{{', r'}}')
196     # Drop tables
197     text = dropNested(text, r'\|', r'\|')
198
199     # Convert wikilinks links to plaintext
200     text = wiki_link.sub(make_anchor_tag, text)
201     # Drop remaining links

```

```

202 text = parametrized_link.sub('', text)
203
204 # Handle external links
205 text = externalLink.sub(r'\1', text)
206 text = externalLinkNoAnchor.sub('', text)
207
208 #Handle text formatting
209 text = bold_italic.sub(r'\1', text)
210 text = bold.sub(r'\1', text)
211 text = italic_quote.sub(r'&quot;\1&quot;', text)
212 text = italic.sub(r'&quot;\1&quot;', text)
213 text = quote_quote.sub(r'\1', text)
214 text = text.replace("'", '').replace('"', '&quot;')
215
216 ##### Process HTML #####
217
218 # turn into HTML
219 text = unescape(text)
220
221 # do it again (&nbsp;)
222 text = unescape(text)
223
224 # Collect spans
225
226 matches = []
227 # Drop HTML comments
228 for m in comment.finditer(text):
229     matches.append((m.start(), m.end()))
230
231 # Drop self-closing tags
232 for pattern in selfclosing_tag_patterns:
233     for m in pattern.finditer(text):
234         matches.append((m.start(), m.end()))
235
236 # Drop ignored tags
237 for left, right in ignored_tag_patterns:
238     for m in left.finditer(text):
239         matches.append((m.start(), m.end()))
240     for m in right.finditer(text):
241         matches.append((m.start(), m.end()))
242
243 # Bulk remove all spans
244 text = drop_spans(matches, text)
245
246 # Cannot use dropSpan on these since they may be nested
247 # Drop discarded elements
248 for pattern in discard_element_patterns:
249     text = pattern.sub('', text)
250
251 # Expand placeholders
252 for pattern, placeholder in placeholder_tag_patterns:
253     index = 1
254     for match in pattern.finditer(text):
255         text = text.replace(match.group(), '%s_%d' % (placeholder, index))
256         index += 1
257
258 #####
259
260 # Drop preformatted
261 # This can't be done before since it may remove tags
262 text = preformatted.sub('', text)
263

```

```

264 # Cleanup text
265 text = text.replace('\t', ' ')
266 text = spaces.sub(' ', text)
267 text = dots.sub('...', text)
268 text = re.sub(u' (,:\.\)\[\]\>)', r'\1', text)
269 text = re.sub(u' \[\(\«' , r'\1', text)
270 text = re.sub(r'\n\W+?\n', '\n', text) # lines with only punctuations
271 text = text.replace(',', '.').replace('.', '. ')
272
273 #Handle section headers, residua etc.
274 page = []
275 headers = {}
276 empty_section = False
277
278 for line in text.split('\n'):
279
280     if not line:
281         continue
282     # Handle section titles
283     m = section.match(line)
284     if m:
285         title = m.group(2)
286         lev = len(m.group(1))
287         if title and title[-1] not in '!?':
288             title += ' '
289         headers[lev] = title
290         # drop previous headers
291         for i in headers.keys():
292             if i > lev:
293                 del headers[i]
294         empty_section = True
295         continue
296     # Handle page title
297     if line.startswith('++'):
298         title = line[2:-2]
299         if title:
300             if title[-1] not in '!?':
301                 title += ' '
302             page.append(title)
303     # handle lists
304     elif line[0] in '*#,:;':
305         continue
306     # Drop residuals of lists
307     elif line[0] in '{| ' or line[-1] in '}':
308         continue
309     # Drop irrelevant lines
310     elif (line[0] == '(' and line[-1] == ')') or line.strip('.-') == '':
311         continue
312     elif len(headers):
313         items = headers.items()
314         items.sort()
315         for (i, v) in items:
316             page.append(v)
317         headers.clear()
318         page.append(line) # first line
319         empty_section = False
320     elif not empty_section:
321         page.append(line)
322
323 text = ''.join(page)
324
325 #Remove quote tags.

```

```

326 text = text.replace("&quot;", '')
327
328 #Get rid of parentheses, punctuation and the like
329 text = re.sub('[^\w\s\d\'-]', '', text)
330 return text

```

A.2.6 Application Examples

This section contains code used for the examples of applications of explicit semantic analysis described in section 3.3.

Trend Discovery and Monitoring

This is the script used to generate the results seen in section 3.3.1.

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import matplotlib as mpl
5  mpl.use('Agg')
6  import matplotlib.pyplot as plt
7  import sys
8  sys.path.insert(0, '../..')
9  import codecs
10 sys.stdout = codecs.getwriter('utf8')(sys.stdout)
11 sys.stderr = codecs.getwriter('utf8')(sys.stderr)
12 from collections import Counter
13 from cunning_linguistics import SemanticAnalyser as SA
14 import json
15 import time
16 import datetime
17 from shared import moerkeroed#, wine, oldhat, nude
18
19 INPUT_FILENAME = 'carlsberg_short.txt'
20 #INPUT_FILENAME = 'carlsberg_filtered_tweets.txt'
21 OUTPUT_DIR = 'animation/'
22 EXT = '.pdf'
23 SHOW = True
24 USE_ONE_IN = 1 # e.g. if 5 only every 5th tweet will be used. 1 for all
25 PROCESS = False #Whether to process data from scratch or read in old data
26
27
28 #INPUT_FILENAME = 'carlsberg_filtered_tweets.txt'
29 MAX_TWEETS = float('inf')
30
31 TOP_N = 10 #Number of best matched concepts to include in analysis
32 IGNORE = ['Carlsberg Group'] #Concepts to ignore
33 #If not empty, track only concepts in this list
34 TRACK = ['Carlsberg Foundation',
35          'Carlsberg Polska',
36          'Carlsberg Srbija',
37          'Old Lions',
38          'Copenhagen',
39          'Carlsberg Laboratory',
40 #          'Raidió Teilifís Éireann',

```



```

41         'Kim Little']
42
43 _display = TOP_N + len(IGNORE)
44
45 _now = time.time()
46
47 def tweet2epoch(t):
48     '''Extracts the time of input tweets creation and returns as datetime.'''
49     epoch = time.mktime(time.strptime(t['created_at'], '%a %b %d %H:%M:%S +0000 %Y'))
50     return epoch
51
52 def epoch2dtindex(epoch, span = 'day'):
53     dt = datetime.datetime.fromtimestamp(epoch)
54     year, week, weekday = datetime.date.isocalendar(dt)
55     day, month = dt.day, dt.month
56     if span == 'week':
57         return "%s-%02d" % (year, week)
58     elif span == 'day':
59         return "%s-%02d-%02d" % (year, month, day)
60
61 def pad_labels(X, limit = 15):
62     '''Limits label length to <limit> characters.'''
63     for i in xrange(len(X)):
64         if len(X[i]) <= limit:
65             continue
66         else:
67             X[i] = X[i][:limit]+"..."
68
69 def process():
70     tweets = []
71     n_read = 0
72     with open(INPUT_FILENAME, 'r') as f:
73         for line in f.readlines():
74             tweet = json.loads(line)
75             n_read += 1
76             if not n_read % USE_ONE_IN == 0:
77                 continue
78             if tweet['retweeted'] or not tweet['lang'] == 'en':
79                 continue
80             else:
81                 tweets.append({'created_at': tweet2epoch(tweet),
82                               'text': tweet['text']})
83             if len(tweets) >= MAX_TWEETS:
84                 break
85             #
86             #
87             #
88
89     # Create a semantic analyser
90     sa = SA()
91
92     data = {} #This will map index yyyy_ww to cobined list of top concepts
93     counter = 0
94     for tweet in tweets:
95         concepts = sa.interpret_text(tweet['text'], display_concepts = _display)
96         filtered = [c for c in concepts if not c in IGNORE][:TOP_N]
97         index = epoch2dtindex(tweet['created_at'])
98         # If tracking, include only tracked concepts
99         if TRACK:
100             filtered = [c for c in filtered if c in TRACK]
101         try:
102             data[index] += filtered

```

```

103     except KeyError:
104         data[index] = filtered
105         counter += 1
106         print "Processed %d of %d tweets..." % (counter, len(tweets))
107
108     #Change data into {index: (sorted topics, their counts)}
109     for index, _list in data.iteritems():
110         # Most often used concepts in entire period corresponding to index
111         d = dict(Counter(_list))
112         top_concepts = sorted(d, key = d.get)[:TOP_N]
113         #If not tracking specific concepts, just save top n concepts and counts
114         if not TRACK:
115             X = sorted(top_concepts)
116             Y = [d[x] for x in X]
117         #If tracking, use NaN for concepts that didn't occur in top n.
118         else:
119             X = TRACK
120             Y = [d[x] if x in top_concepts else float('nan') for x in TRACK]
121         data[index] = (X, Y)
122
123     return data
124
125 if __name__ == '__main__':
126     if PROCESS:
127         data = process()
128
129         with open('processed_data.json', 'w') as f:
130             json.dump(data, f, indent = 4)
131         #
132     else:
133         with open('processed_data.json', 'r') as f:
134             data = json.load(f)
135
136     #=====
137     # Plotting ...
138     #=====
139
140     #Set pylab params to make plots look similar
141     ymin = float('-inf')
142     ymax = float('-inf')
143     for _,Y in data.values():
144         if not Y:
145             continue
146         thismin = min(Y)
147         thismax = max(Y)
148         if thismin < ymin:
149             ymin = thismin
150         if thismax > ymax:
151             ymax = thismax
152
153     ymin = 0 # Don't truncate
154     yair = (ymax - ymin)*0.05
155
156     xair = 0.5
157     xmin = 0
158     xmax = len(TRACK) if TRACK else TOP_N
159
160     for index, datum in data.iteritems():
161         #Enforce axes 'n stuff
162         fig = plt.figure(figsize = (3.25,3))

```

```

165     ax = fig.add_subplot(111)
166     plt.axis((xmin - xair, xmax + xair, ymin - yair, ymax + yair))
167     ax.set_autoscale_on(False)
168
169     X, Y = datum
170     pad_labels(X, limit = 20)
171     plt.xticks([n+0.4 for n in xrange(len(X))], X, rotation = 50,
172               ha='right', fontsize = 8)
173     plt.bar(range(len(X)), Y, color = moerkeroed)
174     plt.gcf().subplots_adjust(bottom = 0.42, left = 0.3)
175     plt.title(index)
176
177     plt.savefig(OUTPUT_DIR+index+EXT, dpi = 600)
178
179     if SHOW:
180         plt.show()
181
182     plt.close()
183
184
185     print "Runtime (m): ", (time.time()-_now)/60

```

Social Media Impact

This is the code used to extract and analyse the data described in section 3.3.2.

```

1  # -*- coding: utf-8 -*-
2  '''This script computes the semantic vector corresponding to an input reference
3  text, loads in tweets contained in files saved in the specified period
4  (change filelist to whatever files contain your tweets), and then computes
5  the cosine similarity of each tweet with the reference text and saves the
6  result.'''
7  from __future__ import division
8
9  import sys
10 sys.path.insert(0, '../..')
11 from cunning_linguistics import SemanticAnalyser
12 import datetime
13 from dateutil.parser import parse
14 import json
15 from matplotlib import pylab as plt
16 import numpy as np
17 import pytz
18 import glob
19 import time
20 import multiprocessing
21
22
23 REFERENCE = 'reference_google.txt'
24 PUBLISHED = datetime.datetime(2015, 06, 17) #Date ref was published
25 EARLY = 16 #Number of days to include around pub date
26 LATE = 42
27 OUTPUT_FILENAME = 'deep_dreams.txt'
28 KEYWORD = 'google'
29 N_JOBS = 8
30

```

```

31 |
32 | #15-06-17
33 | with open(REFERENCE, 'r') as f:
34 |     reference_text = f.read()
35 |
36 | now = time.time()
37 |
38 | epoch = datetime.datetime(1970,1,1, tzinfo = pytz.utc)
39 | def dt2epoch(dt):
40 |     utc_date = dt.astimezone(pytz.utc)
41 |     delta = utc_date - epoch
42 |     return delta.total_seconds()
43 |
44 | def percentof(small, large):
45 |     return str(100*small/large) + "%"
46 |
47 | #Get canonical date string
48 | def timeparse(string):
49 |     dt = parse(string)
50 |     (y,m,d) = (str(dt.year), str(dt.month).zfill(2), str(dt.day).zfill(2))
51 |     return "%s-%s-%s" % (y,m,d)
52 |
53 | with open(REFERENCE, 'r') as f:
54 |     reference_text = f.read()
55 |
56 | #required entries in tweets
57 | ineedthese = ['lang', 'text', 'created_at']
58 |
59 | def worker(filename):
60 |     Y = []
61 |     X = []
62 |     #Make an analyser
63 |     sa = SemanticAnalyser()
64 |     reference_vector = sa.interpretation_vector(reference_text)
65 |     with open(filename, 'r') as f:
66 |         tweets = [json.loads(line) for line in f.readlines()]
67 |         for tweet in tweets:
68 |             if not all(stuff in tweet.keys() for stuff in ineedthese):
69 |                 continue
70 |             if not tweet['lang'] == 'en':
71 |                 continue
72 |             text = tweet['text'].lower()
73 |             if not KEYWORD in text:
74 |                 continue
75 |             t = dt2epoch(parse(tweet['created_at']))
76 |             this_vector = sa.interpretation_vector(text)
77 |             similarity = sa.scalar(this_vector, reference_vector)
78 |             if np.isnan(similarity):
79 |                 continue
80 |             X.append(t)
81 |             Y.append(similarity)
82 |             #
83 |             print "Processed file: ", filename
84 |             d = {'X': X, 'Y': Y}
85 |             return d
86 |
87 | if __name__ == '__main__':
88 |     filelist = set([])
89 |
90 |     for w in np.arange(-EARLY, LATE+1, 1):
91 |         delta = datetime.timedelta(days = w)
92 |         dt = PUBLISHED + delta

```

```
93     prefix = dt.strftime("%Y-%m-%d")
94     pattern = 'tweets/' + prefix + '-data*'
95     filelist.update((glob.glob(pattern)))
96
97     pool = multiprocessing.Pool(processes = N_JOBS)
98     jobs = [pool.apply_async(worker, kwds = {'filename' : fn})
99             for fn in filelist]
100     pool.close() #run
101     pool.join() #Wait for remaining jobs
102
103     #Make sure no children died too early
104     if not all(job.successful() for job in jobs):
105         raise RuntimeError('Some jobs failed.')
106
107     X = []
108     Y = []
109     for d in [j.get() for j in jobs]:
110         X += d['X']
111         Y += d['Y']
112     inds = np.argsort(X)
113     X = [X[i] for i in inds]
114     Y = [Y[i] for i in inds]
115
116     with open(OUTPUT_FILENAME, 'w') as f:
117         json.dump(X, f)
118         f.write('\n')
119         json.dump(Y, f)
120
121     # plt.plot(X, Y)
122     # plt.show()
```


BIBLIOGRAPHY

- [1] Yoshua Bengio and Yves Grandvalet. No Unbiased Estimator of the Variance of K-Fold Cross-Validation. 5:1089–1105, 2004.
- [2] Avrim Blum, Adam Kalai, and John Langford. Beating the Hold-out: Bounds for {K}-fold and Progressive Cross-Validation. *Proceedings of the 12th Annual Conference on Computational Learning Theory*, (c):203–208, 1999.
- [3] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [4] L Breiman. Random forests. *Machine learning*, pages 5–32, 2001.
- [5] Cjc Christopher J C Burges. A Tutorial on Support Vector Machines for Pattern Recognition. *Data Mining and Knowledge Discovery*, 2(2):121–167, 1998.
- [6] Yves Alexandre De Montjoye, Jordi Quoidbach, Florent Robic, and Alex Pentland. Predicting personality using novel mobile phone-based metrics. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, 7812 LNCS:48–55, 2013.
- [7] JM John M Digman. Personality structure: Emergence of the five-factor model. *Annual review of psychology*, 41:417–440, 1990.
- [8] RA Fisher. The use of multiple measurements in taxonomic problems. *Annals of eugenics*, 1936.
- [9] E Gabrilovich and S Markovitch. Computing Semantic Relatedness Using Wikipedia-based Explicit Semantic Analysis. *IJCAI*, 2007.
- [10] Isabelle Guyon. Gene Selection for Cancer Classification. pages 389–422, 2002.

- [11] Marti a. Hearst, Susan T. Dumais, Edgar Osuna, John Platt, and Bernhard Schölkopf. Support vector machines. *IEEE Intelligent Systems and their Applications*, 13(4):18–28, 1998.
- [12] Thomas K Landauer, Peter W. Foltz, and Darrell Laham. An introduction to latent semantic analysis. *Discourse Processes*, 25(2-3):259–284, 1998.
- [13] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. An Introduction to Information Retrieval. *Online*, (c):569, 2009.
- [14] Preslav Nakov. Latent semantic analysis of textual data. *Proceedings of the conference on Computer systems and technologies - CompSysTech '00*, pages 5031–5035, 2000.
- [15] AD Pietersma. Kernel Learning in Support Vector Machines using Dual-Objective Optimization. *Proceedings of the 23rd . . .*, 2011.
- [16] G. Salton, a. Wong, and C. S. Yang. A vector space model for automatic indexing. *Communications of the ACM*, 18(11):613–620, 1975.
- [17] Gerard Salton and Christopher Buckley. Term-weighting approaches in automatic text retrieval. *Information Processing & Management*, 24(5):513–523, 1988.
- [18] AJ Smola and B Schölkopf. *Learning with kernels*. 1998.
- [19] AJ Smola and B Schölkopf. A tutorial on support vector regression. *Statistics and computing*, 2004.
- [20] Philip Wolfe. A duality theorem for nonlinear programming. *Quarterly of applied mathematics*, 19(3):239–244, 1961.
- [21] PA Zandbergen and SJ Barbeau. Positional accuracy of assisted gps data from high-sensitivity gps-enabled mobile phones. *Journal of Navigation*, 2011.