

SPECIALE

**FRA HADES
OF
DOOM**



BJARKE MØNSTED

PRETENTIOUS
QUOTE.

- FAMOUS PERSON, BORN-DIED

FRA HADES OF DOOM

Author	My Name
Advisor	His Name
Co-Advisor	Her Name



Ting, som KU siger der skal stå her

CMOL

Center for Models of Life

Submitted to the University of Great Justice
June 14, 2014

ACKNOWLEDGEMENTS

Thank you! Thank you all!

CONTENTS

I	Wikipedia-based Explicit Semantic Analysis	1
1	ESA	3
1.1	Methods	3
1.1.1	Bag-of-Words	4
1.1.2	Semantic Analysis	4
1.2	Constructing a Semantic Analyser	6
1.2.1	XML Parsing	7
1.2.2	Index Generation	8
1.2.3	Matrix Construction	9
1.3	Applications & Results	10
II	Social Fabric Project	13
2	Data Extraction	15
2.1	Temporal Correlations in Activity	16
2.1.1	Influence of phone calls	17
2.1.2	Influence of GPS activity	20
2.1.3	Influence of Bluetooth signal	23
2.2	Extraction of Input Data	27
2.2.1	Simple Call/Text Data	27
2.2.2	Location Data	28
2.2.3	Time Series Analysis	33
2.2.4	Facebook Data	33
2.2.5	Bluetooth Data	33
2.3	Output Data	35
2.4	Linear Discriminant Analysis	35
2.4.1	A measure of separation for projected Gaussians	36
2.4.2	Optimizing separation	38

3 Psychological Profiling	43
3.1 SVM	43
3.1.1 Obtaining the Maximally Separating Hyperplane	45
3.1.2 Generalizing to the non-linear case	47
3.1.3 Statistical subtleties	53
3.2 Decision Trees & Random Forests	54
3.3 Nearest Neighbour-classifiers	57
3.4 Results	58
A Appendix	59
A.1 Source Code for Explicit Semantic Analysis	60
A.1.1 Parser	60
A.1.2 Index Generator	65
A.1.3 Matrix Builder	68
A.1.4 Library for Computational Linguistics	72
A.1.5 Wikicleaner	76
A.2 Social Fabric-related Code	83
A.2.1 Phonetools	83
A.2.2 Code Communication Dynamics	88
A.2.3 Preprocessing	94
A.2.4 Social Fabric Code	107
A.2.5 Lloyd's Algorithm	122
A.2.6 Smallest Enclosing Circle	125
Bibliography	129

ENGLISH ABSTRACT

WORDS! SOOOOO MANY WORDS!

DANSK SAMMENFATNING

ORD! SAAAAAAAAAAAA MANGE ORD

Part I

WIKIPEDIA-BASED EXPLICIT SEMANTIC ANALYSIS

WIKIPEDIA-BASED EXPLICIT SEMANTIC ANALYSIS



NATURAL language processing has long been both a subject of interest and a source of great challenges in the field of artificial intelligence. The difficulty varies greatly depending with the different language processing tasks; certain problems, such as text categorization, are relatively straightforward to convert to a purely mathematical problem, which in turn can be solved by a computer, whereas other problems, such as computing semantic relatedness, necessitates a deeper understanding of a given text, and thus poses a greater problem. This sections aims firstly to give a brief introduction to some of the most prominent techniques used in language processing in order to explain my chosen method of explicit semantic analysis (ESA), and secondly to explain in detail my practical implementation of an ESA-based text interpretation scheme.

ref

1.1 Methods

This section outlines a few methods used in natural language processing, going into some detail on ESA while touching briefly upon related techniques.

1.1.1 *Bag-of-Words*

ref

An example of a categorization problem is the ‘bag of words’ approach, which has seen use in spam filters. Here, text fragments are treated as unordered collections of words drawn from various bags, which in the case of spam filters would be undesired mails (spam) and desired mails (ham). By analysing large amounts of regular mail and spam, the probability of drawing each of the words constituting a given text from each bag can be computed, and the probability of the given text fragment representing draws from each bag can be computed using Bayesian statistics.

More formally, the text T is represented as a collection of words $T = \{w_1, w_2, \dots, w_n\}$, and the probability of T actually representing draws from bag j is hence

$$P(B_j|T) = \frac{P(T|B_j)P(B_j)}{P(T)}, \quad (1.1)$$

$$= \frac{\prod_i P(w_i|B_j)P(B_j)}{\sum_j \prod_i P(w_i|B_j)P(B_j)}, \quad (1.2)$$

ref

for an arbitrary number of bags labelled by j . This method is simple and powerful whenever a text is expected to fall in one of several discrete categories (such as spam filters or language detection). However, for more complex tasks it proves lucrative to attempt instead to assign some kind of meaning to text fragments rather than to consider them analogous to lottery numbers or marbles. This notion of meaning will be elaborated on shortly, as it varies depending on the method of choice, but the overall idea is to ascribe to words a meaning which depends not only on the word itself, but also on the connection between the word and existing repository of knowledge. The reader may think of this as mimicking the reading comprehension of humans. In itself, the word ‘dog’ for instance, contains a mere 24 bits of information if stored with a standard encoding, yet a human reader immediately associates a rich amount of existing knowledge to the word, such as dogs being mammals, related to wolves, being a common household pet, etc. The objective of both explicit and latent semantic analysis is to establish a high-dimensional ‘concept space’ in which words and text fragments are represented as vectors. The difference between explicit and latent semantic analysis is the method used to obtain said concepts, as explained in the following sections.

1.1.2 Semantic Analysis

Salton et al proposed in their 1975 paper *A Vector Space Model for Automatic Indexing*[15] an approach where words and text fragments are mapped

with a linear transformation to vectors in a high-dimensional concept space,

$$T \rightarrow |V\rangle = \sum_i v_i |i\rangle, v_i \in \mathbb{R}, \quad (1.3)$$

where a similarity measure of two texts can be defined as the inner product of two normalized such vectors,

$$S(V, W) = \langle \hat{V} | \hat{W} \rangle = \frac{\sum_i v_i w_i}{\left(\sum_i v_i^2\right) \left(\sum_i w_i^2\right)}, \quad (1.4)$$

and the cosine quasi-distance can be considered as a measure of semantic distance between texts:

$$D(V, W) = 1 - S(V, W). \quad (1.5)$$

This approach has later seen use in the methods of Latent Semantic Analysis (LSA) and Explicit Semantic Analysis (ESA). Both methods can be said to mimic human cognition in the sense that the transformation from (1.3) is viewed as a mapping of a text fragment to a predefined *concept space* and thus, processing of texts relies heavily on external repositories of knowledge.

The difference between LSA and ESA is how the concept space is established. Although I have used solely ESA for this project, I will give an extremely brief overview of LSA for completeness following (Landauer 1998 [11]). LSA constructs its concept space by first extracting every unique word encountered in a large collection of text corpora and essentially uses the leading eigenvectors (i.e. corresponding to the largest eigenvalues) of the word-word covariance matrix as the basis vectors of its conceptual space. This is the sense in which the concept are latent - rather than interpret text in terms of explicit concepts, such as 'healthcare', LSA would discover correlations between words such as 'doctor', 'surgery' etc. and consider that a latent concept. Owing to the tradeoff between performance and computational complexity, only about 400 such vectors are kept[13]. In psychology, LSA has been proposed as a possible model of fundamental human language acquisition as it provides computers a way of estimating e.g. word-word relatedness (a task which LSA does decently) using nothing but patterns discovered in the language it encounters[11].

In contrast, the concepts in ESA correspond directly to certain parts of the external text corpora one has employed to construct a semantic analyser. Concretely, the matrix playing the role of the reduced covariance

matrix in LSA has columns corresponding to each text corpus used and rows corresponding to individual terms or words, with the value of each matrix element denoting some measure of relatedness between the designated word and concept. I have used the English Wikipedia, so naturally each concept consists of an article, although the process could easily be tuned to be more or less fine-grained and associate instead each concept with e.g. a subsection or a category, respectively. Of course, a wholly different collection of texts could also be used - for instance a version of ESA more suited to compare the style or period of literary works could be constructed using a large collection of literature such as the Gutenberg Project. However, with no prior knowledge of the subject matter of the text to be analysed, Wikipedia seems like a good all-round solution considering its versatility and the massive numbers of volunteers constantly keeping it up to date.

I wish to point out two advantages of ESA over LSA. First, it is more successful, at least using the currently available text corpora and implementation techniques. A standard way of evaluating the performance of a natural language processing program is to measure the correlation between relatedness scores assigned to pairs of words or text fragments by the computer and by human judges. In both disciplines, ESA has outperformed LSA since it was first implemented ([8], p 457).

Second, the concepts employed in ESA are directly understandable by a human reader, whereas the concepts in LSA correspond to the leading moments of the covariance matrix. For example, to test whether the first semantic analyser built by my program behaved reasonably, I fed it a snippet of a news article on CNN with the headline *"In Jerusalem, the 'auto intifada' is far from an uprising"*. This returned an ordered list of top scoring concepts as follows: "Hamas, Second Intifada, Palestinian National Authority, Shuafat, Gaza War (2008-09), Jerusalem, Gaza Strip, Arab Peace Initiative, Yasser Arafat, Israel, West Bank, Temple Mount, Western Wall, Mahmoud Abbas", which seems a very reasonable output.

1.2 Constructing a Semantic Analyser

The process of applying ESA to a certain problem may be considered as the two separate subtasks of first a very computationally intensive construction of the machinery required to perform ESA, followed by the application of said machinery to some collection of texts. For clarity, I'll limit the present section to the details of the former subtask while description its application and results in section 1.3.

The construction itself is divided into three steps which are run in succession to create the desired machinery. The following is a very brief overview of these steps, each of which is elaborated upon in the following subsections.

1. First, a full Wikipedia XML-dump¹ is parsed to a collection of files each of which contains the relevant information on a number of articles. This includes the article contents in plaintext along with some metadata such as designated categories, inter-article link data etc.
2. Then, the information on each concept (article) is evaluated according to some predefined criteria, and concepts thus deemed inferior are purged from the files. Furthermore, two lists are generated and saved, which map unique concepts and words, respectively, to an integer, the combination of which is to designate the relevant row and column in the final matrix. For example, the concept 'Horse' corresponded to column 699221 in my matrix, while the word 'horse' corresponded to row 11533476.
3. Finally, a large sparse matrix containing relevance scores for each word-concept pair is built and, optionally, pruned (a process used to remove 'background noise' from common words as explained in section 1.2.3)

These steps are elaborated upon in the following.

1.2.1 XML Parsing

The Wikipedia dump comes in a rather large (~50GB unpacked for the version I used) XML file which must be parsed to extract each article's contents and relevant information. This file is essentially a very long list of nested fields where the data type in each field is denoted by an XML tag, such as `<text> blabla </text>`. A very simplified example of a field for one Wikipedia article is shown in 1.1. The content of each field has already been sanitised by Wikipedia so that if for instance the symbol '<' is entered into an article, it is instead represented as '<' in the XML file. To this end, I wrote a SAX parser, which processes the dump sequentially to accommodate its large size. When running, the parser walks through the file and sends the various elements it encounters to a suitable processing function depending on the currently open XML tag. For example, when

¹These are periodically released at <http://dumps.wikimedia.org/enwiki/>

```

<page>
  <title>Horse</title>
  <ns>0</ns>
  <id>12</id>
  <revision>
    <id>619093743</id>
    <parentid>618899706</parentid>
    <timestamp>2014-07-30T07:26:05Z</timestamp>
    <contributor>
      <username>Eduen</username>
      <id>7527773</id>
    </contributor>
    <text xml:space="preserve">
      ===Section title===
      [[Image:name_of_image_file|Image caption]]
      Lots of educational text containing, among other things links to [[other article|text to display]].
    </text>
    <sha1>n57mnhtthxpxq1nanak3zhmmmc622</sha1>
    <model>wikitext</model>
    <format>text/x-wiki</format>
  </revision>
</page>

```

Snippet 1.1: A simplified snippet, of a Wikipedia XML dump.

a 'title' tag is encountered, a callback method is triggered which assigns a column number to the article title and adds it to the list of processed articles. For the callback method for processing the 'text' fields, I used a bit of code from the pre-existing Wikiextractor project to remove Wiki markup language (such as links to other articles being displayed with square brackets and the like) from the content. This code is included in section A.1.5. The remainder of the parser is my work, and is included in section A.1.1. Throughout this process, the parser keeps a list of unique words encountered as well as outgoing link information for each article. These lists, along with the article contents, are saved to files each time a set number of articles have been processed. The link information is also kept in a hashmap with target articles as keys and a set articles linking to the target as values. The point of this is to reduce the computational complexity of the link processing as detailed in the following section.

1.2.2 Index Generation

The next step reads in the link information previously saved and adds it as ingoing link information in the respective article content files. The point of this approach is that link information is initially saved as a hashmap so the link going to a given article can be found quickly, rather than having

to search for outgoing links in every other article to determine the ingoing links to each article, which would be of $O(n^2)$ complexity.

Following that, articles with sufficiently few words and/or ingoing/-outgoing links are discarded and index lists for the remaining articles and words are generated to associate a unique row/column number with each word/concept pair. The code performing the step described here is included in section A.1.2.

1.2.3 Matrix Construction

This final step converts the information compiled in the previous steps into a very large sparse matrix. The program allows for this to be done in 'chunks' in order to avoid insane RAM usage. Similarly, the matrix is stored in segments with each file containing a set number of rows in order to avoid loading the entire matrix to interpret a short text.

The full matrix is initially constructed using a DOK (dictionary of keys) sparse format in which the i, j th element simply counts the number of occurrences of word i in the article corresponding to concept j . This is denoted $\text{count}(w_i, c_j)$. The DOK format as a hashmap using tuples (i, j) as keys and the corresponding matrix elements as values and is the fastest format available for element-wise construction. The matrix is subsequently converted to CSR (compressed sparse row) format, which allows faster operations on rows which performs much quicker when computing TF-IDF (term frequency - inverse document frequency) scores and extracting concept vectors from words, i.e. when accessing separate rows corresponding to certain words.

Each non-zero entry is then converted to a TF-IDF score according to

$$T_{ij} = \left(1 + \ln\left(\text{count}(w_i, c_j)\right)\right) \ln\left(\frac{n_c}{df_i}\right), \quad (1.6)$$

where n_c is the total number of concepts and

$$df_i = |\{c_k, w_i \in c_k\}| \quad (1.7)$$

is the number of concepts whose corresponding article contains the i th word. Thus, the first part of (1.6), $1 + \ln(\text{count}(w_i, c_j))$ is the *text frequency* term, as it increases with the frequency of word i in document j . Similarly, $\ln(\frac{n_c}{df_i})$ in (1.6) is the *inverse document frequency* term as it decreases with the frequency of documents containing word i . Thus, the TF-IDF score as somewhat complement to entropy in that it goes to zero as the fraction

Giver det
mening?

of documents containing word i goes to 1, and takes its highest values if word i occurs with high frequency in only a few documents[12]. While (1.6) is not the only expression to have those properties, empirically it tends to achieve superior results in information retrieval[16].

Each row is then L^2 normalized (divided by their Euclidean norm):

$$T_{ij} \rightarrow \frac{T_{ij}}{\sqrt{\sum_i T_{ij}^2}}. \quad (1.8)$$

Finally, each row is pruned to reduce spurious associations between concepts and articles with a somewhat uniform occurrence rate. This was done in practice by following the pragmatic approach of Gabrilovich [8] of sorting the entries of each row, move a sliding window across the entries, truncating when the falloff drops below a set threshold and finally reversing the sorting. The result of this step is the matrix which computes the interpretation vectors as described in 1.1.2. The code is included in section A.1.3.

1.3 Applications & Results

Having constructed a necessary machinery, I wrote a small Python module to provide an easy-to-use interface with the output from the computations described earlier. The code for this is included in section A.1.4. The module consists mainly of a `SemanticAnalyser` class, which loads in the previously mentioned index lists and provides methods for various computations such as estimating the most relevant concepts for a text, determining semantic distance etc. For example, the following code will create a semantic analyser instance and use it to guess the topic of the input string:

```
sa = SemanticAnalyser()
sa.interpret_text("Physicist from Austria known for the theory of relativity")
```

This returns a sorted list of the basis concepts best matching the input string, where the first element is of course 'Albert Einstein'. The `SemanticAnalyser` class contains equally simple methods to interpret a target text file or keyboard input, to calculate the semantic similarity or cosine distance between texts, and to compute interpretations vectors from a text.

The same module contains a `TweetHarvester` class which I wrote in order to obtain a large number of tweets to test the semantic analyser on, as tweets are both numerous and timestamped, which allows investigations

Hearst nævner noget offentligt tilgængeligt Reuters-data som folk øver tekstklassifikation på. Det kunne være ret sjovt. Det kunne være sjovt at lave 'semantisk nearest neighbor'

of the temporal evolution of tweets matching a given search term. The TweetHarvester class provides an equally simply interface - for instances, the 100 most recent tweets regarding a list of companies can be mined and printed by typing

```
terms = ['google', 'carlsberg', 'starbucks']
th = TweetHarvester()
th.mine(terms, 100)
print th.harvested_tweets
```

in addition to actively 'mining' for tweets matching a given query, the class can also passively 'listen' for tweets while automatically saving its held tweets to a generated date-specific filename once a set limit of held tweets is exceeded:

```
th = TweetHarvester(tweets_pr_file = 100)
th.listen(terms)
```

The downloaded tweets are stored as tweet objects which contain a built in method to convert to a JSON-serializable hashmap, an example of which is provided in 1.2. As can be seen in the example, the tweet object contains not only the tweets textual content but also a wide range of metadata such as the hashtags contained in the tweet, users mentioned, time of creation, language etc. Indsæt fine grafer OSV!!!

DO IT NAW!

```

{u'contributors': None,
 u'coordinates': None,
 u'created_at': u'Wed Jun 03 12:16:23 +0000 2015',
 u'entities': {u'hashtags': [],
               u'symbols': [],
               u'urls': [],
               u'user_mentions': [{u'id': 630908729,
                                   u'id_str': u'630908729',
                                   u'indices': [0, 12],
                                   u'name': u'Alexandra White ',
                                   u'screen_name': u'lexywhite86'}]}},
 u'favorite_count': 0,
 u'favorited': False,
 u'geo': None,
 u'id': 606071930282745857L,
 u'id_str': u'606071930282745857',
 u'in_reply_to_screen_name': u'lexywhite86',
 u'in_reply_to_status_id': 605991263129714688L,
 u'in_reply_to_status_id_str': u'605991263129714688',
 u'in_reply_to_user_id': 630908729,
 u'in_reply_to_user_id_str': u'630908729',
 u'is_quote_status': False,
 u'lang': u'en',
 u'metadata': {u'iso_language_code': u'en', u'result_type': u'recent'},
 u'place': None,
 u'retweet_count': 0,
 u'retweeted': False,
 u'source': u'<a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone
 </a>',
 u'text': u'@lexywhite86 if carlsberg did mornings \U0001f602',
 u'truncated': False,
 u'user': {u'contributors_enabled': False,
           u'created_at': u'Thu Apr 05 13:48:00 +0000 2012',
           u'description': u'',
           u'favourites_count': 169,
           u'follow_request_sent': False,
           u'followers_count': 110,
           u'friends_count': 268,
           u'geo_enabled': False,
           u'id': 545986280,
           u'id_str': u'545986280',
           u'lang': u'en',
           u'listed_count': 0,
           u'location': u'',
           u'name': u'Robert Murphy',
           u'notifications': False,
           u'protected': False,
           u'screen_name': u'7Robmurphy',
           u'statuses_count': 1650,
           u'time_zone': None,
           u'url': None,
           u'utc_offset': None,
           u'verified': False}}

```

Snippet 1.2: Example of a downloaded tweet.

Part II

SOCIAL FABRIC PROJECT

PHONE ACTIVITY AND QUANTITATIVE DATA EXTRACTION

THE main objective of part II of this thesis is to investigate behavioural patterns in the phone activities of the participants in the Social Fabric Project, and to predict various traits of the users based only on their phone logs. Throughout the part, I'll provide brief examples of usage for the software I've written to process the large amount of data available and to apply various prediction schemes to it, while the source code itself is included in the appendix.

My first objective was to investigate how various phone activities correlate with each other temporally, i.e. how a given user's probability for e.g. receiving a call increases or decreases around other activities such as moving around physically. This is the topic of section 2.1.

Next, I set out to replicate some recent research results claiming that people's phone activities predict certain psychological traits. In the most general terms, then, the task consists of predicting a collection of numbers or labels denoted Y based on a set of corresponding data points X . The topic of section 2.2 is the extraction of the many-dimensional data points or *feature vectors* X from the phone logs of the participants, while section 2.3 gives a brief description of the psychological traits Y . Finally, section 2.4 derives an often used linear classification method known as Linear Discriminant Analysis or Fisher's Discriminant and provides a discussion of why it fails for the present dataset, which serves to motivate the more sophisticated prediction schemes introduced in chapter 3.

kilder!!!

2.1 Temporal Correlations in Activity

One category of interesting quantities is the predictability of mobile phone behaviour from recorded behaviour at different times, i.e. the influence of certain events deduced from a user's Bluetooth, GPS or call log data on the tendency of some event to happen in the near past or future. A simple example would be to determine how much placing or receiving a call increases or decreases the probability of a user placing or receiving another call in the period following the first call.

This analysis was performed by comparing an 'activity' signal with a 'background' signal in the following fashion: For each user, the time period around each call is sliced into bins and the each of the remaining calls placed in the bin corresponding to the time between the two calls. Once divided by the total number of calls for the user, this is the activity signal. The background is obtained in a similar fashion but comparing each call in a given user's file with calls in the remaining users' files.

This involves repeatedly binning the time around certain events and then determining in which bin to place other events; a situation in which confusion may arise easily and errors may be hard to identify. To accommodate this, I started out by writing a custom array class designed to greatly simplify the binning procedure. This class called features the following:

- Methods to bin the time around a given event and determine determine which bin a given event falls into. This is useful to implement in the class itself as one then avoids having to continually worry about which bin an event fits into, and as it ensures that bin placement errors can only arise in one small piece of code which can then be tested rigorously.
- Attributes that keep track of the number of events that didn't fit into any bin, and of the current centre of the array, which can then be manipulated to move the array and a method to use this to return a normalized list of the bins.

In short, the binarray can be visualized as a collection of enumerated buckets that can be moved so as to center it on some event and then let other events 'drip' into the buckets. The code for this class is included in [A.2.1](#). In general, objects can be converted to byte streams and stored using Python's pickle module, but as that tends to be both slow and insecure, I generally used json to save my objects. This poses a slight problem as some data types, such as tuples, and custom classes in general are not json

serializable. I got around this by writing some recursive helper methods to help store the relevant information about arbitrary nested combinations of some such objects and to help reconstruct said objects again. These are also included in section A.2.1. As an example of usage, the following code constructs a Binarrray, centers it around the present time, and generates a number of timestamps which are then placed in the event. It is then saved to a file using the helper method previously described.

```
from time import time
from random import randint
#Create Binarrray with interval +/- one hour and bin size ten minutes.
ba = Binarrray(interval = 60*60, bin_size = 10*60)
#Center it on the present
now = int(time())
ba.center = now
#Generate some timestamps around the present
new_times = [now + randint(-60*60, 60*60) for _ in xrange(100)]
for tt in new_times:
    ba.place_event(tt)

#Save it
with open('filename.sig', 'w') as f:
    dump(ba, f)
```


This data will be visualized by plotting the relative signal from the activity of some event, such as in- or outgoing calls or texts, over the background (simply A/B) around another type of event hypothesized to trigger the activity.

2.1.1 Influence of phone calls


I first investigated the effects of incoming and outgoing calls as triggers for other phone activities. The call logs were stored in a format where each line represents a hashmap with quantities such as call time or call duration mapping to their corresponding value. Below is an example of one such line, where any personal data have been replaced by a random number or hexadecimal string of equal length.

```
{
  "timestamp": 6335212287,
  "number": "c4bdd708b1d7b82e349780ee1e7875caa600c579",
  "user": "ea42a1dbe422f83b0178d158f154f4",
  "duration": 483,
  "type": 2,
  "id": 45687
}
```

the text logs are similar except for the missing duration entry. Computing the relative signal in Binarrrays centered on each incoming and outgoing




(a) Relative activity of events triggered by incoming calls.




(b) Relative activity of events triggered by outgoing calls.

Figure 2.1: Comparison of the increased activity caused by incoming and outgoing calls over an interval of ± 3 hours around an event with bins of three minutes.

call using bin sizes of three and thirty minutes resulted in the plots shown in figures 2.1 and 2.2, respectively. As the figures clearly show, the all four activities increase significantly for the average user around incoming and outgoing calls.



(a) Relative activity of events triggered by incoming calls.



(b) Relative activity of events triggered by outgoing calls.

Figure 2.2: Comparison of the increased activity caused by incoming and outgoing calls over an interval of ± 12 hours around an event with bins of thirty minutes.

2.1.2 Influence of GPS activity

The raw format of the users' GPS logs looks similar to those of the call and text logs:

```
{
  "timestamp": 8058876274,
  "lon": 6.45051654,
  "user": "0c28e8f4ad9619bca1e5ea4167e10a",
  "provider": "gps",
  "lat": 28.20527041,
  "id": 6429902,
  "accuracy": 39.4
}
```

An analysis similar to that of described in section 2.1.1 was carried out using GPS phone data as triggers. I chose to define a user as being 'active' if they travelled at an average speed of 0.5 m/s between two consecutive GPS log entries, while discarding measurements closely following each other. The reason for this is that the uncertainty on the location measurements could yield false measurements of high average speeds when the measurements are not temporally separated. A lot of the measurements turned out to be grouped somewhat tightly - for instance, approximately 80% of the time intervals were below 100 s. This occurs because the Social Fabric data not only actively records its users' locations with some set interval, but also passively records the location when another app requests it, so when users spend time on apps that need to continually update their position such as Google Maps, a location log entry is written every second. The distribution of intervals between consecutive GPS measurements is shown in figure 2.3. A typical uncertainty on civilian GPS devices is at most 100 m[18], so because I choose to consider a user active if they travel at a mean speed of 0.5 m/s , and based on the time spacings shown in figure 2.3, I chose to discard measurements separated by less than 500 s.

An analysis like that of section 2.1.1 reveals that a user's phone activity is significantly increased around times when they are on the move, as shown in figure 2.4. Note the asymmetry of the signal, especially visible in figure 2.4(a). After a measurement of a user being active, the signal dies off about two and a half hours into the future, whereas it persists much longer into the past. Concretely, this means that people's phone activity (tendency to call or text) becomes uncorrelated with their physical activity after roughly two and a half hours, whereas their tendency to move around is increased for much longer time after calling or texting.

The relative signal in figure 2.4(b) appears to be increasing at around $\pm 24 \text{ h}$, which would seem reasonable assuming people have slightly


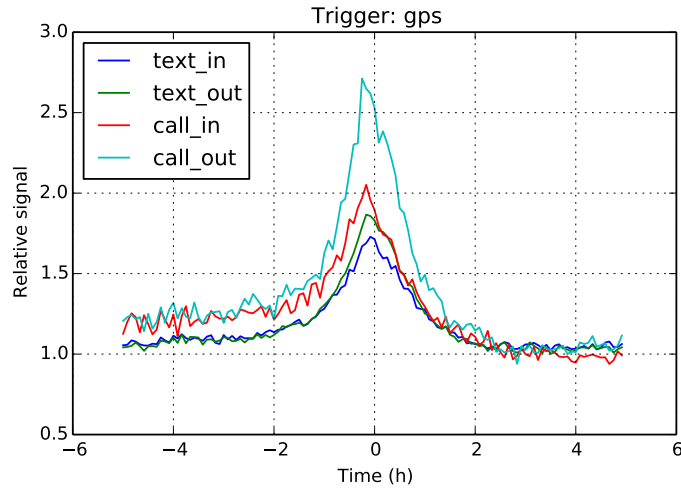
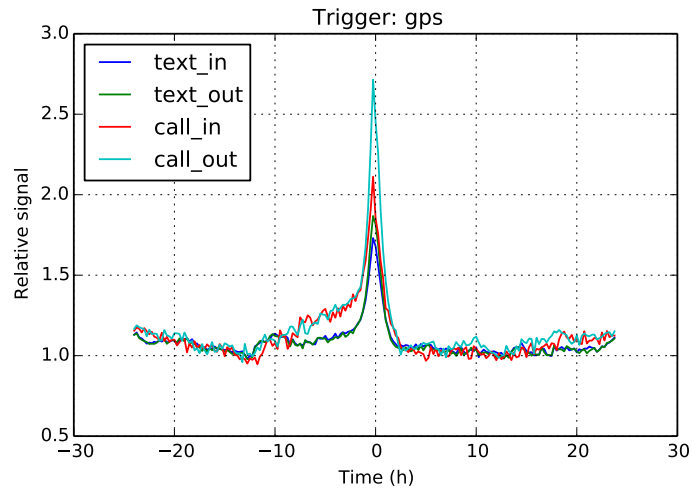


Figure 2.3: Plot of typical temporal spacings between consecutive GPS measurements.

different sleep schedules - if a person is on the move and hence more likely to place a call at time $t = 0$, they're slightly more likely than the general user to be on the move around $t = \pm 24$ h. Figure 2.5 shows the same signal extended to ± 36 h where slight bumps are visible 24 hours before and after activity.



(a) Interval: 5 hours. Bin size: 5 minutes.



(b) Interval: 24 hours. Bin size: 15 minutes.

Figure 2.4: Relative increase of activities triggered by GPS activity.

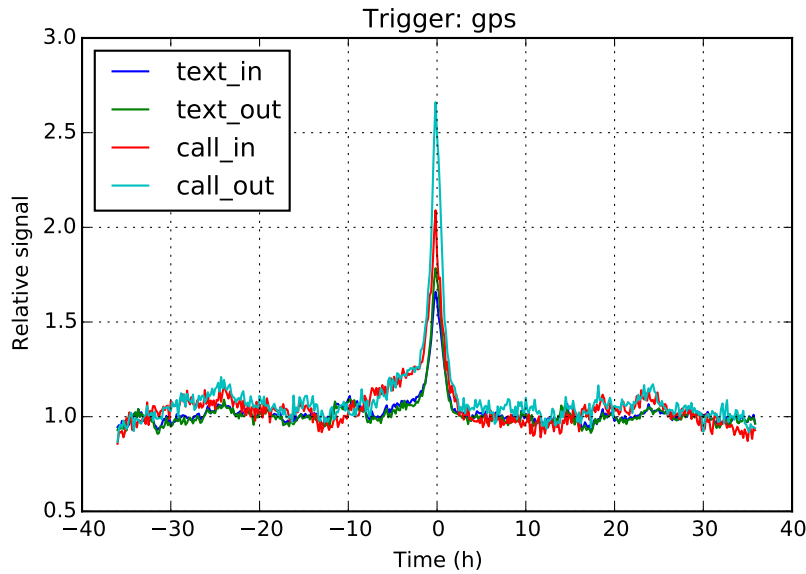


Figure 2.5: GPS-triggered activity increase over an interval of 36 hours using a bin size of 10 minutes.


2.1.3 Influence of Bluetooth signal

The following is a randomized entry in a user's bluetooth log.


```
{
  "name": "d5306a3672b7a0b8f9696d294ec4b731",
  "timestamp": 6870156680,
  "bt_mac": "1f158ae269d69efa5bb4794ee2a0b2dd68bd3a9badfeaf70f258ad3c74b0c09b",
  "class": 1317046,
  "user": "41cdb7ecaaec3d33391ed063e7fa2",
  "rssi": -76,
  "id": 4139043
}
```

The 'bt_mac' entry is the MAC-address of the device which the Bluetooth receiver in the user's phone has registered, so it is reasonable to assume several different MAC addresses occur at several consecutive timestamps. I call the number of repeated MAC addresses needed for a user to be considered social the 'social threshold'. Figures 2.6, 2.7 and 2.8 show the increased activity around times when users were considered social with a threshold of 1, 2 and 4 repeated pings.

Contrary to the previous analyses, phone activities decreased somewhat when users were social. As stated, each of these analyses were fairly similar, I've only explicitly included the code used to extract and save




(a) Interval: 12 hours, Bin size: 10 minutes.




(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 2.6: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as one repeated signal.




(a) Interval: 12 hours, Bin size: 10 minutes.




(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 2.7: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as two repeated signals.



(a) Interval: 12 hours, Bin size: 10 minutes.



(b) Interval: 36 hours, Bin size: 15 minutes.

Figure 2.8: The effect on phone activity of sociality as measured by the user's Bluetooth signal. The threshold used for being considered social as four repeated signals.

Bluetooth data, as well as the code used to load the data and generate figures 2.6 through 2.8. This code is included in section A.2.2.

2.2 Extraction of Input Data

The predictive powers of mobile phone behaviour on the user's psychological profile is currently an area of active research. As part of my thesis work, I have tried to predict the psychological profiles of the SFP participants using various machine learning methods on the available phone logs.

1000 kilder!!!

The software I've written first preprocesses the phone logs to extract various relevant parameters, then collects the parameters and psychological profile scores for each user to serve as input and output, respectively, for the various learning methods. Many of the parameters are chosen following a recent article by de Montjoye et al[5]. The following contains an outline and brief explanation of the extracted parameters.

This section contains a list of the extracted parameters used for psychological profiling along with a brief description of the extraction process when necessary. The preprocessing code is included in section A.2.3.

2.2.1 Simple Call/Text Data

The most straightforward data type is the timestamps from a given user's call/text logs. Six of the parameters used were simply the standard deviation and median of the times between events in the logs for each user's call log, text log, and the combination thereof, excluding time gaps of more than three says on the assumption that it would indicate a user being on vacation or otherwise having a period of telephone inactivity. The entropy S_u of each of the three was also included simply by computing the sum

$$S_u = - \sum_c \frac{n_c}{n_t} \ln_2 \frac{n_c}{n_t}, \quad (2.1)$$

where c denotes a given contact and n_t the total number of interactions, and n_c the number of interactions with the given contact. The number of contacts, i.e. the number of unique phone numbers a given user had contacted by means of calls, texts, and the combination thereof, was also extracted along with the total number of the various kinds of interactions and the contact to interaction ratios. The response rates, defined as the rate of missed calls and incoming texts, respectively, that a given user replied to within an hour, were also determined along with the text

latency defined as the median test response time. Finally the percentage calls and texts that were outgoing was determined as well as the fraction of call interactions that took places during the night, defined as between 22-08.

2.2.2 Location Data

A number of parameters based on the spacial dynamics of the user were also extracted. Among these is the radius of gyration, meaning simply the radius of the smallest enclosing circle enclosing all the registered locations of the user on the given day, and the distance travelled per day. I chose to extract the median and standard deviation of each, filtering out the radii that exceeded 500km so as to keep information about long distance travels in the distance parameter and information about travel within a given region in the radius of gyration parameter.

Cluster Analysis

One parameter which has strong links[5] to psychological traits is the number of locations in which the user typically spends time, and the entropy of their visits to that location. Hence, the task at hand is to identify dense clusters of GPS coordinates for each user. This is a typical example of a task which is very intuitive and quickly approximated by humans, but is extremely computationally expensive to solve exactly. Concretely, the problem of finding the optimal division of n data points into K clusters is formulated as minimizing the 'score' defined as

$$S = \sum_K \sum_{x_n \in C_k} |x_n - c_k|^2, \quad (2.2)$$


where c_k denotes the centroid of the cluster C_k . Each point x_n is assigned to the cluster corresponding to the nearest centroid. The usual way of approaching this problem is to use Lloyd's algorithm, which consists of initializing the centroids randomly, assigning each point to the cluster corresponding to the centroid which is nearest, then moving each centroid to the center of its points and repeating the last two steps until convergence. As this isn't guaranteed to converge on the global minimum of (2.2), the process can be repeated a number of times, keeping only the result with the lowest value of S . I accomplished this by writing a small Python module to perform various variations of Lloyd's algorithm and to produce plots of the resulting clusters. The code is included in section A.2.5.

This allows one to implement Lloyd's algorithm and visualize its result easily, as the code allows automatic plotting of the result from the

algorithm while automatically selecting different colors for the various clusters. As an example, the following code snippet generates 1000 random points, runs Lloyd's algorithm to determine clusters and saves a plot of the results.

```
points = [[random.uniform(-10,10), random.uniform(-10,10)] for _ in xrange(10**3)]
clusters = lloyds(X = points, K = 6, runs = 1)
draw_clusters(clusters = clusters, filename = 'lloyds_example.pdf')
```

This results in the following visualization:



I chose to modify the algorithm slightly on the following basis: Usually, the algorithm takes as its initial centroids a random sample of the data. I'll call this 'sample' initialization. This leads to a greater number of clusters being initialized in the areas with an increased density of data points, meaning that centroids will be highly cluttered at first, 'fighting' over the dense regions of data points then slowly spreading out. A few such iterations are shown in figure 2.9. However, this method is dangerous: The goal is to identify locations in which a user spends much of their time, i.e. in which more than some threshold of their GPS pings originated, and this initialization is likely to 'cut' the more popular locations into several clusters, neither of which contains more data points than the threshold. One example might be the DTU campus, which is a risk of being divided into several locations with too few data points in each, giving the false impression that user doesn't visit the campus that often. To avoid this effect, I implemented another initialization, 'scatter', in which the clusters start out on points select randomly from the entire range of x, y -values in the user's dataset. This turned out to not only solve the problem described above, but also converge much quicker and reach a slightly lower score as define in (2.2). A few such iterations are shown in figure 2.10. The difference in end results for the two methods is exemplified in figure 2.11. While this works great for users who stay in or around Copenhagen, it will cause problems for people who travel a lot. A user who has visited Australia, for instance, will have their initial clusters spread out across

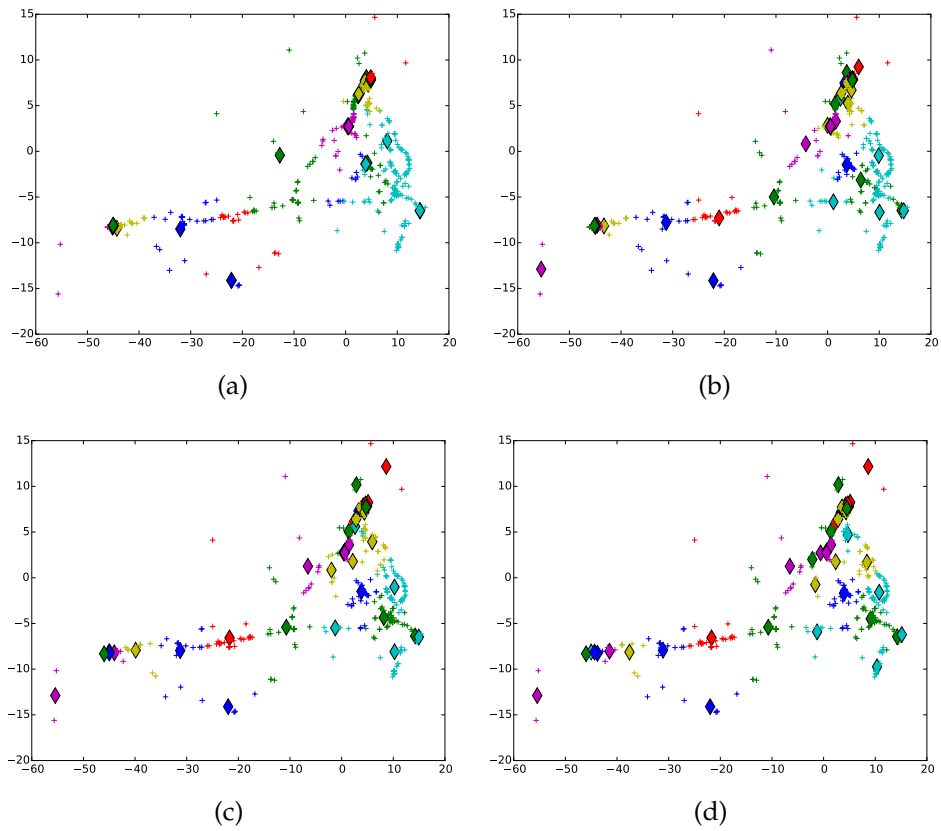


Figure 2.9: A few iterations of Lloyd's algorithm using 'sample' initialization. The axes denote the distance in km to some typical location for the user. Note that clusters are initially cluttered, then slowly creep away from the denser regions.

the globe, and it's highly likely that one them will end up representing all of Denmark. I ended up simply running both versions and keeping the result yielding the highest amount of locations.

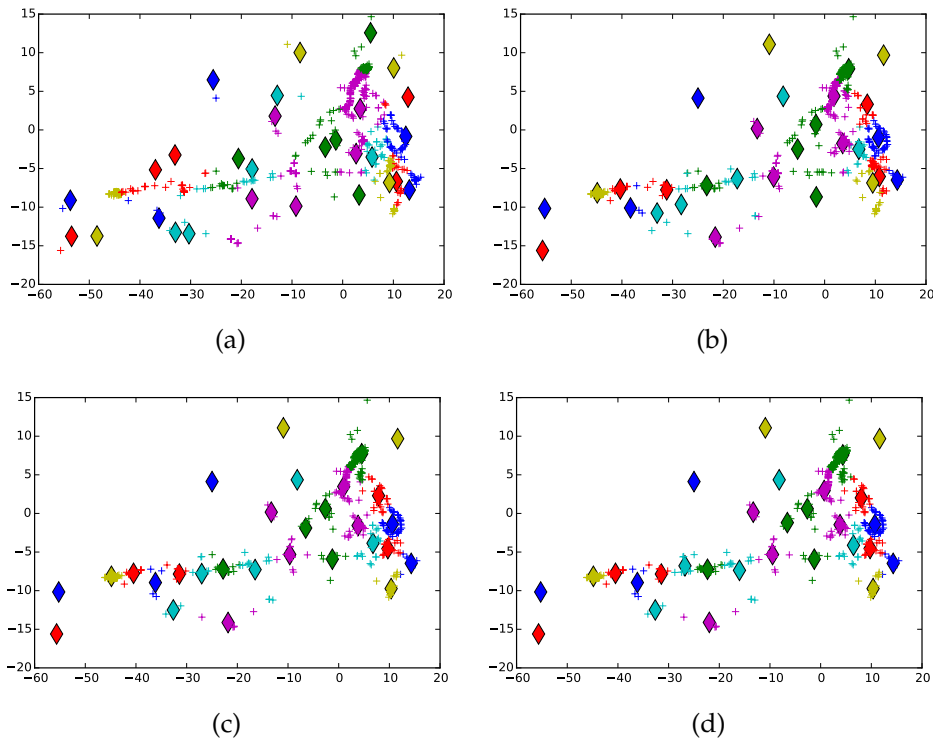




Figure 2.10: A few iterations of Lloyd's algorithm using 'scatter' initialization. The axes denote the distance in km to some typical location for the user. Note that clusters are initially randomly spread across the entire range of x, y -values and converge quickly to a local minimum for (2.2).



(a) Sample initialization.



(b) Scattered initialization.

Figure 2.11: Comparison of the final results of the two initialization methods using 100 initial clusters, a threshold of 5% of the data points before a cluster is considered a popular location and running the algorithm 10 times and keeping the best result. Clusters containing more than 5% of the total amount of data points are in color, whereas the remaining points are black dots.

2.2.3 Time Series Analysis

Another interesting aspect to include is what one somewhat qualitatively might call behavioural regularity - some measure of the degree in which a user's phone activities follow a regular pattern. Quantifying this turns out to take a bit of work. First of all, any user's activity would be expected to closely follow the time of day, so the timestamps of each user's outgoing texts and calls are first converted into 'clock times' meaning simply the time a regular clock in Copenhagen's time zone would display at the given time. This process is fairly painless when using e.g. the UTC time standard, which does not observe daylight saving time (DST), but some subtleties arise in countries that do use DST, as this makes the map from Unix/epoch time to clock time 'almost bijective' - when changing *away* from DST, two consecutive hours of unix time map to the same clock time period (02:00 to 03:00), whereas that same clock period is skipped when changing *to* DST. The most commonly used Python libraries for datetime arithmetic accommodate this by including a `dst` boolean in their datetime objects when ambiguity might arise, however I simply mapped the timestamps to clock times and ignored the fact twice a year, one time bin will artificially contain contributions from one hour too many or few. One resulting histogram is shown in figure 2.12.

Tilføj lidt om
AR-serier når
du har bo-
gen!!!

2.2.4 Facebook Data

Unfortunately, the only available Facebook data was a list of each user's friends, so the only contribution of each user's Facebook log was the number of friends the user had.

2.2.5 Bluetooth Data

I extracted a number of different features from each user's Bluetooth log file. First, I set a threshold for when a given user is considered social, as described in section 2.1.3. I chose to use a threshold of two. I then tried to estimate how much time each user spends in the physical company of others in the following way: for each time stamp in the user's Bluetooth log, I checked if the user was social or not and assumed that this status was the same until the following log entry, unless the delay was more than two hours. The rationale behind this is to avoid skewing the measurements if a user turns off their phone for extended periods of time. Otherwise, e.g. studying with a few friends at DTU, turning off your phone and going on vacation for two weeks would give the false impression that the user




Figure 2.12: Histogram of a user's outgoing calls and texts with a bin size of six hours.

were highly social for a long period of time. I then recorded the fraction of times the user was estimated as being social in this fashion.

Finally, I also wanted some measure of the degrees to which a user's social behaviour follows a pattern. I looked for temporal patterns by fitting AR-series and computing autocorrelation coefficients for each user's social behaviour as described in section 2.2.3. I also chose to compute a 'social entropy' much like (2.1), but weighted by the time the user spends with each acquaintance:

$$E = - \sum_i f_i \ln_2(f_i), \quad (2.3)$$

$$f_i = \frac{\text{time spent with } i}{\sum_j \text{time spent with } j}. \quad (2.4)$$

Note that the denominator of (2.4) is not equal to the total amount of time spent being social, as the contribution from each log entry is multiplied by the number of people present.

2.3 Output Data

The main emphasis of this part of the thesis is on predicting so-called *Big Five* personality traits. This section contains a brief description of those, following[6]. **Extraversion** signifies how extroverted and sociable a person is. People with high extraversion scores are supposed to be more eager to seek the company of others. **Agreeableness** is supposed to be a measure of how sympathetic or cooperative a person is, whereas **conscientiousness** denotes constraint, self discipline, level of organization etc.. **Neuroticism** signify the tendency to experience mood swings, and is complementary to emotional stability. Finally, **Openness**, also called 'openness to experience', or 'inquiring intellect' in earlier works, signifies thoughtfulness, imagination and so on. These five are collectively referred to as the 'big five' or 'OCEAN' after their initials.

In addition to the above, I also had access to a range self-explanatory traits about the participants such as their gender, whether they smoke etc.

2.4 Linear Discriminant Analysis & Gender Prediction

Linear discriminant analysis is basically a dimensionality reduction technique developed by Fisher in 1936 [7] for separating data points into two or more classes. The general idea is to project a collection of data points in n -dimensional variable space, onto the line or hyperplane which maximizes the separation between classes. Representing data points in n -space by vectors denoted x , the objective is to find a vector ω such that separation between the projected data points on it

$$y = \omega^T x \quad (2.5)$$

is maximized.

To break down the derivation of this method, I will first define a convenient distance measure used to optimize the separation between classes, then solve the resulting optimization problem. For clarity, I'll only describe the case of projection of two classes onto one dimension (i.e. using 'line' rather than 'hyperplane' and so on), although the method generalizes easily.

2.4.1 A measure of separation for projected Gaussians

If the projected data points for two classes a and b follow distributions \mathcal{N}_a and \mathcal{N}_b , which are standard Gaussians, $\mathcal{N}_i(x) = \mathcal{N}(x; \mu_i, \sigma_i^2)$, the joint probability distribution for the distance between the projections will be the convolution

$$P(x) = \int_{-\infty}^{\infty} \mathcal{N}_a(y) \cdot \mathcal{N}_b(x - y) dy. \quad (2.6)$$

Computing this for a Gaussian distribution,

$$\mathcal{N}(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}, \quad (2.7)$$

becomes easier with the convolution theorem, which I'll derive in the following.

Denoting convolution by $*$ and Fourier transforms by

$$\mathcal{F}(f) = \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(x) \cdot e^{-i\omega x} dx, \quad (2.8)$$

the convolution theorem is derived as follows:

$$\mathcal{F}(f * g) = \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} \int_{\mathbb{R}^n} f(y) \cdot g(x - y) dy e^{-i\omega x} d\omega, \quad (2.9)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) \int_{\mathbb{R}^n} g(x - y) e^{-i\omega x} dy d\omega, \quad (2.10)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) \int_{\mathbb{R}^n} g(z) e^{-i\omega(z+y)} dz d\omega, \quad (2.11)$$

$$= \frac{1}{(2\pi)^{n/2}} \int_{\mathbb{R}^n} f(y) e^{-i\omega y} \int_{\mathbb{R}^n} g(z) e^{-i\omega z} dz d\omega, \quad (2.12)$$

$$\boxed{\mathcal{F}(f * g) = (2\pi)^{n/2} \mathcal{F}(f) \cdot \mathcal{F}(g)}, \quad (2.13)$$

where the factor in front of the usual form of the theorem $\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$ stems from the convention of using angular frequency in Fourier transforms, as in (2.8), rather than

$$\mathcal{F}(f) = \int_{\mathbb{R}^n} f(x) \cdot e^{-2\pi i v x} dx. \quad (2.14)$$

Using this, the convolution of two Gaussians can be calculated as

$$\mathcal{N}_a * \mathcal{N}_b = (2\pi)^{n/2} \mathcal{F}^{-1}(\mathcal{F}(\mathcal{N}_a) \cdot \mathcal{F}(\mathcal{N}_b)). \quad (2.15)$$

The required Fourier transform can be massaged into a nicer form by displacing the coordinate system and cancelling out terms with odd parity:

$$\begin{aligned}
 \mathcal{F}(\mathcal{N}(x)) &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \cdot e^{-i\omega x} dx, \\
 &= \frac{1}{2\pi\sigma} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} e^{-i\omega(x+\mu)} dx, \\
 &= \frac{1}{2\pi\sigma} e^{-i\omega\mu} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} (\cos(\omega x) + i \sin(\omega x)) dx, \\
 &= \underbrace{\frac{1}{2\pi\sigma} e^{-i\omega\mu}}_a \underbrace{\int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} \cos(\omega x) dx}_{I(\omega)}. \tag{2.16}
 \end{aligned}$$

Noting that $I(\omega)$ reduces to an ordinary Gaussian integral at $\omega = 0$ so $I(0) = \sqrt{2\pi}\sigma$, this can be solved with a cute application of Feynman's trick:

$$\begin{aligned}
 \frac{\partial I}{\partial \omega} &= - \int_{-\infty}^{\infty} x e^{-\frac{x^2}{2\sigma^2}} \sin(\omega x) dx, \\
 &= \int_{-\infty}^{\infty} \sigma^2 \frac{\partial}{\partial x} \left(e^{-\frac{x^2}{2\sigma^2}} \right) \sin(\omega x) dx, \\
 &= \sigma^2 e^{-\frac{x^2}{2\sigma^2}} \sin(\omega x) \Big|_{-\infty}^{\infty} - \omega \int_{-\infty}^{\infty} \sigma^2 e^{-\frac{x^2}{2\sigma^2}} \cos(\omega x) dx, \\
 &= -\omega \sigma^2 I(\omega) \Leftrightarrow \\
 I(\omega) &= C e^{-\sigma^2 \omega^2 / 2}, \\
 I(0) &= C = \sqrt{2\pi}\sigma, \\
 I(\omega) &= \sqrt{2\pi}\sigma e^{-\sigma^2 \omega^2 / 2}.
 \end{aligned}$$

Plugging this into (2.16) gives the result

$$\mathcal{F}(\mathcal{N}) = \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu} e^{-\sigma^2 \omega^2 / 2}. \tag{2.17}$$

This can be used in conjunction with (2.13) to obtain

$$\mathcal{F}(\mathcal{N}_a * \mathcal{N}_b) = \sqrt{2\pi} \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu_a} e^{-\sigma_a^2 \omega^2 / 2} \cdot \frac{1}{\sqrt{2\pi}} e^{-i\omega\mu_b} e^{-\sigma_b^2 \omega^2 / 2}, \tag{2.18}$$

$$= \frac{1}{\sqrt{2\pi}} e^{-i\omega(\mu_a + \mu_b)} e^{-(\sigma_a^2 + \sigma_b^2) \omega^2 / 2}, \tag{2.19}$$

which is recognized as the transform of another Gaussian describing the separation with $\mu_s = \mu_a - \mu_b$ and $\sigma_s^2 = \sigma_a^2 + \sigma_b^2$, so taking the inverse Fourier transformation gives the convolution

$$\mathcal{N}_a * \mathcal{N}_b = \frac{1}{\sqrt{2\pi}\sigma_s} e^{-\frac{(x-\mu_s)^2}{2\sigma_s^2}}. \quad (2.20)$$

Hence, a reasonable measure of the separation of two projected distributions is

$$d = \frac{(\mu_a - \mu_b)^2}{\sigma_a^2 + \sigma_b^2}. \quad (2.21)$$

2.4.2 Optimizing separation

To maximize the separation, the numerator and denominator, respectively, of (2.21) can be rewritten in terms of w in the following way (using $\tilde{\mu}_i$ to denote projected means) and simplified by introducing scattering matrices:

$$(\tilde{\mu}_a - \tilde{\mu}_b)^2 = (w^T (\mu_a - \mu_b))^2, \quad (2.22)$$

$$= w^T (\mu_a - \mu_b) (\mu_a - \mu_b)^T w, \quad (2.23)$$

$$= w^T S_B w, \quad (2.24)$$

and

$$\tilde{\sigma}_i^2 = \sum_{y \in i} \frac{1}{N} (y - \tilde{\mu}_i)^2, \quad (2.25)$$

$$= w^T \sum_{y \in i} (x - \mu_i) (x - \mu_i)^T w, \quad (2.26)$$

$$= w^T S_i w, \quad (2.27)$$

$$\tilde{\sigma}_a^2 + \tilde{\sigma}_b^2 = w^T S_W w, \quad (2.28)$$

having introduced the between-class and within-class scatter matrices S_B and S_W by

$$S_B = (\mu_a - \mu_b) (\mu_a - \mu_b)^T, \quad (2.29)$$

$$S_i = \sum_{y \in i} (x - \mu_i) (x - \mu_i)^T, \quad (2.30)$$

$$S_W = S_a + S_b. \quad (2.31)$$

Hence, the objective is to solve

$$\frac{d}{dw} J(w) = \frac{d}{dw} \left(\frac{w^T S_B w}{w^T S_W w} \right) = 0, \quad (2.32)$$

$$\frac{\frac{d[w^T S_B w]}{dw} w^T S_W w - w^T S_B w \frac{d[w^T S_W w]}{dw}}{(w^T S_W w)^2} = 0, \quad (2.33)$$

$$2S_B w \cdot w^T S_W w - w^T S_B w \cdot 2S_W w = 0, \quad (2.34)$$

$$S_B w - \frac{w^T S_B w \cdot S_W w}{w^T S_W w} = 0, \quad (2.35)$$

$$S_B w - S_W w J(w) = 0, \quad (2.36)$$

$$S_B w = S_W w J(w), \quad (2.37)$$

$$S_W^{-1} S_B w = J(w) w. \quad (2.38)$$

The optimal projection vector w^* which satisfies this is

$$w^* = S_W^{-1} (\mu_a - \mu_b). \quad (2.39)$$

Vær lige
sikker på at
du forstår det
her.

Figure 2.13 shows a visualization of this that I generated by drawing (x, y) points from two random distributions to simulate two distinct classes of points. If the distributions are independent and Gaussian, the projections will also form Gaussian distributions, and the probability of a new point belonging to e.g. class a given its coordinates d can be estimated using Bayesian probability

$$P(a|d) = \frac{P(d|a)P(a)}{P(d|a)P(a) + P(d|b)P(b)}, \quad (2.40)$$

where $P(a)$ and $P(b)$ are simply the prior probabilities for encountering the respective classes, and the conditional probabilities, e.g. $P(d|a)$ are simply given by the value of the projected Gaussian $\mathcal{N}(x'; \tilde{\mu}_a, \tilde{\sigma}_a)$ at the projected coordinate x' . In practise, even when the points are not independent or Gaussian, so that (2.40) is not a precise estimate of the probability of the point representing a given class, the class with the highest posteriori according to (2.40) still often turns out to be a good guess.

This method accurately predicted the gender of 79.8% of the participants, which is not particularly impressive as 77.3% of participants were male, so a classifier that assumes that every participant is male would have a comparable success rate. An immediate source of concern is the assumption of linearity: It is possible that the data is ordered in such a way that it is possible to separate data points fairly well based on




Figure 2.13: Two collections of points drawn from independent Gaussian distributions, representing **class a** and **class b**. If the points are projected onto the **straight line**, which is given by (2.39), the separation between the peaks representing the two classes is maximized.

gender or some psychological trait, just not using a linear classifier. As an extreme example of this, figure 2.14 shows a situation where the points representing one class are grouped together in an 'island' in the middle, isolating them from points representing the remaining class. While it is clear that there's a pattern here, a linear classifier fails to predict classes more precisely than their ratio. Support Vector Machines, or SVMs are another linear classification technique which can be generalized to detect patterns like that in figure 2.14. This is described in section 3.1





Figure 2.14: An example of data points representing **class a** are clearly discernible from those of **class b**, yet a linear Fisher classifier fails to predict the classes more precisely than the ratio of **b** to **a**.

PSYCHOLOGICAL PROFILING & MACHINE LEARNING

ACHINE learning is currently a strong candidate for prediction of psychological profiles from phone data. This chapter describes the application of the quantitative data described in section 2.2 and various machine learning schemes, starting with support vector machines (SVMs).

1000 kilder!!!

Uddyb når
der er flere
modeller.

3.1 Support Vector Machines

The purpose of this section is to introduce SVMs and attempt to apply them to the data obtained in 2.2. The introduction is mainly based on introductory texts by Marti Hearst [10] and Christopher Burges [4]. SVMs in their simplest form (*simplest* meaning using a linear kernel, which I'll explain shortly) can be thought of as a slight variation on the linear classifier described in section 2.4. However, where LDA finds a line such that the distribution of the points representing various classes projected onto the line is maximized, the aim of SVMs is to establish the hyperplane that represents the best possible slicing of the feature space into regions containing only points corresponding to the different classes. A simple example of this is shown in figure 3.1. Using labels ± 1 to denote classes, the problem may be stated as trying to guess the mapping from an N-dimensional data space to classes $f : \mathbb{R}^N \rightarrow \{\pm 1\}$ based on a set of training data in $\mathbb{R}^N \otimes \{\pm 1\}$. I'll describe separately the properties of




Figure 3.1: The same points as those shown in figure 2.13, except points in class a and class b are now pictured along with their maximally separating hyperplane.

this maximally separating hyperplane, how it is obtained, and how the method is generalized to non-linear classification problems as the 'island' illustrated in figure 2.14.

The well-known equation for a plane is obtained by requiring that its normal vector \mathbf{w} be orthogonal to the vector from some point in the plane \mathbf{p} to any point \mathbf{x} contained in it:

$$\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}) = 0. \quad (3.1)$$

The left hand side of (3.1) gives zero for points in the plane and positive or negative values when the point is displaced in the same or opposite direction as the normal vector, respectively. Hence, $\text{sign}(\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}))$ may be taken as the decision function. It is clear from (3.1) that the normal vector may be scaled without changing the actual plane (of course the decision function is inverted if a negative value is chosen), so \mathbf{w} is usually rescaled such that

$$\mathbf{w} \cdot (\mathbf{x} - \mathbf{p}) = \mathbf{w} \cdot \mathbf{x} + b = \pm 1, \quad (3.2)$$

for the points that are closest to the separating plane. Those points located on the margin are encircled in figure 3.1. In general then, the meaning of

the sign and magnitude of

$$\mathbf{w} \cdot \mathbf{x} + b \quad (3.3)$$

will be the predicted class and a measure of prediction confidence, respectively, for new data points. Finally, note that \mathbf{w} can be expanded in terms of the data points that are on the margin in figure 3.1 as

$$\mathbf{w} = \sum_i v_i \mathbf{x}_i, \quad (3.4)$$

these \mathbf{x}_i , the position vectors of the margin points in data space, are the 'support vectors' that lend their name to the method.

3.1.1 Obtaining the Maximally Separating Hyperplane

Assuming first that it is possible to slice the data space into two regions that contain only points corresponding to one class each, and that the plane's normal vector has already been rescaled according to (3.2), the following inequalities hold:

$$\begin{aligned} \mathbf{x}_i \cdot \mathbf{w} + b &\geq 1, y_i = +1, \\ \mathbf{x}_i \cdot \mathbf{w} + b &\leq -1, y_i = -1. \end{aligned} \quad (3.5)$$

Multiplying by y_i , both simply become

$$y_i (\mathbf{x}_i \cdot \mathbf{w} + b) - 1 \geq 0. \quad (3.6)$$

The separation between the two margins shown with dashed lines in figure 3.1 is $2/|\mathbf{w}|$, so the Lagrangian

$$L = \frac{1}{2} |\mathbf{w}|^2 - \sum_i \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{w} + b - 1), \quad (3.7)$$

must be minimized with the constraints

$$\alpha_i \geq 0, \quad (3.8)$$

$$\frac{\partial L}{\partial \alpha_i} = 0. \quad (3.9)$$

A result from convex optimization theory known as Wolfe Duality[17] states that one may instead maximize the above Lagrangian subject to

$$\nabla_w L = \frac{\partial L}{\partial b} = 0, \quad (3.10)$$

ikke 100% tryg
her...

which gives conditions

$$\mathbf{w} = \sum_j \alpha_j y_j \mathbf{x}_j, \quad (3.11)$$

$$\sum_j \alpha_j y_j = 0. \quad (3.12)$$

These can be plugged back into (3.7) to obtain

$$L_D = \frac{1}{2} \sum_i \sum_j \alpha_i y_i \alpha_j y_j \mathbf{x}_i \cdot \mathbf{x}_j - \sum_i \alpha_i y_i \left(\mathbf{x}_i \cdot \sum_j \alpha_j y_j \mathbf{x}_j + b \right) + \sum_i \alpha_i, \quad (3.13)$$

$$L_D = -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i \mathbf{x}_j + \sum_i \alpha_i. \quad (3.14)$$

A problem with this is that eqs 3.5 can only be satisfied in the completely separable case, although it is easy to imagine an example in which a classifier performs well but not flawlessly on the training set. For instance, if two points, one from each class, in figure 3.1 were permuted, the classifier shown in the plot would still do a very good job, but eqs. 3.5 would not be satisfiable, causing the method to fail. This is remedied by introducing slack variables

$$\mathbf{x}_i \cdot \mathbf{w} + b \geq 1 - \xi_i, \quad y_i = +1, \quad (3.15)$$

$$\mathbf{x}_i \cdot \mathbf{w} + b \leq -(1 - \xi_i), \quad y_i = -1, \quad (3.16)$$

$$\xi_i \geq 0, \quad (3.17)$$

The above can be rewritten exactly as previously, except another set of non-negative Lagrange multipliers are added to (3.7) to ensure positivity of the ξ_i . Also a cost term, usually $C \cdot \sum_i \xi_i$, is added to the Lagrangian to keep separability high:

$$L = \frac{1}{2} |\mathbf{w}|^2 + C \cdot \sum_i \xi_i - \sum_i \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{w} + b - 1 + \xi_i) - \sum_i \mu_i \xi_i. \quad (3.18)$$

This results in the same dual Lagrangian L_D as before, but with an upper bound on the α_i :

$$0 \leq \alpha_i \leq C. \quad (3.19)$$

The values of the slack variables ξ_i and the cost term C are typically decided by performing a 'grid search' in which the performance is evaluated for each possibly combination of the parameters and the optimal combination




Figure 3.2: Heat map of the result of a grid search over the parameter space of the variables ξ and C . The colors signify the mean squared error of a support vector regression problem, which is closely related to the classification problem described in the present section. Note how overfitting damages performance for high values of C .

used in the final classifier. Figure 3.2 shows a heat map of the results of a grid search over the parameters in question. The main point to be emphasized here is that the training data \mathbf{x}_i only enter into the dual Lagrangian of (3.14) as inner products. This is essential when extending the SVM model to nonlinear cases, which is the subject of the following section.

3.1.2 Generalizing to the non-linear case

The fact that the data \mathbf{x}_i only occur as inner products in (3.14) makes one way of generalizing to non-linearly separable datasets straightforward: Referring back to figure 2.14, one might imagine bending the plane containing the data points by curling the edges outwards in a third dimension after which a two-dimensional plane could separate the points very well. In general, this means applying some mapping

$$\Phi : \mathbb{R}^l \rightarrow \mathbb{R}^h, \quad h > l, \quad (3.20)$$

to the \mathbf{x}_i (l and h are for low and high, respectively). For example, one could look for a mapping such that the new inner product becomes

$$\Phi(\mathbf{x}_i) \cdot \Phi(\mathbf{x}_j) = (\mathbf{x}_i \cdot \mathbf{x}_j)^2. \quad (3.21)$$

I'll describe the components of each vector separately, so I'm going to change notation to let the subscripts denote coordinates and using \mathbf{x} and \mathbf{y} as two arbitrary feature vectors, where the latter shouldn't be confused with the class label used earlier. As an example, in two dimensions the above becomes

$$(\mathbf{x} \cdot \mathbf{y})^2 = \left(\sum_{i=1}^2 x_i y_i \right)^2 = x_1^2 y_1^2 + 2x_1 y_1 x_2 y_2 + x_2^2 y_2^2, \quad (3.22)$$

meaning that one possibility for Φ is

$$\Phi : \mathbf{x} \mapsto \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1 x_2 \\ x_2^2 \end{pmatrix} \quad (3.23)$$

This can be generalized to d -dimensional feature vectors and to taking the n' th power rather than the square using the multinomial theorem:

$$\left(\sum_{i=1}^d x_i \right)^n = \sum_{\sum_{i=1}^d k_i = n} \frac{n!}{\prod_{l=1}^d k_l!} \prod_{j=1}^d x_j^{k_j}, \quad (3.24)$$

where the subscript $\sum_{i=1}^d k_i = n$ simply means that the sum goes over any combination of d non-negative integers k_i that sum to n . I wish to rewrite this slightly for two reasons: to simplify the notation in order to make a later proof more manageable, and to help quantify how quickly the number of dimensions in the output space grows to motivate a trick to avoid these explicit mappings.

As stated, the sum on the RHS of (3.24) runs over all combinations of d integers which sum to n . This can be simplified by introducing a function K , which simply maps

$$K : n, d \mapsto \left\{ \{k\} \in \mathbb{N}^d \mid \sum_{i=1}^d k_i = n \right\}, \quad (3.25)$$

and denoting each of those collections $\{k\}_i$ so each of the coefficients in (3.24) can be written

$$\frac{n!}{\prod_{i=1}^d k_i!} = C_{\{k\}}. \quad (3.26)$$

Dobbeltjek om notation er fornuftig.

Then, (3.24) becomes

$$\left(\sum_{i=1}^d x_i \right)^n = \sum_{K(n,d)} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} \quad (3.27)$$

To show how quickly the dimensions of the required embedding space grows, note that the dimension is equal to the number of terms in the sum above, i.e.

$$\dim(\mathbb{R}^h) = |K(n, d)| = \left| \left\{ \{k\} \in \mathbb{N}^d \left| \sum_{i=1}^d k_i = n \right. \right\} \right|, \quad (3.28)$$

which can be computed using a nice trick known from enumerative combinatorics.

Consider the case where $n = 5$ and $d = 3$. $K(5, 3)$ then contains all sets of 3 integers summing to 5, such as 1, 3, 1 or 0, 1, 4. Each of these can be uniquely visualized as 5 unit values distributed into 3 partitions in the following fashion:

$$\begin{array}{c} \circ \mid \circ \circ \circ \mid \circ, \\ \mid \circ \mid \circ \circ \circ \circ, \end{array}$$

and so on. It should be clear that you need n \circ -symbols and $d - 1$ \mid separators. The number of possible such combinations, and hence the dimensionality of the embedding space, is then

$$\binom{n + d - 1}{n} = \frac{(n + d - 1)!}{n!(d - 1)!}. \quad (3.29)$$

This number quickly grows to be computationally infeasible, which motivates one to look for a way to compute the inner product in the embedded space without performing the explicit mapping itself. This is the point of the so-called 'kernel trick', which I'll introduce in the following.

The idea of the kernel trick is that since only the inner products between feature vectors in the embedded space are required, one might as well look for some function K of the original feature vectors which gives the same scalar as the inner product in the embedded space, i.e.

$$K(\mathbf{x}, \mathbf{y}) = \Phi(\mathbf{x}) \cdot \Phi(\mathbf{y}). \quad (3.30)$$

In the polynomial case treated above, the correspondence between the kernel function $K(\mathbf{x}, \mathbf{y})$ and the explicit mapping Φ is straightforward:

$$K(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y})^n, \quad (3.31)$$

$$\Phi(\mathbf{x}) = \sum_{K(n,d)} \sqrt{C_{\{k\}}}} \prod_{j=1}^d x_j^{k_j}, \quad (3.32)$$

so that (3.30) is true by the multinomial theorem and the above considerations. However, situations arise in which the explicit mapping Φ isn't directly obtainable, and the correspondence of the kernel function to inner products in higher dimensional spaces is harder to demonstrate. This is the subject of the following section.

Radial Basis Functions

One commonly used kernel function is the RBF, or radial basis function, kernel:

$$K(\mathbf{x}, \mathbf{y}) = e^{|\mathbf{x}-\mathbf{y}|^2/2\sigma}. \quad (3.33)$$

Burges [4] shows that the polynomial kernel is valid, so I'll show how the argument extends to the RBF kernel in the following.

find en eller
anden kilde

Mercer's condition states that for a kernel function $K(\mathbf{x}, \mathbf{y})$, there exists a corresponding Hilbert space \mathcal{H} and a mapping Φ as specified earlier, iff any L^2 -normalizable function $g(\mathbf{x})$ satisfies

$$\int K(\mathbf{x}, \mathbf{y}) g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \geq 0. \quad (3.34)$$

This can be shown by rewriting (3.33) as

$$K(\mathbf{x}, \mathbf{y}) = e^{(\mathbf{x}-\mathbf{y}) \cdot (\mathbf{x}-\mathbf{y})/2\sigma} = e^{|\mathbf{x}|^2/2\sigma} e^{|\mathbf{y}|^2/2\sigma} e^{-\mathbf{x} \cdot \mathbf{y}/\sigma}, \quad (3.35)$$

and expanding the last term in $(\mathbf{x} \cdot \mathbf{y})$ as

$$e^{-\mathbf{x} \cdot \mathbf{y}/\sigma} = \sum_{i=0}^{\infty} \frac{(-1)^i}{i! \sigma^i} (\mathbf{x} \cdot \mathbf{y})^i, \quad (3.36)$$

but using (3.27) on the dot product gives

$$(\mathbf{x} \cdot \mathbf{y})^i = \left(\sum_{j=1}^d x_j y_j \right)^i = \sum_{K(i,d)} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} y_j^{k_j} \quad (3.37)$$

so the Taylor expansion becomes

$$e^{-\mathbf{x} \cdot \mathbf{y} / \sigma} = \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \prod_{j=1}^d x_j^{k_j} y_j^{k_j}, \quad (3.38)$$

which can be plugged back into (3.35) to yield

$$K(\mathbf{x}, \mathbf{y}) = \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} e^{|\mathbf{x}|^2 / 2\sigma} e^{|\mathbf{y}|^2 / 2\sigma} \prod_{j=1}^d x_j^{k_j} y_j^{k_j}. \quad (3.39)$$

The underlying reason for these algebraic shenanigans is that (3.39) is clearly separable so that the integral in (3.34) from Mercer's condition becomes

$$\int K(\mathbf{x}, \mathbf{y}) g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \quad (3.40)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \int_{\mathbb{R}^{2d}} e^{|\mathbf{x}|^2 / 2\sigma} e^{|\mathbf{y}|^2 / 2\sigma} \prod_{j=1}^d x_j^{k_j} y_j^{k_j} g(\mathbf{x}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \quad (3.41)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \left(\int_{\mathbb{R}^d} e^{|\mathbf{x}|^2 / 2\sigma} \prod_{j=1}^d x_j^{k_j} g(\mathbf{x}) d\mathbf{x} \right) \cdot \left(\int_{\mathbb{R}^d} e^{|\mathbf{y}|^2 / 2\sigma} \prod_{j=1}^d y_j^{k_j} g(\mathbf{y}) d\mathbf{y} \right) \quad (3.42)$$

$$= \sum_{i=0}^{\infty} \sum_{K(i,d)} \frac{(-1)^i}{i! \sigma^i} C_{\{k\}} \left(\int_{\mathbb{R}^d} e^{|\mathbf{x}|^2 / 2\sigma} \prod_{j=1}^d x_j^{k_j} g(\mathbf{x}) d\mathbf{x} \right)^2 \quad (3.43)$$

$$\geq 0. \quad (3.44)$$

Hence, radial basis functions satisfy Mercer's condition and the kernel described above can be plugged into the dual Lagrangian from (3.14) to obtain

$$L_D = -\frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j e^{|\mathbf{x}_i - \mathbf{x}_j|^2 / 2\sigma} + \sum_i \alpha_i, \quad (3.45)$$

which must be maximized subject to the same constraints as earlier. The concrete optimization procedure is complicated and already implemented in most machine learning libraries, so I choose not to go into details with that, but instead to demonstrate the effectiveness of the RBF kernel approach on the non-linear-separable points that were generated earlier. figure 3.3 shows the points again, along with the *decision frontier* i.e. the curve which separates regions in which points are classified into separate




Figure 3.3: The ‘island’ scenario of figure 2.14 revisited. The points representing **class a** and **class b** have been mapped to a higher-dimensional space in which it is possible to construct a separating hyperplane whose **decision frontier** is also shown.

classes. The danger of overfitting should be clear from figure 3.3. If the cost of misclassification C and the sharpness of the RBFs, usually denoted by $\gamma = 2/\sigma$ are set sufficiently high, the algorithm will simply end up with a tiny decision boundary around every training point of class a, resulting in flawless classification on the training set, but utter failure on new data. The typical way of evaluating this is to perform k -fold validation, meaning that the available data is k equal parts and the SVM is consecutively trained on $k - 1$ parts and tested on the remaining. A variant of this, which my code uses, is stratified k -fold validation, which only differs in that the data is partitioned so as to keep the ratio between the different classes in each parts as close to equal as possible.

The γ parameter is often fixed by performing a grid search similar to that discussed earlier. Figure 3.4 shows the resulting heat map from a grid search.




Figure 3.4: Result of a grid search for the optimal combination of values for the cost parameter C and the sharpness γ of the Gaussian kernel function used.

3.1.3 Statistical subtleties

An important note should be made here about some often neglected subtleties relating to uncertainties. Physicists often deal with measurements that can be assumed to be independently drawn from a normal distribution $\mathcal{N}(x_i; \mu, \sigma^2)$ due to the central limit theorem. With a large number of measurements n , the standard deviation of a sample

$$\sigma^2 = \frac{1}{N} \sum_i^N (x_i - \mu)^2, \quad (3.46)$$

converges as $N \rightarrow \infty$ to the maximum likelihood, minimum variance unbiased estimator for the true variance of the underlying distribution with unknown mean

$$\hat{\sigma}^2 = \frac{1}{N-1} \sum_i^N (x_i - \mu)^2. \quad (3.47)$$

The standard deviation σ and the width of the underlying gaussian $\hat{\sigma}^2$ can then often be used interchangeably. This tempts some people into the

questionable habit of always assuming that the sample standard deviance can be used as the 68% confidence interval of their results.

When using a K-fold validation scheme, the performance scores for the various folds cannot be assumed to be independently drawn from an underlying distribution, as the test set of one fold is used in the training sets of the remaining folds. In fact, it has been shown [1] that there is no unbiased estimator for the variance of the performance estimated using K-fold validation. However, as K-fold validation is more effective than keeping the test, and training data separate, which can be shown using Jensen's inequality along with some basic properties of expectation values [2], I'll mostly use K-fold regardless. As the standard deviation still provides a qualitative measure of the consistency of the model's performance, I'll still use the sample STD in a usual fashion, such as error bars, unless otherwise is specified, but the reader should keep in mind that these do not indicate precise uncertainties whenever K-fold validation has been involved.

3.2 Decision Trees & Random Forests

Another popular machine learning scheme is that of random forests, which consist of an ensemble of decision trees. A decision tree is a very intuitive method for classification problems which can be visualized as a kind of flow chart in the following fashion. As usual, the problem consists of a set of feature vectors \mathbf{x}_i and a set of corresponding class labels y_i . A decision tree then resembles a flowchart starting at the root of the tree, at each node splitting into branches and finally branching into leaves at which all class labels should be identical. At each node, part of the feature vector is used to split the dataset into parts. This resembles the 'twenty questions' game, in which one participant thinks of a famous person and another attempts to guess who it is by asking a series of yes/no-questions, each one splitting the set of candidates in two parts. In this riddle game and in decision tree learning, there are good and bad questions (asking whether the person was born on March 14th, 1879 is a very bad first question, for instance). There are several ways of quantifying how 'good' a yes/no-question, corresponding to a partitioning of the dataset, is.

One metric for this is the Gini Impurity Index I_G , which is computed by summing over each class label:

$$I_G = \sum_i f_i(1 - f_i) = 1 - \sum_i f_i^2, \quad (3.48)$$

where f_i denotes the fraction of the set that consists of class y_i . Using this as a metric, the best partitioning is the one which results in the largest drop in total Gini impurity following a branching. Another metric is the information gain measured by comparing the entropy before a split with a weighted average of the entropy in the groups resulting from the split. Denoting the fractions of the various classes in the parent group, i.e. before splitting, by f_i and the two child groups by a_i and b_i , the information gain is

$$I_E = - \sum_i f_i \log_2 f_i + \frac{n_a}{N} \sum_i a_i \log_2 a_i + \frac{n_b}{N} \sum_i b_i \log_2 b_i. \quad (3.49)$$

However, if too many such nodes are added to a decision tree, overfitting, i.e. extreme accuracies on training data but poor performance on new data, becomes a problem. This can be remedied by instead predicting with a majority vote, or averaging in the case of regression problems, from an ensemble of randomized decision trees called a random forest. The main merits of random forests are their accuracy and ease of use, and their applications as auxiliary methods in other machine learning schemes, which I'll elaborate on shortly.

The individual trees in a random forest are grown using a randomly selected subset of the training data for each tree. The data used to construct a given tree is referred to as 'in bag', whereas the remaining training data is referred to as 'out of bag' (OOB) for the given tree. At each node, a set number of features is randomly selected and the best possible branching, cf. the above considerations, is determined. The only parameters that must be tweaked manually are the number of trees in the forest, number of features to include in each branching, and the maximum tree depth. While other variables such as the metric for determining branching quality as described above, may be customized, those aren't essential to achieve a decent predictor, which is robust in regard to both overfitting and irrelevant parameters.[3]

There doesn't seem to be a single universally accepted way of adjusting these parameters, so I chose a somewhat pragmatic approach of simply checking how well various choices for each parameter performed on a randomly selected trait. For instance, figure 3.5 shows how well a random forest predicted the tertiles of participants' extroversion as a function of the fraction of available features each tree was allowed to include in each branching. This was done using a large number of trees ($n = 1000$) and using each of the two metrics described earlier. The number of features used per split doesn't seem to have any significant effect on performance, and as the entropy metric seems to perform as well or slightly better than Gini impurity, I decided to stick to that. A similar plot of the performance




Figure 3.5: Performance of a random forest with 1000 decision trees using various fractions of the available features in each branching using both the **entropy** and the **Gini impurity** metric to determine the optimal branching. The number of features seems not to play a major role, and the **entropy** metric seems to perform slightly better in general.

of various numbers of decision trees in the forest is shown in figure 3.6. The performance seems to stagnate around 100 trees, and remain constant after that, so I usually used at least 500 trees to make sure to get the optimal performance, as runtime wasn't an issue.

The robustness to irrelevant features and overfitting described earlier also plays a role in the application of random forests in conjunction with other schemes. SVMs as described in section 3.1 can be sensitive to irrelevant data[14]. There exist off-the-shelf methods, such as recursive feature elimination (RFE)[9], for use with linear SVMs, but to my knowledge, there is no 'standard' way to eliminate irrelevant features when using a non-linear kernel. However, it is possible to use a random forest approach to obtain the relative importance of the various features and then use only the most important ones in another machine learning scheme which is less tolerant to the inclusion of irrelevant data. The relative importance of feature j can be estimated by first constructing a random forest and evaluating its performance s , then randomly permuting the values of feature j across the training sample and measure the damage it does to




Figure 3.6: Example of a random forest performance versus number of decision trees. Performance seems to increase steadily until about 100 trees, then stagnate.

the performance of the forest by comparing with the permuted score s_p . The ratio of the mean to standard deviation of those differences:

$$w_j = \frac{\langle s - s_p \rangle}{\text{std}(s - s_p)} \quad (3.50)$$

Random forests also provide a natural measure of similarity between data points. Given data points i and j these can be plugged into all their OOB decision trees, or a random subset thereof, and the fraction of the attempts in which both end up at the same leaf can be taken as a measure of similarity. This can be used to generate a proximity matrix for the data points, and it can be used as a metric for determining the nearest neighbours of a new point in conjunction with a simple nearest neighbour classifier.

3.3 Nearest Neighbour-classifiers

Skriv en masse om den smart random forest NN-model.

Do iiit!




Figure 3.7: Comparison of performance of models using **random forest** and **support vector regression** with a **baseline model** which always predicts the mean of the training sample. The y axis shows the mean error of each model and the error bars show the 95-percentile around the median scores obtained by running on 1000 bootstrap samples.

3.4 Results

HARJ



APPENDIX

A.1 Source Code for Explicit Semantic Analysis

This section contains code pertaining to part I of the thesis.

A.1.1 Parser

```

1  # -*- coding: utf-8 -*-
2  '''Parses a full Wikipedia XML-dump and saves to files containing
3  a maximum of 1000 articles.
4  In the end, each file is saved as a JSON file containing entries like:
5  {{
6      'concept':
7      {
8          'text': <article contents>,
9          'links_in' : Set of links TO the article in question,
10         'links_out' : Set of links FROM the article in question,
11     }
12 }
13 Although links_in is added by the generate_indices script.
14 Also saved are dicts for keeping track of word and concept indices when
15 building a large sparse matrix for the semantic interpreter.
16 The file structure is like {'word blah' : index blah}'''
17
18 import re
19 import xml.sax as SAX
20 import wikicleaner
21 import os
22 import glob
23 import shared
24 import sys
25
26 DEFAULT_FILENAME = 'medium_wiki.xml'
27
28 def canonize_title(title):
29     # remove leading whitespace and underscores
30     title = title.strip(' _')
31     # replace sequences of whitespace and underscore chars with a single space
32     title = re.compile(r'[\s_]+').sub(' ', title)
33     # remove forbidden characters
34     title = re.sub('[?/\*\\"\'`]', '', title)
35     return title.title()
36
37 # Import shared parameters
38 from shared import extensions, temp_dir
39
40 # Cleanup
41 for ext in extensions.values():
42     for f in glob.glob(temp_dir + '*' + ext):
43         os.remove(f)
44
45 def filename_generator(folder):
46     '''Generator for output filenames'''
47     if not os.path.exists(folder):
48         os.makedirs(folder)
49     count = 0
50     while True:
51         filename = folder + "content" + str(count)

```



```

52         count += 1
53         yield filename
54
55     make_filename = filename_generator(temp_dir)
56
57     #Format {right title : redirected title }, e.g. {because: ([cuz, cus])}
58     redirects = {}
59
60     #Minimum number of links/words required to keep an article.
61     from shared import min_links_out, min_words
62
63     #Open log file for writing and import logging function
64     logfile = open(os.path.basename(__file__) + '.log', 'w')
65     log = shared.logmaker(logfile)
66
67     class WikiHandler(SAX.ContentHandler):
68         '''ContentHandler class to process XML and deal with the WikiText.
69         It works basically like this:
70         It traverses the XML file, keeping track of the type of data being read and
71         adding any text to its input buffer. When event handlers register a page
72         end, the page content is processed, the processed content is placed in the
73         output buffer, and the input buffer is flushed.
74         Whenever a set number of articles have been processed, the output buffer is
75         written to a file. The point of this approach is to
76         limit memory consumption.'''
77
78     def __init__(self):
79         SAX.ContentHandler.__init__(self)
80         self.current_data = None
81         self.title = ''
82         self.input_buffer = []
83         self.output_buffer = {}
84         self.article_counter = 0
85         self.links = []
86         self.categories = []
87         self.redirect = None
88         self.verbose = False
89         #Harvest unique words here
90         self.words = set([])
91         #keeps track of ingoing article links. format {to : set([from])}
92         self.linkhash = {}
93
94     def flush_input_buffer(self):
95         '''Deletes info on the currently processed article.
96         This is called when a page end event is registered.'''
97         self.input_buffer = []
98         self.current_data = None
99         self.title = ''
100         self.links = []
101         self.categories = []
102         self.redirect = None
103
104     def flush_output_buffer(self):
105         '''Flushes data gathered so far to a file and resets.'''
106         self.output_buffer = {}
107         self.words = set([])
108         self.linkhash = {}
109
110     def startElement(self, tag, attrs):
111         '''EventHandler for element start - keeps track of current datatype.'''
112         self.current_data = tag
113         #Informs the parser of the redirect destination of the article

```

```

114     if tag == "redirect":
115         self.redirect = attrs['title']
116         return None
117
118 def endElement(self, name):
119     '''Eventhandler for element end. This causes the parser to process
120     its input buffer when a pageend is encountered.'''
121     #Process content after each page
122     if name == 'page':
123         self.process()
124     #Write remaining data at EOF.
125     elif name == 'mediawiki':
126         self.writeout()
127
128 def characters(self, content):
129     '''Character event handler. This simply passes any raw text from an
130     article field to the input buffer and updates title info.'''
131     if self.current_data == 'text':
132         self.input_buffer.append(content)
133     elif self.current_data == 'title' and not content.isspace():
134         self.title = content
135
136 def process(self):
137     '''Process input buffer contents. This converts wikilanguage to
138     plaintext, registers link information and checks if content has
139     sufficient words and outgoing links (ingoing links can't be checked
140     until the full XML file is processed).'''
141
142     #Ignore everything else if article redirects
143     if self.redirect:
144         self.flush_input_buffer()
145         return None
146         global redirects
147         try:
148             redirects[self.title].add(self.redirect)
149         except KeyError:
150             redirects[self.title] = set([self.redirect])
151         self.flush_input_buffer()
152         return None
153
154     #Redirects handled - commence processing
155     print "processing: "+self.title.encode('utf8')
156     #Combine buffer content to a single string
157     text = ''.join(self.input_buffer).lower()
158
159     #Find and process link information
160     link_regexp = re.compile(r'\[[([^\]]*)\]')
161     links = re.findall(link_regexp, text) #grap stuff like [[<something>]]
162     #Add links to the parsers link hash
163     for link in links:
164         #Check if link matches a namespace, e.g. ' file :something.png'
165         if any([ns+' ':' in link for ns in wikicleaner.namespaces]):
166             continue #Proceed to next link
167         #Namespaces done, so remove any colons:
168         link = link.replace(':', '')
169         if not link:
170             continue #Some noob could've written an empty link...
171         #remove chapter designations/displaytext - keep article title
172         raw = re.match(r'([^\|\\#]*)', link).group(0)
173         title = canonize_title(raw)
174         #note down that current article has outgoing link to 'title'
175         self.links.append(title)

```

```

176         #also note that 'title' has incoming link from here
177         try:
178             self.linkhash[title].add(self.title) #maps target->sources
179         except KeyError:
180             self.linkhash[title] = set([self.title])
181
182     #Disregard current article if it contains too few links
183     if len(self.links) < min_links_out:
184         self.flush_input_buffer()
185         return None
186
187     #Cleanup text
188     text = wikicleaner.clean(text)
189     article_words = text.split()
190
191     #Disregard article if it contains too few words
192     if len(article_words) < min_words:
193         self.flush_input_buffer()
194         return None
195
196     #Update global list of unique words
197     self.words.update(set(article_words))
198
199     #Add content to output buffer
200     output = {
201         'text': text,
202         #Don't use category info for now
203         #'categories': self.categories,
204         'links_out': self.links
205     }
206     self.output_buffer[self.title] = output
207     self.article_counter += 1
208
209     #Flush output buffer to file
210     if self.article_counter%1000 == 0:
211         self.writeout()
212
213     #Done, flushing buffer
214     self.flush_input_buffer()
215     return None
216
217 def writeout(self):
218     '''Writes output buffer contents to file'''
219     #Generate filename and write to file
220     filename = make_filename.next()
221     #Write article contents to file
222     with open(filename+extensions['content'], 'w') as f:
223         shared.dump(self.output_buffer, f)
224
225     #Store wordlist as files
226     with open(filename+extensions['words'], 'w') as f:
227         shared.dump(self.words, f)
228
229     #Store linkhash in files
230     with open(filename+extensions['links'], 'w') as f:
231         shared.dump(self.linkhash, f)
232
233     if self.verbose:
234         log("wrote "+filename)
235
236     #Empty output buffer
237     self.flush_output_buffer()

```

```
238         return None
239
240     if __name__ == "__main__":
241         if len(sys.argv) == 2:
242             file_to_parse = sys.argv[1]
243         else:
244             file_to_parse = DEFAULT_FILENAME
245
246         #Create and configure content handler
247         test = WikiHandler()
248         test.verbose = True
249
250         #Create a parser and set handler
251         ATST = SAX.make_parser()
252         ATST.setContentHandler(test)
253
254         #Let the parser walk the file
255         log("Parsing started...")
256         ATST.parse(file_to_parse)
257         log("...Parsing done!")
258
259         #Attempt to send notification that job is done
260         if shared.notify:
261             try:
262                 shared.pushme(sys.argv[0]+' completed.')
263             except:
264                 log("Job's done. Push failed.")
265
266     logfile.close()
```

A.1.2 Index Generator

```

1  # -*- coding: utf-8 -*-
2  '''This finishes preprocessing of the output from the XML parser.
3  This script reads in link data and removes from the content files those
4  concepts that have too few incoming links. Information on incoming links
5  is saved to each content file.
6  Finally, index maps for words and approved concepts are generated and saved.'''
7
8  from __future__ import division
9  import glob
10 import gc
11 import shared
12 import os
13 import sys
14
15 logfile = open(os.path.basename(__file__)+'.log', 'w')
16 log = shared.logmaker(logfile)
17
18 #Import shared parameters
19 from shared import extensions, temp_dir, min_links_in, matrix_dir
20
21 def listchopper(l):
22     '''Generator to chop lists into chunks of a predefined length'''
23     n = shared.link_chunk_size
24     ind = 0
25     while ind < len(l):
26         yield l[ind:ind+n]
27         ind += n
28
29 def main():
30     #Import shared parameters and verify output dir exists
31     if not os.path.exists(temp_dir):
32         raise IOError
33
34     #=====
35     #   Read in link data and update content files accordingly
36     #=====
37
38     #Get list of files containing link info and chop it up
39     linkfiles = glob.glob(temp_dir + '*' + extensions['links'])
40     linkchunks = listchopper(linkfiles)
41
42     linkfiles_read = 0
43     for linkchunk in linkchunks:
44         #Hash mapping each article to a set of articles linking to it
45         linkhash = {}
46
47         for filename in linkchunk:
48             with open(filename, 'r') as f:
49                 newstuff = shared.load(f)
50                 #Add link info to linkhash
51                 for target, sources in newstuff.iteritems():
52                     try:
53                         linkhash[target].update(set(sources))
54                     except KeyError:
55                         linkhash[target] = set(sources)
56
57     #Log status

```

```

58     linkfiles_read += 1
59     log("Read " + filename + " - " +
60         str(100*linkfiles_read/len(linkfiles))[:4] + " % of link data.")
61
62     log("Chunk finished - updating content files")
63     #Update concept with newly read link data
64     contentfiles = glob.glob(temp_dir + '*' + extensions['content'])
65     contentfiles_read = 0
66     for filename in contentfiles:
67         #Read file. Content is like {'article title': {'text': blah}}
68         with open(filename, 'r') as f:
69             content = shared.load(f)
70
71         #Search linkhash for links going TO concept
72         for concept in content.keys():
73             try:
74                 sources = linkhash[concept]
75             except KeyError:
76                 sources = set([]) #Missing key => zero incoming links
77
78             #Update link info for concept
79             try:
80                 content[concept]['links_in'] = set(content[concept]['links_in'])
81                 content[concept]['links_in'].update(sources)
82             except KeyError:
83                 content[concept]['links_in'] = sources
84
85             #Save updated content
86             with open(filename, 'w') as f:
87                 shared.dump(content, f)
88
89             contentfiles_read += 1
90             if contentfiles_read % 100 == 0:
91                 log("Fixed " + str(100*contentfiles_read/len(contentfiles))[:4]
92                     + "% of content files")
93         pass #Proceed to next link chunk
94
95     #=====
96     # Finished link processing
97     # Remove unworthy concepts and combine concept/word lists.
98     #=====
99
100    #What, you think memory grows on trees?
101    del linkhash
102    gc.collect()
103
104    #Set of all approved concepts
105    concept_list = set([])
106
107    #Purge inferior concepts (with insufficient incoming links)
108    for filename in contentfiles:
109        #Read in content file
110        with open(filename, 'r') as f:
111            content = shared.load(f)
112
113        for concept in content.keys():
114            entry = content[concept]
115            if 'links_in' in entry and len(entry['links_in']) >= min_links_in:
116                concept_list.add(concept)
117            else:
118                del content[concept]
119

```

```
120         with open(filename, 'w') as f:
121             shared.dump(content, f)
122
123     log("Links done - saving index files")
124
125     #Make sure output dir exists
126     if not os.path.exists(matrix_dir):
127         os.makedirs(matrix_dir)
128
129     #Generate and save a concept index map. Structure: {concept : index}
130     concept_indices = {n: m for m, n in enumerate(concept_list)}
131     with open(matrix_dir+'concept2index.ind', 'w') as f:
132         shared.dump(concept_indices, f)
133
134     #Read in all wordlists and combine them.
135     words = set([])
136     for filename in glob.glob(temp_dir + '*' + extensions['words']):
137         with open(filename, 'r') as f:
138             words.update(shared.load(f))
139
140     #Generate and save a word index map. Structure: {word : index}
141     word_indices = {n: m for m, n in enumerate(words)}
142     with open(matrix_dir+'word2index.ind', 'w') as f:
143         shared.dump(word_indices, f)
144
145     log("Wrapping up.")
146     #Attempt to notify that job is done
147     if shared.notify:
148         try:
149             shared.pushme(sys.argv[0]+' completed.')
150         except:
151             log("Job's done. Push failed.")
152
153     logfile.close()
154
155 if __name__ == '__main__':
156     main()
```

A.1.3 Matrix Builder

```

1  # -*- coding: utf-8 -*-
2  '''Builds a huge sparse matrix of Term frequency/Inverse Document Frequency
3  (TFIDF) of the previously extracted words and concepts.
4  First a matrix containing simply the number of occurrences of word i in the
5  article corresponding to concept j is build (in DOK format as that is faster
6  for iterative construction), then the matrix is converted to sparse row format
7  (CSR), TFIDF values are computed, each row is normalized and finally pruned.'''
8
9  from __future__ import division
10 import scipy.sparse as sps
11 import numpy as np
12 from collections import Counter
13 import glob
14 import shared
15 import sys
16 import os
17
18 def percentof(small, large):
19     return str(100*small/large) + "%"
20
21 logfile = open(os.path.basename(__file__)+'.log', 'w')
22 log = shared.logmaker(logfile)
23
24 #import shared parameters
25 from shared import (extensions, matrix_dir, prune, temp_dir, column_chunk_size,
26                     row_chunk_size, datatype)
27
28 def main():
29     #Cleanup
30     for f in glob.glob(matrix_dir + '/*'+extensions['matrix']):
31         os.remove(f)
32
33     #Set pruning parameters
34     window_size = shared.window_size
35     cutoff = shared.cutoff
36
37     #Read in dicts mapping words and concepts to their respective indices
38     log("Reading in word/index data")
39     word2index = shared.load(open(matrix_dir+'word2index.ind', 'r'))
40     concept2index = shared.load(open(matrix_dir+'concept2index.ind', 'r'))
41     log("...Done!")
42
43     #=====
44     #     Construct count matrix in small chunks
45     #=====
46
47     #Count words and concepts
48     n_words = len(word2index)
49     n_concepts = len(concept2index)
50
51     #Determine matrix dimensions
52     matrix_shape = (n_words, n_concepts)
53
54     #Allocate sparse matrix. Dict-of-keys should be faster for iterative
55     #construction. Convert to csr for fast row operations later.
56     mtx = sps.dok_matrix(matrix_shape, dtype = datatype)
57

```



```

58 def matrix_chopper(matrix, dim):
59     '''Generator to split a huge matrix into small submatrices, which can
60     then be stored in individual files.
61     This is handy both when constructing the matrix (building the whole
62     matrix without saving to files in the process takes about 50 gigs RAM),
63     and when applying it, as this allows one to load only the submatrix
64     relevant to a given word.'''
65     ind = 0
66     counter = 0
67     rows = matrix.get_shape()[0]
68     while ind < rows:
69         end = min(ind+dim, rows)
70         #Return pair of submatrix number and the submatrix itself
71         yield counter, sps.vstack([matrix.getrow(i)\
72                                 for i in xrange(ind, end)], format = 'csr')
73         counter += 1
74         ind += dim
75
76 def writeout():
77     '''Saves the matrix as small submatrices in separate files.'''
78     for n, submatrix in matrix_chopper(mtx, row_chunk_size):
79         filename = matrix_dir+str(n)+extensions['matrix']
80         #Update submatrix if it's already partially calculated
81         log("Writing out chunk %s" % n)
82         try:
83             with open(filename, 'r') as f:
84                 submatrix = submatrix + shared.mload(f)
85             #
86         except IOError:
87             pass #File doesn't exist yet, so no need to change mtx
88
89         #Dump the submatrix to file
90         with open(filename, 'w') as f:
91             shared.mdump(submatrix, f)
92     return None
93
94 log("Constructing matrix.")
95 filelist = glob.glob(temp_dir + '*' + extensions['content'])
96 files_read = 0
97 for filename in filelist:
98     with open(filename, 'r') as f:
99         content = shared.load(f)
100
101     #Loop over concepts (columns) as so we don't waste time with rare words
102     for concept, entry, in content.iteritems():
103         #This is the column index (concept w. index j)
104         j = concept2index[concept]
105
106         #Convert concept 'countmap' like so: {word : n}
107         wordmap = Counter(entry['text'].split()).iteritems()
108
109         #Add them all to the matrix
110         for word, count in wordmap:
111             #Find row index of the current word
112             i = word2index[word]
113
114             #Add the number of times word i occurs in concept j to the matrix
115             mtx[i,j] = count
116         #
117     #Update file count
118     files_read += 1
119     log("Processed content file no. %s of %s - %s"

```

```

120         %(files_read, len(filelist)-1, percentof(files_read, len(filelist))))
121
122     if files_read % column_chunk_size == 0:
123         mtx = mtx.tocsr()
124         writeout()
125         mtx = sps.dok_matrix(matrix_shape)
126     #
127
128     #Convert matrix to CSR format and write to files.
129     mtx = mtx.tocsr()
130     writeout()
131
132     #=====
133     # Count matrix/matrices constructed – computing TF-IDF
134     #=====
135
136     log("Done - computing TF-IDF")
137
138     #Grap list of matrix files (containing the submatrices from before)
139     matrixfiles = glob.glob(matrix_dir + "*" + extensions['matrix'])
140     words_processed = 0 #for logging purposes
141
142     for filename in matrixfiles:
143         with open(filename, 'r') as f:
144             mtx = shared.mload(f)
145
146         #Number of words in a submatrix
147         n_rows = mtx.get_shape()[0]
148
149         for w in xrange(n_rows):
150             #Grap non-zero elements from the row corresponding to word w
151             row = mtx.data[mtx.indptr[w]:mtx.indptr[w+1]]
152             if len(row) == 0:
153                 continue
154
155             #Make a vectorized function to convert a full row to TF-IDF
156             f = np.vectorize(lambda m_ij: (1+np.log(m_ij))*
157                             np.log(n_concepts/len(row)))
158
159             #Map all elements to TF-IDF and update matrix
160             row = f(row)
161
162             #Normalize the row
163             assert row.dtype.kind == 'f' #Non floats round to zero w/o warning
164             normfact = 1.0/np.linalg.norm(row)
165             row *= normfact
166
167             #Start inverted index pruning
168             if prune:
169                 #Number of documents containing w
170                 n_docs = len(row)
171
172                 #Don't prune if the windows exceeds the array bounds (duh)
173                 if window_size < n_docs:
174
175                     #Obtain list of indices such that row[index] is sorted
176                     indices = np.argsort(row)[::-1]
177
178                     #Generate a sorted row
179                     sorted_row = [row[index] for index in indices]
180
181                     #Go through sorted row and truncate when pruning condition is met

```

```

182         for i in xrange(n_docs-window_size):
183             if sorted_row[i+window_size] >= cutoff*sorted_row[i]:
184                 #Truncate, i.e. set the remaining entries to zero
185                 sorted_row[i:] = [0]*(n_docs-i)
186                 break
187             else:
188                 pass
189
190         #Unsort to original positions
191         for i in xrange(n_docs):
192             row[indices[i]] = sorted_row[i]
193
194         #Update matrix
195         mtx.data[mtx.indptr[w] : mtx.indptr[w+1]] = row
196
197         #Log it
198         words_processed += 1
199         if words_processed % 10**3 == 0:
200             log("Processing word %s of %s - %s" %
201                 (words_processed, n_words,
202                  percentof(words_processed, n_words)))
203
204         #Keep it sparse – no need to store zeroes
205         mtx.eliminate_zeros()
206         with open(filename, 'w') as f:
207             shared.mdump(mtx, f)
208
209         log("Done!")
210
211         #Notify that the job is done
212         if shared.notify:
213             try:
214                 shared.pushme(sys.argv[0]+' completed.')
215             except:
216                 log("Job's done. Push failed.")
217
218         logfile.close()
219         return None
220
221 if __name__ == '__main__':
222     main()

```

A.1.4 Library for Computational Linguistics

```

1  # -*- coding: utf-8 -*-
2  '''Small module for computational linguistics applied to Twitter.
3  The main classes are a TweetHarvester, which gathers data from Twitters' API,
4  and a SemanticAnalyser, which relies on the previously constructed TFIDF
5  matrices.'''
6
7  from __future__ import division
8  from scipy import sparse as sps
9  from collections import Counter
10 from numpy.linalg import norm
11 import re
12 import shared
13 import tweepy
14 from datetime import date
15 import json
16 import time
17 import sys
18 import codecs
19 from pprint import pprint
20 sys.stdout = codecs.getwriter('utf8')(sys.stdout)
21 sys.stderr = codecs.getwriter('utf8')(sys.stderr)
22
23 #=====
24 # This stuff defines a twitter 'harvester' for downloading Tweets
25 #=====
26
27 #Import credentials for accessing Twitter API
28 from supersecretstuff import consumer_key, consumer_secret, access_token, access_token_secret
29 auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
30 auth.set_access_token(access_token, access_token_secret)
31
32 class listener(tweepy.StreamListener):
33     '''Listener class to access Twitter stream.'''
34     #What to do with a tweet (override later)
35     def process(self, content):
36         print content
37         return None
38
39     def on_status(self, status):
40         self.process(status)
41         return True
42
43     def on_error(self, status):
44         print status
45
46 # Exception to be raised when the Twitter API messes up. Happens occasionally.
47 class IncompleteRead(Exception):
48     pass
49
50 class TweetHarvester(object):
51     '''Simple class to handle tweet harvest.
52     Harvest can be performed actively or passively, i.e. using the 'mine'
53     method to gather a fixed number of tweets or using the 'listen' method
54     to stream tweets matching a given search term.
55     Harvested tweets are sent to the process method which by default simply
56     stores them inside the object.'''
57

```

```

58 def __init__(self, max_tweets=-1, verbose = False, tweets_pr_file = 10*5):
59     #Set parameters
60     self.max_tweets = max_tweets #-1 for unlimited stream
61     self.verbose = verbose
62     self.tweets_pr_file = tweets_pr_file
63
64     #Internal parameters to keep track of harvest status
65     self.files_saved = 0
66     self.harvested_tweets = []
67     self.current_filename = 0
68     self.current_date = date.today()
69
70 def filename_maker(self):
71     #Update counter and date if neccessary
72     if not self.current_date == date.today():
73         self.current_date = date.today()
74         self.current_filename = 0
75     else:
76         pass #Date hasn't changed. Proceed.
77     filename = str(self.current_date) + "-data%s.json" % self.current_filename
78     self.current_filename += 1
79     return filename
80
81 #Simple logging function
82 def log(self, text):
83     string = text+" at "+time.asctime()+"\n"
84     if self.verbose:
85         print string
86     with open('tweetlog.log', 'a') as logfile:
87         logfile.write(string)
88     #Must return true so I can log errors without breaking the stream.
89     return True
90
91 def listen(self, search_term):
92     #Make a listener
93     listener = tweepy.StreamListener()
94     #Override relevant methods.
95     listener.on_status = self.process
96     listener.on_error = lambda status_code: self.log("Error: "+status_code)
97     listener.on_timeout = lambda: self.log("Timeout.")
98
99     twitterStream = tweepy.Stream(auth, listener)
100     twitterStream.filter(track=search_term)
101
102 def mine(self, search_term, n = None):
103     '''Mine a predefined number of tweets using input search word'''
104     if n == None:
105         n = self.max_tweets
106
107     api = tweepy.API(auth)
108     tweets = tweepy.Cursor(api.search, q=search_term).items(n)
109     for tweet in tweets:
110         self.process(tweet)
111
112 def process(self, tweet):
113     self.harvested_tweets.append(tweet)
114     if self.verbose:
115         print "Holding %s tweets." % len(self.harvested_tweets)
116
117     #Write to file if buffer is full
118     if len(self.harvested_tweets) == self.tweets_pr_file:
119         self.writeout()

```

```

120
121     #Check if limit has been reached (returning false cuts off listener)
122     return not (len(self.harvested_tweets) == self.max_tweets)
123
124     def writeout(self):
125         filename = self.filename_maker()
126         with open(filename, 'w') as outfile:
127             outfile.writelines([json.dumps(t._json) + "\n"
128                                for t in self.harvested_tweets])
129
130         self.harvested_tweets = []
131         self.files_saved += 1
132         #Log event
133         s = "Saved %s files" % self.files_saved
134         self.log(s)
135
136
137     #=====
138     # Defines stuff to analyse text using an already constructed interpretation
139     # matrix.
140     #=====
141
142     from shared import matrix_dir, row_chunk_size, extensions
143
144     class SemanticAnalyser(object):
145         '''Analyser class using Explicit Semantic Analysis (ESA) to process
146         text fragments. It can compute semantic (pseudo) distance and similarity,
147         as well'''
148         def __init__(self, matrix_filename = 'matrix.mtx', display_concepts = 20):
149             #Number of top concepts to display
150             self.display_concepts = display_concepts
151
152             #Hashes for word and concept indices
153             with open(matrix_dir + 'word2index.ind', 'r') as f:
154                 self.word2index = shared.load(f)
155             with open(matrix_dir + 'concept2index.ind', 'r') as f:
156                 self.concept2index = shared.load(f)
157             self.index2concept = {i : c for c, i in self.concept2index.iteritems()}
158
159             #Count number of words and concepts
160             self.n_words = len(self.word2index)
161             self.n_concepts = len(self.concept2index)
162
163         def clean(self, text):
164             text = re.sub('[^w\s\d\`-]', '', text)
165             text = text.lower()
166
167             return text
168
169         def interpretation_vector(self, text):
170             '''Converts a text fragment string into a row vector where the i'th
171             entry corresponds to the total TF-IDF score of the text fragment
172             for concept i'''
173
174             #Remove mess (quotes, parentheses etc) from text
175             text = self.clean(text)
176
177             #Convert string to hash like {'word': no. of occurrences}
178             countmap = Counter(text.split()).iteritems()
179
180             #Interpretation vector to be returned
181             result = sps.csr_matrix((1, self.n_concepts), dtype = float)

```

```

182
183 #Add word count in the correct position of the vector
184 for word, count in countmap:
185     try:
186         ind = self.word2index[word]
187         #Which file to look in
188         file_number = int(ind/row_chunk_size)
189         filename = matrix_dir+str(file_number)+extensions['matrix']
190
191         #And which row to extract
192         row_number = ind % row_chunk_size
193
194         #Do it! Do it naw!
195         with open(filename, 'r') as f:
196             temp = shared.mload(f)
197             result = result + count*temp.getrow(row_number)
198     except KeyError:
199         pass #No data on this word -> discard
200
201 #Done. Return row vector as a 1x#concepts CSR matrix
202 return result
203
204 def interpret_text(self, text):
205     '''Attempts to guess the core concepts of the given text fragment'''
206     #Compute the interpretation vector for the text fragment
207     vec = self.interpretation_vector(text)
208
209     #Magic, don't touch
210     top_n = vec.data.argsort()[len(vec.data)-1-self.display_concepts:-1]
211
212     #List top scoring concepts and their TD-IDF
213     concepts = [self.index2concept[vec.indices[i]] for i in top_n]
214     return concepts
215 #
216 #     scores = [vec.data[i] for i in top_n]
217 #     #Return as dict {concept: score}
218 #     return dict(zip(concepts, scores))
219
220 def interpret_file(self, filename):
221     with open(filename, 'r') as f:
222         data = self.clean(f.read())
223     return self.interpret_text(data)
224
225 def interpret_input(self):
226     text = raw_input("Enter text fragment: ")
227     topics = self.interpret_text(text)
228     print "Based on your input, the most probable topics of your text are:"
229     print topics[:self.display_concepts]
230
231 def compare_texts(self, text1, text2):
232     '''Determines cosine similarity between input texts.
233     Returns float in [0,1]'''
234
235     #Determine interpretation vectors
236     v1 = self.interpretation_vector(text1)
237     v2 = self.interpretation_vector(text2)
238
239     #Compute their inner product and make sure it's a scalar
240     dot = v1.dot(v2.transpose())
241     assert dot.shape == (1,1)
242
243     if dot.data:
244         scal = dot.data[0]

```

```

244     else:
245         scal = 0    #Empty sparse matrix means zero
246
247         #Normalize and return
248         sim = scal/(norm(v1.data)*norm(v2.data))
249         return sim
250
251     def cosine_distance(self, text1, text2):
252         return 1-self.compare_texts(text1, text2)
253
254     if __name__ == '__main__':
255         th = TweetHarvester(verbose=True, max_tweets=10)
256         th.mine('carlsberg', n=10)
257         temp = [t._json for t in th.harvested_tweets if t._json['lang'] == 'en']
258         js = temp[4]
259         with open('tweet_example.json', 'w') as f:
260             pprint(js, stream=f)
261
262     # if len(sys.argv) > 1:
263     #     fn = sys.argv[1]
264     # else:
265     #     fn = 'interpret_me.txt'
266     #     with open(fn, 'r') as f:
267     #         data = f.read()
268     #     #
269     #     data = sa.clean(data)
270     #     guesses = sa.interpret_text(data)
271     #
272     # if len(sys.argv) > 2:
273     #     output_filename = sys.argv[2]
274     # else:
275     #     output_filename = 'guesses.txt'
276     #     with open(output_filename, 'w') as f:
277     #         for line in guesses:
278     #             f.write(line.encode('utf8'))
279     #         f.write('\n')

```

A.1.5 Wikicleaner

```

1  # -*- coding: utf-8 -*-
2  import re
3  from htmlentitydefs import name2codepoint
4
5  namespaces = set(['help', 'file talk', 'module', 'topic', 'mediawiki',
6  'wikipedia talk', 'file', 'user talk', 'special', 'category talk', 'category',
7  'media', 'wikipedia', 'book', 'draft', 'book talk', 'template', 'help talk',
8  'timedtext', 'mediawiki talk', 'portal talk', 'portal', 'user', 'module talk',
9  'template talk', 'education program talk', 'education program',
10 'timedtext talk', 'draft talk', 'talk'])
11
12 def dropNested(text, openDelim, closeDelim):
13     '''Helper function to match nested expressions which may cause problems
14     example: {{something something {{something else}} and something third}}
15     cannot be easily matched with a regexp to remove all occurrences.
16     Copied from the WikiExtractor project.'''
17     openRE = re.compile(openDelim)
18     closeRE = re.compile(closeDelim)

```



```

19 # partition text in separate blocks { } { }
20 matches = [] # pairs (s, e) for each partition
21 nest = 0 # nesting level
22 start = openRE.search(text, 0)
23 if not start:
24     return text
25 end = closeRE.search(text, start.end())
26 next = start
27 while end:
28     next = openRE.search(text, next.end())
29     if not next: # termination
30         while nest: # close all pending
31             nest -= 1
32             end0 = closeRE.search(text, end.end())
33             if end0:
34                 end = end0
35             else:
36                 break
37         matches.append((start.start(), end.end()))
38         break
39     while end.end() < next.start():
40         # { } {
41         if nest:
42             nest -= 1
43             # try closing more
44             last = end.end()
45             end = closeRE.search(text, end.end())
46             if not end: # unbalanced
47                 if matches:
48                     span = (matches[0][0], last)
49                 else:
50                     span = (start.start(), last)
51                 matches = [span]
52                 break
53             else:
54                 matches.append((start.start(), end.end()))
55                 # advance start, find next close
56                 start = next
57                 end = closeRE.search(text, next.end())
58                 break # { }
59         if next != start:
60             # { { }
61             nest += 1
62 # collect text outside partitions
63 res = ''
64 start = 0
65 for s, e in matches:
66     res += text[start:s]
67     start = e
68 res += text[start:]
69 return res
70
71 def unescape(text):
72     '''Removes HTML or XML character references and entities
73     from a text string.
74     @return nice text'''
75 def fixup(m):
76     text = m.group(0)
77     code = m.group(1)
78     return text
79     try:
80         if text[1] == "#": # character reference

```

```

81         if text[2] == "x":
82             return unichr(int(code[1:], 16))
83         else:
84             return unichr(int(code))
85     else: # named entity
86         return unichr(name2codepoint[code])
87 except UnicodeDecodeError:
88     return text # leave as is
89
90 return re.sub("&#?(\w+);", fixup, text)
91
92 def drop_spans(matches, text):
93     """Drop from text the blocks identified in matches"""
94     matches.sort()
95     res = ''
96     start = 0
97     for s, e in matches:
98         res += text[start:s]
99         start = e
100     res += text[start:]
101     return res
102
103 ###Compile regexps for text cleanup:
104 #Construct patterns for elements to be discarded:
105 discard_elements = set([
106     'gallery', 'timeline', 'noinclude', 'pre',
107     'table', 'tr', 'td', 'th', 'caption',
108     'form', 'input', 'select', 'option', 'textarea',
109     'ul', 'li', 'ol', 'dl', 'dt', 'dd', 'menu', 'dir',
110     'ref', 'references', 'img', 'imagemap', 'source'
111 ])
112 discard_element_patterns = []
113 for tag in discard_elements:
114     pattern = re.compile(r'<\s*%s\b[^>]*>.*?<\s*/\s*%s>' % (tag, tag), re.DOTALL | re.IGNORECASE)
115     discard_element_patterns.append(pattern)
116
117 #Construct patterns to recognize HTML tags
118 selfclosing_tags = set(['br', 'hr', 'noabr', 'ref', 'references'])
119 selfclosing_tag_patterns = []
120 for tag in selfclosing_tags:
121     pattern = re.compile(r'<\s*%s\b[/]?/*\s*>' % tag, re.DOTALL | re.IGNORECASE)
122     selfclosing_tag_patterns.append(pattern)
123
124 #Construct patterns for tags to be ignored
125 ignored_tags = set([
126     'a', 'b', 'big', 'blockquote', 'center', 'cite', 'div', 'em',
127     'font', 'h1', 'h2', 'h3', 'h4', 'hiero', 'i', 'kbd', 'nowiki',
128     'p', 'plaintext', 's', 'small', 'span', 'strike', 'strong',
129     'sub', 'sup', 'tt', 'u', 'var',
130 ])
131 ignored_tag_patterns = []
132 for tag in ignored_tags:
133     left = re.compile(r'<\s*%s\b[^>]*>' % tag, re.IGNORECASE)
134     right = re.compile(r'<\s*/\s*%s>' % tag, re.IGNORECASE)
135     ignored_tag_patterns.append((left, right))
136
137 #Construct patterns to recognize math and code
138 placeholder_tags = {'math': 'formula', 'code': 'codice'}
139 placeholder_tag_patterns = []
140 for tag, repl in placeholder_tags.items():
141     pattern = re.compile(r'<\s*%s(\s*| [^>]+?)>.*?<\s*/\s*%s>' % (tag, tag), re.DOTALL | re.IGNORECASE)
142     placeholder_tag_patterns.append((pattern, repl))

```

```

143
144 #HTML comments
145 comment = re.compile(r'<!--.*?-->', re.DOTALL)
146
147 #Wikilinks
148 wiki_link = re.compile(r'\[[^\[\]\?]*(?:\|([^\[\]\?])?\]\](\w*)')
149 parametrized_link = re.compile(r'\[[^\[\]\?]*\]\]')
150
151 #External links
152 externalLink = re.compile(r'\[w+.*? (.*)\]')
153 externalLinkNoAnchor = re.compile(r'\[w+[\&]]*\]')
154
155 #Bold/italic text
156 bold_italic = re.compile(r''''''([^\[\]\?])''''')
157 bold = re.compile(r''''(.*)''''')
158 italic_quote = re.compile(r''''\''''')
159 italic = re.compile(r''''([^\[\]\?])''''')
160 quote_quote = re.compile(r''''(.*)''''')
161
162 #Spaces
163 spaces = re.compile(r' {2,}')
164
165 #Dots
166 dots = re.compile(r'\.{4,}')
167
168 #Sections
169 section = re.compile(r'^(=+)\s*(.*)\s*\1')
170
171 # Match preformatted lines
172 preformatted = re.compile(r'^ .*$', re.MULTILINE)
173
174 #Wikilinks
175 def make_anchor_tag(match):
176     '''Recognizes links and returns only their anchor. Example:
177     <a href="www.something.org">Link text</a> -> Link text'''
178     link = match.group(1)
179     colon = link.find(':')
180     if colon > 0 and link[colon] not in namespaces:
181         return ''
182     trail = match.group(3)
183     anchor = match.group(2)
184     if not anchor:
185         if link[colon] in namespaces:
186             return '' #Don't keep stuff like "category: shellfish "
187         anchor = link
188     anchor += trail
189     return anchor
190
191 def clean(text):
192     '''Outputs an article in plaintext from its format in the raw xml dump.'''
193     # Drop transclusions (template, parser functions)
194     # See: http://www.mediawiki.org/wiki/Help:Templates
195     text = dropNested(text, r'{{', r'}}')
196     # Drop tables
197     text = dropNested(text, r'{\|', r'\|}')
198
199     # Convert wikilinks links to plaintext
200     text = wiki_link.sub(make_anchor_tag, text)
201     # Drop remaining links
202     text = parametrized_link.sub('', text)
203
204     # Handle external links

```

```

205 text = externalLink.sub(r'\1', text)
206 text = externalLinkNoAnchor.sub('', text)
207
208 #Handle text formatting
209 text = bold_italic.sub(r'\1', text)
210 text = bold.sub(r'\1', text)
211 text = italic_quote.sub(r'&quot;\1&quot;', text)
212 text = italic.sub(r'&quot;\1&quot;', text)
213 text = quote_quote.sub(r'\1', text)
214 text = text.replace("'", '').replace('"', '&quot;')
215
216 ##### Process HTML #####
217
218 # turn into HTML
219 text = unescape(text)
220
221 # do it again (&nbsp;)
222 text = unescape(text)
223
224 # Collect spans
225
226 matches = []
227 # Drop HTML comments
228 for m in comment.finditer(text):
229     matches.append((m.start(), m.end()))
230
231 # Drop self-closing tags
232 for pattern in selfclosing_tag_patterns:
233     for m in pattern.finditer(text):
234         matches.append((m.start(), m.end()))
235
236 # Drop ignored tags
237 for left, right in ignored_tag_patterns:
238     for m in left.finditer(text):
239         matches.append((m.start(), m.end()))
240     for m in right.finditer(text):
241         matches.append((m.start(), m.end()))
242
243 # Bulk remove all spans
244 text = drop_spans(matches, text)
245
246 # Cannot use dropSpan on these since they may be nested
247 # Drop discarded elements
248 for pattern in discard_element_patterns:
249     text = pattern.sub('', text)
250
251 # Expand placeholders
252 for pattern, placeholder in placeholder_tag_patterns:
253     index = 1
254     for match in pattern.finditer(text):
255         text = text.replace(match.group(), '%s_%d' % (placeholder, index))
256         index += 1
257
258 #####
259
260 # Drop preformatted
261 # This can't be done before since it may remove tags
262 text = preformatted.sub('', text)
263
264 # Cleanup text
265 text = text.replace('\t', ' ')
266 text = spaces.sub(' ', text)

```

```

267 text = dots.sub('...', text)
268 text = re.sub(u' (,:\.\)\[\]\>)', r'\1', text)
269 text = re.sub(u' \([\(\<)\> ', r'\1', text)
270 text = re.sub(r'\n\W+?\n', '\n', text) # lines with only punctuations
271 text = text.replace(',', ',').replace('.', '. ')
272
273 #Handle section headers, residua etc.
274 page = []
275 headers = {}
276 empty_section = False
277
278 for line in text.split('\n'):
279
280     if not line:
281         continue
282     # Handle section titles
283     m = section.match(line)
284     if m:
285         title = m.group(2)
286         lev = len(m.group(1))
287         if title and title[-1] not in '!?':
288             title += '.'
289         headers[lev] = title
290         # drop previous headers
291         for i in headers.keys():
292             if i > lev:
293                 del headers[i]
294         empty_section = True
295         continue
296     # Handle page title
297     if line.startswith('++'):
298         title = line[2:-2]
299         if title:
300             if title[-1] not in '!?':
301                 title += '.'
302         page.append(title)
303     # handle lists
304     elif line[0] in '*#;,:':
305         continue
306     # Drop residuals of lists
307     elif line[0] in '{| ' or line[-1] in '}':
308         continue
309     # Drop irrelevant lines
310     elif (line[0] == '(' and line[-1] == ')') or line.strip('.-') == '':
311         continue
312     elif len(headers):
313         items = headers.items()
314         items.sort()
315         for (i, v) in items:
316             page.append(v)
317             headers.clear()
318         page.append(line) # first line
319         empty_section = False
320     elif not empty_section:
321         page.append(line)
322
323 text = ''.join(page)
324
325 #Remove quote tags.
326 text = text.replace(""", '')
327
328 #Get rid of parentheses, punctuation and the like

```

```
329 | text = re.sub('[^\w\s\d\'\-]', '', text)
330 | return text
```

A.2 Social Fabric-related Code

This section contains the code referred to in part II of the thesis.

A.2.1 Phonetools

```

1  # -*- coding: utf-8 -*-
2  """
3  Created on Tue Dec 02 15:12:29 2014
4
5  @author: Bjarke
6  """
7
8  import datetime
9  from pkg_resources import resource_filename
10 from ast import literal_eval as LE
11 import numpy as np
12 import json
13 import os
14
15 def _global_path(path):
16     '''Helper method to ensure that data files belonging to the social_fabric
17     module are available both when importing the module and when running
18     individual parts from it for testing.
19     usage: Always use _global_path(somefile) rather than somefile'''
20     if __name__ == '__main__':
21         return path
22     else:
23         return resource_filename('social_fabric', path)
24
25 def make_filename(prefix, interval, bin_size, ext=''):
26     '''Returns a filename like "call_out-intervalsize-binsize_N", where N is
27     an int so files aren't accidentally overwritten.
28     I've been known to do that.'''
29     stem = "%s-int%s-bin%s" % (prefix, interval, bin_size)
30     n = 0
31     ext = '.' + ext
32     attempt = stem + ext
33     while os.path.isfile(attempt):
34         n += 1
35         attempt = stem + "_" + str(n) + ext
36     return attempt
37
38 def unix2str(unixtime):
39     '''Converts timestamp to datetime object'''
40     dt = datetime.datetime.fromtimestamp(unixtime)
41     return str(dt)
42
43 #Converts bluetooth MAC-addresses to users
44 with open(_global_path('user_mappings/bt_to_user.txt'), 'r') as f:
45     bt2user = LE(f.read())
46
47 #Converts phone 'number' code to users
48 with open(_global_path('user_mappings/phonenumbers.txt'), 'r') as f:
49     number2user = LE(f.read())
50
51 #Converts IDs from psychological profile to users

```

```

52 | with open(_global_path('user_mappings/user_mapping.txt'), 'r') as f:
53 |     psych2user = {}
54 |     #This is in tab-separated values, for some reason.
55 |     for line in f.read().splitlines():
56 |         (ID, usercode) = line.split('\t')
57 |         assert ID.startswith('user_')
58 |         psych2user[ID] = usercode
59 |
60 | #Converts users to info on their psych profiles
61 | with open(_global_path('user_mappings/user2profile.json'), 'r') as f:
62 |     user2profile = json.load(f)
63 |
64 | def is_valid_call(call_dict):
65 |     '''Determine whether an entry is 'valid', i.e. make sure user isn't
66 |     calling/texting themselves, which people apparently do...'''
67 |     caller = call_dict['user']
68 |     try:
69 |         receiver = number2user[call_dict['number']]
70 |     except KeyError:
71 |         receiver = None
72 |     return caller != receiver
73 |
74 | def readncheck(path):
75 |     '''Reads in all valid call info from the file at path'''
76 |     try:
77 |         with open(path, 'r') as f:
78 |             raw = [LE(line) for line in f.readlines()]
79 |     except IOError:
80 |         return [] #no file :(
81 |     #File read. Return proper calls
82 |     return [call for call in raw if is_valid_call(call)]
83 |
84 | class Binarray(list):
85 |     '''Custom array type to automatically bin the time around a set center
86 |     and place elements in each bin.
87 |     The array can be centered using Binarray.center = <some time>.
88 |     After centering, timestamps can be placed in bins around the center with
89 |     Binarray.place_event(<some other time>).'''
90 |
91 |     def __init__(self, interval = 3*60*2, bin_size = 3*60, center = None,
92 |                  initial_values = None):
93 |         '''
94 |         Args:
95 |         -----
96 |         interval : int
97 |             Total number of seconds covered by the Binarray.
98 |
99 |         bin_size : int
100 |             Width of each bin measured in seconds. The total interval must be
101 |             an integer multiplum of the bin size.
102 |
103 |         center : int
104 |             Where to center the Binarray. If the array is centered at time t,
105 |             any event placed in it will placed in a bin depending on how long
106 |             before or after t the event occurred.
107 |
108 |         initial_values : list
109 |             List of values to start the Binarray with. Default is zeroes.'''
110 |
111 |
112 |     #Make sure interval is an integer multiplum of bins
113 |     if not interval % bin_size == 0:

```



```

114         suggest = interval - interval % bin_size
115         error = "Interval isn't an integer multiple of bin size. \
116             Consider changing interval to %s." % suggest
117         raise ValueError(error)
118
119     #Set parameters
120     self.bin_size = bin_size
121     self.interval = interval
122     self.size = 2*int(interval/bin_size)
123     self.centerindex = int(self.size//2)
124     self.center = center
125     #Keep track of how many events missed the bins completely
126     self.misses = 0
127     #Call parent constructor
128     if not initial_values:
129         startlist = [0]*self.size
130     else:
131         if not len(initial_values) == self.size:
132             msg = '''Array of start value must have length %d. Tried to
133                 instantiate with length of %d.''' % (self.size,
134                 len(initial_values))
135             raise ValueError(msg)
136         startlist = initial_values
137     super(Binarray, self).__init__(startlist)
138
139
140     def place_event(self, position):
141         '''Places one count in the appropriate bin if event falls within
142         <interval> of <center>. Returns True on success.'''
143         if self.center == None:
144             raise TypeError('Center must be set!')
145         delta = position - self.center
146         #Check if event is outside current interval
147         if np.abs(delta) >= self.interval:
148             self.misses += 1
149             return False
150         #Woo, we're in the correct interval
151         index = int(delta//self.bin_size) #relative to middle of array
152         self[self.centerindex + index] += 1
153         return True
154
155     def normalized(self):
156         '''Returns a normalized copy of the array's contents.'''
157         events = sum(self) + self.misses
158         #Use numpy vectorized function for increased speed.
159         f = np.vectorize(lambda x: x*1.0/events if events else 0)
160         return f(self)
161
162     def _todict(self):
163         '''Helper method to allow dumping to JSON format.'''
164         attrs = ['misses', 'interval', 'bin_size']
165         d = {att: self.__getattr__(att) for att in attrs}
166         d['values'] = list(self)
167         d['type'] = 'binarray'
168         return d
169
170
171     def _dumphelper(obj):
172         '''Evil recursive helper method to convert various nested objects to
173         a JSON-serializeable format.
174         This should only be called by the dump method!'''
175         if isinstance(obj, Binarray):

```

```

176     d = obj._todict()
177     return _dumpher(d)
178 elif isinstance(obj, tuple):
179     hurrayimhelping = [_dumpher(elem) for elem in obj]
180     return {'type': 'tuple', 'values': hurrayimhelping}
181 elif isinstance(obj, dict):
182     temp = {'type': 'dict'}
183     contents = [{'key': _dumpher(key), 'value': _dumpher(value)}
184                 for key, value in obj.iteritems()]
185     temp['contents'] = contents
186     return temp
187 #Do nothing if obj is an unrecognized type. Let JSON raise errors.
188 else:
189     return obj
190
191 def _hook(obj):
192     '''Evil recursive object hook method to reconstruct various nested
193     objects from a JSON dump.
194     This should only be called by the load method!'''
195     if isinstance(obj, (unicode, str)):
196         try:
197             return _hook(LE(obj))
198         except ValueError: #happens for simple strings that don't need eval
199             return obj
200     elif isinstance(obj, dict):
201         if not 'type' in obj:
202             raise KeyError('Missing type info')
203         if obj['type'] == 'dict':
204             contents = obj['contents']
205             d = {_hook(e['key']): _hook(e['value']) for e in contents}
206             #Make sure we also catch nested expressions
207             if 'type' in d:
208                 return _hook(d)
209             else:
210                 return d
211         elif obj['type'] == 'binarray':
212             instance = Binarray(initial_values = obj['values'],
213                                bin_size = obj['bin_size'],
214                                interval = obj['interval'])
215             instance.misses = obj['misses']
216             for key, val in obj.iteritems():
217                 if key == 'values':
218                     continue
219                 instance.__setattr__(key, val)
220             return instance
221
222         elif obj['type'] == 'tuple':
223             #Hook elements individually, then convert back to tuple
224             restored = [_hook(elem) for elem in obj['values']]
225             return tuple(restored)
226         else:
227             temp = {}
228             for k, v in obj.iteritems():
229                 k = _hook(k)
230                 temp[k] = _hook(v)
231             return temp
232     #
233     #Do nothing if obj is an unrecognized type
234     else:
235         return obj
236
237 def load(file_handle):

```

```
238     '''Reads in json serialized nested combinations of dicts, binarrays
239     and tuples.'''
240     temp = json.load(file_handle, encoding='utf-8')
241     return _hook(temp)
242
243 def dump(obj, file_handle):
244     '''json serializes nested combinations of dicts, binarrays
245     and tuples.'''
246     json.dump(_dumphelper(obj), file_handle, indent=4, encoding='utf-8')
247
248
249 if __name__ == '__main__':
250     from time import time
251     from random import randint
252     #Create Binarray with interval +/- one hour and bin size ten minutes.
253     ba = Binarray(interval = 60*60, bin_size = 10*60)
254     #Center it on the present
255     now = int(time())
256     ba.center = now
257     #Generate some timestamps around the present
258     new_times = [now + randint(-60*60, 60*60) for _ in xrange(100)]
259     for tt in new_times:
260         ba.place_event(tt)
261
262     #Save it
263     with open('filename.sig', 'w') as f:
264         dump(ba, f)
265
266     print ba
```

A.2.2 Code Communication Dynamics

Code for extracting Bluetooth signals

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import os
5  import glob
6  import itertools
7  import random
8  from ast import literal_eval as LE
9  from social_fabric.phonetools import readncheck, Binarray, dump, make_filename
10
11  #=====
12  # Parameters
13  #=====
14
15  #Grap all the files we need to read
16  userfile_path = "userfiles/" #linux
17
18
19  #Interval of interest and bin size (seconds)
20  interval = 12*60**2
21  bin_size = 10*60
22
23  #How many other users contribute to background for a given user
24  number_of_background_samples = 1
25
26  #Number of users to analyse. Use None or 0 to include everyone.
27  max_users = 2
28
29  #Number of repeated signals required to be considered social
30  social_threshold = 2
31
32  #=====
33  # Functions
34  #=====
35
36  def call_type(n):
37      if n == 1:
38          return 'call_in'
39      elif n == 2:
40          return 'call_out'
41      else:
42          return None
43
44  def text_type(n):
45      if n == 1:
46          return 'text_in'
47      elif n == 2:
48          return 'text_out'
49      else:
50          return None
51
52
53  def user2social_times(user):
54      '''Converts user code to a list of times when the user was social i.e. had
55      two or more repeated Bluetooth signals.'''

```

```

56 with open(userfile_path+user+"/bluetooth_log.txt", 'r') as f:
57     raw = [LE(line) for line in f.readlines()]
58     #List of times when user was social
59     social_times = []
60     #Temporary variables
61     current_time = 0
62     previous_time = 0
63     current_users = []
64     previous_users = []
65     for ind in xrange(len(raw)-1):
66         signal = raw[ind]
67         new_time = signal['timestamp']
68         #Check if line represents a new signal and if so, update values
69         if new_time != current_time:
70             #Determine if previous signal was social and append to results
71             overlap = set(previous_users).intersection(set(current_users))
72             if len(overlap) >= social_threshold:
73                 social_times.append(previous_time)
74             #Update variables
75             previous_time = current_time
76             previous_users = current_users
77             current_users = []
78             current_time = new_time
79             new_user = signal['name']
80             if new_user == '-1' or not new_user:
81                 continue
82             else:
83                 current_users.append(new_user)
84             #
85     return social_times
86
87 #=====
88 # Time to crunch some numbers
89 #=====
90
91 user_folders = [f for f in glob.glob(userfile_path+"/*")
92                 if os.path.isfile(f+"/call_log.txt")
93                 and os.path.isfile(f+"/sms_log.txt")
94                 and os.path.isfile(f+"/bluetooth_log.txt")]
95 users = [folder.split(userfile_path)[1] for folder in user_folders]
96
97 if not users:
98     raise IOError('Found no users. Check userfile path.')
99
100 if max_users:
101     users = users[:max_users]
102
103 trigger = 'bluetooth'
104 events = ['call_out', 'call_in', 'text_in', 'text_out']
105 pairs = [p for p in itertools.product([trigger], events)]
106
107 activity = {p: Binarray(interval, bin_size) for p in pairs}
108 background = {p: Binarray(interval, bin_size) for p in pairs}
109
110
111 #Read in data
112 call_data = {user: readncheck(userfile_path+user+"/call_log.txt")
113              for user in users}
114 text_data = {user: readncheck(userfile_path+user+"/sms_log.txt")
115              for user in users}
116
117 count = 0

```

```

118
119 for user in users:
120     count += 1
121     print "Analyzing user %s out of %s. Code: %s" % (count, len(users), user)
122
123     #Get user data
124     user_calls = call_data[user]
125     user_texts = text_data[user]
126     user_social_times = user2social_times(user)
127
128     if not user_social_times:
129         continue
130
131     #Get background data
132     others = []
133     other_social_times = []
134     while len(others) < number_of_background_samples:
135         temp = random.choice(users)
136         if not (temp in others or temp == user):
137             newstuff = user2social_times(temp)
138             if not newstuff:
139                 continue
140             others.append(temp)
141             other_social_times += newstuff
142
143     #Determine the interval in which we have data on the current user
144     first = min(user_social_times)
145     last = max(user_social_times)
146
147     #=====
148     # #Establish activity signal
149     #=====
150     for time in user_social_times:
151         for e in events:
152             activity[(trigger, e)].center = time
153
154         for user_call in user_calls:
155             event = call_type(user_call['type'])
156             if not event:
157                 continue
158             time = user_call['timestamp']
159             activity[(trigger, event)].place_event(time)
160
161         for user_text in user_texts:
162             event = text_type(user_text['type'])
163             if not event:
164                 continue
165             time = user_text['timestamp']
166             activity[(trigger, event)].place_event(time)
167
168     #=====
169     # #Establish background signal
170     #=====
171     for other_time in other_social_times:
172         #Reposition the relevant binarrays
173         if not first <= other_time <= last:
174             continue
175         for e in events:
176             background[(trigger, e)].center = other_time
177
178     #Determine call background
179     for user_call in user_calls or []:

```

```

180         event = call_type(user_call['type'])
181         if not event:
182             continue
183         time = user_call['timestamp']
184         background[(trigger,event)].place_event(time)
185
186         #Determine text background
187         for user_text in user_texts or []:
188             event = text_type(user_text['type'])
189             if not event:
190                 continue
191             time = user_text['timestamp']
192             background[(trigger,event)].place_event(time)
193         #
194     #
195 #
196
197 #=====
198 # Done. Save signals
199 #=====
200
201 #Make a filename for the output file
202 filename = make_filename(prefix = trigger, interval=interval,
203                          bin_size=bin_size, ext = 'json')
204 with open(filename, 'w') as f:
205     dump((activity, background), f)
206
207 print "Saved to "+filename

```

Code for loading and plotting Bluetooth data

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import glob
5  import numpy as np
6  import matplotlib.pyplot as plt
7
8  from social_fabric.phonetools import (load, make_filename)
9
10 #=====
11 display = True
12 #These are just default - they're updated when data is read
13 interval = 12*60*2
14 bin_size = 10*60
15 #=====
16
17
18 #Numpy-compliant method to convert activity and background to relative signal
19 get_signal = np.vectorize(lambda act, back: act * 1.0/back if back else 0)
20
21 def read_data(filename):
22     '''Reads in data to plot. Updates the interval and bin sizes parameters.
23     Returns data to be plotted as a hashmap with the following structure:
24     {trigger : {event : signal}}.'''
25     data = {}
26     with open(filename, 'r') as f:
27         (a,b) = load(f)
28         first = True
29         for key in a.keys():
30             #Update interval and bin info
31             if first:
32                 global interval
33                 global bin_size
34                 interval = a[key].interval
35                 bin_size = a[key].bin_size
36                 first = False
37
38             (trigger, event) = key
39             act = a[key].normalized()
40             back = b[key].normalized()
41             signal = get_signal(list(act), list(back))
42             if not trigger in data:
43                 data[trigger] = {}
44             data[trigger][event] = signal
45     return data
46
47
48 def make_plot(trigger, signals):
49     '''Generate a plot of relative user activity signal.'''
50     legendstuff = []
51     for event, vals in signals.iteritems():
52         t = np.arange(-interval, interval, bin_size)
53         t = map(lambda x: x*1.0/3600, t)
54         legendstuff.append(event)
55         s = list(vals)
56         plt.plot(t, s)
57     plt.legend(legendstuff, loc='upper left')
58     plt.xlabel('Time (h)')
59     plt.ylabel('Relative signal')
60     plt.title("Trigger: " + trigger)
61     plt.grid(True)

```



```
62     saveas = make_filename(trigger, interval, bin_size, ext = 'pdf')
63     plt.savefig(saveas)
64     print "Saved to "+saveas
65     if display:
66         plt.show()
67
68 if __name__ == '__main__':
69     # Read in data
70     filenames = glob.glob('bluetooth-int43200-bin600.json')
71     for filename in filenames:
72         data = read_data(filename)
73         for trigger, signals in data.iteritems():
74             make_plot(trigger, signals)
75     #
```

A.2.3 Preprocessing

```

1  # -*- coding: utf-8 -*-
2  from __future__ import division
3
4  import os
5  import glob
6  import numpy as np
7  from ast import literal_eval as LE
8  from social_fabric.phonetools import readncheck, user2profile
9  from social_fabric.secrets import pushme
10 from social_fabric.smallestenclosingcircle import make_circle
11 from social_fabric import lloyds
12 from collections import Counter
13 import math
14 from datetime import datetime, timedelta
15 import pytz
16 import json
17 from statsmodels.tsa import ar_model
18 import multiprocessing
19
20 #=====
21 # Parameters
22 #=====
23
24 #Grap all the files we need to read
25 #userfile_path = "c:\\ userfiles\\" #@Windows
26 userfile_path = "/lscr_paper/amoellga/Data/Telefon/userfiles/" #linux
27
28 #Filename to save output to.
29 output_filename = 'data.json'
30
31 #Number of users to analyse. Use None or 0 to include everyone.
32 max_users = None
33
34 #Specific user codes to analyze for debugging purposes. Empty means include all
35 exclusive_users = []
36 #Den her driller: '28b76d7b7879d364321f164df5169f'
37
38 #Conversion factors from degrees to meters (accurate around Copenhagen)
39 longitude2meters = 111319
40 latitude2meters = 110946
41
42 #In meters. These improve convergence of the stochastic SEC algorithm.
43 x_offset = 1389425.2238223257
44 y_offset = 6181209.0059678229
45
46 required = ['bluetooth_log.txt', 'call_log.txt', 'facebook_log.txt',
47            'gps_log.txt', 'sms_log.txt']
48
49 user_folders = [f for f in glob.glob(userfile_path+"*")
50                if all(os.path.isfile(f+"/"+s) for s in required)]
51 users = [folder.split(userfile_path)[1] for folder in user_folders]
52
53 if not users:
54     raise IOError('Found no users. Check userfile path.')
55
56 if exclusive_users:
57     users = list(set(exclusive_users).intersection(users))

```

```

58
59 if max_users:
60     users = users[:max_users]
61
62 # Number of wallclock hours pr bin when fitting autoregressive series
63 hours_pr_ar_bin = 6
64 assert 24%hours_pr_ar_bin == 0
65
66 # Number of hours pr bin when compute daily rythm entropy
67 hours_pr_daily_rythm_bin = 1
68 assert 24%hours_pr_daily_rythm_bin == 0
69
70 # Time zone information
71 cph_tz = pytz.timezone('Europe/Copenhagen')
72
73 n_jobs = 16 #maximum number of processors to use.
74
75 #Which data kinds to include
76 include_calls = True
77 include_ar = True
78 include_gps = True
79 include_network = False #Allow geolocation from network data
80 include_bluetooth = True
81 include_facebook = True
82
83 #Whether to include not-a-number values in final output
84 allow_nan = True
85
86 #Threshold values to discard users with insufficient data
87 minimum_number_of_texts = 10
88 minimum_number_of_calls = 5
89 minimum_number_of_gps_points = 100
90 minimum_number_of_facebook_friends = 1
91
92 #N_pings required to be considered social
93 bluetooth_social_threshold = 2
94
95 #Whether to output plots of cluster analysis
96 plot_clusters = False
97
98 #=====
99 # Define helper methods
100 #=====
101
102 def get_distance(p, q):
103     return math.sqrt(sum([(p[i]-q[i])**2 for i in xrange(len(p))]))
104
105 def is_sorted(l):
106     return all([l[i+1] >= l[i] for i in xrange(len(l)-1)])
107
108 def get_entropy(event_list):
109     '''Takes a list of contacts from in/ or outgoing call/text events and
110     computes its entropy. event_list must be simply a list of user codes
111     corresponding to events.'''
112     n = len(event_list)
113     counts = Counter(event_list)
114     ent = sum([-v/n * math.log(v/n, 2) for v in counts.values()])
115     return ent
116
117 def next_time(dt, deltahours):
118     '''Accepts a datetime object and returns the next datetime object at which
119     the 'hour' count modulo deltahours is zero.

```

```

120     For example, deltahours = 6 gives the next time clock time is
121     0, 6, 12 or 18.
122     Sounds simple but is pretty annoying due to daylight saving time and so on,
123     so take care not to mess with this.'''
124     base = datetime(dt.year, dt.month, dt.day, dt.hour)
125     interval = timedelta(hours = deltahours - dt.hour%deltahours)
126     naive_guess = base + interval
127     return cph_tz.localize(naive_guess)
128
129 def next_midnight(dt):
130     '''Takes a datetime object and returns dt object of following midnight'''
131     base = datetime(dt.year, dt.month, dt.day)
132     naive = base + timedelta(days = 1)
133     return cph_tz.localize(naive)
134
135 def epoch2dt(timestamp):
136     '''Converts unix timestamp into a pytz timezone-aware datetimeobject.'''
137     utc_time = datetime.utcnow().fromtimestamp(timestamp)
138     smart_time = cph_tz.fromutc(utc_time)
139     return smart_time
140
141 def sort_dicts_by_key(dictlist, key):
142     '''Takes a list of dicts and returns the same list sorted by its
143     key-entries.'''
144     decorated = [(d[key], d) for d in dictlist]
145     decorated.sort()
146     return [d for (k, d) in decorated]
147
148 def get_autocovar_coefficient(X, lag):
149     '''Returns the autocovariance coefficient for the input series at the
150     input lag.'''
151     mu = np.mean(X)
152     temp = sum((X[i] - mu)*(X[i+lag]-mu) for i in xrange(len(X)-lag))
153     return temp*1.0/len(X)
154
155 def get_autocorrelation_coefficients(series, lags):
156     '''Determines the autocorrelation coefficients of input series at
157     each of the input lags. Uses .r_k = c_k/c_0.
158     Accepts a list of lags or an int in which case it returns lags up to
159     and including the input.'''
160     if isinstance(lags, int):
161         lags = range(lags+1)
162     c0 = get_autocovar_coefficient(series, 0)
163     inv = 1.0/c0
164     return [inv*get_autocovar_coefficient(series, lag) for lag in lags]
165
166 def make_time_series(dts, hours_pr_ar_bin):
167     '''Takes a sorted list of datetime objects and converts to a time series
168     where each entry denotes the number of events in the corresponding bin.'''
169     first_time = next_time(dts[0], hours_pr_ar_bin)
170     for i in xrange(len(dts)):
171         if dts[i] >= first_time:
172             dts = dts[i:]
173             break
174     #
175     last_time = next_time(first_time, hours_pr_ar_bin)
176     time_series = []
177
178     summer = 0
179     for dt in dts:
180         while not first_time <= dt < last_time:
181             time_series.append(summer)

```

```

182         summer = 0
183         first_time = last_time
184         last_time = next_time(last_time, hours_pr_ar_bin)
185         summer += 1
186     return time_series
187
188
189 def timestamps2daily_entropy(timestamps, hours_pr_bin):
190     '''Constructs a histogram of hour-values of the input timestamps
191     and computes its entropy.'''
192     if not 24%hours_pr_bin == 0:
193         raise ValueError("24 must be divisible by hours_pr_bin.")
194     bins = {}
195     for timestamp in timestamps:
196         hour = epoch2dt(timestamp).hour
197         _bin = int(hour/hours_pr_bin)
198         try:
199             bins[_bin] += 1
200         except KeyError:
201             bins[_bin] = 1
202     #
203     total = sum(bins.values())
204     entropy = sum([-v/total*math.log(v/total, 2) for v in bins.values()])
205     return entropy
206
207 #Make sure we have a blank file to write to.
208 open(output_filename, 'w').close()
209 assert os.stat(output_filename).st_size == 0 #Check that it worked.
210
211 '''Main processing method. This is written as a separate method to allow
212 easy multiprocessing.'''
213 def process_user(user, user_counter):
214     msg = "Processing user %s of %s: %s" % (user_counter, len(users), user)
215     print msg
216     if user_counter % 100 == 0:
217         pushme(msg)
218
219     #Dict to hold all data extracted on current user
220     data = {}
221
222     #Try to load user psych profile- Discard user if they're not in file
223     try:
224         profile = user2profile[user]
225     except KeyError:
226         print "No psychological data on user."
227         return None # No questionnaire data. Shouldn't happen.
228
229     #Read in calls/texts
230     calls = readncheck(userfile_path+user+"/call_log.txt")
231
232     if len(calls) < minimum_number_of_calls:
233         print "too few calls"
234         return None
235     texts = readncheck(userfile_path+user+"/sms_log.txt")
236     if len(texts) < minimum_number_of_texts:
237         print "too few texts"
238         return None
239
240     #Extract list of times for calls, texts and combination
241     call_times = sorted([d['timestamp'] for d in calls])
242     text_times = sorted([d['timestamp'] for d in texts])
243     call_text_times = sorted(call_times + text_times)

```

```

244
245 #Get calls from the first three months
246 tmin = call_times[0]
247 tmax = call_times[0] + 60*60*24*30*3
248 early_calls = [c['number'] for c in calls if tmin<=c['timestamp']<=tmax]
249
250 #Repeat for texts
251 tmin = text_times[0]
252 tmax = text_times[0] + 60*60*24*30*3
253 early_texts = [t['body'] for t in texts if tmin<=t['timestamp']<=tmax]
254
255 #Get number of unique contacts for first 3 months and append to data.
256 uniques = len(set(early_calls + early_texts))
257 data['n_contacts_first_three_months'] = uniques
258
259 #Compute daily entropy for calls and texts
260 data['call_daily_entropy'] = timestamps2daily_entropy(call_times,
261                                                         hours_pr_daily_rhythm_bin)
262 data['text_daily_entropy'] = timestamps2daily_entropy(text_times,
263                                                         hours_pr_daily_rhythm_bin)
264
265 #Compute median and std for call durations
266 call_durations = [c['duration'] for c in calls if not c['duration'] == 0]
267 data['call_duration_med'] = np.median(call_durations)
268 data['call_duration_std'] = np.std(call_durations)
269
270 #=====
271 # Crunch time-series info
272 #=====
273 if include_ar:
274     #Grab a sorted list of only times of events caused by user
275     outgoing_stuff = sorted([d['timestamp'] for d in calls
276                             if d['type'] == 2] + [d['timestamp']
277                                                  for d in texts if d['type'] == 2])
278
279     n_params = int(24*7/hours_pr_ar_bin + 1) #1 week plus 1 extra bin
280
281     #Fit time series and extract parameters
282     try:
283         #Convert into timezone aware datetime objects
284         dts = [epoch2dt(timestamp) for timestamp in outgoing_stuff]
285
286         time_series = make_time_series(dts, hours_pr_ar_bin)
287
288         model = ar_model.AR(time_series)
289         result = model.fit(n_params)
290         #Grab parameters from fitted model
291         params = result.params[1:]
292         while len(params) < n_params:
293             params.append(float('nan'))
294     except:
295         if not allow_nan:
296             return None
297         else:
298             params = [float('nan') for _ in xrange(n_params)]
299
300     #Append AR-coefficients to user data
301     count = 0
302     for par in params:
303         count += 1
304         name = "outgoing_activity_AR_coeff_" + str(count)
305         data[name] = par

```

```

306 #Get autocorrelation coefficients as well
307
308 try:
309     accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
310 except:
311     if not allow_nan:
312         return None
313     accs = [float('nan') for _ in xrange(n_params + 1)]
314
315 #Append autocorrelation coefficients to user data
316 for i in xrange(len(accs)):
317     name = "outgoing_activity_acc_" + str(i)
318     data[name] = accs[i]
319
320 #Repeat with incoming signals. Might be interesting.
321 incoming_stuff = sorted([d['timestamp'] for d in calls
322                        if d['type'] == 1] + [d['timestamp']
323                        for d in texts if d['type'] == 1])
324
325 try:
326     # Convert into timezone aware datetime objects
327     dts = [epoch2dt(timestamp) for timestamp in incoming_stuff]
328
329     time_series = make_time_series(dts, hours_pr_ar_bin)
330     model = ar_model.AR(time_series)
331     result = model.fit(n_params)
332     params = result.params[1:]
333     while len(params) < n_params:
334         params.append(float('nan'))
335 except:
336     if not allow_nan:
337         return None
338     else:
339         params = [float('nan') for _ in xrange(n_params)]
340
341 # Name each of them and append to user data
342 count = 0
343 for par in params:
344     count += 1
345     name = "incoming_activity_AR_coeff_" + str(count)
346     data[name] = par
347
348 #Get autocorrelation coefficients as well
349 try:
350     accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
351 except:
352     if not allow_nan:
353         return None
354     accs = [float('nan') for _ in xrange(n_params + 1)]
355
356 for i in xrange(len(accs)):
357     name = "incoming_activity_acc_" + str(i)
358     data[name] = accs[i]
359
360
361 #=====
362 #   Crunch call/text info
363 #=====
364
365 if include_calls:
366     #Add values to temporary data map.
367     d = {'call': call_times, 'text': text_times, 'ct': call_text_times}

```

```

368     for label, times in d.iteritems():
369         timegaps = [times[i+1] - times[i] for i in xrange(len(times)-1)]
370         timegaps = filter(lambda x: x < 259200, timegaps) #3 days, tops
371         data[label+'_iet_med'] = np.median(timegaps)
372         data[label+'_iet_std'] = np.std(timegaps)
373
374     #Generate lists of the contact for each text/call event
375     call_numbers = [call['number'] for call in calls]
376     text_numbers = [text['address'] for text in texts]
377     ct_numbers = call_numbers + text_numbers
378
379     #Compute entropy and add to data
380     data['call_entropy'] = get_entropy(call_numbers)
381     data['text_entropy'] = get_entropy(text_numbers)
382     data['ct_entropy'] = get_entropy(ct_numbers)
383
384     #Compute contact list info
385     call_contacts = Counter([c['number'] for c in calls
386                             if c['type'] == 2]).keys()
387     text_contacts = Counter([t['address'] for t in texts
388                             if t['type'] == 2]).keys()
389
390     #Grap number of contacts
391     n_call_contacts = len(call_contacts)
392     n_text_contacts = len(text_contacts)
393     n_ct_contacts = len(set(call_contacts).union(set(text_contacts)))
394
395     #Add to data map
396     data['n_call_contacts'] = n_call_contacts
397     data['n_text_contacts'] = n_text_contacts
398     data['n_ct_contacts'] = n_ct_contacts
399
400     #Compute and add contact/interaction ratio (cir)
401     data['call_cir'] = n_call_contacts/len(calls)
402     data['text_cir'] = n_text_contacts/len(texts)
403     data['ct_cir'] = n_ct_contacts/(len(calls) + len(texts))
404
405     #Add data on number of interactions
406     data['n_calls'] = len(calls)
407     data['n_texts'] = len(texts)
408     data['n_ct'] = len(calls + texts)
409
410     #Determine percentage of calls/texts that were initiated by user.
411     initiated_calls = len([c for c in calls if c['type'] == 2])
412     data['call_percent_initiated'] = initiated_calls/len(calls)
413     initiated_texts = len([t for t in texts if t['type'] == 2])
414     data['text_percent_initiated'] = initiated_texts/len(texts)
415
416     #Determine call response rate.
417     with open(userfile_path+user+"/call_log.txt", 'r') as f:
418         all_calls = [LE(line) for line in f.readlines()]
419     #Make sure the call data is sorted
420     if not is_sorted([c['timestamp'] for c in all_calls]):
421         all_calls = sort_dicts_by_key(all_calls, 'timestamp')
422
423     '''Check for unanswered called that are replied to within an hour.
424     This is performed in the following fashion: iterate through all the
425     calls. If a call is unanswered, add it to "holding" list. If a call
426     from holding matches the current call, it counts as a reply.
427     If the time of the current call is more than hour after a held call,
428     it is discarded'''
429     missed = 0
430     replied = 0

```



```

430 holding = []
431 for call in all_calls:
432     if call['type']==3 or call['type']==1 and call['duration']==0:
433         holding.append(call)
434         missed += 1
435     else:
436         for held_call in holding:
437             #Drop calls that have been held for too long
438             if call['timestamp'] - held_call['timestamp'] > 3600:
439                 holding.remove(held_call)
440             #Check if given call is a response
441             elif (call['type'] == 2
442                   and call['number'] == held_call['number']):
443                 holding.remove(held_call)
444                 replied += 1
445             #
446         #
447     #
448 data['call_response_rate'] = replied/(missed+replied) if replied else 0
449
450 #Determine text response rate
451 missed = 0
452 replied = 0
453 holding = []
454 response_times = []
455 if not is_sorted([t['timestamp'] for t in texts]):
456     texts = sort_dicts_by_key(texts, 'timestamp')
457
458 for text in texts:
459     #Make sure incoming text is not from a user already held
460     if text['type'] == 1:
461         if not holding or all([text['address'] !=
462                               t['address'] for t in holding]):
463             #It's good - append it
464             holding.append(text)
465             missed += 1
466         #
467     else:
468         for held_text in holding:
469             if text['timestamp'] - held_text['timestamp'] > 3600:
470                 holding.remove(held_text)
471             #Check if text counts as reply
472             elif (text['type'] == 2
473                   and text['address'] == held_text['address']):
474                 holding.remove(held_text)
475                 replied += 1
476                 dt = text['timestamp'] - held_text['timestamp']
477                 response_times.append(dt)
478             #
479         #
480     #
481 data['text_response_rate'] = replied/(missed+replied) if replied else 0
482 data['text_latency'] = np.median(response_times)
483
484 #Check percentage of calls taken place in during the night
485 count = 0
486 for call in calls:
487     hour = epoch2dt(call['timestamp']).hour
488     if not (8 <= hour <22):
489         count += 1
490     #
491 data['call_night_activity'] = count/len(calls)

```

[illegible]

```

554         # Reset counters and update bins
555         current_points = []
556         early_day = late_day
557         late_day = next_midnight(late_day)
558         current_points.append(datum)
559
560
561     data['radius_of_gyration_med'] = np.median(radii)
562     data['radius_of_gyration_std'] = np.std(radii)
563     data['travel_med'] = np.median(distances)
564     data['travel_std'] = np.std(distances)
565
566     #Run Lloyd's algorithm to identify clusters
567     #Determine which points are stationary – less movement than 100m
568     stationary_data = []
569     try:
570         for i in xrange(1,len(gps_data)-1):
571             a,b,c = tuple([gps_data[ind]['point'] for ind in
572                             [i-1, i, i+1]])
573             if (get_distance(a, b) < 0.1 and get_distance(b, c) < 0.1):
574                 stationary_data.append(gps_data[i])
575
576             #
577             initial_clusters = 50
578             threshold_percent = 0.05
579             points = [elem['point'] for elem in stationary_data]
580
581             minimum_points = int(threshold_percent*len(points))
582
583             clusters_scatter = lloyds.lloyds(points, initial_clusters, runs=3,
584                                             init='scatter')
585
586             clusters_sample = lloyds.lloyds(points, initial_clusters, runs=3,
587                                             init='sample')
588
589             #Determine most succesful method
590             locs = lambda c: [p for p in c.values() if len(p)>=minimum_points]
591             if len(locs(clusters_scatter)) >= len(locs(clusters_sample)):
592                 method = 'scatter'
593                 best = clusters_scatter
594             else:
595                 method = 'sample'
596                 best = clusters_sample
597
598             # Number of places which survive cutoff (true = 1, so just sum)
599             n_places = sum([len(pl) >= minimum_points for pl in best.values()])
600
601             #Output plots of clustering
602             if plot_clusters:
603                 if not os.path.isdir('pics'):
604                     os.mkdir('pics')
605                 for ext in ['.pdf', '.png']:
606                     filename = 'pics/'+user+"_"+method+ext
607                     lloyds.draw_clusters(clusters = best,
608                                         threshold = minimum_points,
609                                         show = False,
610                                         filename = filename)
611
612             except:
613                 if not allow_nan:
614                     return None
615                 n_places = float('nan')
616             data['n_places'] = n_places

```

```

616
617 #Compute location entropy and add to data
618 try:
619     n_points = sum(len(location) for location in best.values())
620     data['location_entropy'] = sum([-len(p)/n_points*math.log(len(p)/n_points)
621                                   for p in best.values()])
622 except ValueError:
623     data['location_entropy'] = float('nan')
624
625 '''Guess where people live. Probably where they spend weeknights...
626 It's important to avoid selection bias here (people probably turn off
627 their phone when sleeping at home but not while partying at DTU, which
628 means fewer data points at their actual home).
629 This is rectified by excluding points that aren't logged monday to
630 thursday and only recording one 'late' or 'early' data point pr date.
631 These points are labelled 'weird' and are used to determine the user's
632 home.'''
633 try:
634     weirdpoints = []
635     latedays = []
636     earlydays = []
637     for datum in stationary_data:
638         now = datum['smarttime']
639         #ignore weekends
640         if now.weekday() > 3:
641             continue
642         thisdate = (now.year, now.month, now.day)
643         if now.hour >= 20 and not thisdate in latedays:
644             weirdpoints.append(datum['point'])
645             latedays.append(thisdate)
646         elif now.hour <= 7 and not thisdate in earlydays:
647             weirdpoints.append(datum['point'])
648             earlydays.append(thisdate)
649
650     best_score = 0
651     home = None
652     for key, val in best.iteritems():
653         score = len(set(val).intersection(weirdpoints))
654         if score > best_score:
655             home = key
656             best_score = score
657     #
658     # Estimate how much user spends at home
659     ordered_gps = sort_dicts_by_key(gps_data, 'timestamp')
660     is_home = lambda p: get_distance(home, p) <= 0.200
661     time_home = 0
662     time_away = 0
663     for i in xrange(len(ordered_gps)-1):
664         a = ordered_gps[i]
665         b = ordered_gps[i+1]
666         dt = b['timestamp'] - a['timestamp']
667         if dt > 7200:
668             continue
669         elif is_home(a['point']) and is_home(b['point']):
670             time_home += dt
671         elif (not is_home(a['point'])) and (not is_home(b['point'])):
672             time_away += dt
673     #
674     data['home_away_time_ratio'] = time_home/time_away
675 except:
676     if not allow_nan:
677         return None

```

```

678     data['home_away_time_ratio'] = float('nan')
679
680
681 #=====
682 # Facebook data
683 #=====
684
685 if include_facebook:
686     with open(userfile_path+user+'/facebook_log.txt', 'r') as f:
687         n = len(f.readlines())
688         if n < minimum_number_of_facebook_friends:
689             return None
690         data['number_of_facebook_friends'] = n
691
692 #=====
693 # Bluetooth data
694 #=====
695
696 if include_bluetooth:
697     with open(userfile_path+user+"/bluetooth_log.txt", 'r') as f:
698         raw = [LE(line) for line in f.readlines()]
699         #Make sure data is sorted chronologically
700         if not is_sorted([entry['timestamp'] for entry in raw]):
701             raw = sort_dicts_by_key(raw, 'timestamp')
702         #List of times when user was social
703         social_times = []
704         total_social_time = 0
705         total_time = 0
706         #maps from each other user encountered to time spend with said user
707         friend2time_spent = {}
708         #Temporary variables
709         current_time = 0
710         previous_time = 0
711         current_users = []
712         previous_users = []
713         for signal in raw:
714             new_time = signal['timestamp']
715             #Check if line represents a new signal and if so, update values
716             if new_time != current_time:
717                 dt = new_time - current_time
718                 #Determine number of pings
719                 overlap = set(previous_users).intersection(set(current_users))
720                 if len(overlap) >= bluetooth_social_threshold:
721                     social_times.append(previous_time)
722                     if dt <= 7200:
723                         total_social_time += dt
724                         total_time += dt
725                         for friend in overlap:
726                             try:
727                                 friend2time_spent[friend] += dt
728                             except KeyError:
729                                 friend2time_spent[friend] = dt
730
731                     elif dt <= 7200:
732                         total_time += dt
733                     #Update variables
734                     previous_time = current_time
735                     previous_users = current_users
736                     current_users = []
737                     current_time = new_time
738             new_user = signal['name']
739             if new_user == '-1' or not new_user:

```

```

740         continue
741     else:
742         current_users.append(new_user)
743     #
744     # Add fraction of time spent social to output
745     data['fraction_social_time'] = total_social_time/total_time
746     # Compute social entropy
747     normfac = 1.0/sum(friend2time_spent.values())
748     ent = sum(-t*normfac*math.log(t*normfac)
749             for t in friend2time_spent.values())
750     data['social_entropy'] = ent
751
752     data['bluetooth_daily_entropy'] = timestamps2daily_entropy(social_times,
753                                                             hours_pr_daily_rythm_bin)
754
755     #Ensure time span is suficient to make a time series
756     if not (social_times[-1] - social_times[0] > 24*3600*7
757            +1+3600*hours_pr_ar_bin):
758         return None
759
760     #Fit AR-series and append parameters to output
761     try:
762         dts = [epoch2dt(timestamp) for timestamp in social_times]
763         time_series = make_time_series(dts, hours_pr_ar_bin)
764         model = ar_model.AR(time_series)
765         n_params = int(24*7/hours_pr_ar_bin + 1) #1 week plus 1 extra bin
766         result = model.fit(maxlag = None, ic = None)
767         params = result.params#[1:]
768         while len(params) < n_params:
769             params.append(float('nan'))
770     except:
771         if not allow_nan:
772             return None
773         params = [float('nan') for _ in xrange(n_params)]
774
775     count = 0
776     for par in params:
777         count += 1
778         name = "bluetooth_activity_AR_coeff_" + str(count)
779         data[name] = par
780
781     # Compute autocorrelation coeffs and append to output
782     try:
783         accs = get_autocorrelation_coefficients(time_series, n_params)[:1]
784     except:
785         if not allow_nan:
786             return None
787         accs = [float('nan') for _ in xrange(n_params + 1)]
788     for i in xrange(len(accs)):
789         name = "bluetooth_activity_acc_" + str(i)
790         data[name] = accs[i]
791
792     #=====
793     # Wrap up user
794     #=====
795
796     # Double check thata doesn't containing nan values
797     if any(np.isnan(value) for value in data.values()) and not allow_nan:
798         return None #Discard user due to insufficient data
799
800     #Collect results
801     final = {'user': user, 'data': data, 'profile': profile}

```

```

802     return final
803
804
805 if __name__ == '__main__':
806     user_counter = 0
807     #Make job queue
808     pool = multiprocessing.Pool(processes = n_jobs)
809     jobs = []
810     for user in users:
811         user_counter += 1
812         args = {'user': user, 'user_counter': user_counter}
813         jobs.append(pool.apply_async(process_user, kwds = args))
814
815     pool.close() #run
816     pool.join() #Wait for remaining jobs
817
818     #Write results
819     with open(output_filename, 'a') as f:
820         for job in jobs:
821             result = job.get()
822             if not result:
823                 continue
824             json.dump(result, f)
825             f.write("\n")
826
827
828     #Done.
829     pushme("Data extraction done.")

```

A.2.4 Social Fabric Code

```

1  # -*- coding: utf-8 -*-
2  """This module aims to allow sharing of some common methods and settings
3  when testing and tweaking various machine learning schemes.
4  Always import settings and the like from here!"""
5
6  from __future__ import division
7  import abc
8  from collections import Counter
9  import itertools
10 import json
11 import math
12 import matplotlib.colors as mcolors
13 import matplotlib.pyplot as plt
14 import matplotlib.patches as mpatches
15 import multiprocessing
16 import numpy as np
17 import random
18 from scipy.sparse import dok_matrix
19 from sklearn import svm
20 from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
21 from sklearn.cross_validation import (cross_val_score, LeaveOneOut, KFold,
22                                     StratifiedKFold)
23
24 import sys
25 from time import time
26 import traceback

```

```

27 | oldhat = (35/256,39/256,135/256)
28 | nude = (203/256,150/256,93/256)
29 | wine = (110/256,14/256,14/256)
30 | moerkeroed = (156/256,30/256,36/256)
31 |
32 | def _make_colormap(seq):
33 |     """Return a LinearSegmentedColormap
34 |     seq: a sequence of floats and RGB-tuples. The floats should be increasing
35 |     and in the interval (0,1).
36 |     """
37 |     seq = [(None,) * 3, 0.0] + list(seq) + [1.0, (None,) * 3]
38 |     cdict = {'red': [], 'green': [], 'blue': []}
39 |     for i, item in enumerate(seq):
40 |         if isinstance(item, float):
41 |             r1, g1, b1 = seq[i - 1]
42 |             r2, g2, b2 = seq[i + 1]
43 |             cdict['red'].append([item, r1, r2])
44 |             cdict['green'].append([item, g1, g2])
45 |             cdict['blue'].append([item, b1, b2])
46 |     return mcolors.LinearSegmentedColormap('CustomMap', cdict)
47 |
48 | color_map = _make_colormap([oldhat, moerkeroed, 0.33, moerkeroed, nude, 0.67, nude])
49 |
50 | big_five = ['openness', 'conscientiousness', 'extraversion', 'agreeableness',
51 |             'neuroticism']
52 |
53 | #_default_features = ["n_texts",
54 | #                     "ct_iet_std ",
55 | #                     "call_cir ",
56 | #                     "call_entropy",
57 | #                     "text_cir ",
58 | #                     "n_calls ",
59 | #                     "text_latency ",
60 | #                     "call_outgoing",
61 | #                     "fraction_social_time",
62 | #                     "text_outgoing",
63 | #                     "call_iet_std ",
64 | #                     "n_text_contacts",
65 | #                     "call_night_activity ",
66 | #                     "call_iet_med",
67 | #                     "outgoing_activity_AR_coeff_2",
68 | #                     "text_entropy",
69 | #                     "ct_cir ",
70 | #                     "text_response_rate",
71 | #                     "n_ct_contacts",
72 | #                     "social_entropy",
73 | #                     "n_call_contacts",
74 | #                     "n_ct",
75 | #                     "text_iet_std ",
76 | #                     "ct_iet_med",
77 | #                     "ct_entropy",
78 | #                     "text_iet_med",
79 | #                     "call_response_rate",
80 | #                     "number_of_facebook_friends"]
81 |
82 | _default_features = ['call_iet_med', 'text_iet_med', 'social_entropy',
83 |                     'call_entropy', 'travel_med', 'n_places', 'text_latency',
84 |                     'call_night_activity']
85 |
86 | def split_ntiles(values, n):
87 |     """Determines the values that separate the input list into n equal parts.
88 |     this is a generalization of the notion of median (in the case n = 2) or

```



```

89     quartiles (n=4).
90     Usage: ntiles([5,6,7], 2) gives [6] for instance.'''
91     result = []
92     for i in xrange(1,n):
93         percentile = 100/n * i
94         result.append(np.percentile(values, percentile,
95                                   interpolation='linear'))
96     return result
97
98 def determine_ntile(value, ntiles):
99     '''Determines which n-tile the input value belongs to.
100     Usage: determine_ntile([7,9,13], 10) gives 2 (third quartile).
101     This uses zero indexing so data split into e.g. quartiles will give results
102     like 0,1,2,3 - NOT 1,2,3,4.'''
103     #Check if value is outside either extreme, meaning n-tile 1 or n.
104     if value >= ntiles[-1]:
105         return len(ntiles) #Remember the length is n-1
106     elif value < ntiles[0]:
107         return 0 #Values was in the first n-tile
108     # Define possible region and search for where value is between two elements
109     left = 0
110     right = len(ntiles)-2
111     #Keep checking th middle of the region and updating region
112     ind = (right + left)//2
113     while not ntiles[ind] <= value < ntiles[ind + 1]:
114         #Check if lower bound tile is on the left
115         if value < ntiles[ind]:
116             right = ind - 1
117         else:
118             left = ind + 1
119         ind = (right + left)//2
120     # Being between ntiles 0 and 1, means second n-tile and so on.
121     return ind + 1
122
123 def assign_labels(Y, n):
124     '''Accepts a list and an int n and returns a list of discrete labels
125     corresponding to the ntile each original y-value was in.'''
126     ntiles = split_ntiles(Y, n)
127     labels = [determine_ntile(y, ntiles) for y in Y]
128     return labels
129
130 def normalize_data(list_):
131     '''Normalizes input data to the range [-1, 1]'''
132     lo, hi = min(list_), max(list_)
133     if lo == hi:
134         z = len(list_)*[0]
135         return z
136     else:
137         return [2*(val-lo)/(hi - lo) - 1 for val in list_]
138
139
140 def read_data(filename, trait, n_classes = None, normalize = True,
141               features='default', interpolate = True):
142     '''This reads in a preprocessed datafile, splits psych profile data into
143     n classes if specified, filters desired psychological traits and
144     features and returns as a tuple (X,Y, indexdict), which can be fed to a
145     number of off-the-shelf ML schemes.
146     If trait=='Sex', female and male are converted to 0 and 1, respectively.
147     indexdict maps each element of the feature vectors to their label, as in
148     {42 : 'distance_travelled_pr_day'} etc.
149
150     Args:

```

```

151     filename : str
152         Name of the file containing the data.
153     trait : str
154         The psychological trait to extract data on.
155     n_classes : int
156         Number of classes to split data into. Default is None,
157         i.e. just keep the decimal values. Ignored if trait == 'Sex', as data
158         only has two discreet values.
159     normalize : bool
160         Whether to hard normalize data to [-1, 1].
161     features : str/list
162         Which features to read in. Can also be 'all'
163         or 'default', meaning the ones I've pragmatically found to be
164         reasonable.
165
166     interpolate : bool:
167         Whether to replace NaN's with the median value of
168         the feature in question.'''
169     if trait == 'sex':
170         n_classes = None
171     #Read in the raw data
172     with open(filename, 'r') as f:
173         raw = [json.loads(line) for line in f.readlines()]
174     #Get list of features to be included - everything if nothing's specified
175     included_features = []
176     if features == 'default':
177         included_features = _default_features
178     elif features == 'all':
179         included_features = raw[0]['data'].keys() if features == 'all' else features
180     else:
181         included_features = features
182
183     #Remove any features that only have NaN values.
184     for i in xrange(len(included_features)-1,-1,-1):
185         feat = included_features[i]
186         if all(math.isnan(line['data'][feat]) for line in raw):
187             del included_features[i]
188
189
190     # ----- Handle feature vectors -----
191     # Dict mapping indices to features
192     indexdict = {ind: feat for ind, feat in enumerate(included_features)}
193     # N_users x N_features array to hold data
194     rows = len(raw)
195     cols = len(included_features)
196     X = np.ndarray(shape = (rows, cols))
197     for i in xrange(rows):
198         line = raw[i]
199         #Construct data matrix
200         for j, feat in indexdict.iteritems():
201             val = line['data'][feat]
202             X[i, j] = val
203     #
204     #Replace NaNs with median values
205     if interpolate:
206         for j in xrange(cols):
207             #Get median of feature j
208             med = np.median([v for v in X[:,j] if not math.isnan(v)])
209             if math.isnan(med):
210                 raise ValueError('Feature %s contains only NaN's and should
211                                     have been removed.' % indexdict[j])
212         for i in xrange(rows):

```

```

213         if math.isnan(X[i,j]):
214             X[i,j] = med
215
216     if normalize:
217         for j in xrange(cols):
218             col = X[:,j]
219             X[:,j] = normalize_data(col)
220
221     # ----- Handle class info -----
222     trait_values = []
223     for line in raw:
224         #Add value of psychological trait
225         psych_trait = line['profile'][trait]
226         if trait == 'Sex':
227             if psych_trait == 'Female':
228                 psych_trait = 0
229             elif psych_trait == 'Male':
230                 psych_trait = 1
231             else:
232                 raise ValueError('My code is binary gender normative, sorry.')
233         trait_values.append(psych_trait)
234     Y = []
235     if n_classes == None:
236         Y = trait_values
237     else:
238         ntiles = split_ntiles(trait_values, n_classes)
239         Y = [determine_ntile(tr, ntiles) for tr in trait_values]
240
241     return (X, Y, indexdict)
242
243
244 def plot_stuff(input_filename, output_filename=None, color=moerkeroed):
245     with open(input_filename, 'r') as f:
246         d = json.load(f)
247         x = d['x']
248         y = d['y']
249         yerr = d['mean_stds']
250         plt.plot(x,y, color=color, linestyle='dashed',
251                 marker='o')
252         plt.errorbar(x, y, yerr=yerr, linestyle="None", marker="None",
253                     color=color)
254     if output_filename:
255         plt.savefig(output_filename)
256
257 def get_TPRs_and_FPRs(X, Y, forest = None, verbose = False):
258     '''Accepts a list of feature vectors and a list of labels and returns a
259     tuple of true positive and false positive rates (TPRs and FPRs,
260     respectively) for various confidence thresholds.'''
261     kf = LeaveOneOut(n=len(Y))
262
263     results = []
264     thresholds = []
265
266     counter = 0
267     for train, test in kf:
268         counter += 1
269         if counter % 10 == 0 and verbose:
270             print "Testing on user %s of %s..." % (counter, len(Y))
271
272         result = {}
273         train_data = [X[i] for i in train]
274         train_labels = [Y[i] for i in train]

```

```

275     test_data = [X[i] for i in test]
276     test_labels = [Y[i] for i in test]
277
278     if not forest:
279         forest = RandomForestClassifier(
280             n_estimators = 1000,
281             n_jobs=-1,
282             criterion='entropy')
283
284     forest.fit(train_data, train_labels)
285     result['prediction'] = forest.predict(test_data)[0]
286     result['true'] = test_labels[0]
287     confidences = forest.predict_proba(test_data)[0]
288     result['confidences'] = confidences
289     thresholds.append(max(confidences))
290
291     results.append(result)
292
293     #ROC curve stuff – false and true positive rates
294     TPRs = []
295     FPRs = []
296
297     unique_thresholds = sorted(list(set(thresholds)), reverse=True)
298
299     for threshold in unique_thresholds:
300         tn = 0
301         fn = 0
302         tp = 0
303         fp = 0
304         for result in results:
305             temp = result['prediction']
306             if temp == 1 and result['confidences'][1] >= threshold:
307                 pred = 1
308             else:
309                 pred = 0
310             if pred == 1:
311                 if result['true'] == 1:
312                     tp += 1
313                 else:
314                     fp += 1
315             #
316             elif pred == 0:
317                 if result['true'] == 0:
318                     tn += 1
319                 else:
320                     fn += 1
321             #
322         #
323         TPRs.append(tp/(tp + fn))
324         FPRs.append(fp/(fp + tn))
325     return (TPRs, FPRs)
326
327 def make_roc_curve(TPRs, FPRs, output_filename = None):
328     '''Accepts a list of true and false positive rates (TPRs and FPRs,
329     respectively) and generates a ROC-curve.'''
330     predcol = moerkeroed
331     basecol = oldhat
332     fillcol = nude
333
334     fig = plt.figure()
335     ax = fig.add_subplot(1,1,1)
336

```

```

337 TPRs = [0] + TPRs + [1]
338 FPRs = [0] + FPRs + [1]
339
340 area = 0.0
341 for i in xrange(len(TPRs)-1):
342     dx = FPRs[i+1] - FPRs[i]
343     y = 0.5*(TPRs[i] + TPRs[i+1])
344     under_curve = dx*y
345     baseline = dx*0.5*(FPRs[i] + FPRs[i+1])
346     area += under_curve - baseline
347
348 baseline = FPRs
349 ax.fill_between(x = FPRs, y1 = TPRs, y2 = baseline, color = fillcol,
350                interpolate = True, alpha=0.8)
351 ax.plot(baseline, baseline, color = basecol, linestyle = 'dashed',
352        linewidth = 1.0, label = 'Baseline')
353 ax.plot(FPRs, TPRs, color=predcol, linewidth = 1,
354        label = 'Prediction')
355
356 plt.xlabel('False positive rate.')
357 plt.ylabel('True positive rate.')
358
359 handles, labels = ax.get_legend_handles_labels()
360 hest = mpatches.Patch(color=fillcol)
361
362 labels += ['Area = %.3f' % area]
363 handles += [hest]
364 ax.legend(handles, labels, loc = 'lower right')
365 # plt.legend(handles = [tp_line, base])
366 if output_filename:
367     plt.savefig(output_filename)
368 plt.show()
369
370 def rank_features(X, Y, forest, indexdict, limit = None):
371     '''Ranks the features of a given dataset and classifier.
372     indexdict should be a map from indices to feature names like
373     {0 : 'average_weight'} etc.
374     if limit is specified, this method returns only the top n ranking features.
375     Returns a dict like {'feature name' : (mean importances, std)}.'''
376     importances = forest.feature_importances_
377     stds=np.std([tree.feature_importances_ for tree in forest.estimators_],
378                axis=0)
379     indices = np.argsort(importances)[::-1]
380     if limit:
381         indices = indices[:limit]
382     d = {indexdict[i]: (importances[i], stds[i]) for i in indices}
383     return d
384
385 def check_performance(X, Y, clf, strata = None):
386     '''Checks forest performance compared to baseline.'''
387     N_samples = len(Y)
388     #Set up validation indices
389     if not strata: #Do leave-one-out validation
390         skf = KFold(N_samples, n_folds=N_samples, shuffle = False)
391     else: #Do stratified K-fold
392         skf = StratifiedKFold(Y, n_folds=strata)
393
394     #Evaluate classifier performance
395     scores = []
396     for train, test in skf:
397         train_data = [X[ind] for ind in train]
398         train_labels = [Y[ind] for ind in train]

```

```

399     test_data = [X[ind] for ind in test]
400     test_labels = [Y[ind] for ind in test]
401
402     #Check performance of input forest
403     clf.fit(train_data, train_labels)
404     score = clf.score(test_data, test_labels)
405     scores.append(score)
406
407     #Compute baseline
408     most_common_label = max(Counter(Y).values())
409     baseline = float(most_common_label)/N_samples
410
411     #Compare results with prediction baseline
412     score_mean = np.mean(scores)/baseline
413     score_std = np.std(scores)/baseline
414
415     return (score_mean, score_std)
416
417 def check_regressor(X, Y, reg, strata = None):
418     '''Checks the performance of a regressor against mean value baseline.'''
419     N_samples = len(Y)
420     #Set up validation indices
421     if not strata: #Do leave-one-out validation
422         skf = KFold(N_samples, n_folds=N_samples,
423                     shuffle = False)
424     else: #Do stratified K-fold
425         skf = StratifiedKFold(Y, n_folds=strata)
426
427     #Evaluate performance
428     model_abs_errors = []
429     baseline_abs_errors = []
430     for train, test in skf:
431         train_data = [X[ind] for ind in train]
432         train_labels = [Y[ind] for ind in train]
433         test_data = [X[ind] for ind in test]
434         test_labels = [Y[ind] for ind in test]
435
436         #Check performance of input forest
437         reg.fit(train_data, train_labels)
438         base = np.mean(train_labels)
439         for i in xrange(len(test_data)):
440             pred = reg.predict(test_data[i])
441             true = test_labels[i]
442             model_abs_errors.append(np.abs(pred - true))
443             baseline_abs_errors.append(np.abs(base - true))
444         #
445     return (np.mean(model_abs_errors), np.mean(baseline_abs_errors))
446
447 class _RFNN(object):
448     __metaclass__ = abc.ABCMeta
449     '''Abstract class for random forest nearest neighbor predictors.
450     This should never be instantiated.'''
451
452     def __init__(self, forest, n_neighbors):
453         self.forest = forest
454         self.n_neighbors = n_neighbors
455         self.X = None
456         self.Y = None
457
458     def fit(self, X, Y):
459         '''Fits model to training data.
460

```

```

461     Args
462     -----
463     X : List
464         List of training feature vectors.
465
466     Y : List
467         List of training labels or values to be predicted.'''
468     if not len(X) == len(Y):
469         raise ValueError("Training input and output lists must have "
470                          "same length.")
471     if not self.n_neighbors <= len(X):
472         raise ValueError("Fewer data points than neighbors.")
473
474     self.forest.fit(X, Y)
475     self.X = X
476     self.Y = Y
477
478     def _rf_similarity(self, a, b):
479         '''Computes a similarity measure for two points using a trained random
480          forest classifier.'''
481         if self.X == None or self.Y == None:
482             raise NotImplementedError("Model has not been fitted to data yet.")
483
484         #Feature vectors must be single precision.
485         a = np.array([a], dtype = np.float32)
486         b = np.array([b], dtype = np.float32)
487         hits = 0
488         tries = 0
489         for estimator in self.forest.estimators_:
490             tries += 1
491             tree = estimator.tree_
492             # Check whether the points end up on the same leaf for this tree
493             if tree.apply(a) == tree.apply(b):
494                 hits += 1
495             #
496         return hits/tries
497
498     def find_neighbors(self, point):
499         '''Determine the n nearest neighbors for the given point.
500          Returns a list of n tuples like (yval, similarity).
501          The tuples are sorted descending by similarity.'''
502         if self.X == None or self.Y == None:
503             raise NotImplementedError("Model has not been fitted to data yet.")
504
505         #Get list of tuples like (y, similarity) for the n 'nearest' points
506         nearest = [(None, float('-infinity')) for _ in xrange(self.n_neighbors)]
507         for i in xrange(len(self.X)):
508             similarity = self._rf_similarity(self.X[i], point)
509             # update top n list if more similar than the furthest neighbor
510             if similarity > nearest[-1][1]:
511                 nearest.append((self.Y[i], similarity))
512                 nearest.sort(key = lambda x: x[1], reverse = True)
513                 del nearest[-1]
514             #
515         return nearest
516
517     #Mandatory methods – must be overridden
518     @abc.abstractmethod
519     def predict(self, point):
520         pass
521
522     @abc.abstractmethod

```

```

523     def score(self, X, Y):
524         pass
525
526     def _reservoir_sampler(start = 1):
527         '''Generator of the probabilities need to do reservoir sampling. The point
528         it that this can be used to iterate through a list, discarding each element
529         for the following element with probability P_n and ending up with a random
530         element from the list.'''
531         n = start
532         while True:
533             p = 1/n
534             r = random.uniform(0,1)
535             if r < p:
536                 yield True
537             else:
538                 yield False
539             n += 1
540
541     class RFNNClassifier(_RFNN):
542         '''Random Forest Nearest Neighbor Classifier.
543
544         Parameters
545         ----
546         n_neighbors : int
547             Number of neighbors to consider.
548
549         forest : RandomForestClassifier
550             The forest which will provide a
551             distance measure on which determine nearest neighbors.
552
553         weighting : str
554             How to weigh the votes of different neighbors.
555             'equal' means each neighbor has an equivalent vote.
556             'linear' mean votes are weighed by their similarity to the input point.
557         '''
558     def predict(self, point):
559         '''Predicts the label of a given point.'''
560         neighbortuples = self.find_neighbors(point)
561         if self.weighting == 'equal':
562             #Simple majority vote. Select randomly if it's a tie.
563             predictions = [t[0] for t in neighbortuples]
564             best = 0
565             winner = None
566             switch = _reservoir_sampler(start = 2)
567             for label, votes in Counter(predictions).iteritems():
568                 if votes > best:
569                     best = votes
570                     winner = label
571                     switch = _reservoir_sampler(start = 2)
572                 elif votes == best:
573                     if switch.next():
574                         winner = label
575                 else:
576                     pass
577             else:
578                 pass
579             #
580             return winner
581
582         #Weigh votes by their similarity to the input point
583         elif self.weighting == 'linear':
584             #The votes are weighted by their similarity

```



```

585         d = {}
586         for yval, similarity in neighbortuples:
587             try:
588                 d[yval] += similarity
589             except KeyError:
590                 d[yval] = similarity
591         best = float('-infinity')
592         winner = None
593         for k, v in d.iteritems():
594             if v > best:
595                 best = v
596                 winner = k
597             else:
598                 pass
599         return winner
600
601     def score(self, X, Y):
602         if not len(X) == len(Y):
603             raise ValueError("Training data and labels must have same length.")
604
605         hits = 0
606         n = len(X)
607         for i in xrange(n):
608             pred = self.predict(X[i])
609             if pred == Y[i]:
610                 hits += 1
611         #
612         return hits/n
613
614
615     def __init__(self, forest = None, n_neighbors = 3, weighting = 'equal',
616                 n_jobs = 1):
617         #Make sure we have a forest * classifier *
618         if forest == None:
619             forest = RandomForestClassifier(n_estimators = 1000,
620                                           criterion = 'entropy',
621                                           n_jobs = n_jobs)
622         if not isinstance(forest, RandomForestClassifier):
623             raise TypeError("Forest must be a classifier")
624
625         self.weighting = weighting
626
627         #Call parent constructor
628         super(RFNNClassifier, self).__init__(forest, n_neighbors)
629
630 class RFNNRegressor(_RFNN):
631     '''Random Forest Nearest Neighbor Regressor.
632
633     Parameters
634     ----
635     n_neighbors : int
636         Number of neighbors to consider.
637
638     forest : RandomForestRegressor
639         The forest which will provide a
640         distance measure on which determine nearest neighbors.
641
642     weighting : str
643         How to weigh the votes of different neighbors.
644         'equal' means each neighbor has an equivalent weight.
645         'linear' mean votes are weighed by their similarity to the input point.
646     '''

```

```

647 def predict(self, point):
648     # lists of the y vaues and similarities of nearest neighbors
649     neighbortuples = self.find_neighbors(point)
650     yvals, similarities = zip(*neighbortuples)
651
652     # Weigh each neighbor y value equally is that's how we roll
653     if self.weighting == 'equal':
654         weight = 1.0/len(yvals)
655         result = 0.0
656         for y in yvals:
657             result += y*weight
658         return result
659     #
660     # Otherwise, weigh neighbors by similarity
661     elif self.weighting == 'linear':
662         weight = 1.0/(sum(similarities))
663         result = 0.0
664         for i in xrange(len(yvals)):
665             y = yvals[i]
666             similarity = similarities[i]
667             result += y*similarity*weight
668         return result
669
670
671 def score(self, X, Y):
672     if not len(X) == len(Y):
673         raise ValueError("X and Y must be same length.")
674     errors = [Y[i] - self.predict(X[i]) for i in xrange(len(X))]
675     return np.std(errors)
676
677
678 def __init__(self, forest = None, n_neighbors = 3, weighting = 'equal'):
679     #Check forest type.
680     if forest == None:
681         forest = RandomForestRegressor(n_estimators = 1000, n_jobs = -1)
682     if not isinstance(forest, RandomForestRegressor):
683         raise TypeError("Must use Random Forest Regressor to initialize.")
684     # Set params
685     self.weighting = weighting
686     # Done. Call parent constructor
687     super(RFNNRegressor, self).__init__(forest, n_neighbors)
688
689
690 class _BaselineRegressor(object):
691     '''Always predicts the mean of the training set.'''
692     def __init__(self, guess=None):
693         self.guess = guess
694     def fit(self, xtrain, ytrain):
695         '''Find the average of input list of target values and guess on that
696         from now on.'''
697         self.guess = np.mean(ytrain)
698     def predict(self, x):
699         return self.guess
700
701
702 class _BaselineClassifier(object):
703     '''Always predicts the most common label in the training set'''
704     def __init__(self, guess=None):
705         self.guess = guess
706     def fit(self, xtrain, ytrain):
707         '''Find the most common label and guess on that from now on.'''
708         countmap = Counter(ytrain)

```

```

709         best = 0
710         for label, count in countmap.iteritems():
711             if count > best:
712                 best = count
713                 self.guess = int(label)
714         #
715         #
716     def predict(self, x):
717         return self.guess
718
719
720
721
722 def _worker(X, Y, score_type, train_percentage, classifier, clf_args, n_groups,
723            replace):
724     '''Worker method for parallelizing bootstrap evaluations.'''
725     #Create bootstrap sample
726     try:
727         rand = np.random.RandomState() #Ensures PRNG works in children
728         indices = rand.choice(xrange(len(X)), size = len(X), replace = replace)
729         rand.randint
730         xsample = [X[i] for i in indices]
731         ysample = [Y[i] for i in indices]
732         #Create regressor if we're doing regression
733         if classifier == 'RandomForestRegressor':
734             clf = RandomForestRegressor(**clf_args)
735         elif classifier == 'SVR':
736             clf = svm.SVR(**clf_args)
737         elif classifier == 'baseline_mean':
738             clf = _BaselineRegressor()
739
740         #Create classifier and split dataset into labels
741         elif classifier == 'RandomForestClassifier':
742             clf = RandomForestClassifier(**clf_args)
743             ysample = assign_labels(ysample, n_groups)
744         elif classifier == 'SVC':
745             clf = svm.SVC(**clf_args)
746             ysample = assign_labels(ysample, n_groups)
747         elif classifier == 'baseline_most_common_label':
748             clf = _BaselineClassifier()
749             ysample = assign_labels(ysample, n_groups)
750         #Fail if none of the above classifiers were specified
751         else:
752             raise ValueError('Regressor or classifier not defined.')
753         #Generate training and testing set
754         cut = int(train_percentage*len(X))
755         xtrain = xsample[:cut]
756         ytrain = ysample[:cut]
757         xtest = xsample[cut:]
758         ytest = ysample[cut:]
759
760         #Fit the classifier or regressor
761         clf.fit(xtrain, ytrain)
762
763         #Compute score and append to output list
764         if score_type == 'mse':
765             scores = [(ytest[i] - clf.predict(xtest[i]))**2
766                       for i in xrange(len(xtest))]
767             return np.mean(scores)
768
769         elif score_type == 'fraction_correct':
770             n_correct = sum([ytest[i] == int(clf.predict(xtest[i]))

```

```

771         for i in xrange(len(ytest)))
772         score = n_correct/len(ytest)
773         return score
774     elif score_type == 'over_baseline':
775         #Get score
776         score = sum([ytest[i] == int(clf.predict(xtest[i]))
777                        for i in xrange(len(ytest))])
778         #Get baseline
779         baselineclf = _BaselineClassifier()
780         baselineclf.fit(xtrain, ytrain)
781         baseline = sum([ytest[i] == baselineclf.predict(xtest[i])
782                        for i in xrange(len(ytest))])
783         return score/baseline
784
785     #Fail if none of the above performance metrics were specified
786     else:
787         raise ValueError('Score type not defined.')
788
789     #Job's done!
790     return None
791 except:
792     raise Exception(''.join(traceback.format_exception(*sys.exc_info())))
793
794
795 def bootstrap(X, Y, classifier, score_type = 'mse', train_percentage = 0.8,
796              clf_args = {}, iterations = 1000, n_groups = 3, n_jobs = 1,
797              replace = True):
798     '''Performs bootstrap resampling to evaluate the performance of some
799     classifier or regressor. Note that this takes the *complete dataset* as
800     arguments as well as arguments specifying which predictor to use and which
801     function to estimate the distribution of.
802     This seems to be the most straightforward generalizable implementation
803     which can be parallelized, as passing e.g. the scoring function directly
804     clashed with the mechanisms implemented to work around the GIL for
805     multiprocessing for obscure reasons.
806
807     Parameters:
808     -----
809     X : list
810         All feature vectors in the complete dataset.
811
812     Y : list
813         All 'true' labels or output values in the complete dataset.
814
815     classifier : str
816         Which classifier to use to predict the test set. Allowed values:
817         'RandomForestRegressor', 'baseline_mean', 'SVR',
818         'RandomForestClassifier', 'SVC', 'baseline_most_common_label'
819
820     score_type : str
821         String signifying which function to estimate the distribution of.
822         Allowed values: 'mse', 'fraction_correct', 'over_baseline'
823
824     train_percentage : float
825         The percentage [0:1] of each bootstrap sample to be used for training.
826
827     clf_args : dict
828         optional arguments to the constructor method of the regressor/classifier.
829
830     iterations : int
831         Number of bootstrap samples to run.
832

```

```

833     n_jobs : int
834         How many cores (maximum) to use.
835
836     replace : bool
837         Whether to sample with replacement when obtaining the bootstrap samples.
838     '''
839
840     if not len(X) == len(Y):
841         raise ValueError("X and Y must have equal length.")
842
843     #Arguments to pass to worker processes
844     d = {'X' : X, 'Y' : Y, 'train_percentage' : train_percentage,
845         'classifier' : classifier, 'clf_args' : clf_args,
846         'score_type' : score_type, 'n_groups' : n_groups,
847         'replace' : replace}
848
849     #Make job queue
850     pool = multiprocessing.Pool(processes = n_jobs)
851     jobs = [pool.apply_async(_worker, kwds = d) for _ in xrange(iterations)]
852     pool.close() #run
853     pool.join() #Wait for remaining jobs
854
855     #Make sure no children died too early
856     if not all(job.successful() for job in jobs):
857         raise RuntimeError('Some jobs failed.')
858
859     return [j.get() for j in jobs]
860
861 def get_correlations(X, Y):
862     '''Given a list of feature vectors X and labels or values Y, returns a list
863     of correlation coefficients for each dimension of the feature vectors.'''
864     n_feats = len(X[0])
865     correlations = []
866     for i in xrange(n_feats):
867         temp = np.corrcoef([x[i] for x in X], Y)
868         correlation = temp[0,1]
869         if math.isnan(correlation):
870             correlation = 0
871         correlations.append(correlation)
872     return correlations
873
874
875 def make_kernel(correlations = None, gamma = 1.0, threshold = 0.0):
876     '''Returns a weighted radial basis function (WRBF) kernel.'''
877     def kernel(x,y, *args, **kwargs):
878         if correlations == None:
879             _corrs = np.ones(shape = (len(x)), dtype = np.float64)
880         else:
881             _corrs = correlations
882         d = len(_corrs) #number of features
883         #Strong (above threshold) correlations
884         strong = [np.abs(c) if np.abs(c) >= threshold else 0.0
885                 for c in _corrs]
886         normfactor = 1.0/np.sqrt(sum([e**2 for e in strong]))
887         #Metric to compute distance between points
888         metric = dok_matrix((d,d), dtype = np.float64)
889         for i in xrange(d):
890             metric[i,i] = strong[i]*normfactor
891         #
892         result = np.zeros(shape = (len(x), len(y)))
893         for i in xrange(len(x)):
894             for j in xrange(len(y)):

```

```

895         dist = x[i] - y[j]
896         result[i,j] = np.exp(-gamma*np.dot(dist,dist))
897     return result
898 return kernel
899
900 if __name__ == '__main__':
901     pass
902     X, Y, ind_dict = read_data('../data.json', trait = 'openness',
903                               features = ['call_iet_med', 'text_iet_med', 'social_entropy', 'call_entropy', '←
904                               travel_med', 'n_places', 'text_latency', 'call_night_activity'],
905                               n_classes = 3
906                               )
907     # X = [[1,2,7,0],[3,1,6,0.01],[6,8,1,0],[10,8,2,0.01]]
908     # Y = [1,1,0,0]
909     cut = int(0.8*len(X))
910
911     xtrain = X[:cut]
912     ytrain = Y[:cut]
913     xtest = X[cut:]
914     ytest = Y[cut:]
915
916     corrs = get_correlations(X, Y)
917     print corrs
918
919     C = 70
920     gamma = 3.75
921
922     kernel = make_kernel(corrs, 0.05)
923
924     clf = svm.SVC(kernel = kernel)
925     clf.fit(xtrain, ytrain)
926
927     hits = 0
928
929     for i in xrange(len(xtest)):
930         if clf.predict(xtest[i]) == ytest[i]:
931             hits += 1
932     #
933     print 100.0*hits/len(ytest)
934
935     for i in xrange(len(corrs)):
936         print ind_dict[i], corrs[i]
937
938
939
940 # print len(X[0])
941 ## print i
942 # print [el for el in i.values() if 'init' in el]

```

A.2.5 Lloyd's Algorithm

```

1 # -*- coding: utf-8 -*-
2 from __future__ import division
3
4 import numpy as np
5 import matplotlib

```

```

6 matplotlib.use('Agg') #ugly hack to allow plotting from terminal
7 import matplotlib.pyplot as plt
8 import random
9 from copy import deepcopy
10
11 def _dist(p,q):
12     return sum([(p[i]-q[i])**2 for i in xrange(len(p))])
13
14 def _lloyds_single_run(X, K, max_iterations, init):
15     # Initialize with a subset of the data
16     if init == 'sample':
17         initials = random.sample(X, K)
18     # Or initialize with random points across the same range as data
19     elif init == 'scatter':
20         vals = zip(*X)
21         xmin = min(vals[0])
22         xmax = max(vals[0])
23         ymin = min(vals[1])
24         ymax = max(vals[1])
25         initials = [(random.uniform(xmin, xmax),
26                     random.uniform(ymin, ymax)) for _ in xrange(K)]
27     # Or yell RTFM at user
28     else:
29         raise ValueError('Invalid initialization mode!')
30
31     #Construct hashmap mapping integers up to K to centroids
32     centroids = dict(enumerate(initials))
33     converged = False
34     iterations = 0
35
36     while not converged and iterations < max_iterations:
37         clusters = {i: [] for i in xrange(K)}
38         #Make sure clusters and centroids have identical keys, or we're doomed.
39         assert set(clusters.keys()) == set(centroids.keys())
40         prev_centroids = deepcopy(centroids)
41
42         ### STEP ONE -update clusters
43         for x in X:
44             #Check distances to all centroids
45             bestind = -1
46             bestdist = float('inf')
47             for ind, centroid in centroids.iteritems():
48                 dist = _dist(x, centroid)
49                 if dist < bestdist:
50                     bestdist = dist
51                     bestind = ind
52             #
53             clusters[bestind].append(x)
54
55         ### STEP TWO -update centroids
56         for ind, points in clusters.iteritems():
57             if not points:
58                 pass #Cluster's empty - nothing to update
59             else:
60                 centroids[ind] = np.mean(points, axis = 0)
61
62         ### We're converged when all old centroids = new centroids.
63         converged = all([_dist(prev_centroids[k], centroids[k]) == 0
64                         for k in xrange(K)])
65         iterations += 1
66         #
67     return {tuple(centroids[i]): clusters[i] for i in xrange(K)}

```

```

68
69 def lloyds(X, K, runs = 1, max_iterations = float('inf'), init = 'sample'):
70     '''Runs Lloyd's algorithm to identify K clusters in the dataset X.
71     X is a list of points like [[x1,y1],[x2,y2]---].
72     Returns a hash of centroids mapping to points in the corresponding cluster.
73     The objective is to minimize the sum of distances from each centroid to
74     the points in the corresponding cluster. It might only converge on a local
75     minimum, so the configuration with the lowest score (sum of distances) is
76     returned.
77     init denotes initialization mode, which can be 'sample', using a randomly
78     select subset of the input data, or 'scatter', using random points selected
79     from the same range as the data as initial centroids.
80
81     Parameters
82     -----
83     X : array_like
84         list of points. 2D example: [[3,4],[3.4, 7.2], ...]
85
86     K : int
87         Number of centroids
88
89     runs : int
90         Number of times to run the entire algorithm. The result with the lowest
91         score will be returned.
92
93     max_iterations : int or float
94         Number of steps to allow each run. Default if infinit, i.e. the algorithm
95         runs until it's fully converged.
96
97     init : str
98         Initialization mode. 'sample' means use a random subset of the data as
99         starting centroids. 'scatter' means place starting centroids randomly in
100         the entire x-y range of the dataset.
101
102     Returns
103     -----
104     result : dict
105         A dictionary in which each key is a tuple of coordinated corresponding to
106         a centroid, and each value is a list of points belonging to that cluster.
107         '''
108
109     record = float('inf')
110     result = None
111     for _ in xrange(runs):
112         clusters = _lloyds_single_run(X, K, max_iterations = max_iterations,
113                                     init = init)
114         #Determine how good the clusters came out
115         score = 0
116         for centroid, points in clusters.iteritems():
117             score += sum([_dist(centroid, p) for p in points or []])
118         if score < record:
119             result = clusters
120             record = score
121         #
122     return result
123
124
125 def _makecolor():
126     i = 0
127     cols = ['b', 'g', 'r', 'c', 'm', 'y']
128     while True:
129         yield cols[i]

```



```

130         i = (i+1)%len(cols)
131
132
133     def draw_clusters(clusters, threshold=0, show=True, filename=None):
134         '''Accepts a dict mapping cluster centroids to cluster points and makes
135         a color-coded plot of them. Clusters containing fewer points than the
136         threshold are plotted in black.'''
137         colors = _makecolor()
138         plt.figure()
139         for centroid, points in clusters.iteritems():
140             if not points:
141                 continue
142             if len(points) < threshold:
143                 style = ['k,']
144             else:
145                 color = colors.next()
146                 style = [color+'+']
147                 #Plot centroids
148                 x,y = centroid
149                 plt.plot(x,y, color = color, marker = 'd', markersize = 12)
150                 #plot points
151                 plt.plot(*(zip(*points))+style))
152             if filename:
153                 plt.savefig(filename, bbox_inches = 'tight')
154             if show:
155                 plt.show()
156
157
158 if __name__ == '__main__':
159     points = [[random.uniform(-10,10), random.uniform(-10,10)] for _ in xrange(10*3)]
160     clusters = lloyds(X = points, K = 6, runs = 1)
161     draw_clusters(clusters = clusters, filename = 'lloyds_example.pdf')

```

A.2.6 Smallest Enclosing Circle

```

1  # -*- coding: utf-8 -*-
2  #
3  # Smallest enclosing circle
4  #
5  # Copyright (c) 2014 Project Nayuki
6  # http://www.nayuki.io/page/smallest-enclosing-circle
7  #
8  # This program is free software: you can redistribute it and/or modify
9  # it under the terms of the GNU General Public License as published by
10 # the Free Software Foundation, either version 3 of the License, or
11 # (at your option) any later version.
12 #
13 # This program is distributed in the hope that it will be useful,
14 # but WITHOUT ANY WARRANTY; without even the implied warranty of
15 # MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
16 # GNU General Public License for more details.
17 #
18 # You should have received a copy of the GNU General Public License
19 # along with this program (see COPYING.txt).
20 # If not, see <http://www.gnu.org/licenses/>.
21 #
22

```

```

23 import math, random
24
25
26 # Data conventions: A point is a pair of floats (x, y). A circle is a triple of floats (center x, center y, radius).
27
28 #
29 # Returns the smallest circle that encloses all the given points. Runs in expected O(n) time, randomized.
30 # Input: A sequence of pairs of floats or ints, e.g. [(0,5), (3.1,-2.7)].
31 # Output: A triple of floats representing a circle.
32 # Note: If 0 points are given, None is returned. If 1 point is given, a circle of radius 0 is returned.
33 #
34 def make_circle(points):
35     '''Accepts list of points as tuples and returns (x, y, r).'''
36     # Convert to float and randomize order
37     shuffled = [(float(p[0]), float(p[1])) for p in points]
38     random.shuffle(shuffled)
39
40     # Progressively add points to circle or recompute circle
41     c = None
42     for (i, p) in enumerate(shuffled):
43         if c is None or not _is_in_circle(c, p):
44             c = _make_circle_one_point(shuffled[0:i + 1], p)
45     return c
46
47
48 # One boundary point known
49 def _make_circle_one_point(points, p):
50     c = (p[0], p[1], 0.0)
51     for (i, q) in enumerate(points):
52         if not _is_in_circle(c, q):
53             if c[2] == 0.0:
54                 c = _make_diameter(p, q)
55             else:
56                 c = _make_circle_two_points(points[0:i + 1], p, q)
57     return c
58
59
60 # Two boundary points known
61 def _make_circle_two_points(points, p, q):
62     diameter = _make_diameter(p, q)
63     if all(_is_in_circle(diameter, r) for r in points):
64         return diameter
65
66     left = None
67     right = None
68     for r in points:
69         cross = _cross_product(p[0], p[1], q[0], q[1], r[0], r[1])
70         c = _make_circumcircle(p, q, r)
71         if c is None:
72             continue
73         elif cross > 0.0 and (left is None or _cross_product(p[0], p[1], q[0], q[1], c[0], c[1]) > _cross_product(←
74             p[0], p[1], q[0], q[1], left[0], left[1])):
75             left = c
76         elif cross < 0.0 and (right is None or _cross_product(p[0], p[1], q[0], q[1], c[0], c[1]) < _cross_product(←
77             p[0], p[1], q[0], q[1], right[0], right[1])):
78             right = c
79     return left if (right is None or (left is not None and left[2] <= right[2])) else right
80
81 def _make_circumcircle(p0, p1, p2):
82     # Mathematical algorithm from Wikipedia: Circumscribed circle
83     ax = p0[0]; ay = p0[1]

```

```

83     bx = p1[0]; by = p1[1]
84     cx = p2[0]; cy = p2[1]
85     d = (ax * (by - cy) + bx * (cy - ay) + cx * (ay - by)) * 2.0
86     if d == 0.0:
87         return None
88     x = ((ax * ax + ay * ay) * (by - cy) + (bx * bx + by * by) * (cy - ay) + (cx * cx + cy * cy) * (ay - by)) / d
89     y = ((ax * ax + ay * ay) * (cx - bx) + (bx * bx + by * by) * (ax - cx) + (cx * cx + cy * cy) * (bx - ax)) / d
90     return (x, y, math.hypot(x - ax, y - ay))
91
92
93 def _make_diameter(p0, p1):
94     return ((p0[0] + p1[0]) / 2.0, (p0[1] + p1[1]) / 2.0, math.hypot(p0[0] - p1[0], p0[1] - p1[1]) / 2.0)
95
96
97 _EPSILON = 1e-12
98
99 def _is_in_circle(c, p):
100     return c is not None and math.hypot(p[0] - c[0], p[1] - c[1]) < c[2] + _EPSILON
101
102
103 # Returns twice the signed area of the triangle defined by (x0, y0), (x1, y1), (x2, y2)
104 def _cross_product(x0, y0, x1, y1, x2, y2):
105     return (x1 - x0) * (y2 - y0) - (y1 - y0) * (x2 - x0)
106
107 #pts = [(0.0, 0.0), (6.0, 8.0)]
108 #
109 #test = make_circle(pts)
110 #
111 #print test

```


BIBLIOGRAPHY

- [1] Yoshua Bengio and Yves Grandvalet. No Unbiased Estimator of the Variance of K-Fold Cross-Validation. 5:1089–1105, 2004.
- [2] Avrim Blum, Adam Kalai, and John Langford. Beating the Hold-out: Bounds for {K}-fold and Progressive Cross-Validation. *Proceedings of the 12th Annual Conference on Computational Learning Theory*, (c):203–208, 1999.
- [3] L Breiman. Random forests. *Machine learning*, pages 5–32, 2001.
- [4] Cjc Christopher J C Burges. A Tutorial on Support Vector Machines for Pattern Recognition. *Data Mining and Knowledge Discovery*, 2(2):121–167, 1998.
- [5] Yves Alexandre De Montjoye, Jordi Quoidbach, Florent Robic, and Alex Pentland. Predicting personality using novel mobile phone-based metrics. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, 7812 LNCS:48–55, 2013.
- [6] JM John M Digman. Personality structure: Emergence of the five-factor model. *Annual review of psychology*, 41:417–440, 1990.
- [7] RA Fisher. The use of multiple measurements in taxonomic problems. *Annals of eugenics*, 1936.
- [8] E Gabrilovich and S Markovitch. Computing Semantic Relatedness Using Wikipedia-based Explicit Semantic Analysis. *IJCAI*, 2007.
- [9] Isabelle Guyon. Gene Selection for Cancer Classification. pages 389–422, 2002.
- [10] Marti a. Hearst, Susan T. Dumais, Edgar Osuna, John Platt, and Bernhard Schölkopf. Support vector machines. *IEEE Intelligent Systems and their Applications*, 13(4):18–28, 1998.

- [11] Thomas K Landauer, Peter W. Foltz, and Darrell Laham. An introduction to latent semantic analysis. *Discourse Processes*, 25(2-3):259–284, 1998.
- [12] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schutze. An Introduction to Information Retrieval. *Online*, (c):569, 2009.
- [13] Preslav Nakov. Latent semantic analysis of textual data. *Proceedings of the conference on Computer systems and technologies - CompSysTech '00*, pages 5031–5035, 2000.
- [14] AD Pietersma. Kernel Learning in Support Vector Machines using Dual-Objective Optimization. *Proceedings of the 23rd . . .*, 2011.
- [15] G. Salton, a. Wong, and C. S. Yang. A vector space model for automatic indexing. *Communications of the ACM*, 18(11):613–620, 1975.
- [16] Gerard Salton and Christopher Buckley. Term-weighting approaches in automatic text retrieval. *Information Processing & Management*, 24(5):513–523, 1988.
- [17] Philip Wolfe. A duality theorem for nonlinear programming. *Quarterly of applied mathematics*, 19(3):239–244, 1961.
- [18] PA Zandbergen and SJ Barbeau. Positional accuracy of assisted gps data from high-sensitivity gps-enabled mobile phones. *Journal of Navigation*, 2011.